



CSE 124/224 Week 9 Discussion

Patrick Liu



Overview and Goals

- Implement a subset of DynamoDB
 - No consistent hashing or joining/leaving
- Understand causality using vector clocks
- Learn to implement and debug basic distributed systems

Vector Clocks


- Methods
 - LessThan
 - Implements $<$ from lecture
 - Equals
 - Implements $=$ from lecture
 - Combine
 - Changes the current vector clock such that all clocks in the argument list as well as the clock this is called on are \leq the resulting vector clock
 - Concurrent
 - For vector clocks A and B, $A < B$ and $B < A$ are both false

Vector Clocks

- Internal representation
 - Must be able to associate nodeID with version
 - Choices: Fixed size vector clock, or dynamic size vector clock
 - Fixed-size: We know cluster_size, so we know how big vector clocks can get
 - Dynamic size: Add elements whenever you increment for a nodeID that isn't tracked yet. Need logic for if one vector clock tracks one nodeID, and another one doesn't.

Put(on one node)


- I recommend thinking about Put on a single node separately from a Put to W nodes
 - Easy to compartmentalize
- Put on a single server does:
 - Retrieve (Context, value) pairs associated with this key
 - Check for causality
 - Keep only those (Context, value) pairs that have no causal descendant
 - If new Context == one of existing Contexts, keep existing value



"x": (["A":1], "old"),
(["B":1], "concurrent")




"x",["A": 2], "value"




"x": (["A":1], "old"),
(["B":1], "concurrent")



"x",["A": 2], "value"



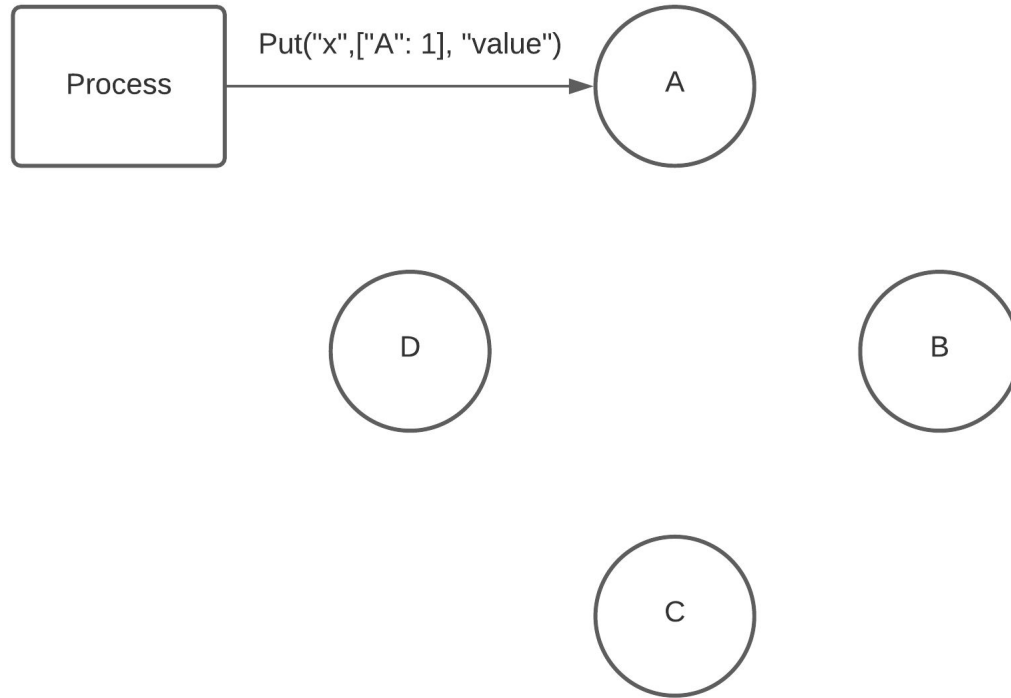
"x": (["A":2], "value"),
(["B":1], "concurrent")







"x": (["A":2, "value"),
(["B":1, "concurrent")

Put

- Put will do the following things:
 - Increment the incoming Context's version associated with this node
 - Store the (Context, value) pair in local key-value store
 - Try to store the the same (Context, value) pair at $W - 1$ other nodes, in the order of this node's preference list.
 - Return true if everything went well, otherwise return false







Put("X", ["A": 2], "value")










Put("x", ["A": 2], "value")



Put("x", ["A": 2], "value")







Get(from one node)

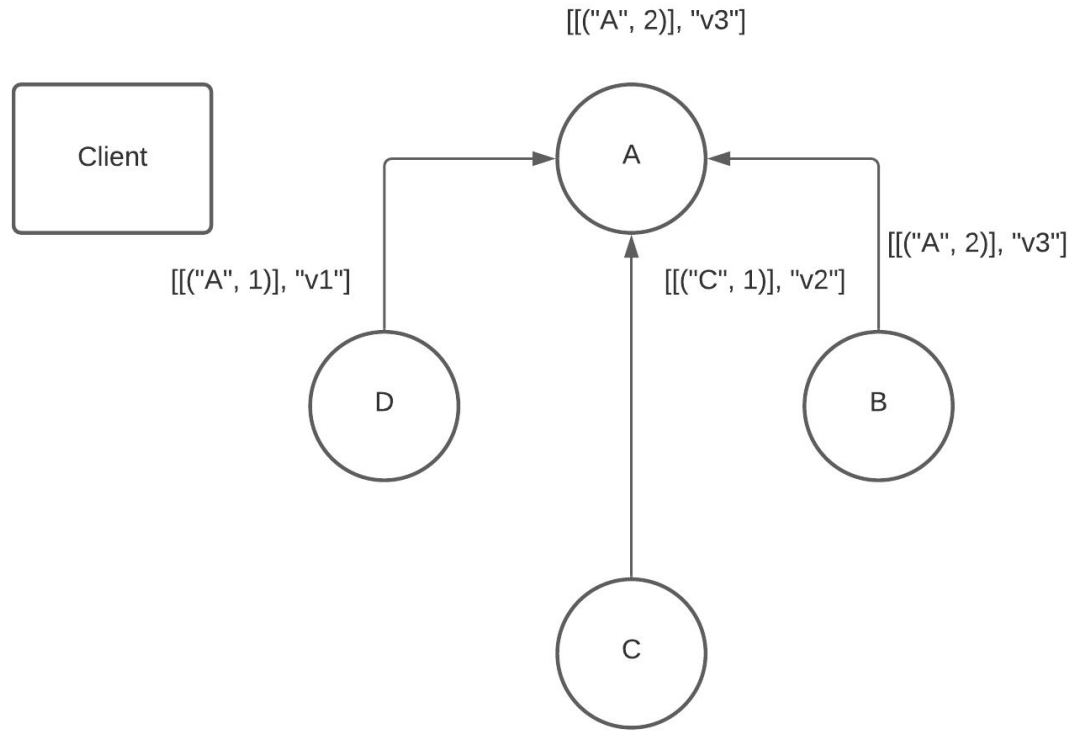
- Simply return from local key/value store
- No causality checks required(why?)

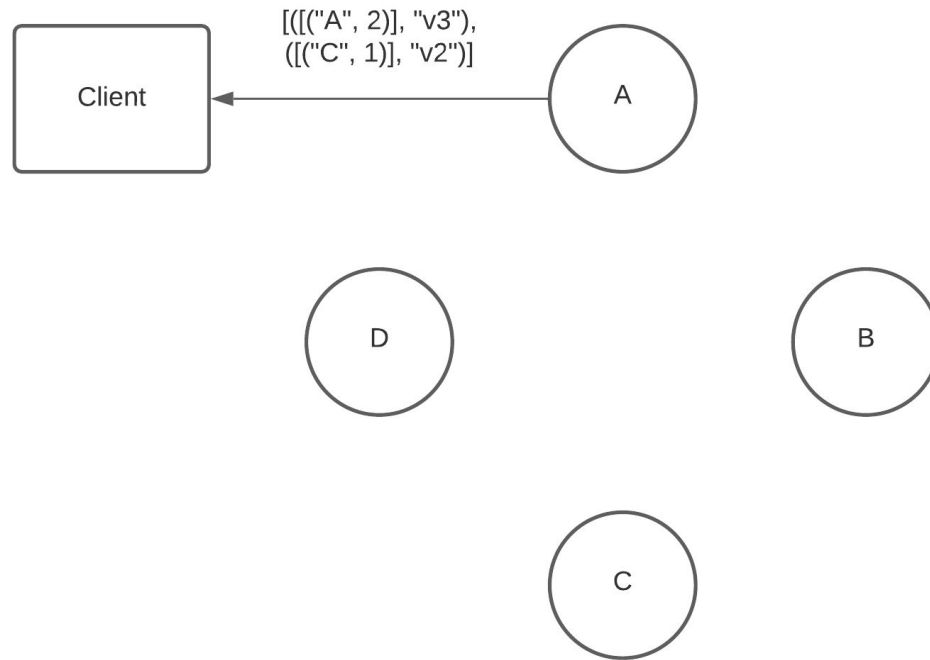
Get(from multiple nodes)

- Get from local storage, and $R - 1$ additional nodes
- Return only those elements that have no causal descendants.
 - This is mostly the same operation as in Put







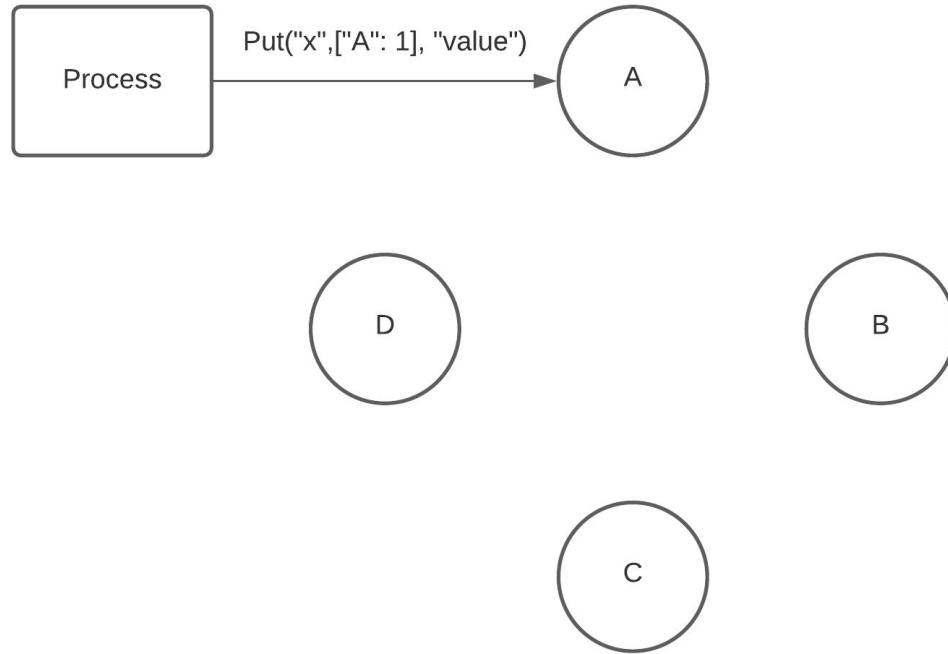


Gossip





- Replicates this node's key/value store to all other nodes.
- Each individual key replication should look pretty similar to Put to a single node
- Not two-way
 - The node running Gossip will not update its own store

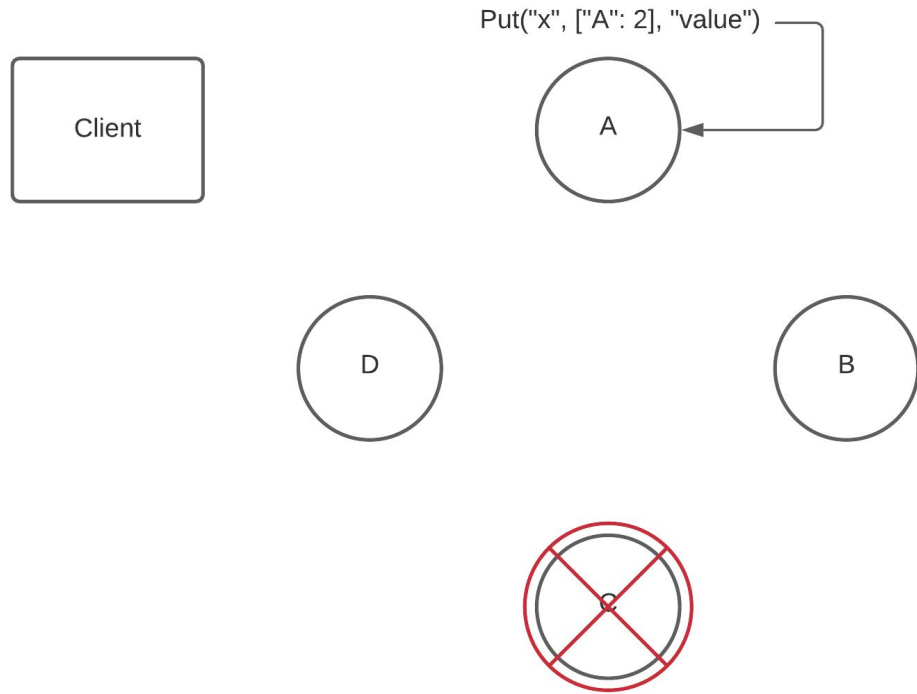
Crash

- Makes this node unresponsive to other RPC queries for some amount of seconds
- How to implement?
 - Likely have to keep some kind of crash state.
 - Crash state changes in only one place, but is checked everywhere.






Put("x", ["A": 2], "value")

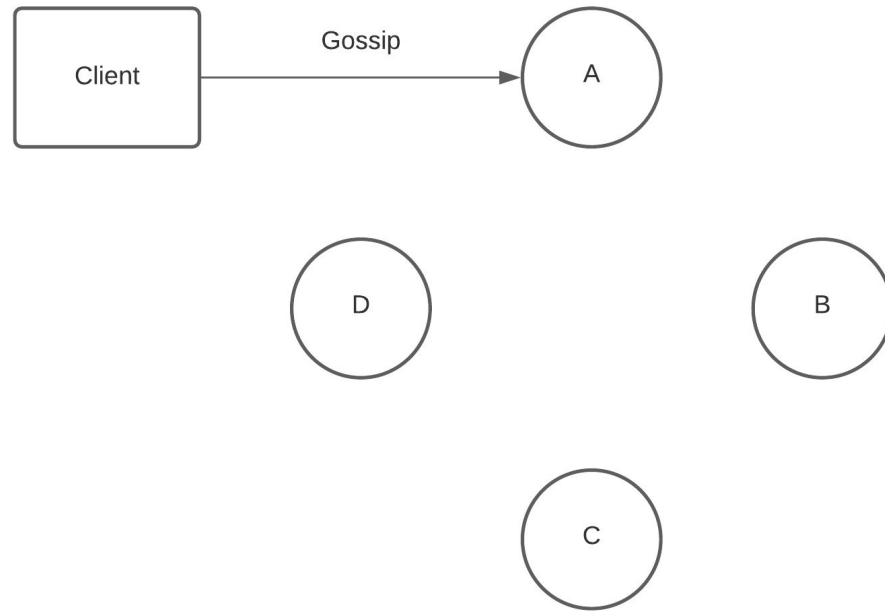


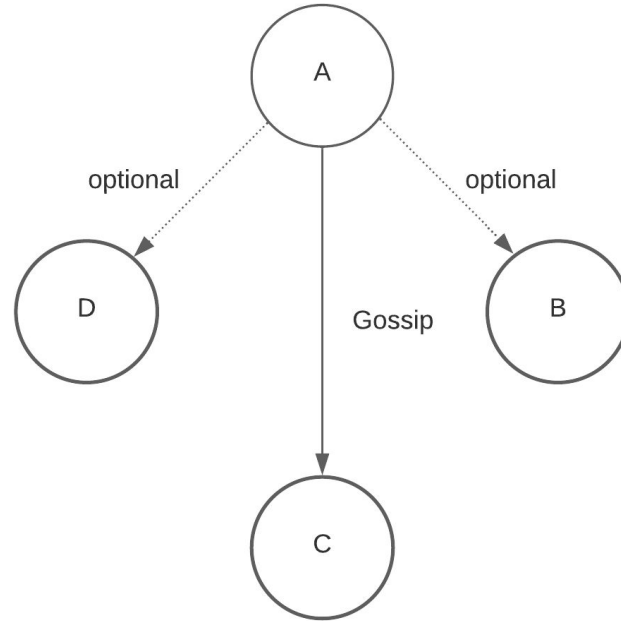


Put("x", ["A": 2], "value")









FAQ

- Failures in DynamoCoordinator
 - Most likely server isn't spun up when DynamoCoordinator tries to send preference list
 - Try adding a small wait
- Can I add struct fields/functions?
 - Yes! In fact, you will most likely need to do so.

Tips

- Don't worry about $W > 1$, $R > 1$ for your first iteration. Ensure your implementation works for $W = 1$, $R = 1$, then add functionality for $W > 1$, $R > 1$
- Use unit testing extensively. TDD is perhaps helpful here.
- Pay attention to your logic for determining causality: most of your bugs will likely occur here.