Containers

Guest Lecture

December 8th, 2020

UC San Diego

Shelby Thomas Ph.D.

 @realshelbyt

What can containers do?

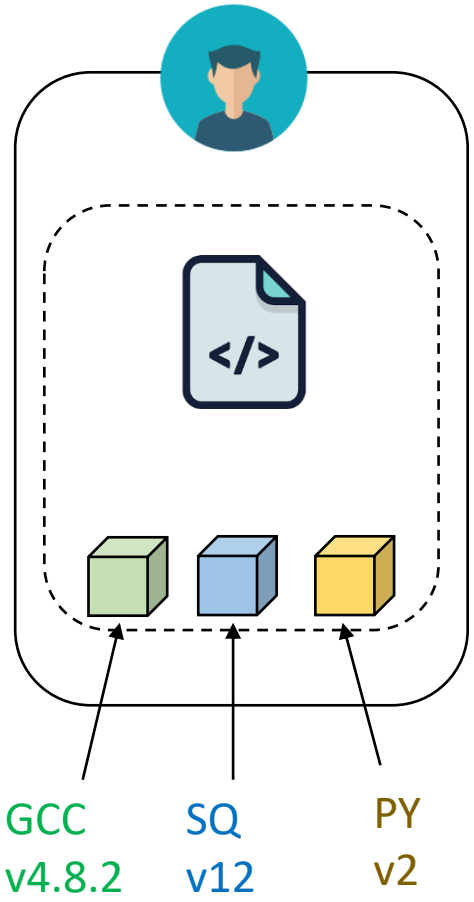

What can containers do?

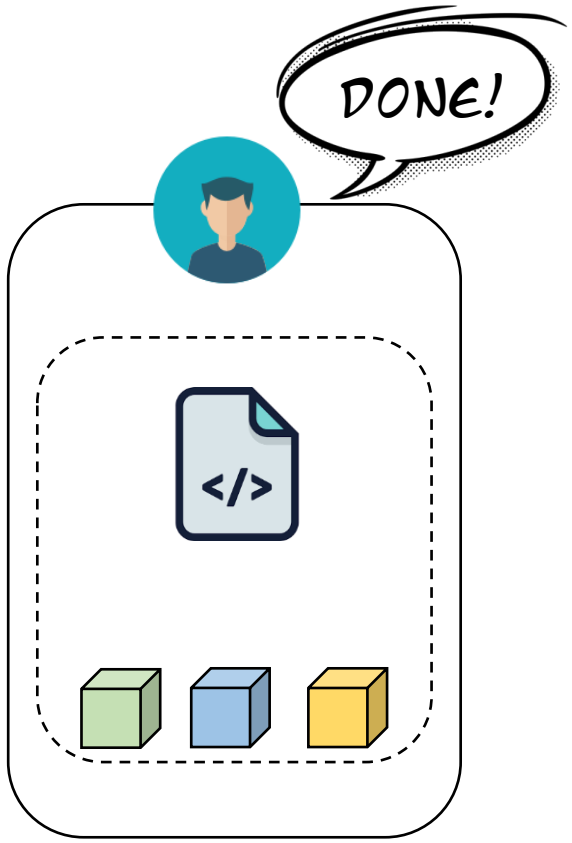

What can containers do?

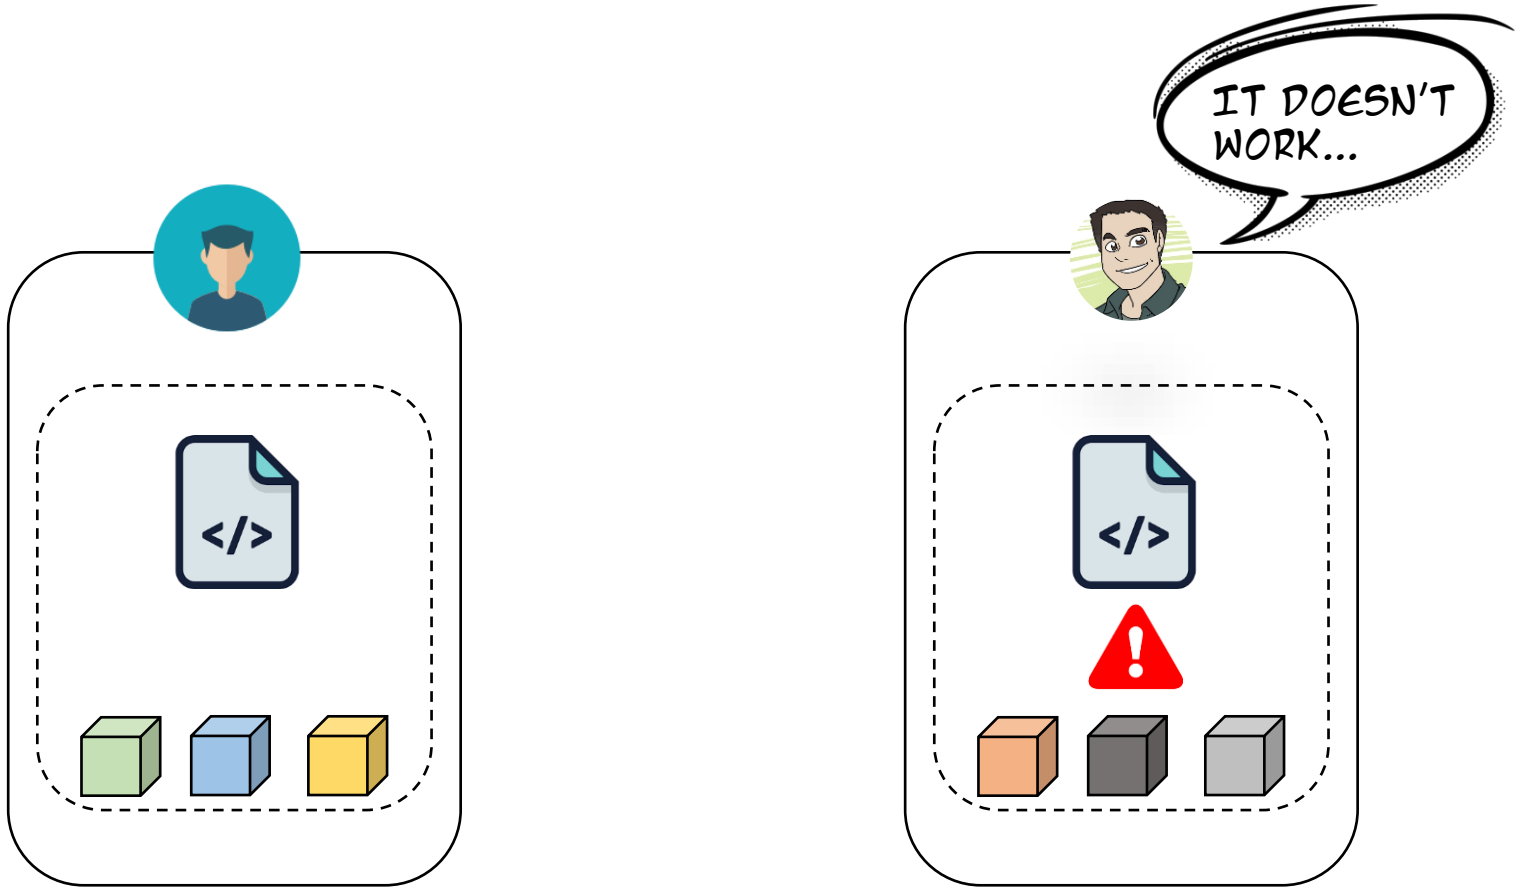

What can containers do?

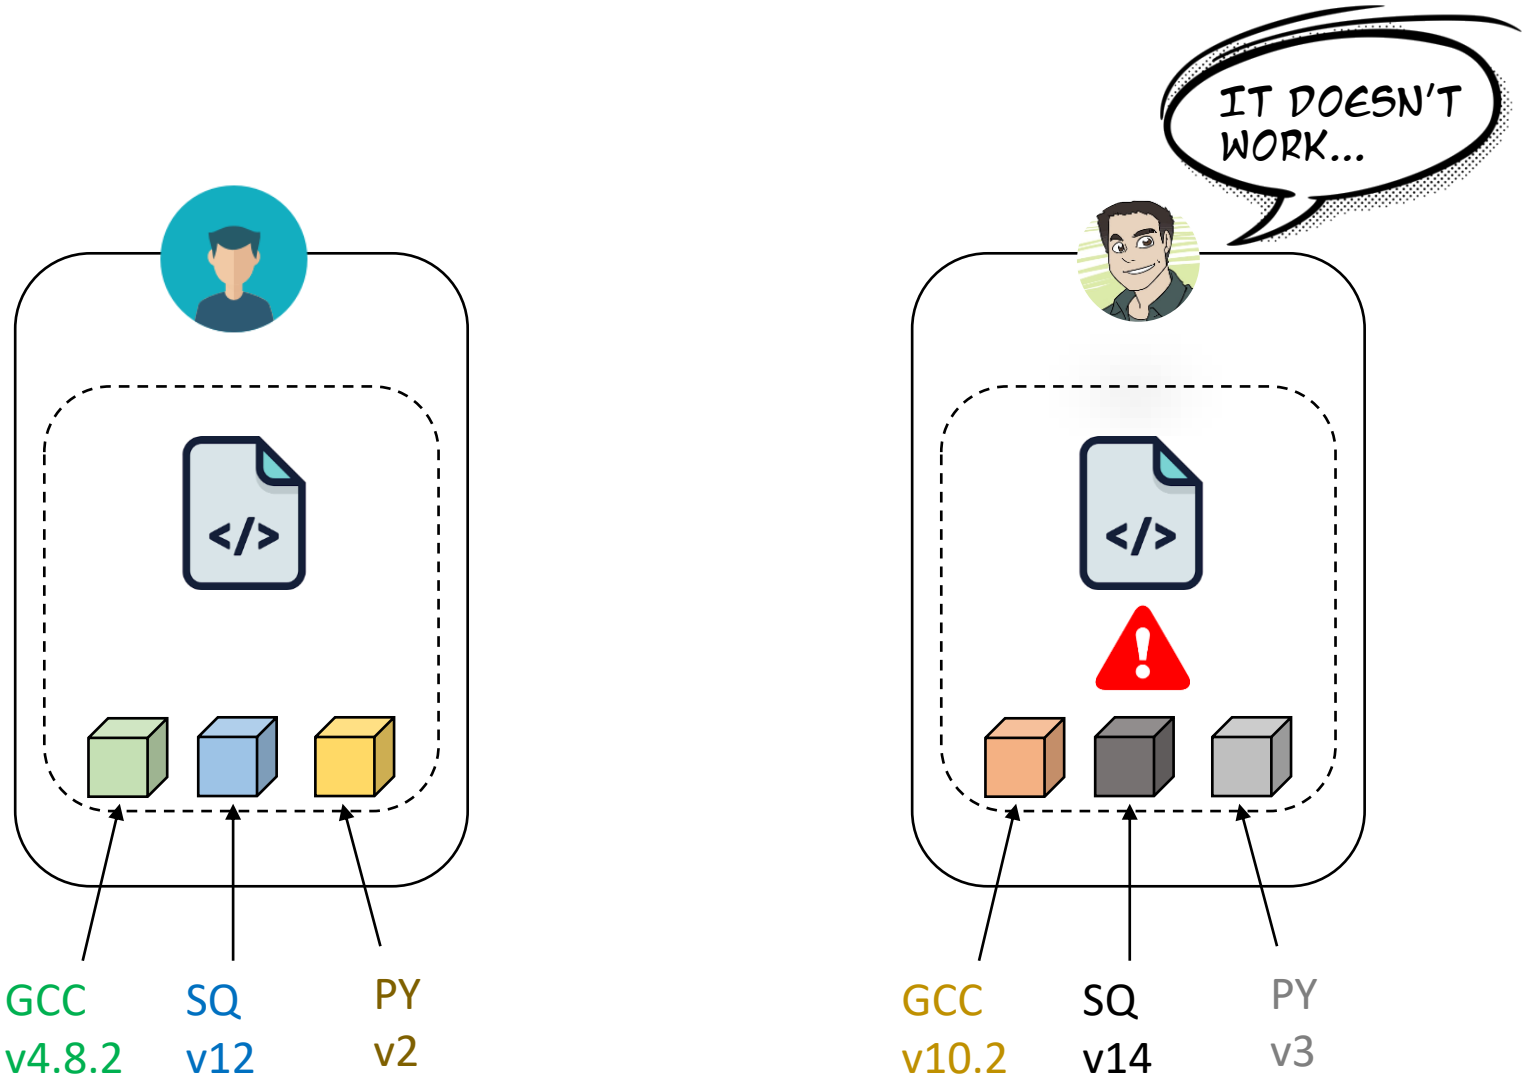

Scenario One

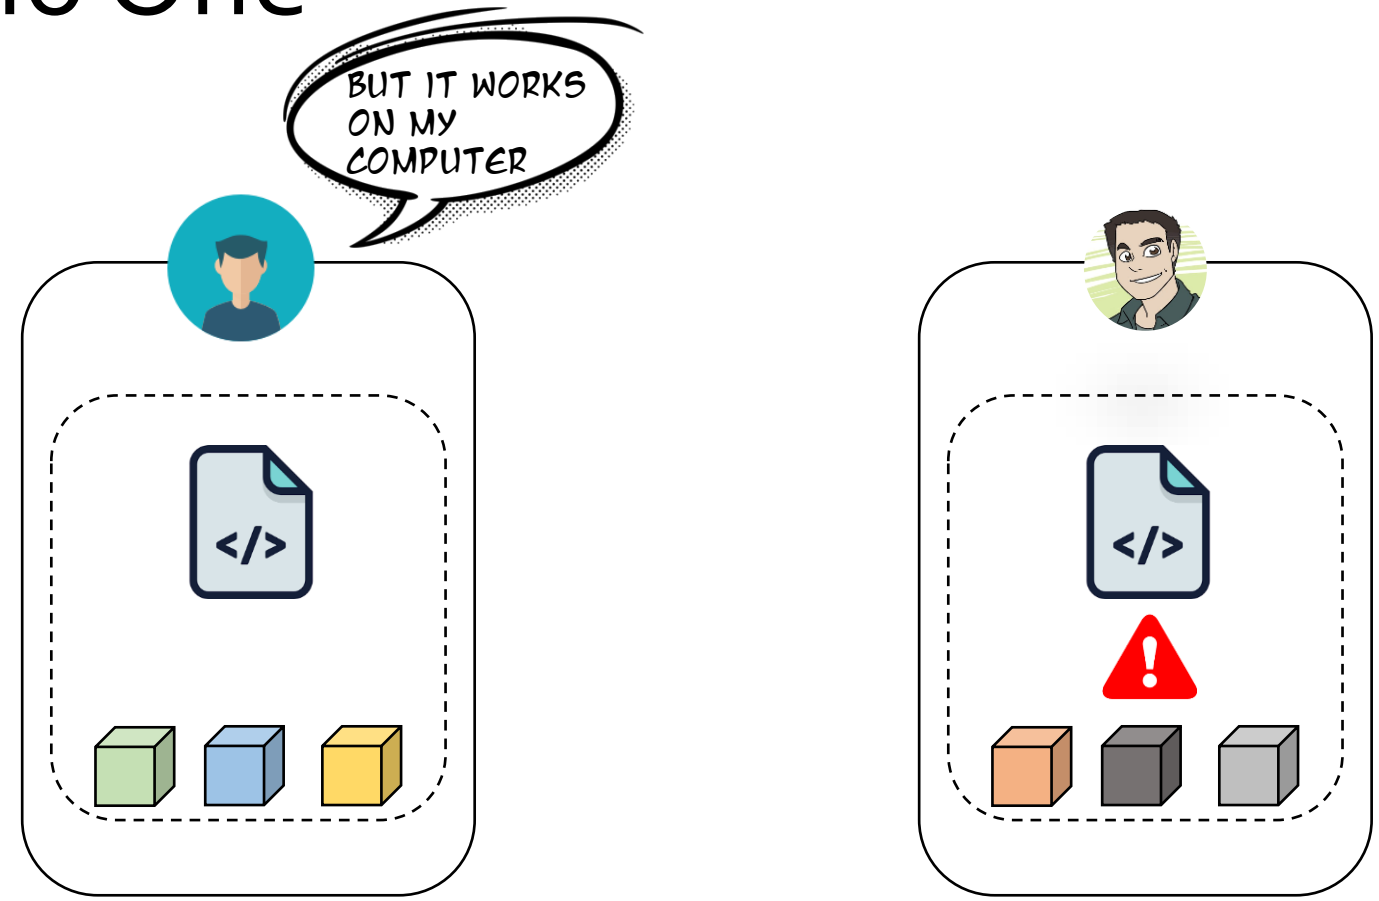

Scenario One

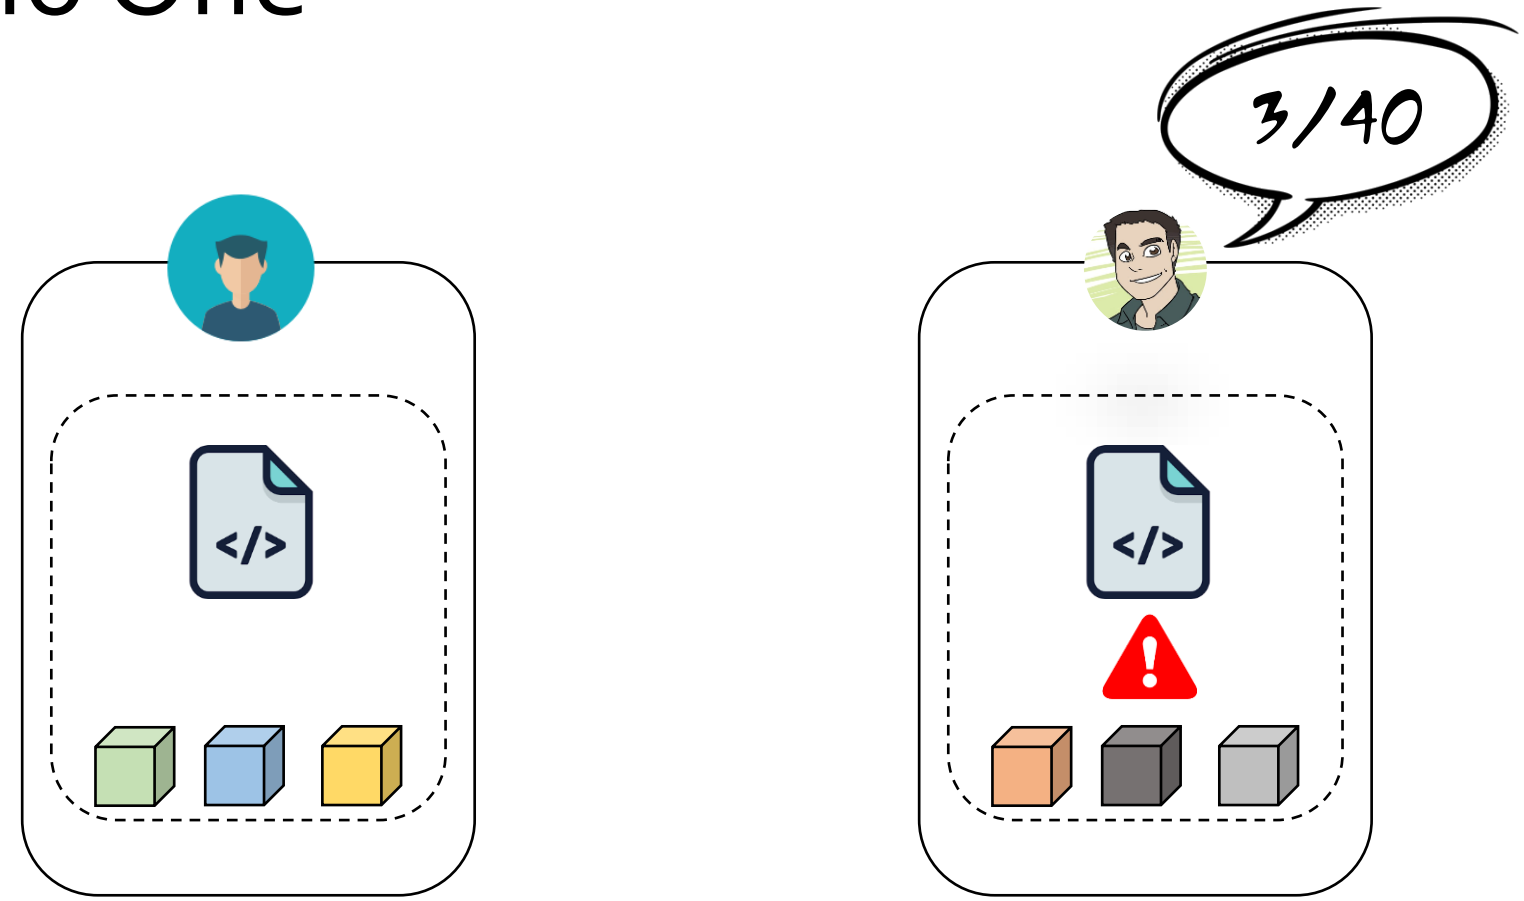

Scenario Two

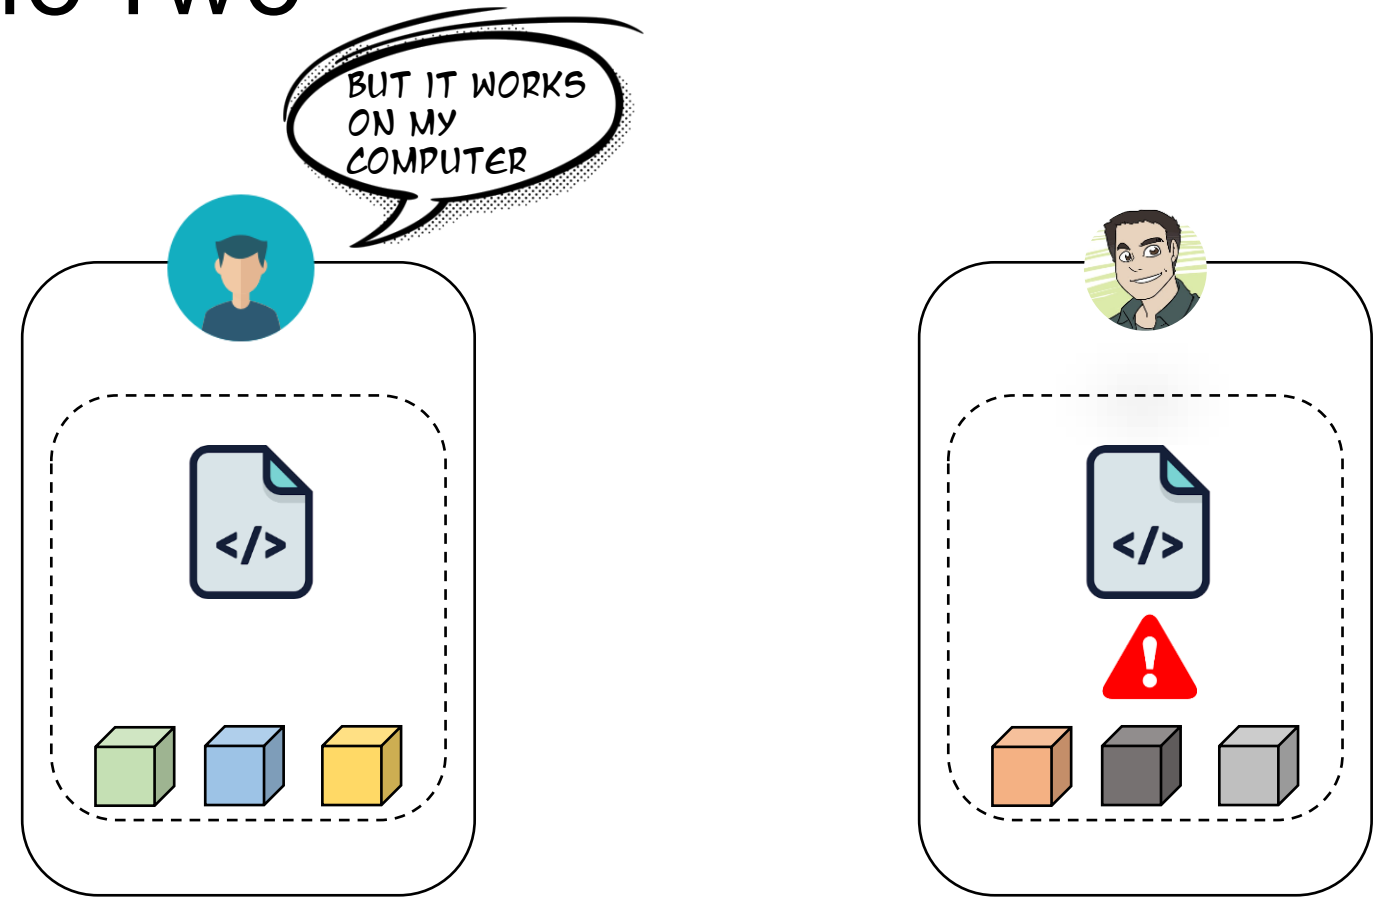

Scenario Two

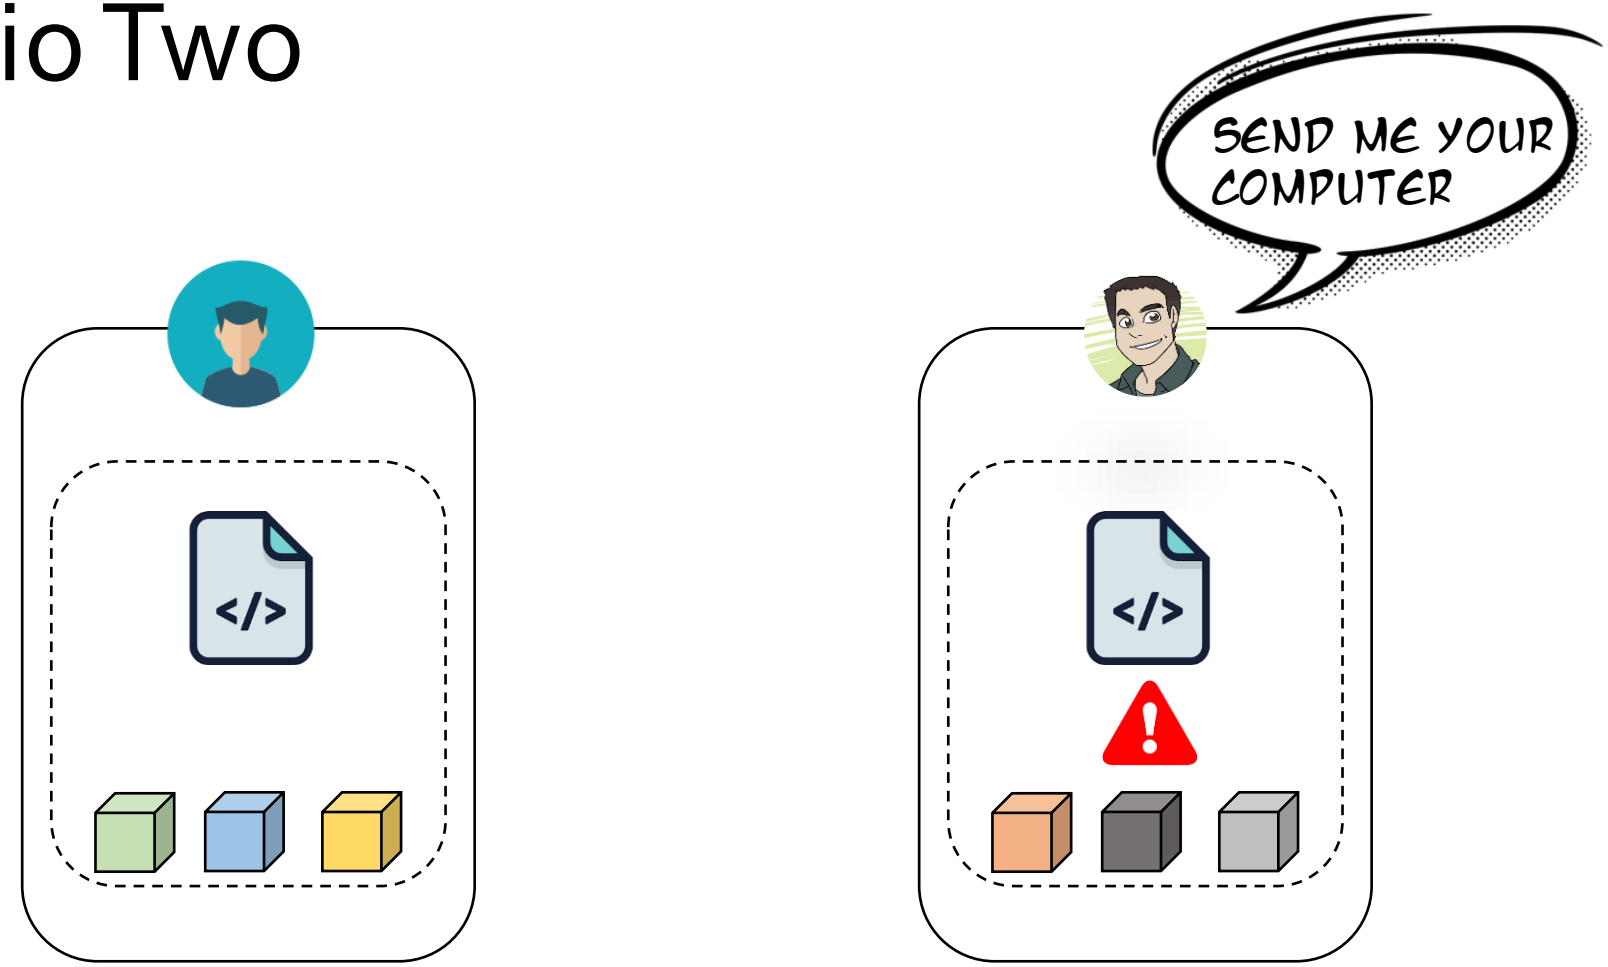

Scenario Two

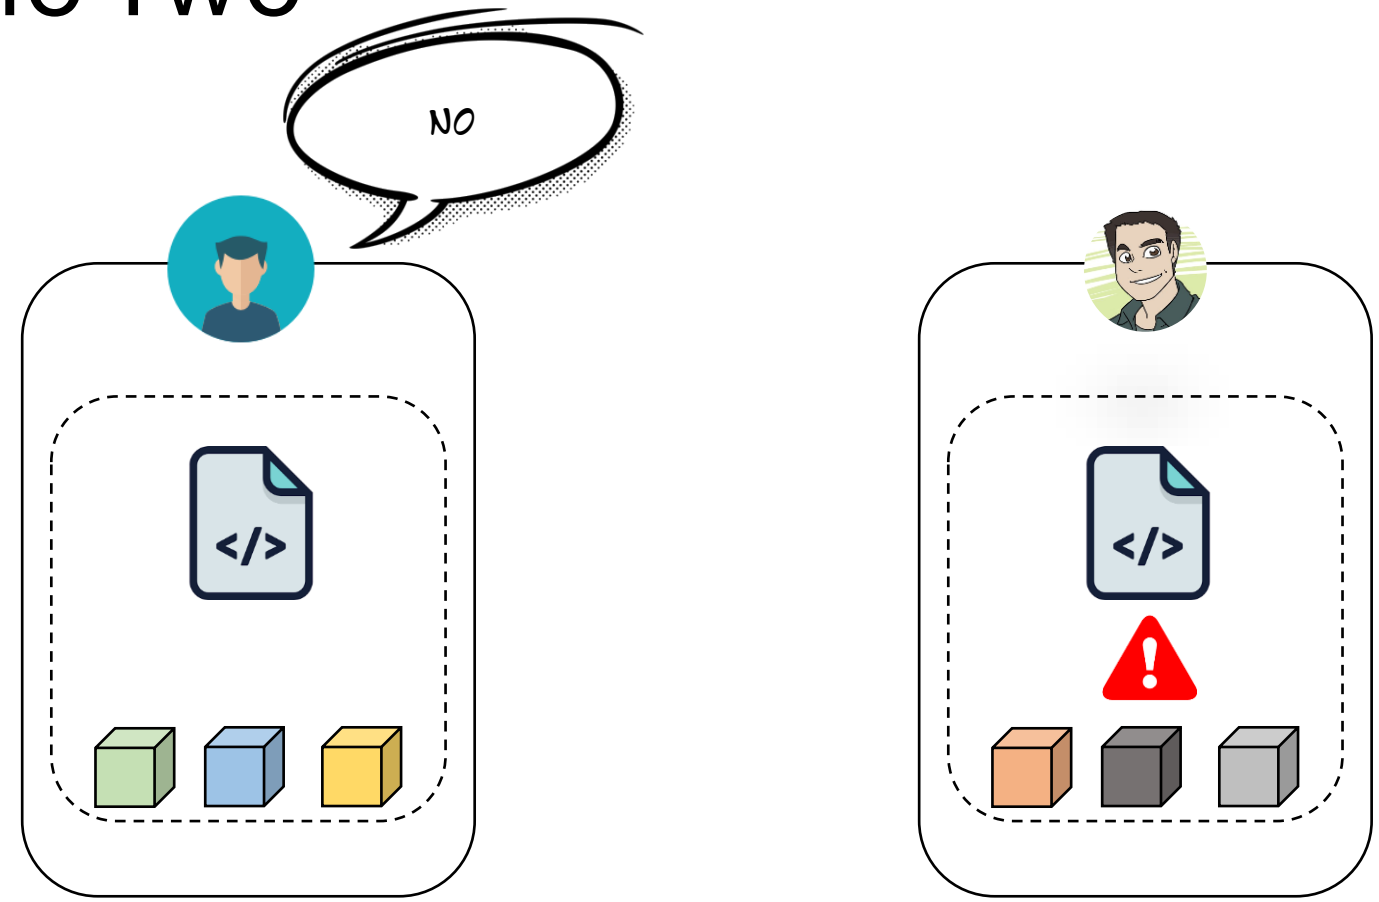

Scenario Two

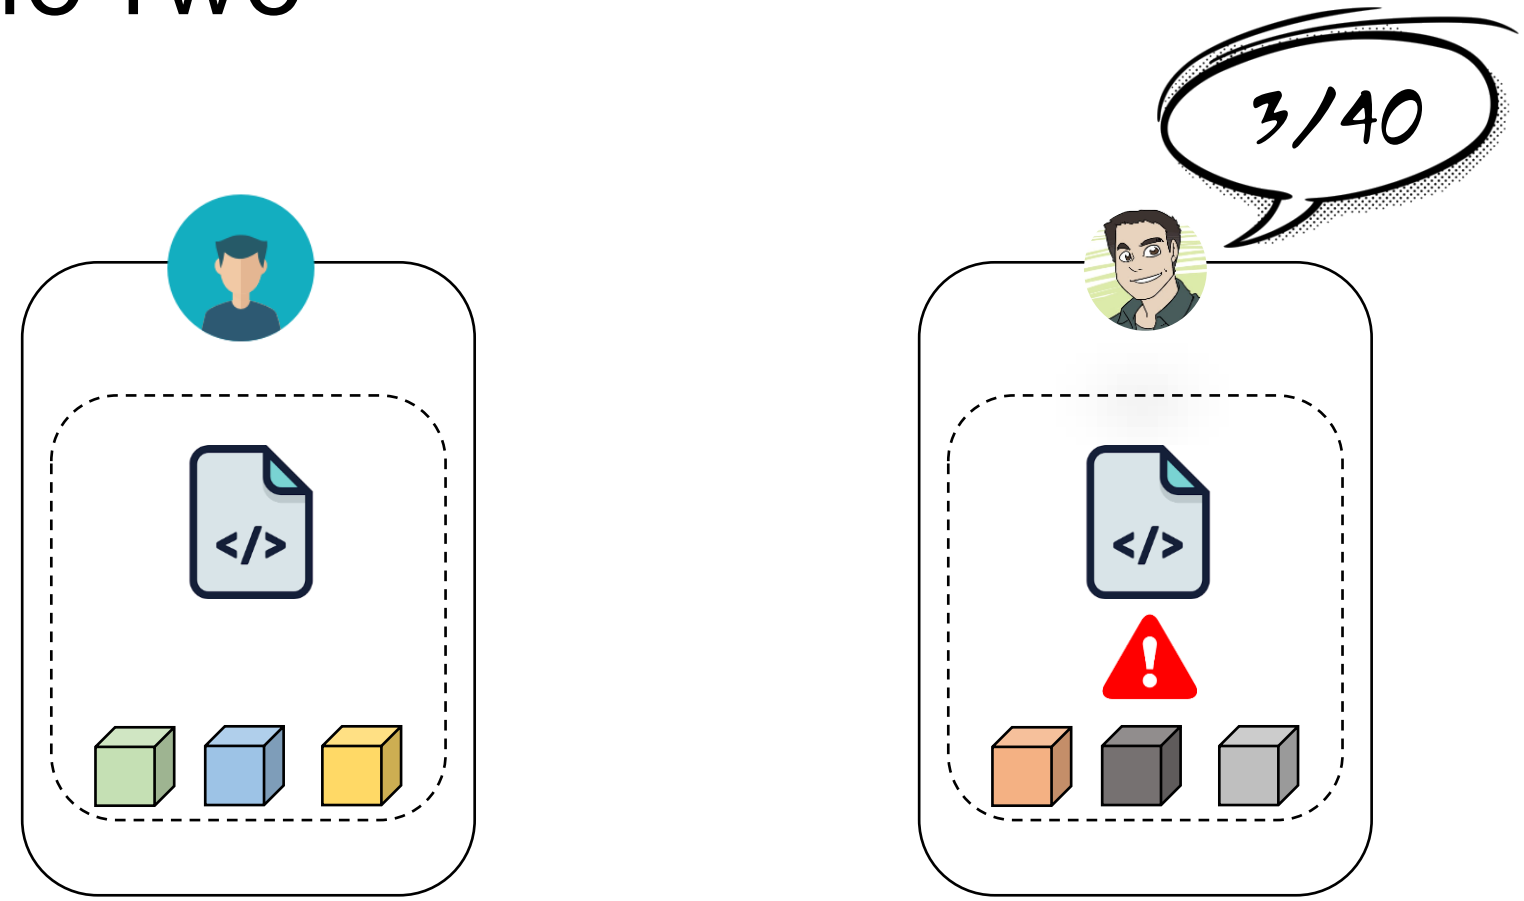

Scenario Three

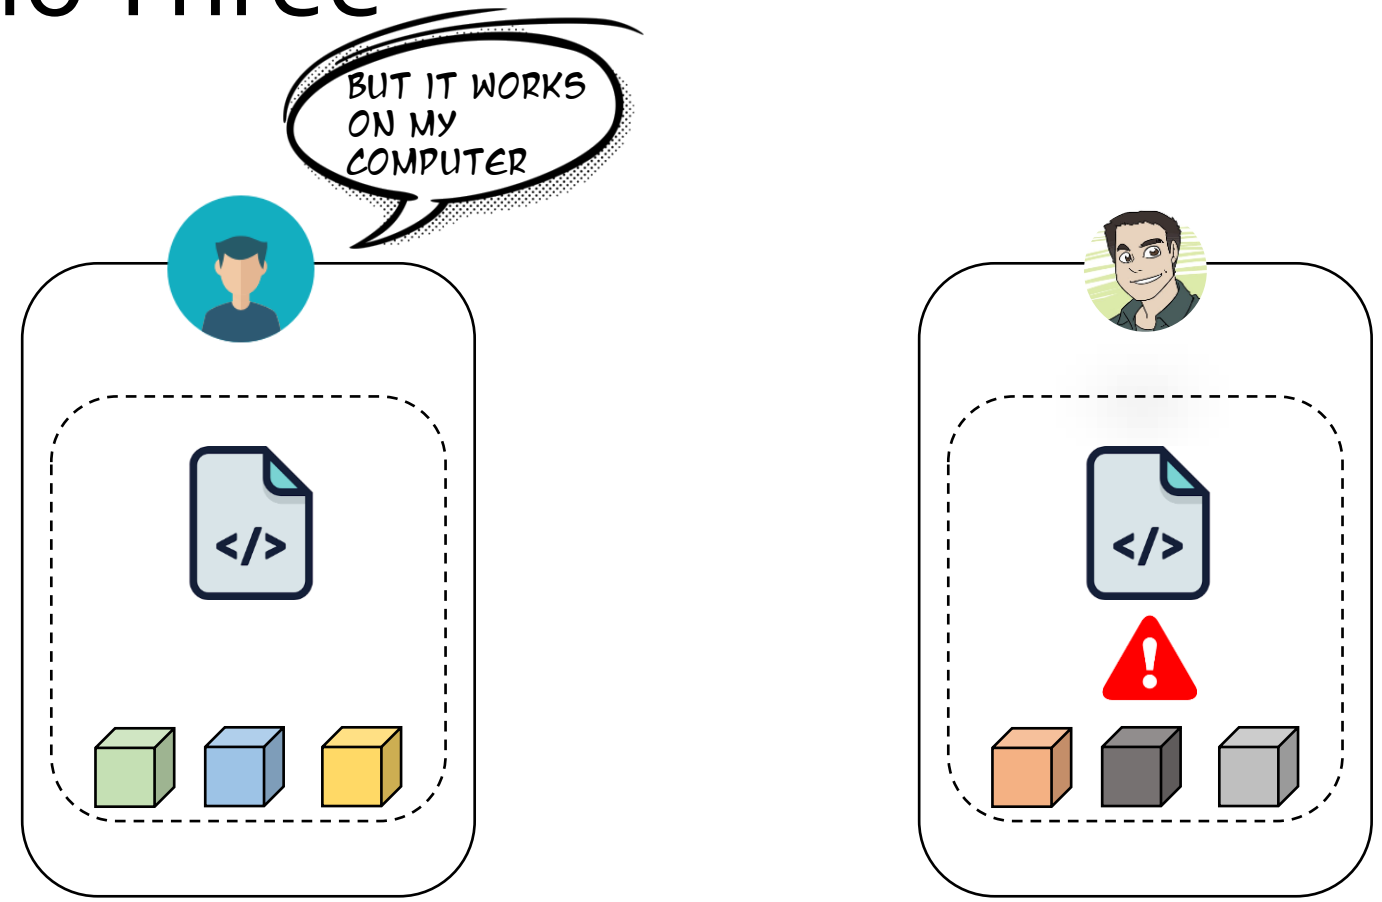

Scenario Three

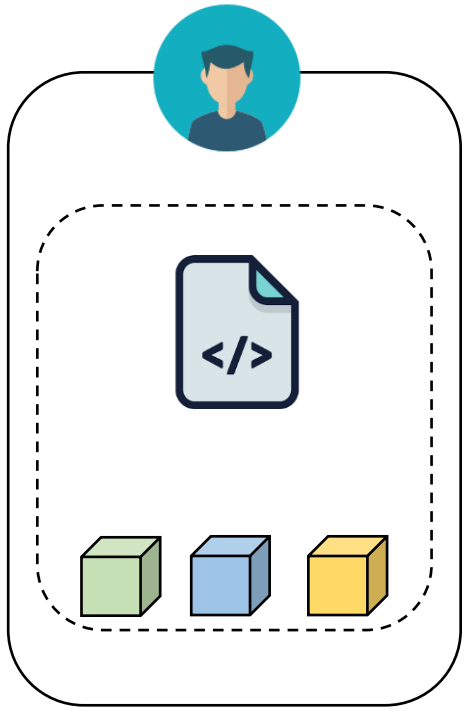

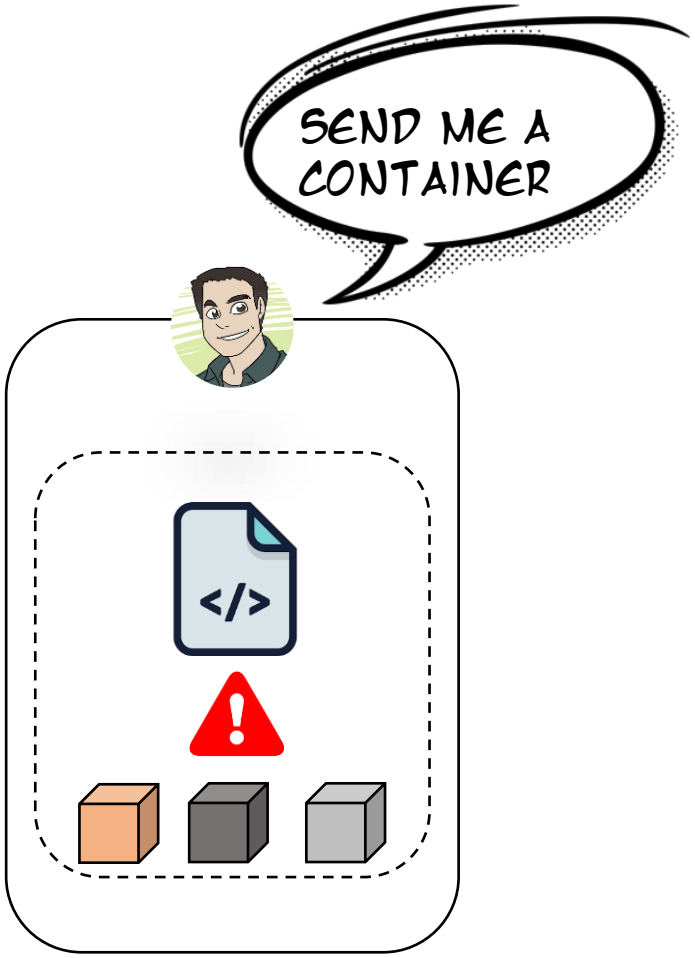

Scenario Three

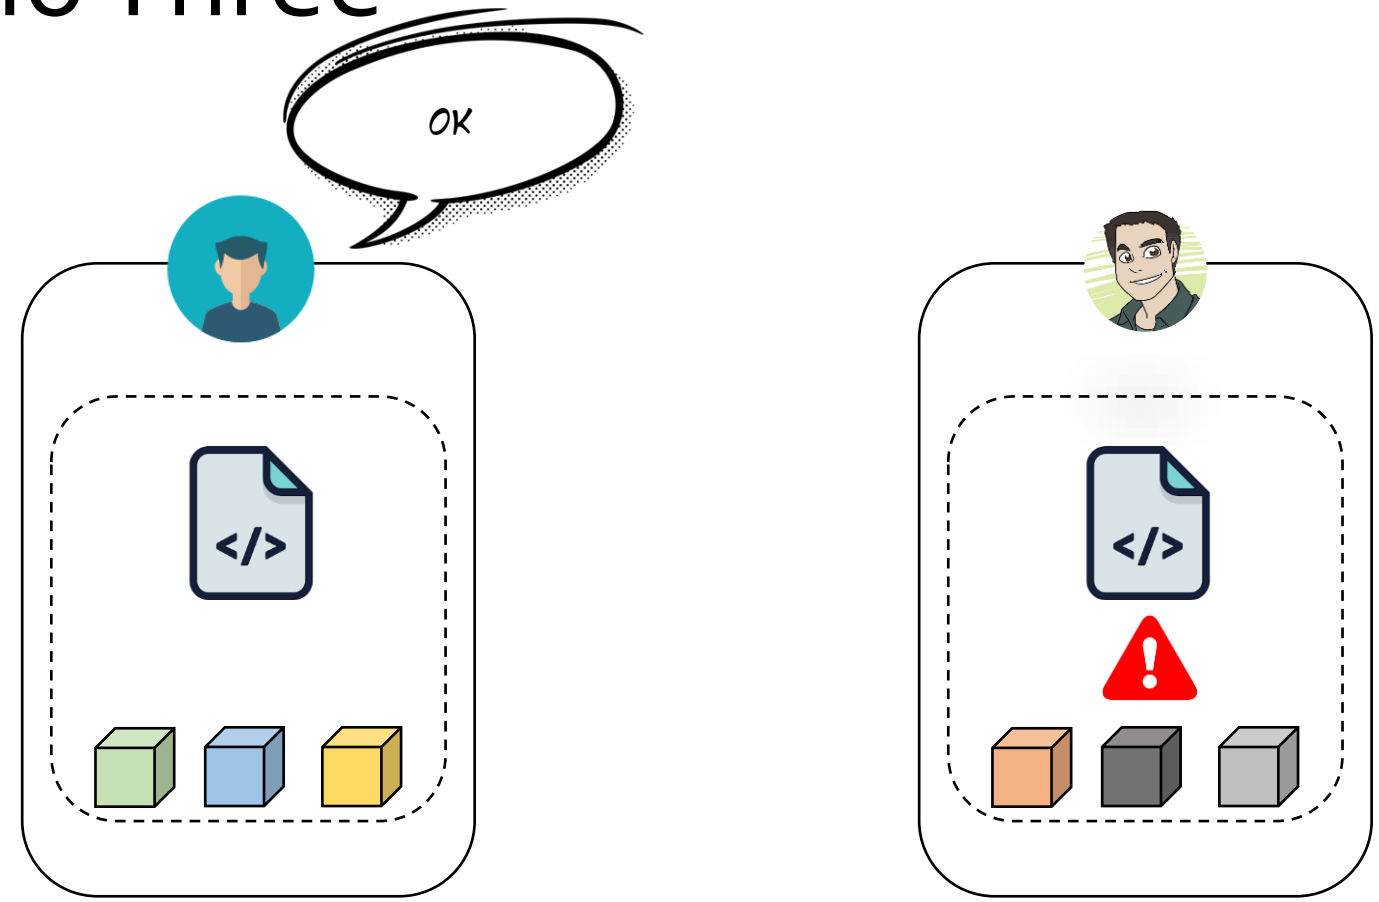

Scenario Three

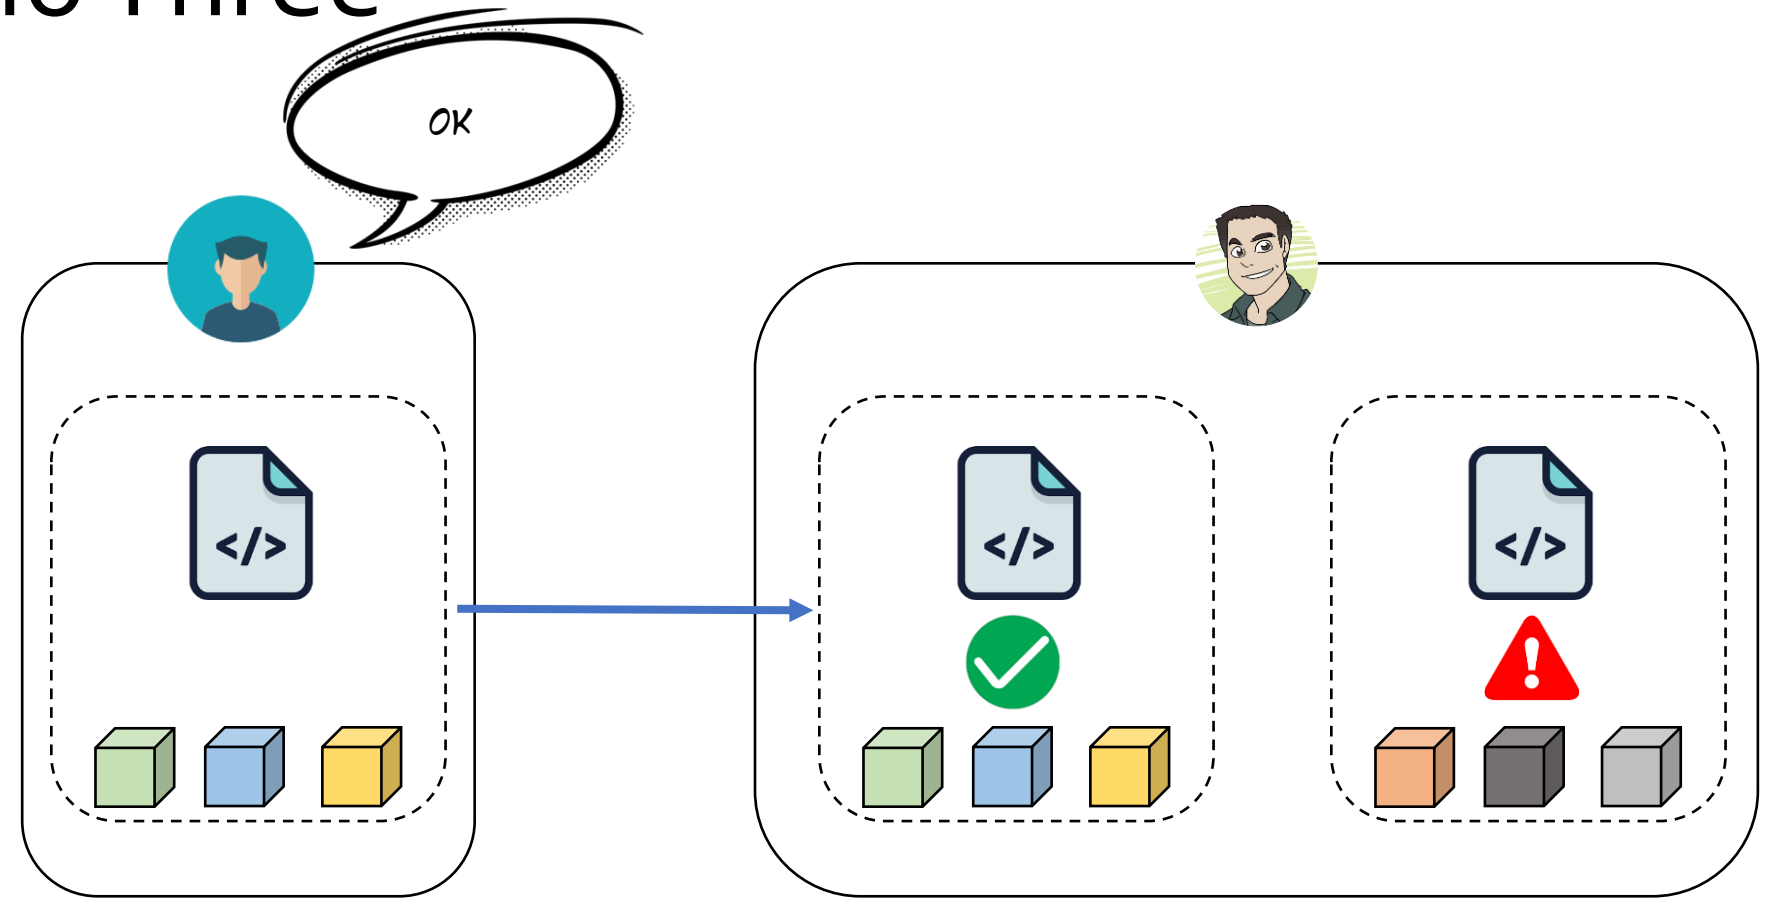

Multiple Versions and Dependencies Living Side by Side

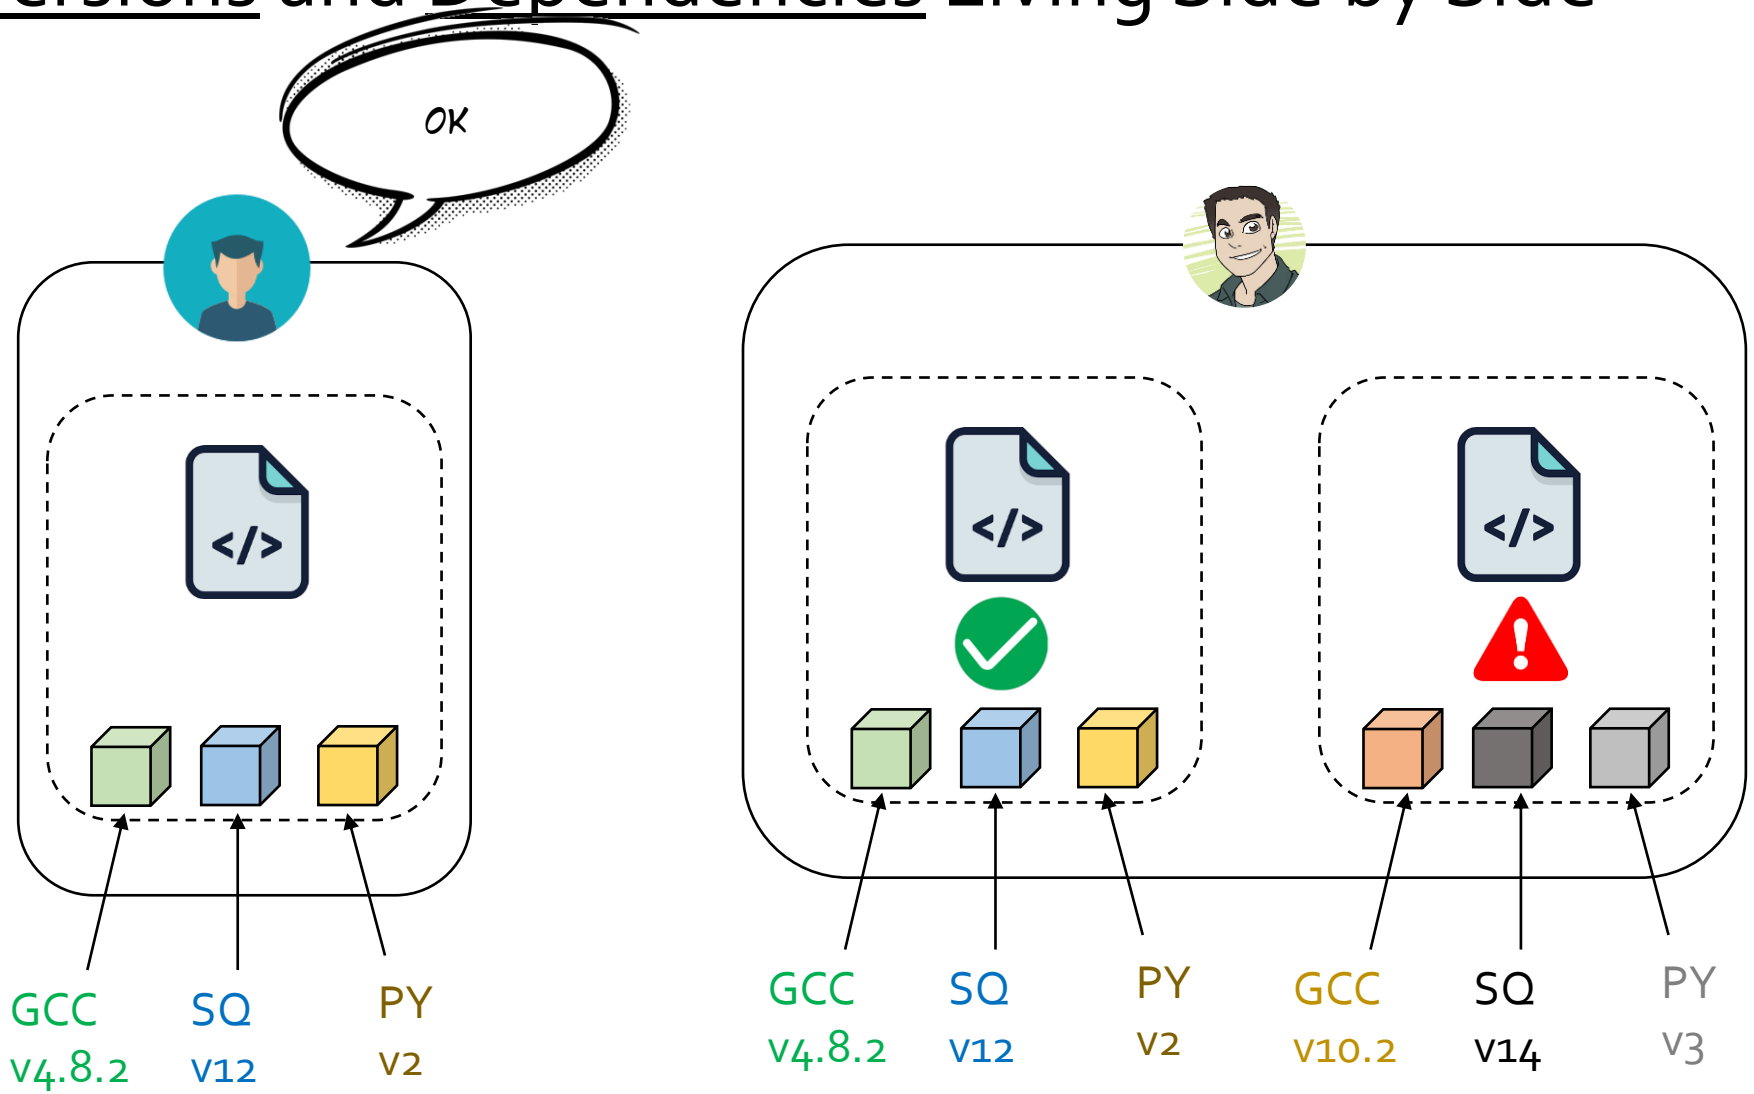

What is a container?

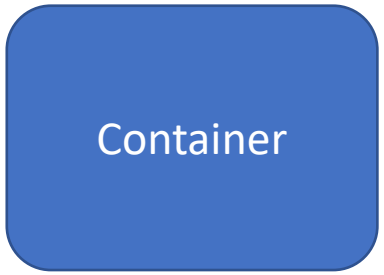

What is a container?

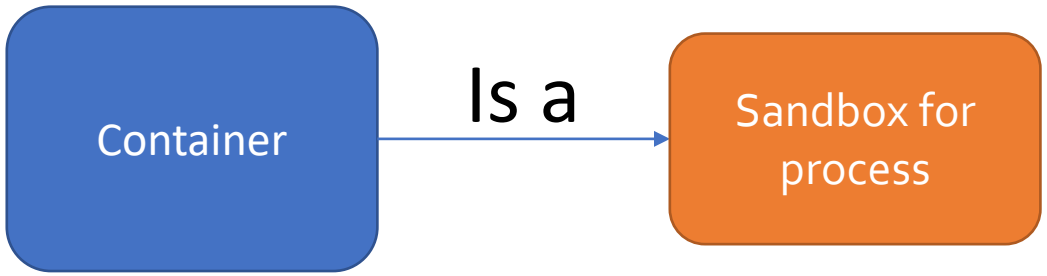

What is a container?

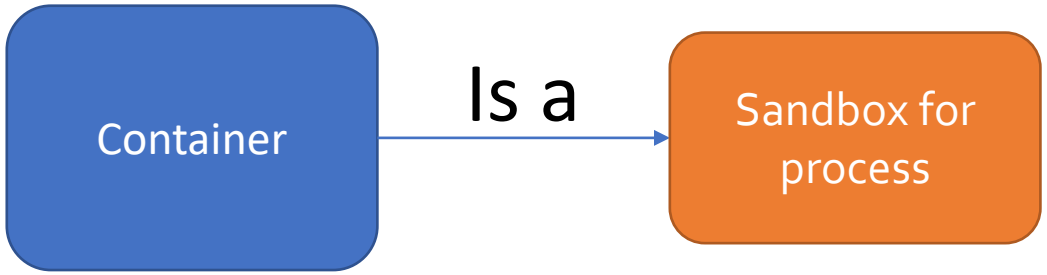

Containers are a type of virtualization for Linux Kernel

How does the kernel build this sandbox?

What kernel capabilities might a process (program) need?

What kernel capabilities might a process (program) need?

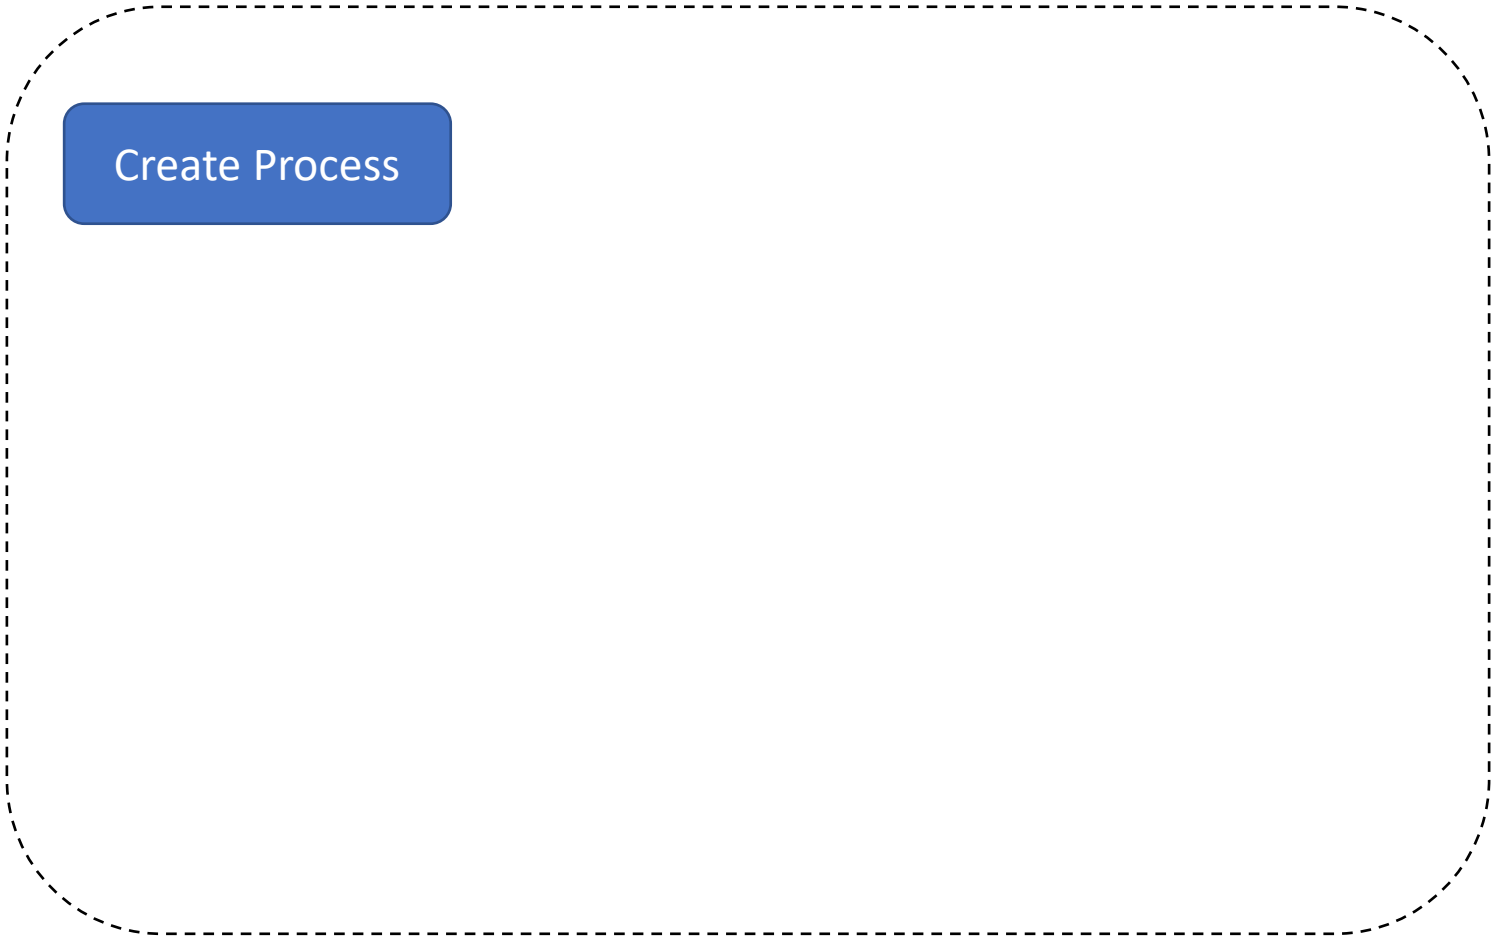

What kernel capabilities might a process (program) need?

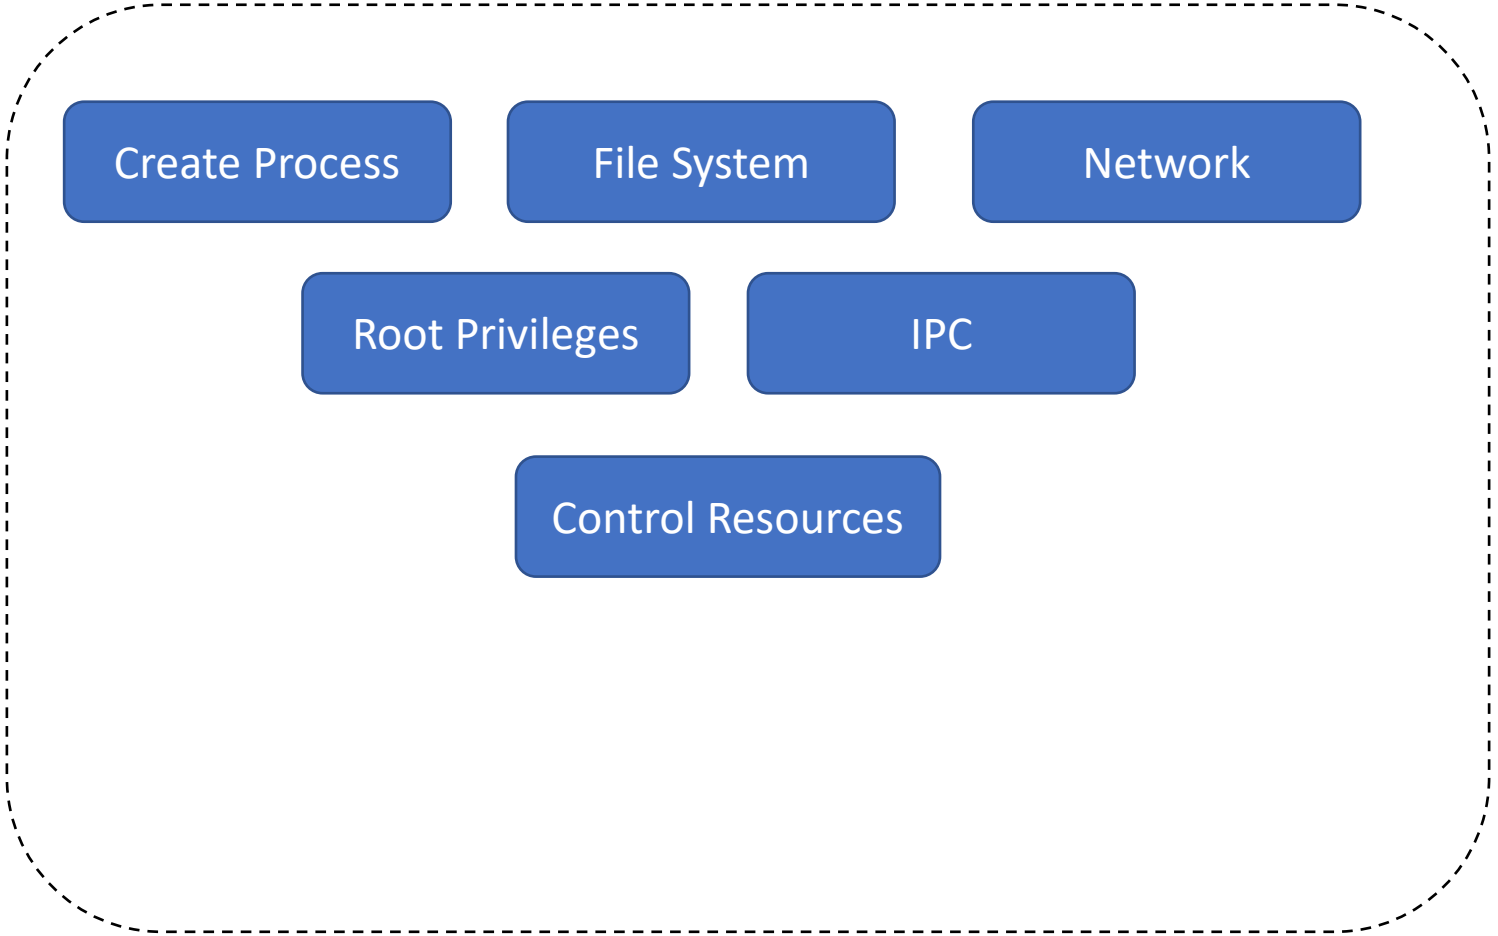

The capabilities are called *namespaces*

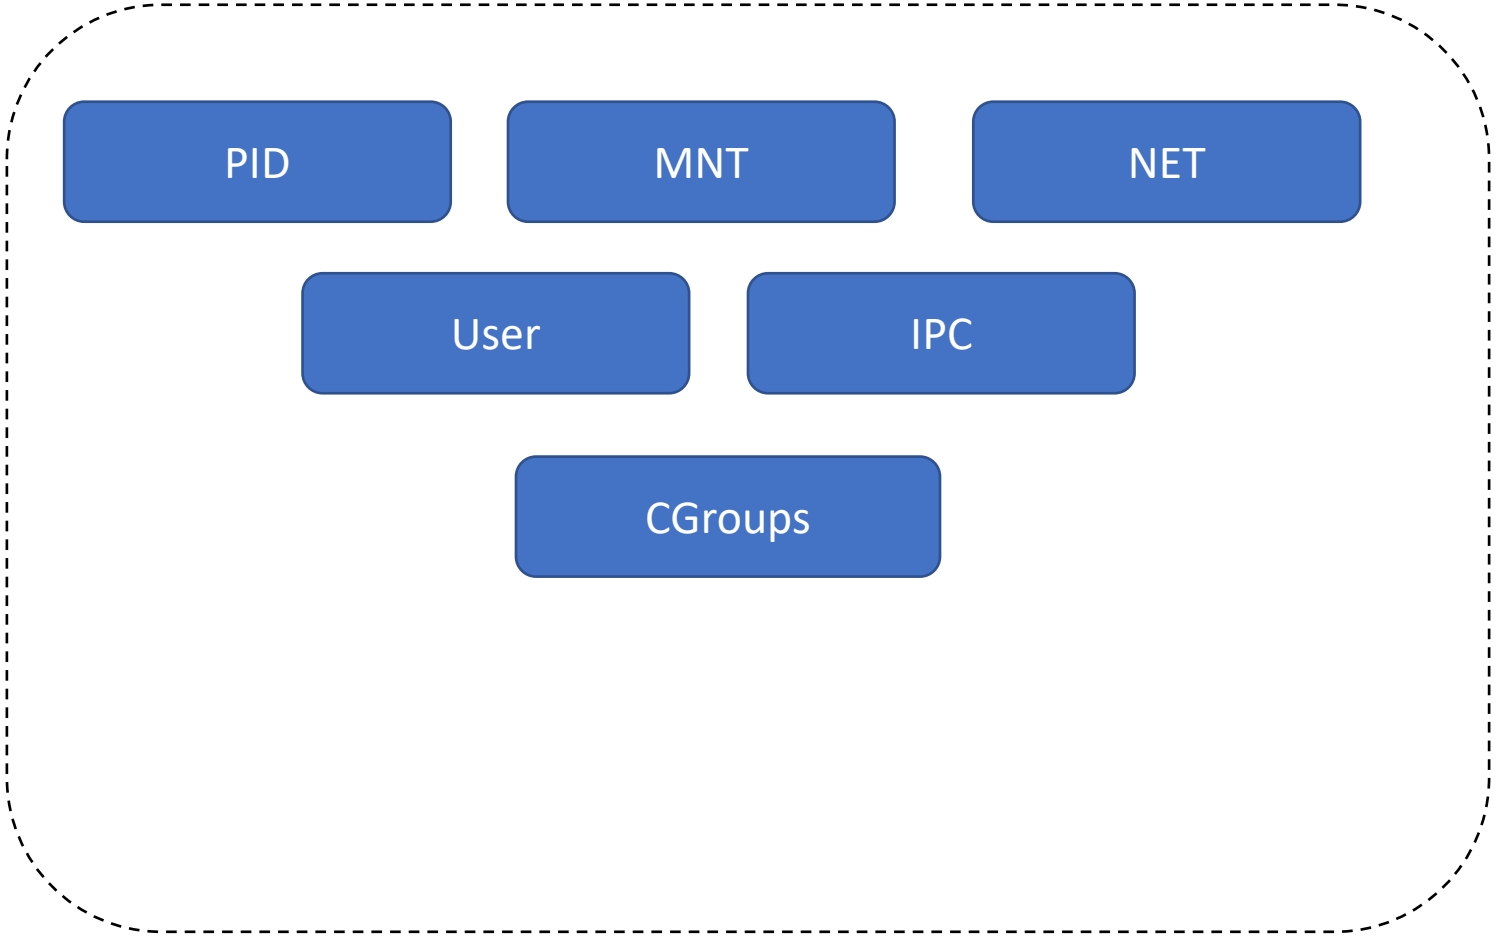

New containers mean new namespaces

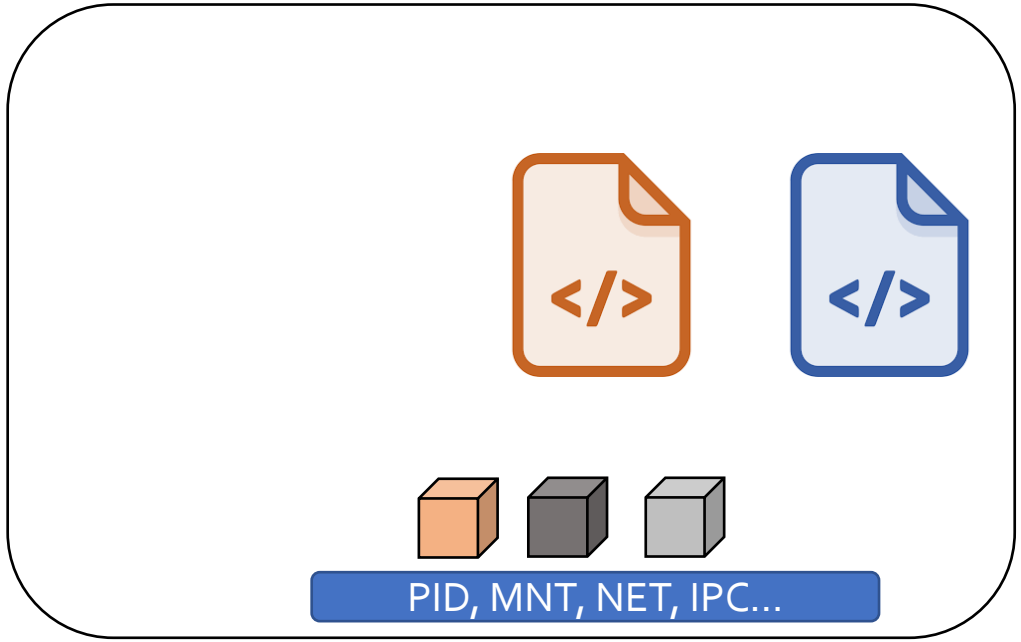

New containers mean new namespaces

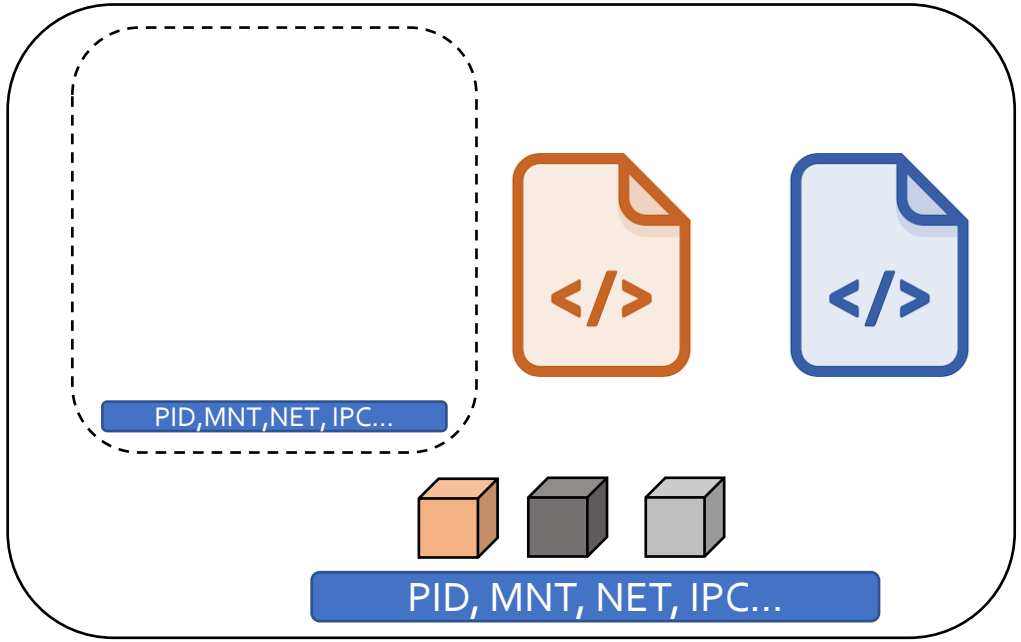

New containers mean new namespaces

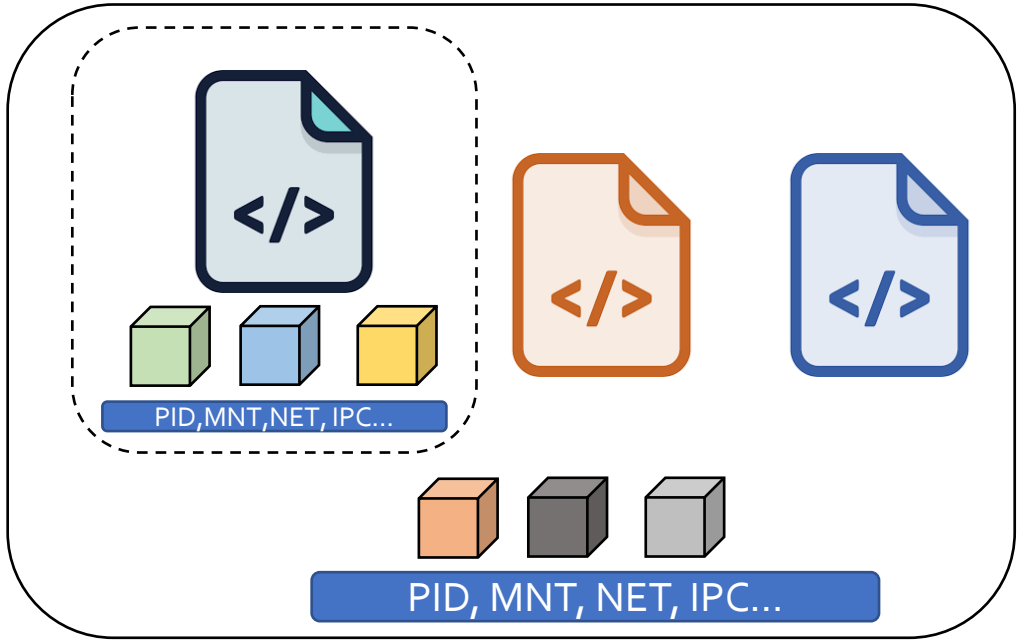

Review

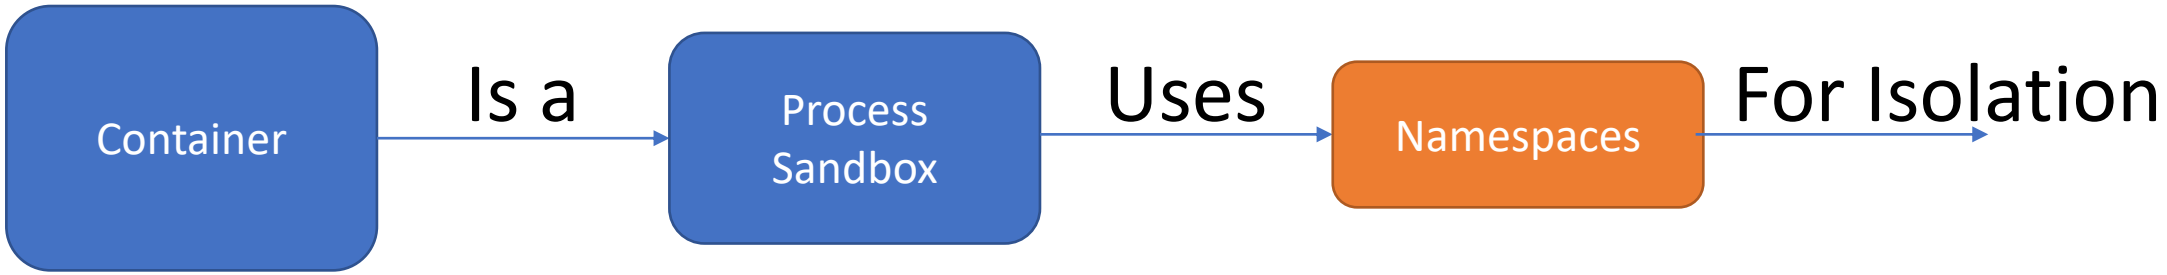
Container
Is a
Process
Sandbox
Uses
Namespaces
For Isolation

Network Namespaces

Closer look at the network namespace

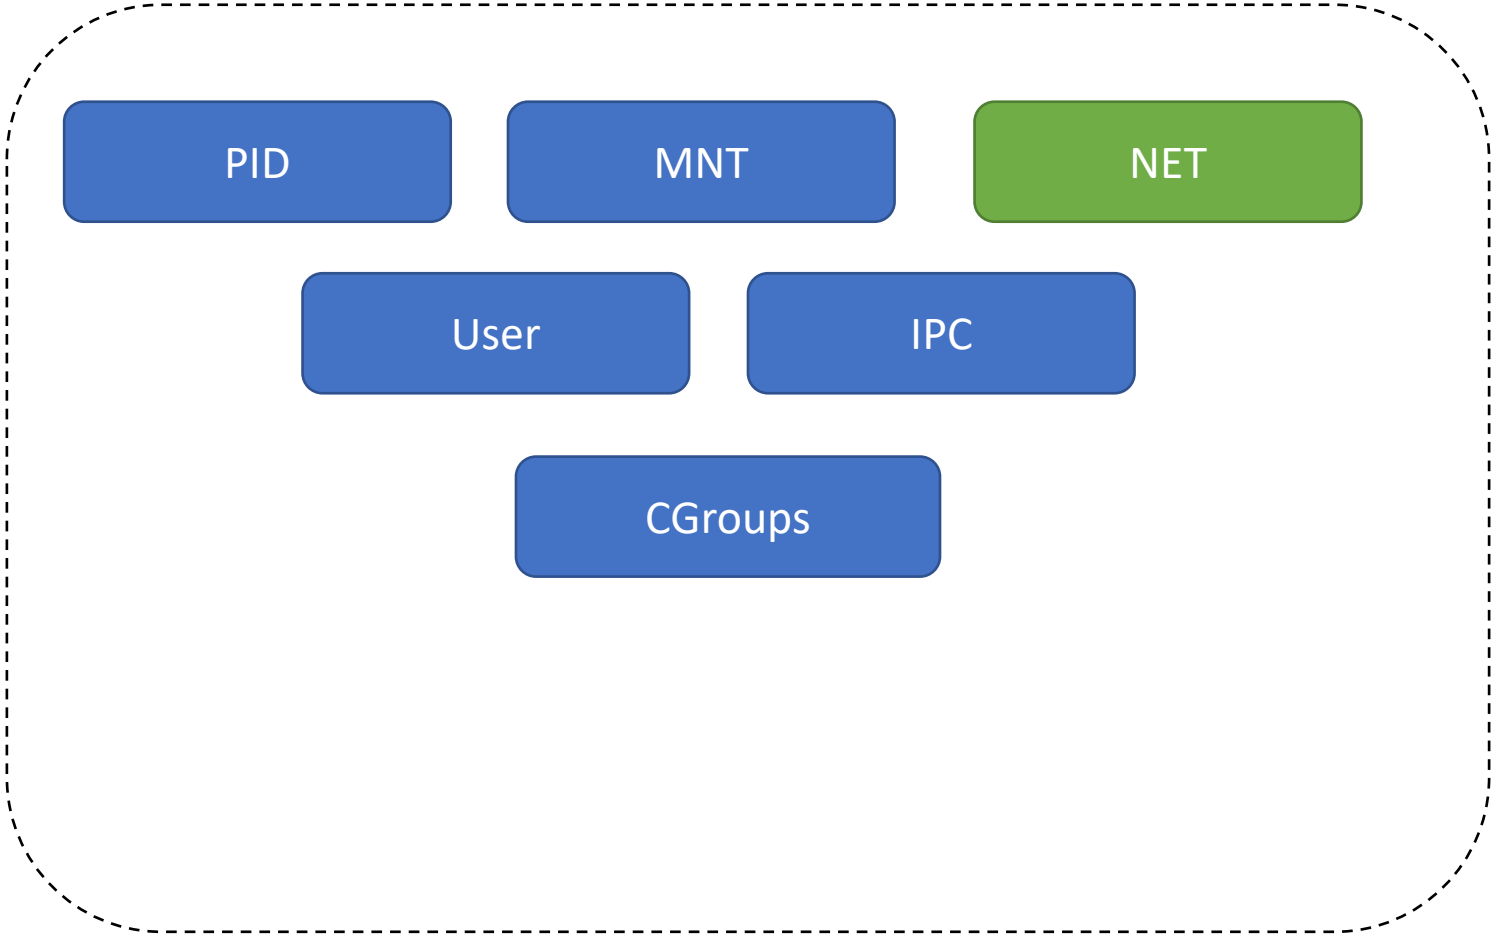

What is a network?

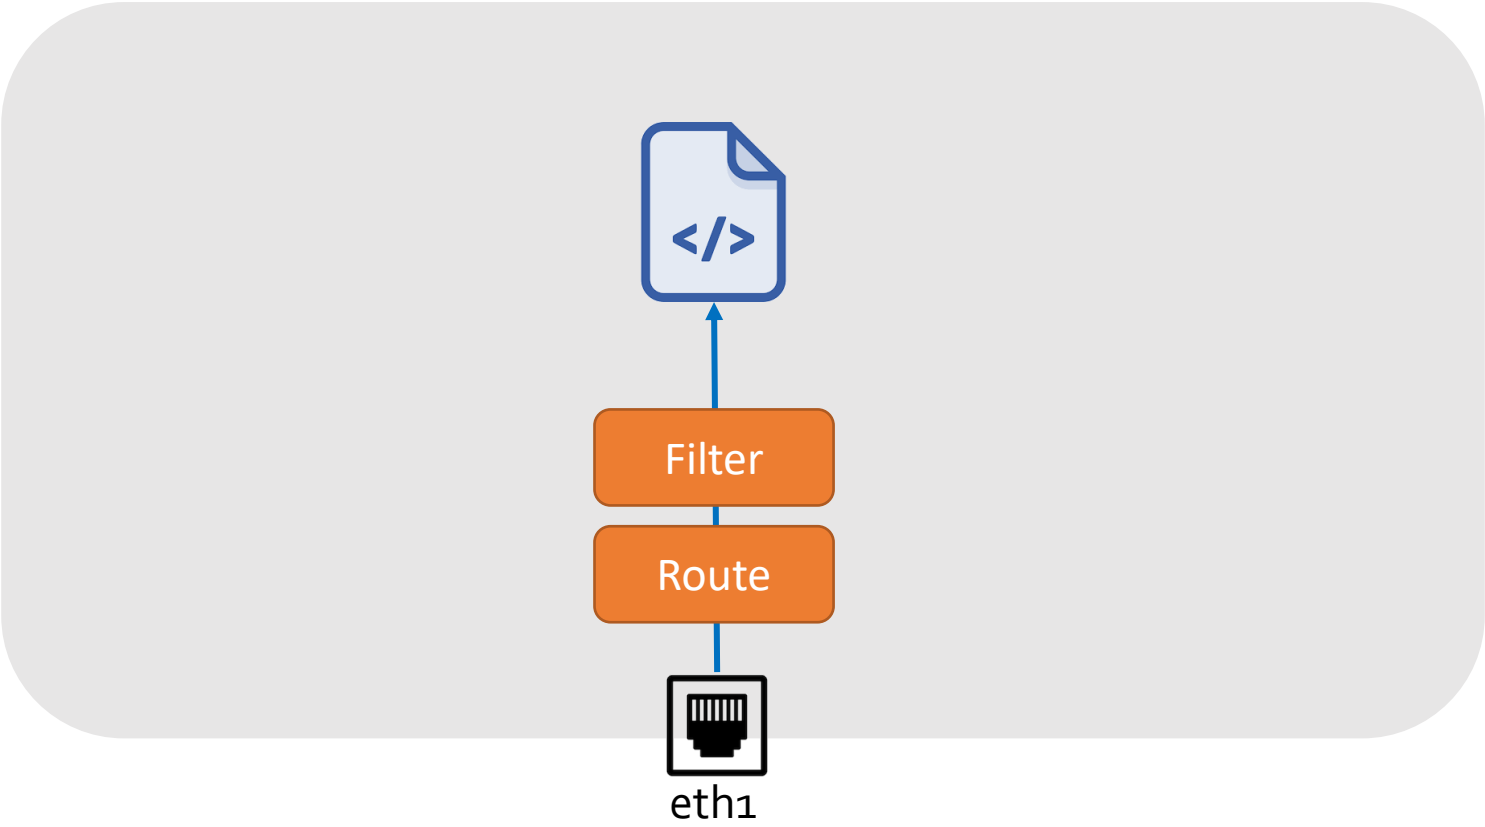

Network namespace isolation

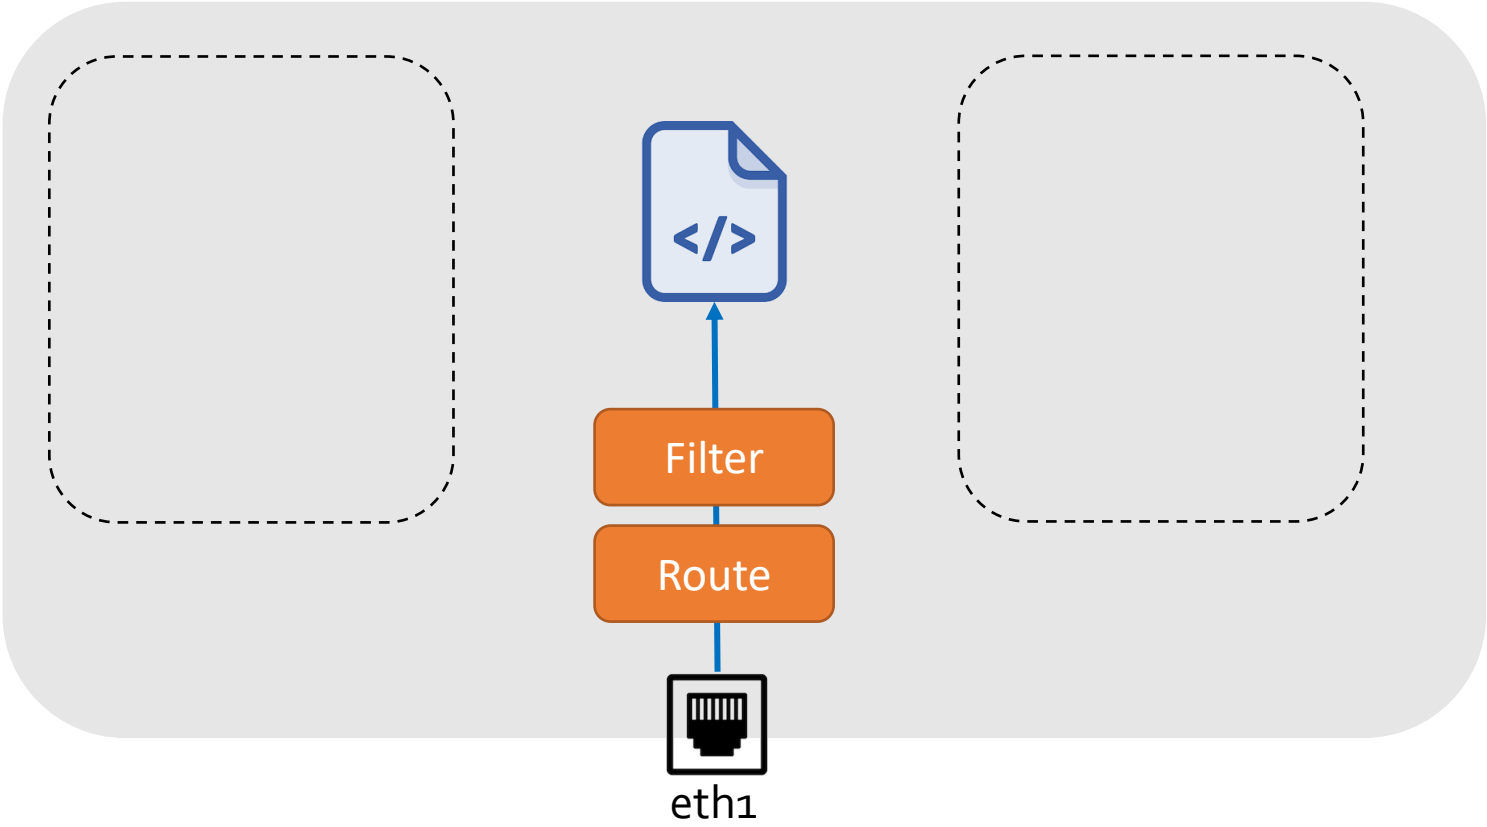

Network namespace isolation

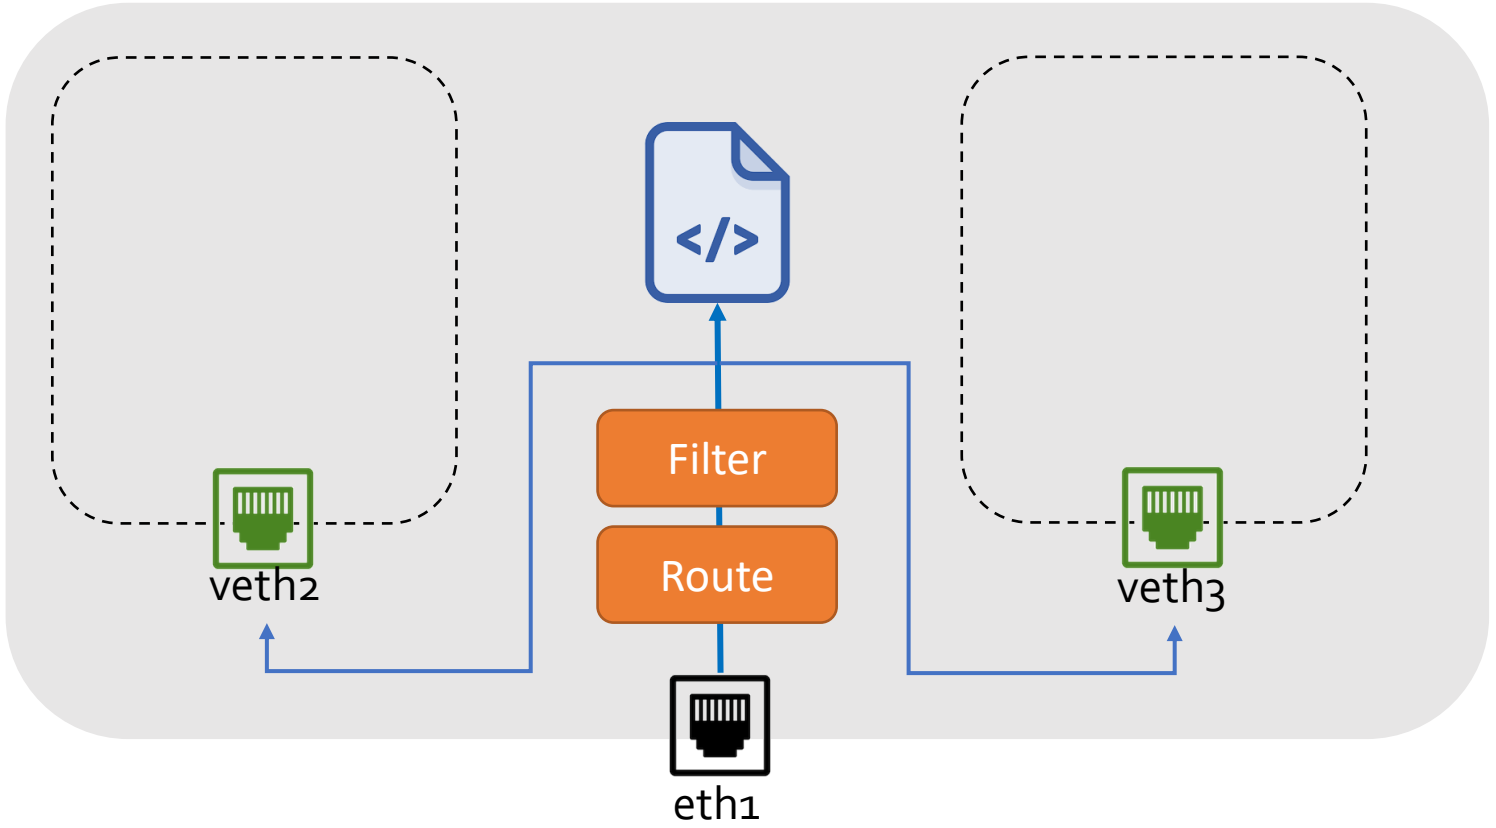

Network namespace isolation

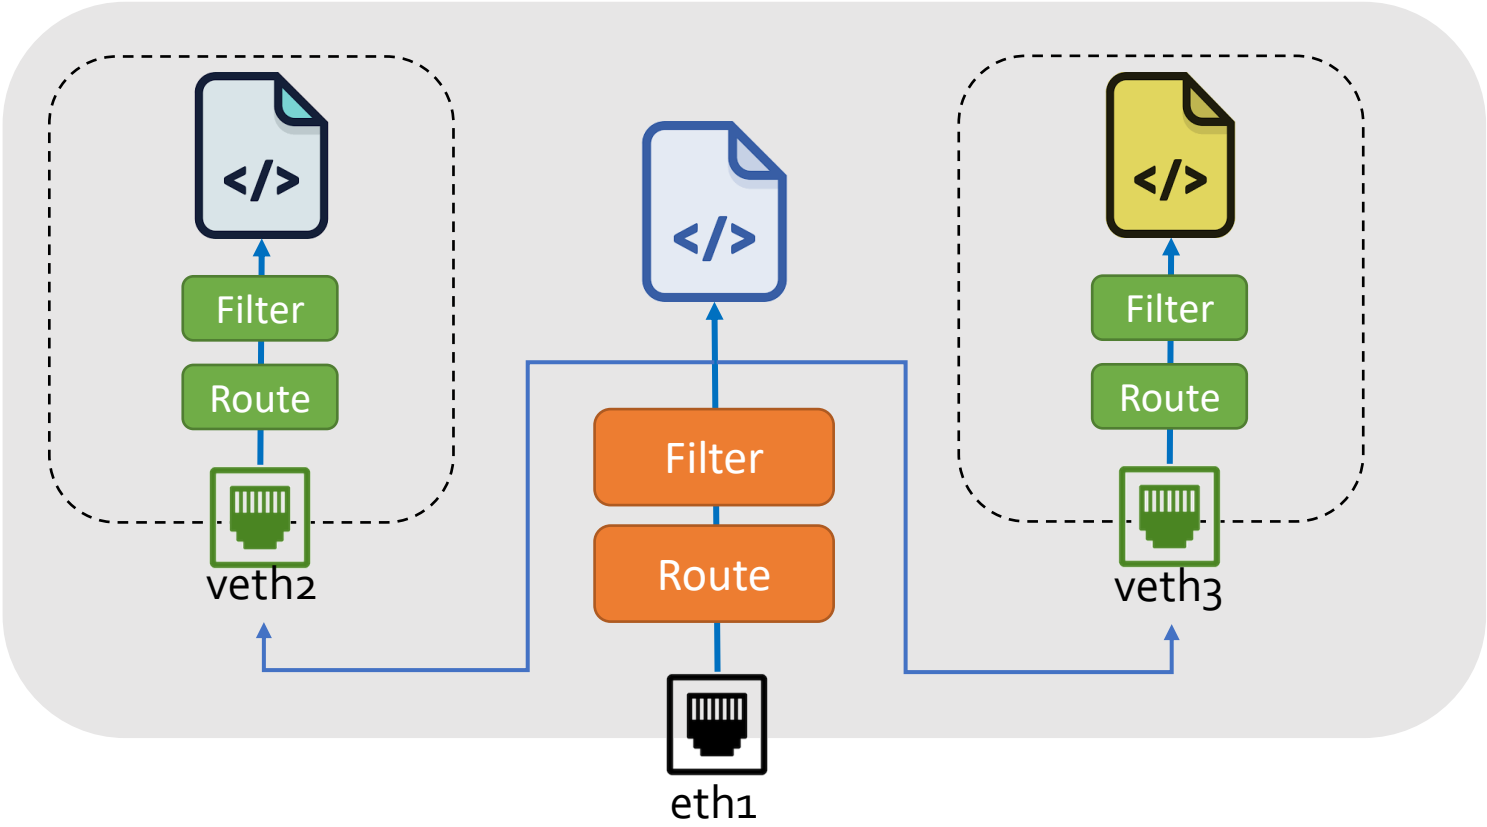

Operational Benefit: Easy to experiment with rules because they can be tied to the container

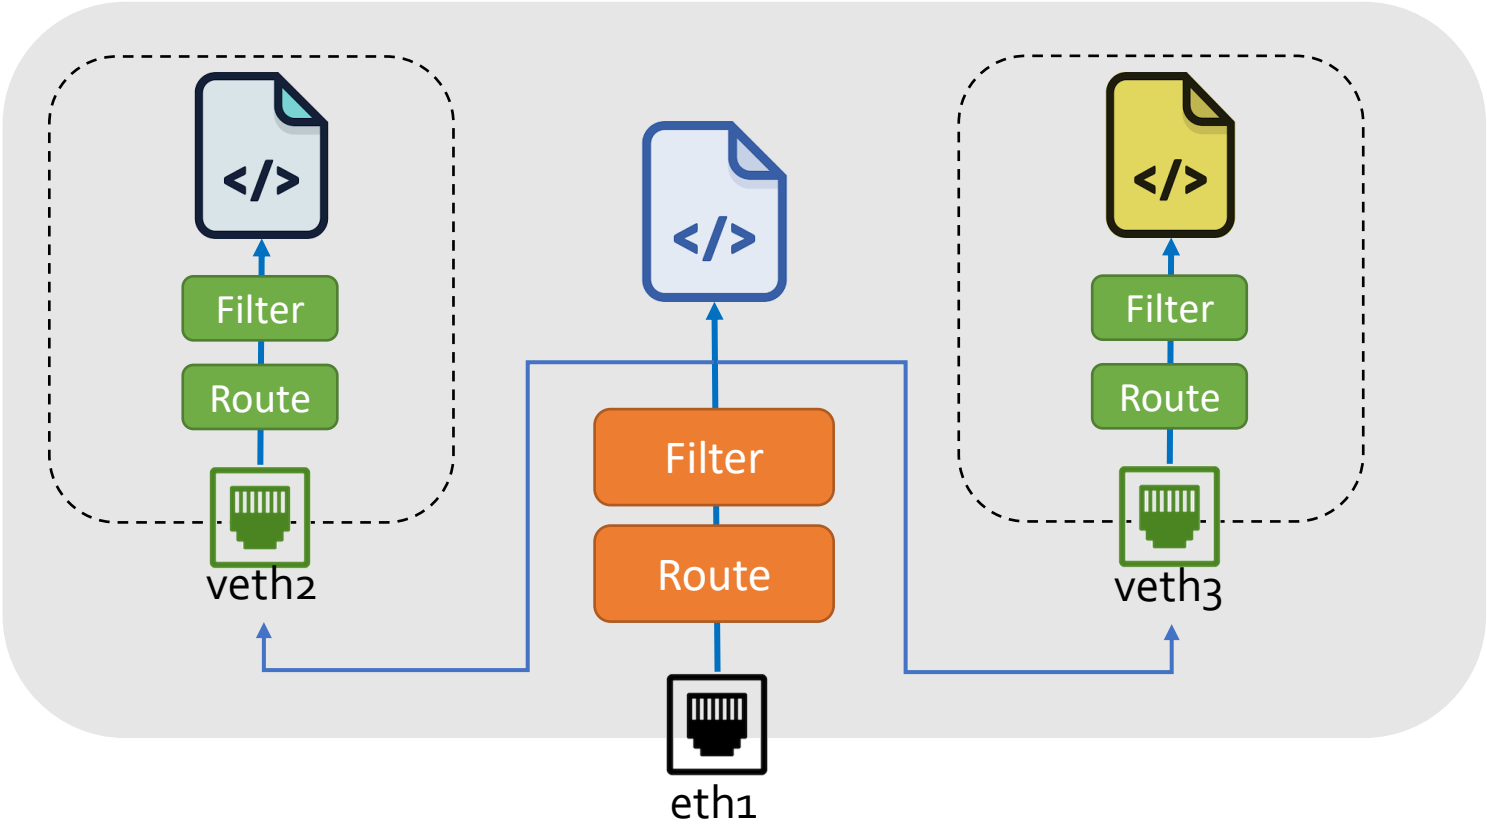

Operational Benefit: Easy to experiment with rules because they can be tied to the container

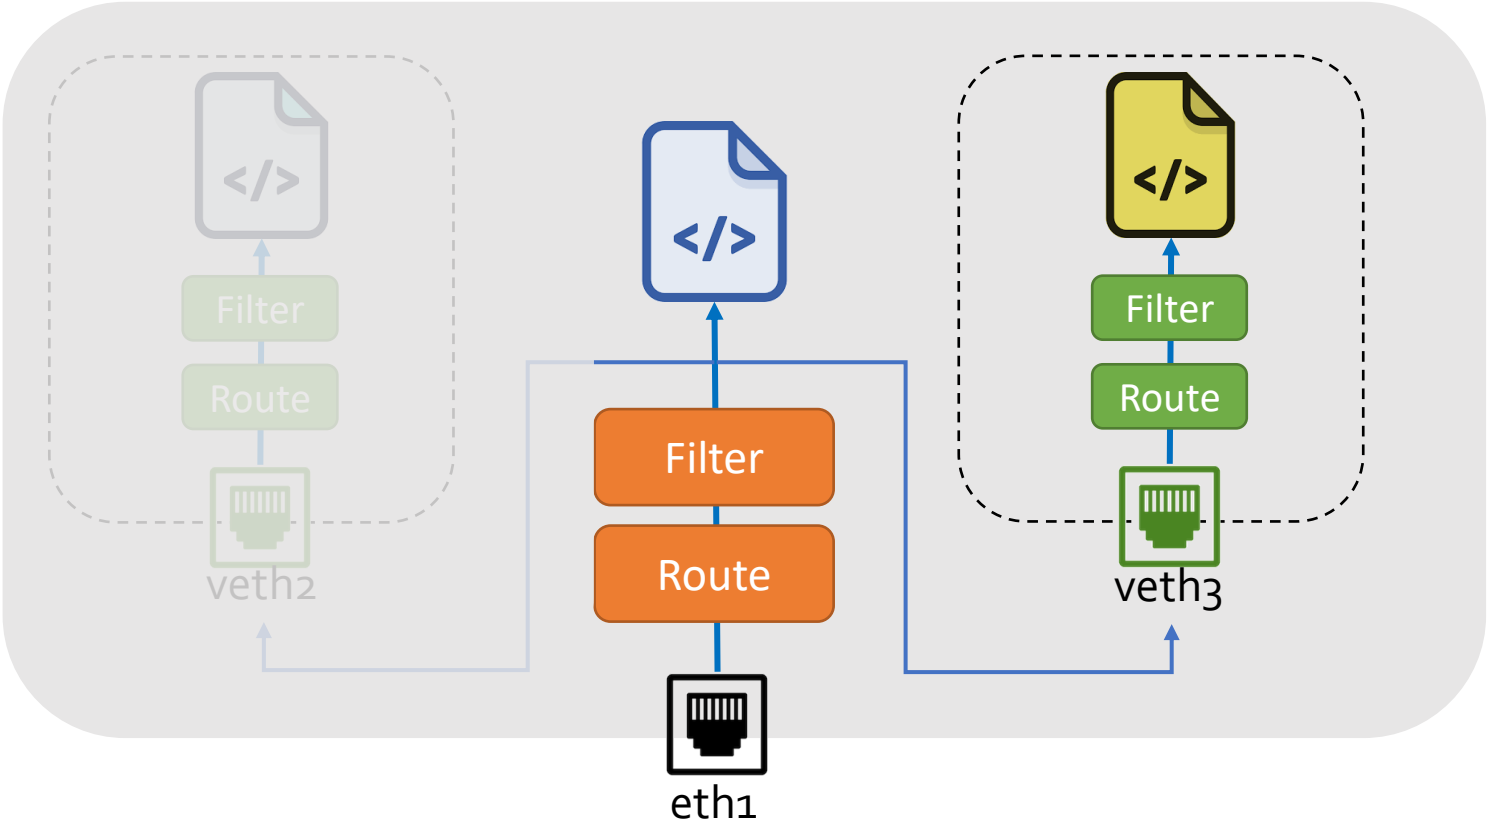

Container Network is often managed through a virtual bridge/switch

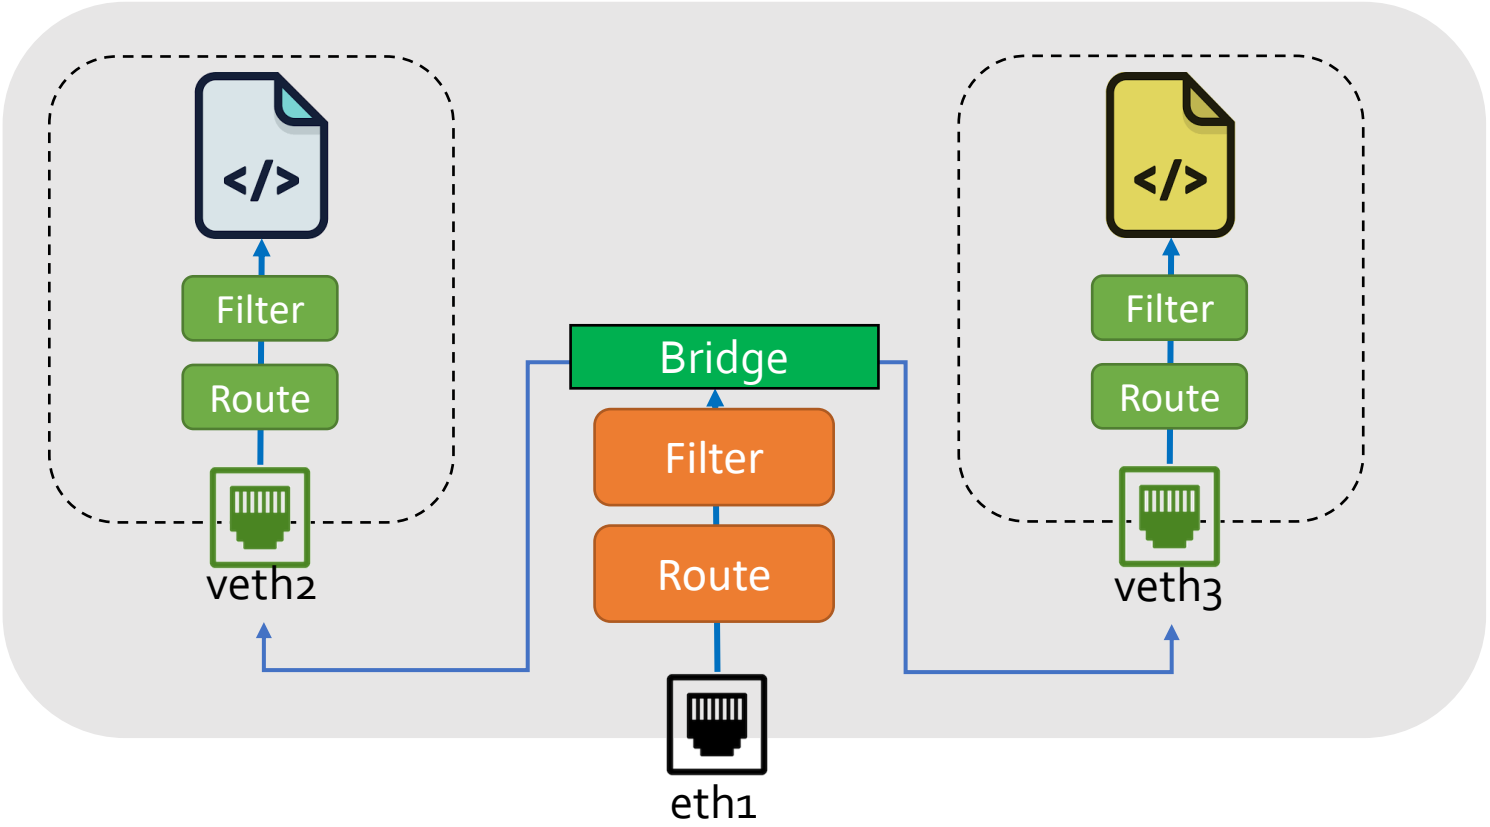

Other namespaces provide isolation in similar ways

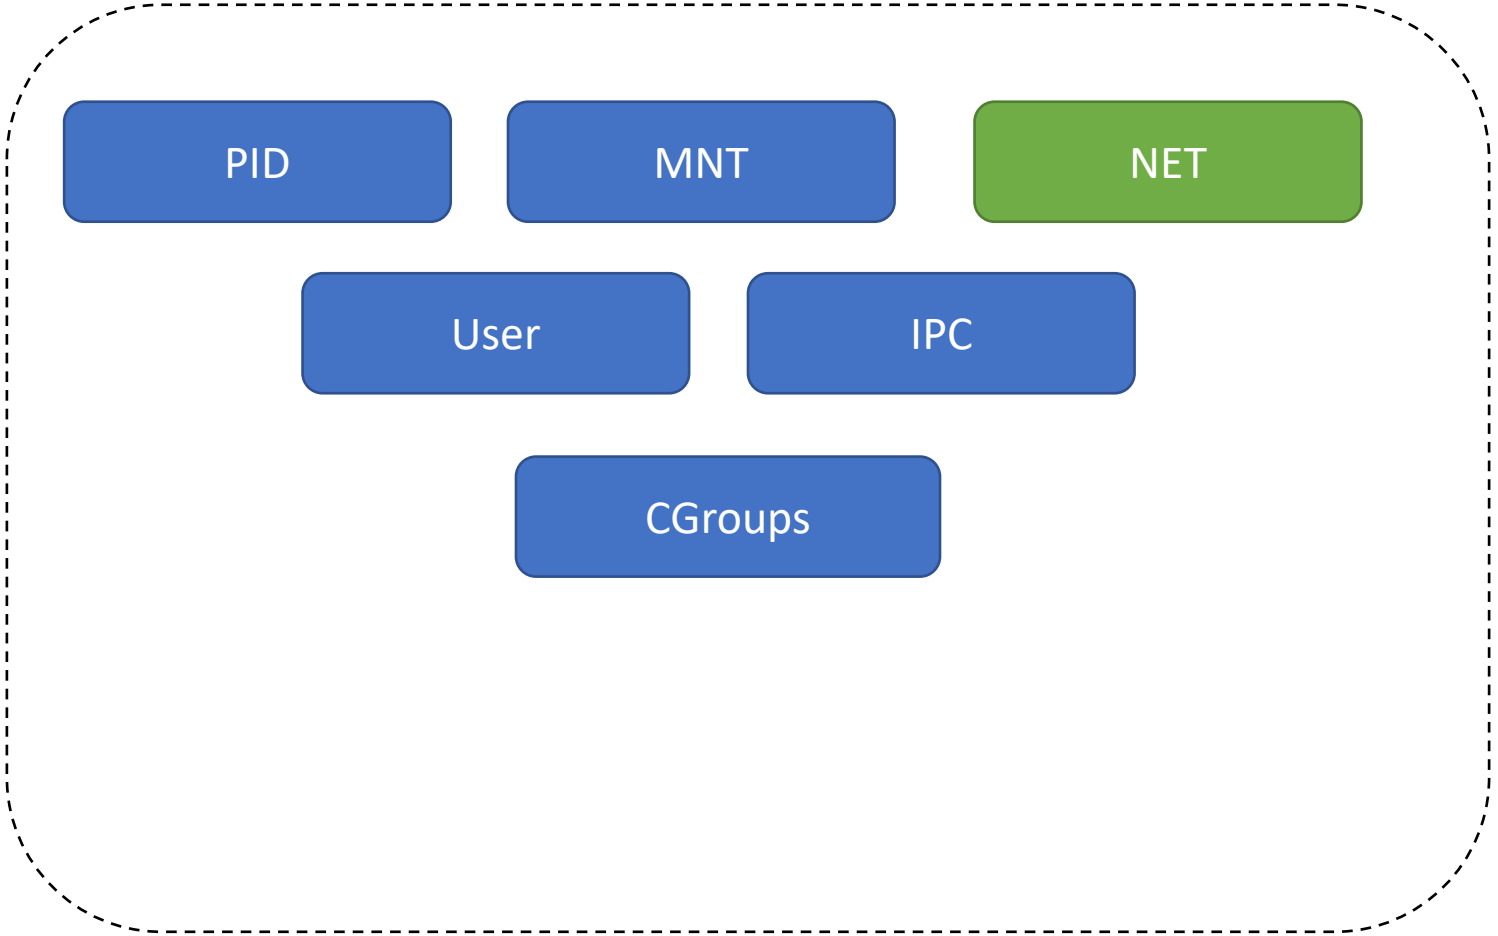

Review

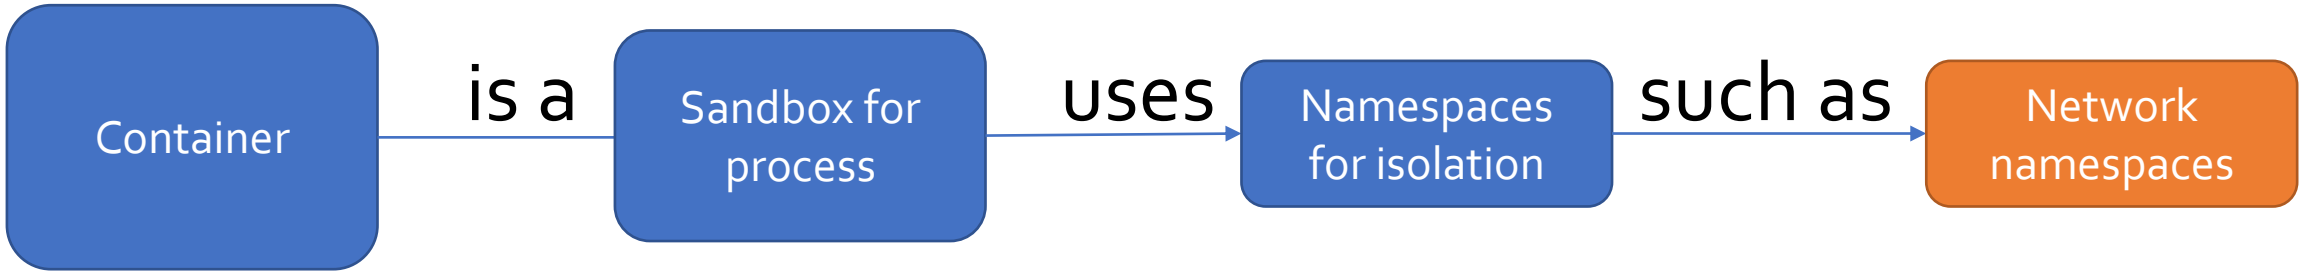
Container
is a
Sandbox for process
uses
Namespaces for isolation
such as
Network namespaces

Container Platforms and Docker

These are all kernel capabilities
i.e. you can build a container right now without additional software

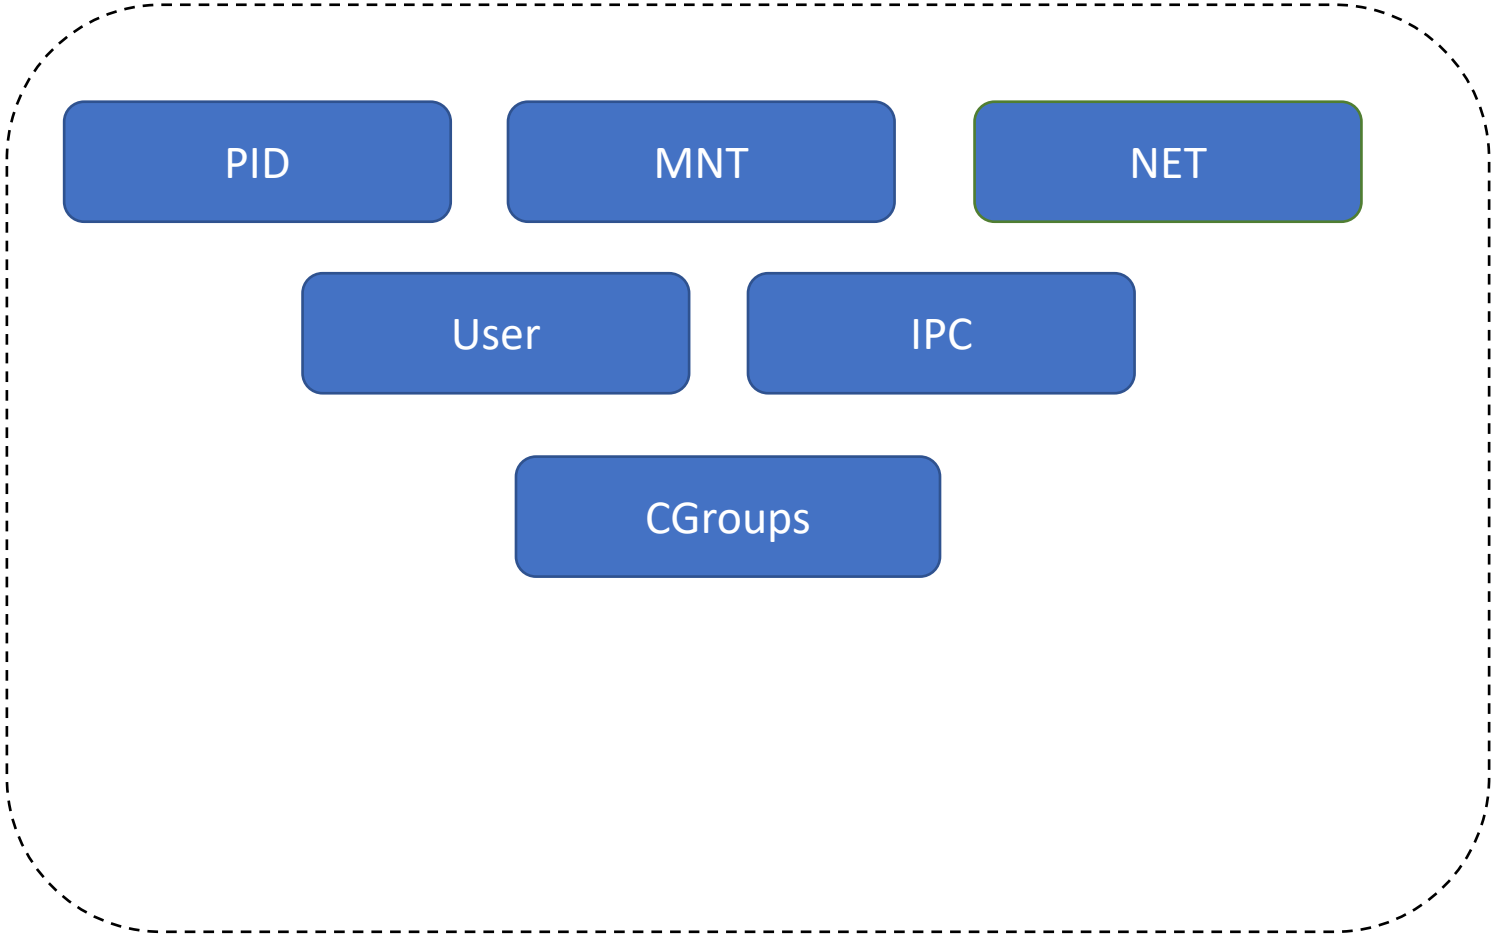

Struggles of Creating Namespaces

```
$ ip link add veth0 type veth peer name veth1
$ ip link set veth1 netns pa1
$ ip netns exec pa1 ip addr add 10.1.1.1/24 dev veth1
$ ip netns exec pa1 ip link set dev veth1 up
```

Repeat for additional isolation and you've created *one container*

```
$ ip link add veth0 type veth peer name veth1
$ ip link set veth1 netns pa1
$ ip netns exec pa1 ip addr add 10.1.1.1/24 dev veth1
$ ip netns exec pa1 ip link set dev veth1 up
```

Scenario Three

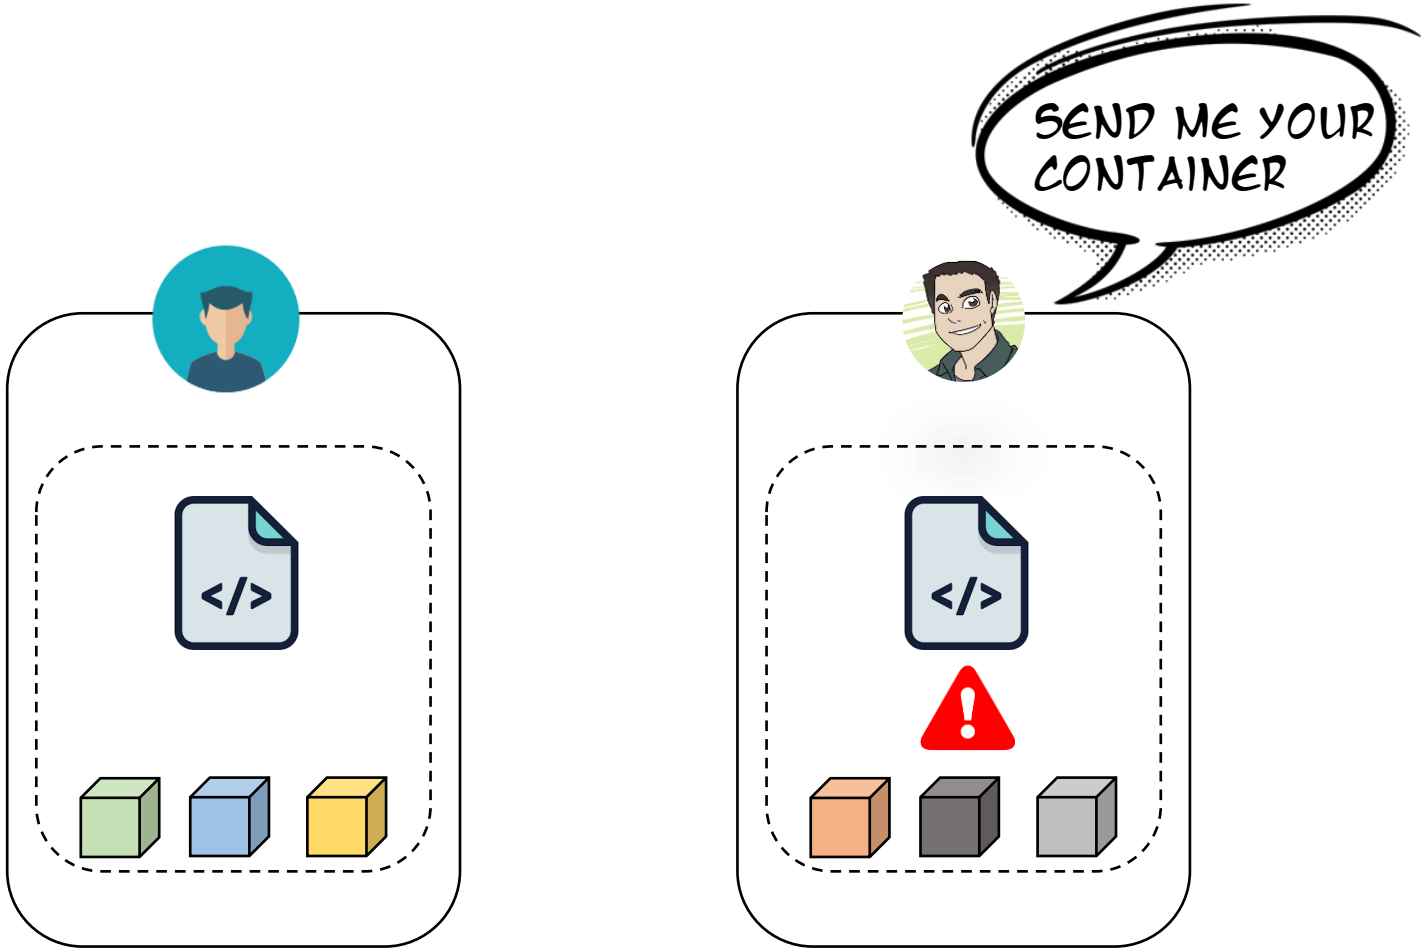

Scenario Three

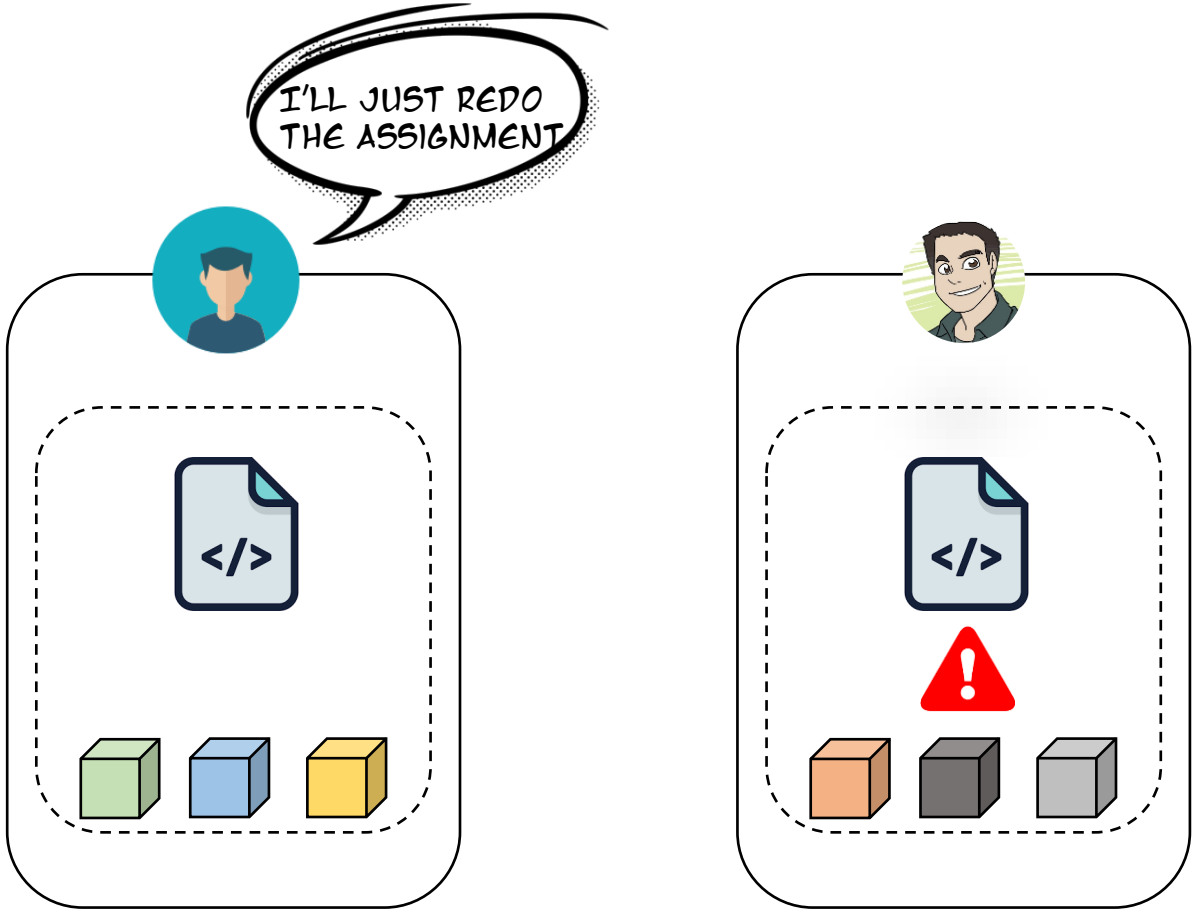

Container software makes it easy to manage containers and namespaces

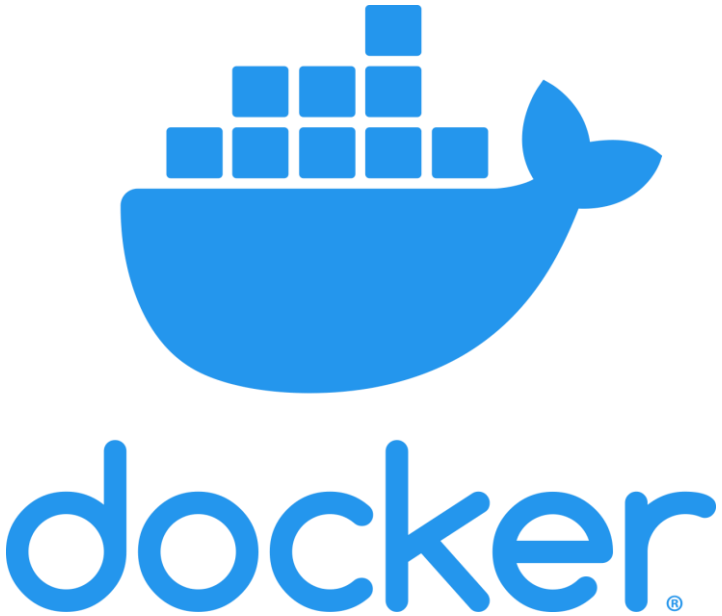

Docker abstracts underlying namespaces

Create Dockerfile

```
FROM ubuntu
RUN apt-get update
RUN apt-get install mysql:5.6.50 -y
COPY /usr/shelbyt/cse224/pa1/
/usr/cse224/pa1/
```

Docker abstracts underlying namespaces

Create Dockerfile

```
FROM ubuntu
RUN apt-get update
RUN apt-get install mysql:5.6.50 -y
COPY /usr/shelbyt/cse224/pa1/
/usr/cse224/pa1/
```

Build an Image

```
docker build –t  pa1 .
```

Docker abstracts underlying namespaces

Create Dockerfile

```
FROM ubuntu
RUN apt-get update
RUN apt-get install mysql:5.6.50 -y
COPY /usr/shelbyt/cse224/pa1/
/usr/cse224/pa1/
```

Build an Image

```
docker build –t  pa1 .
```

Run

```
docker run –it pa1
% (pa1-shell):
```

Docker abstracts underlying namespaces

Create Dockerfile

```
FROM ubuntu
RUN apt-get update
RUN apt-get install mysql:5.6.50 -y
COPY /usr/shelbyt/cse224/pa1/
/usr/cse224/pa1/
```

Build an Image

```
docker build –t  pa1 .
```

Run

```
docker run –it pa1
% (pa1-shell):
```

Isolated shell, no manual configuration of namespaces needed →

Scenario Three

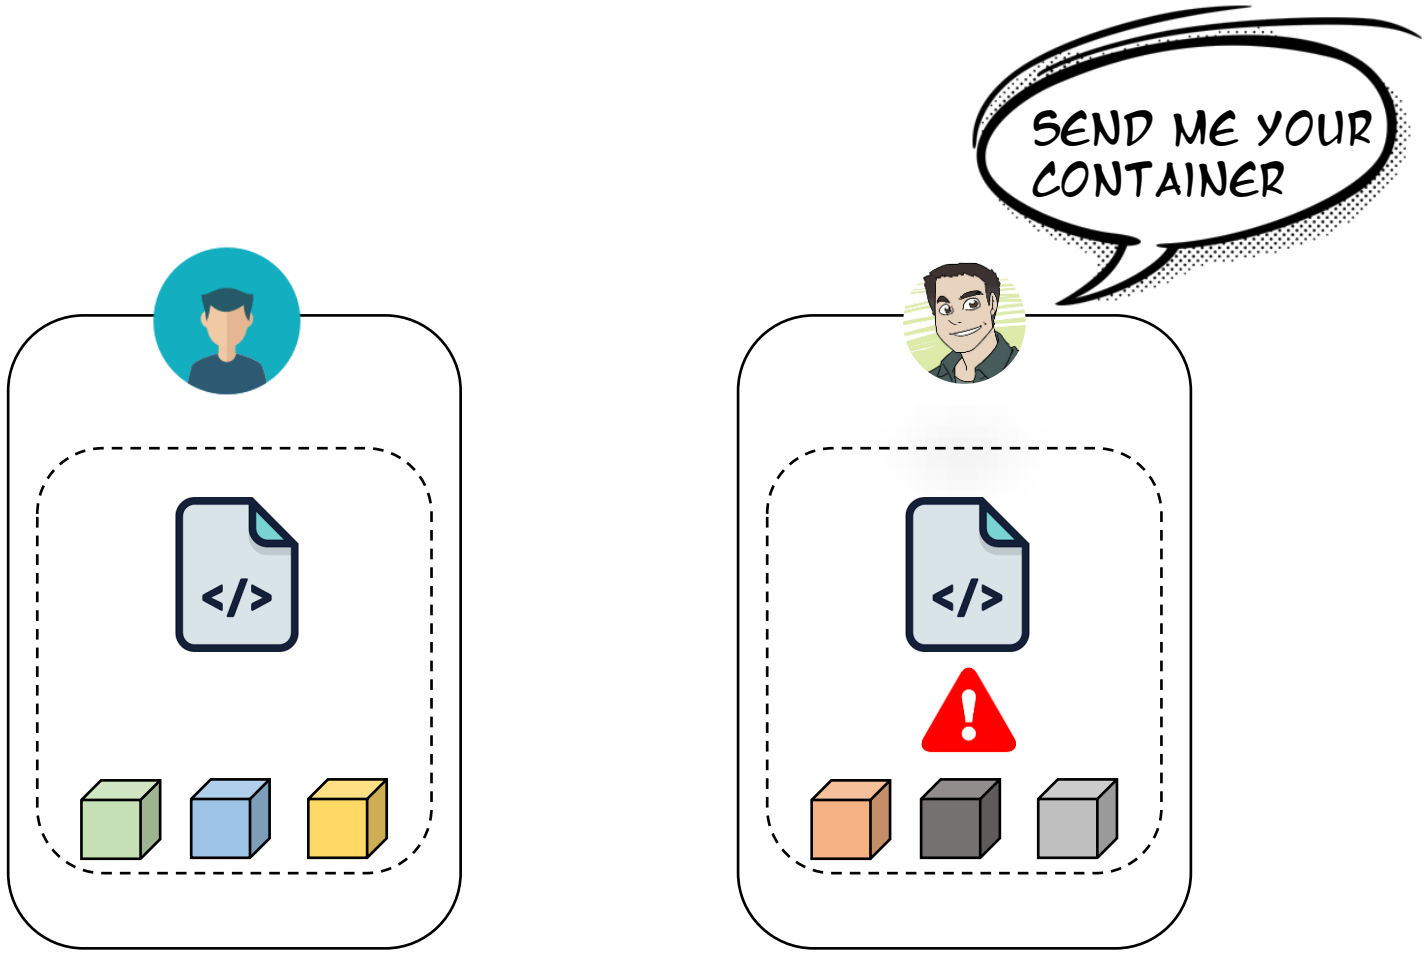

Scenario Three

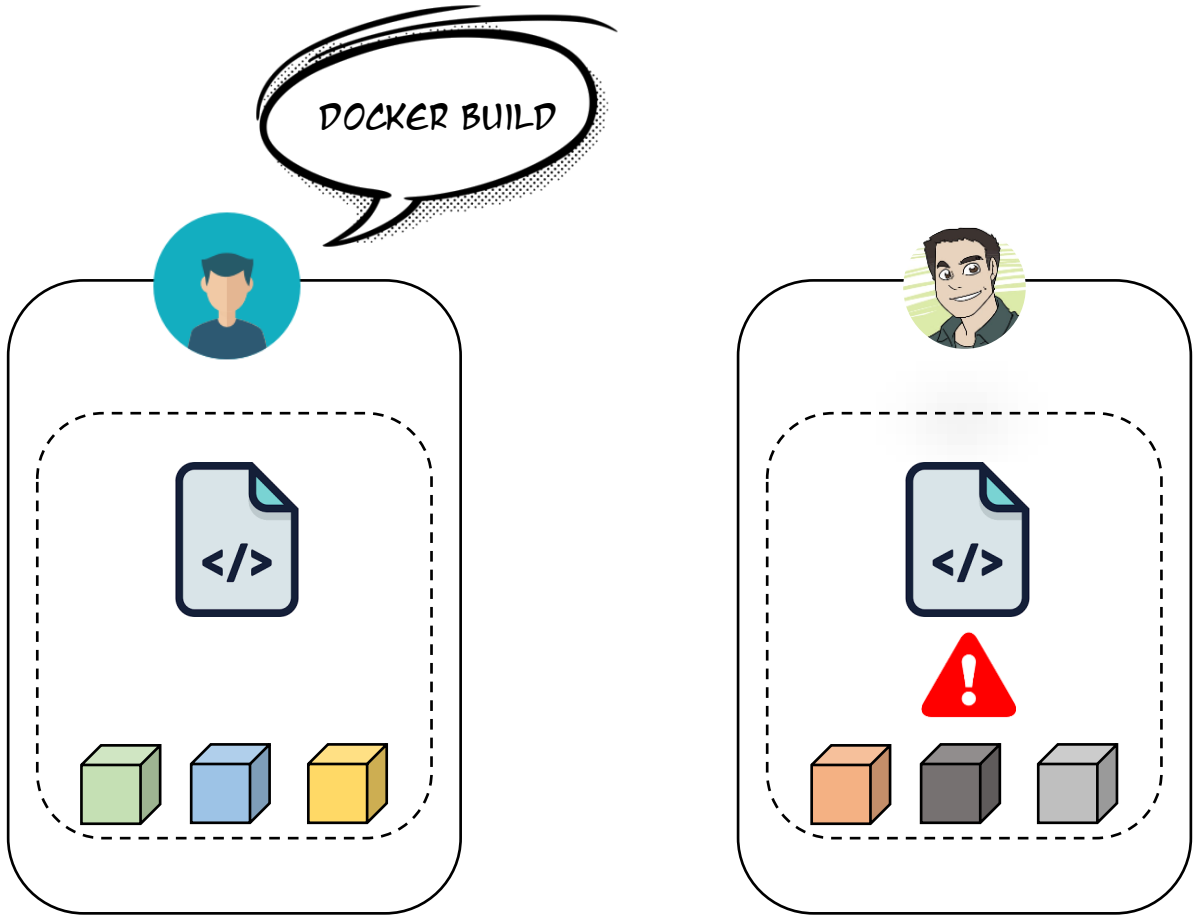

Just send the image

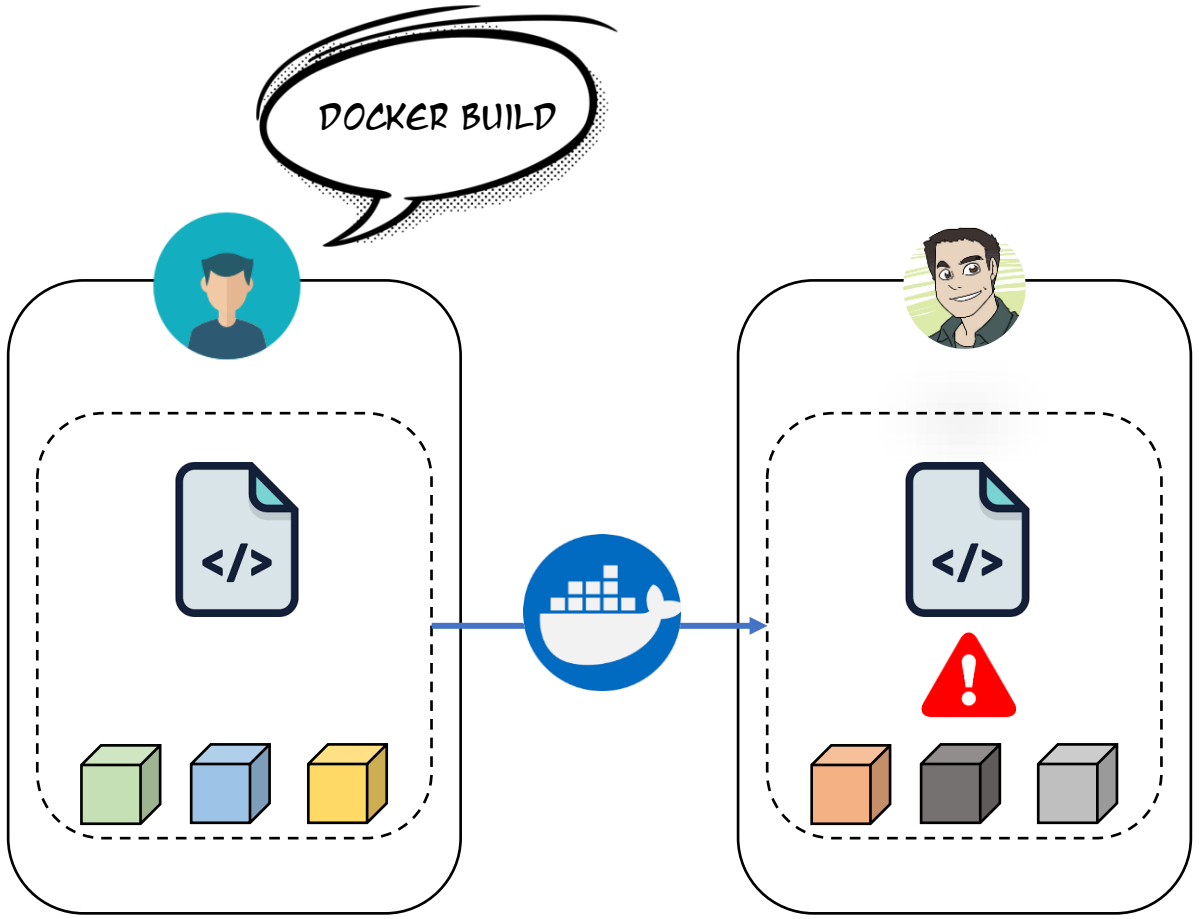

Just send the image

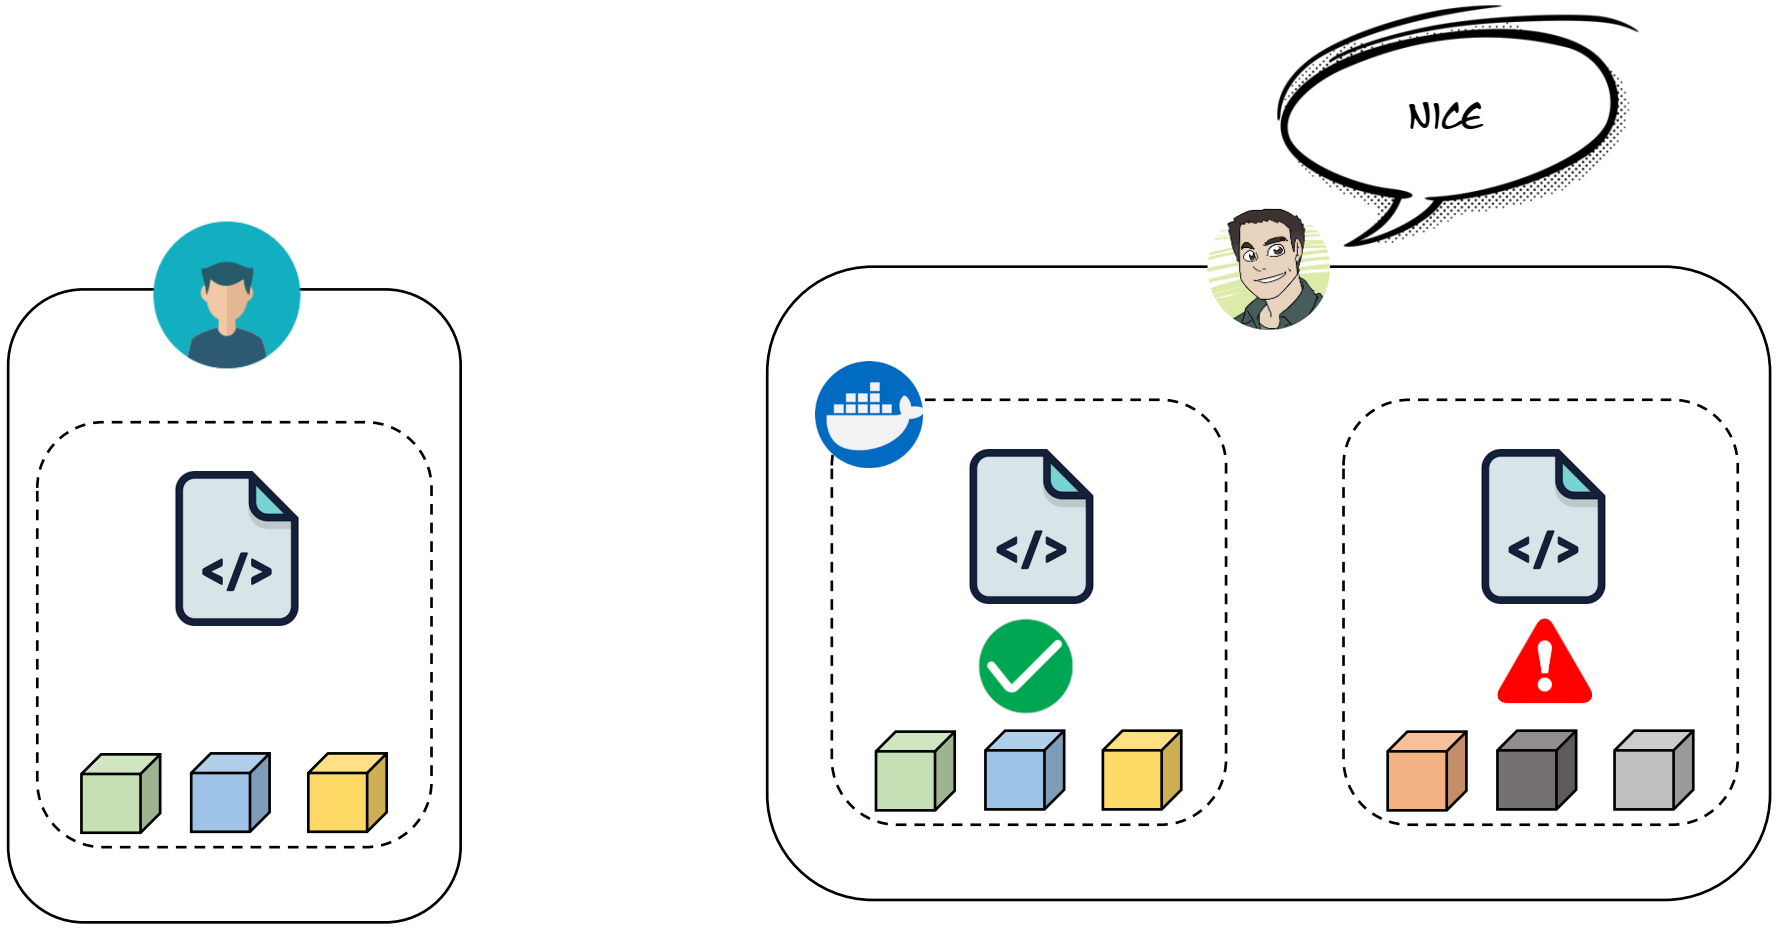

Review

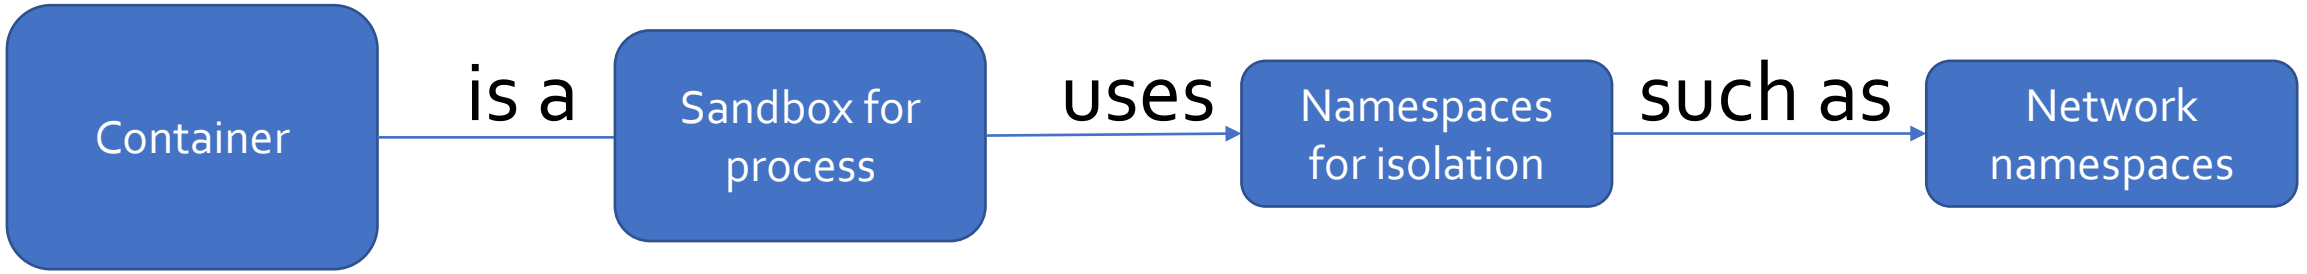
Container
is a
Sandbox for process
uses
Namespaces for isolation
such as
Network namespaces

Review

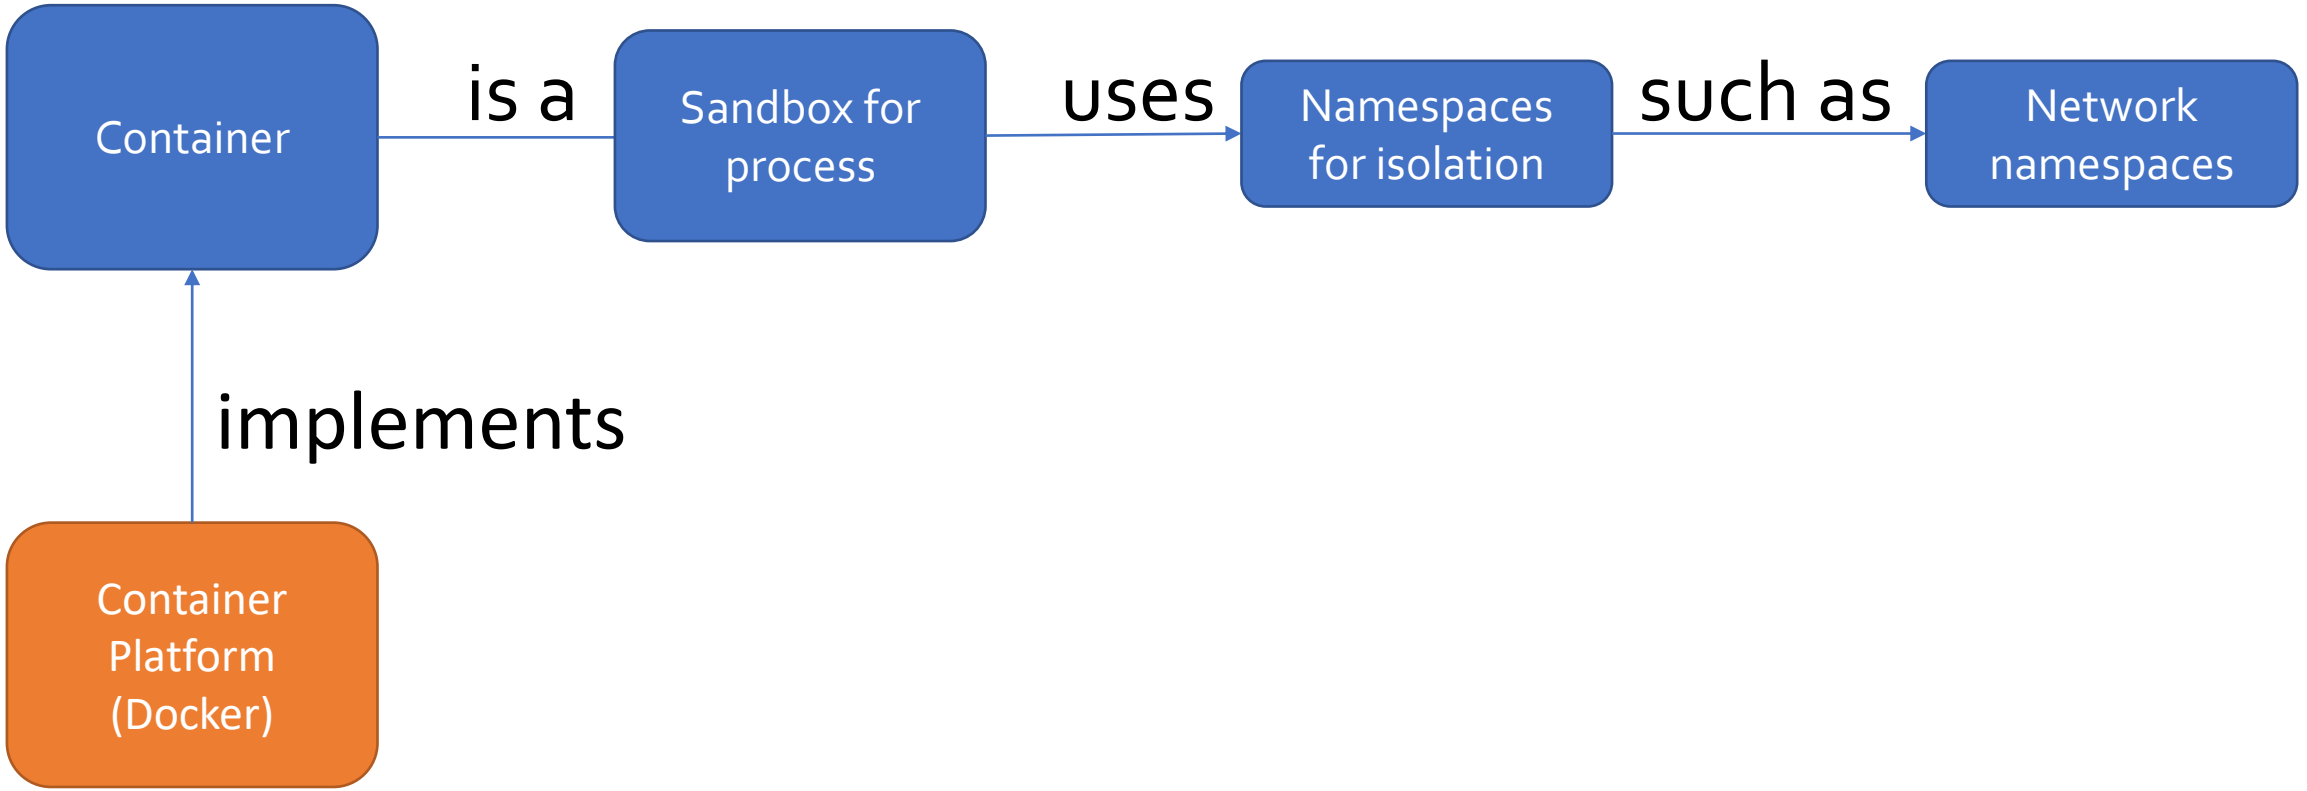
Container
is a
Sandbox for process
uses
Namespaces for isolation
such as
Network namespaces
implements
Container Platform (Docker)

Virtual Machine vs. Containers

Why not use a Virtual Machine?

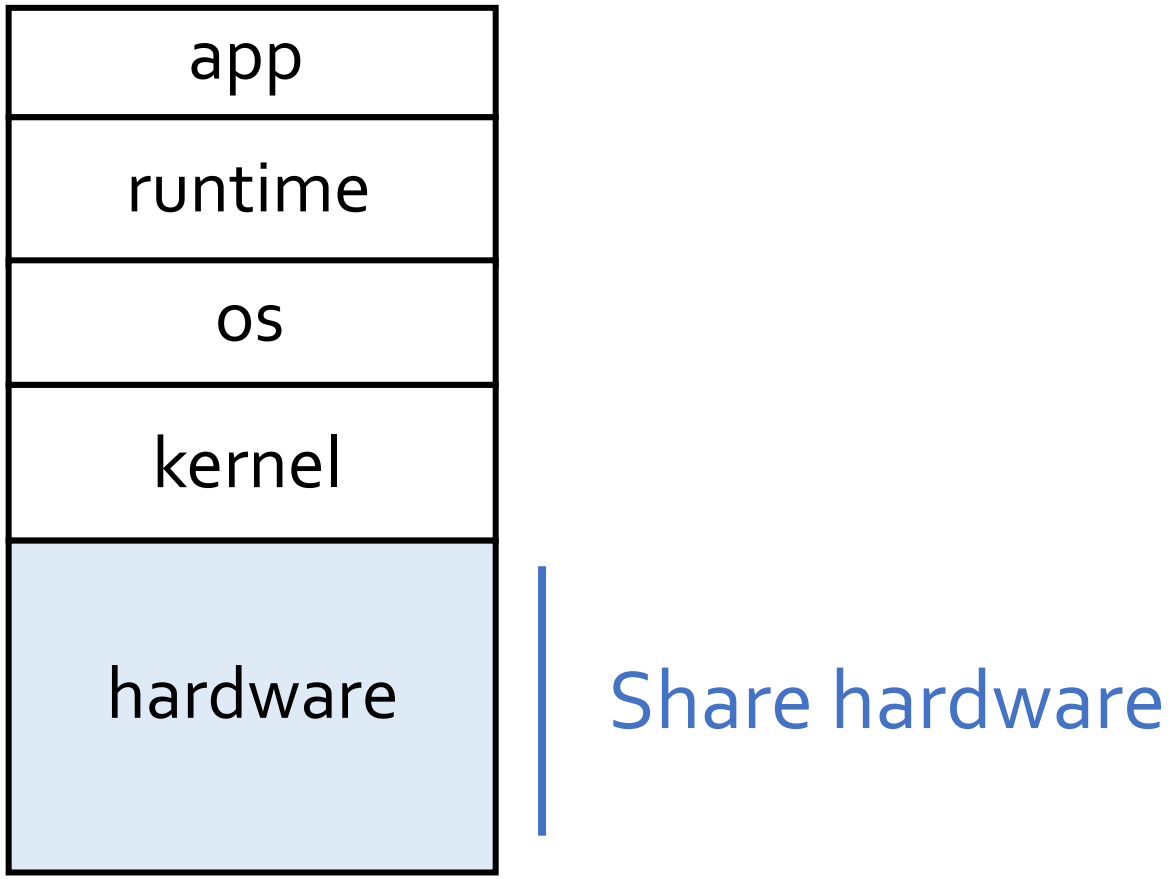

Different Isolation Points

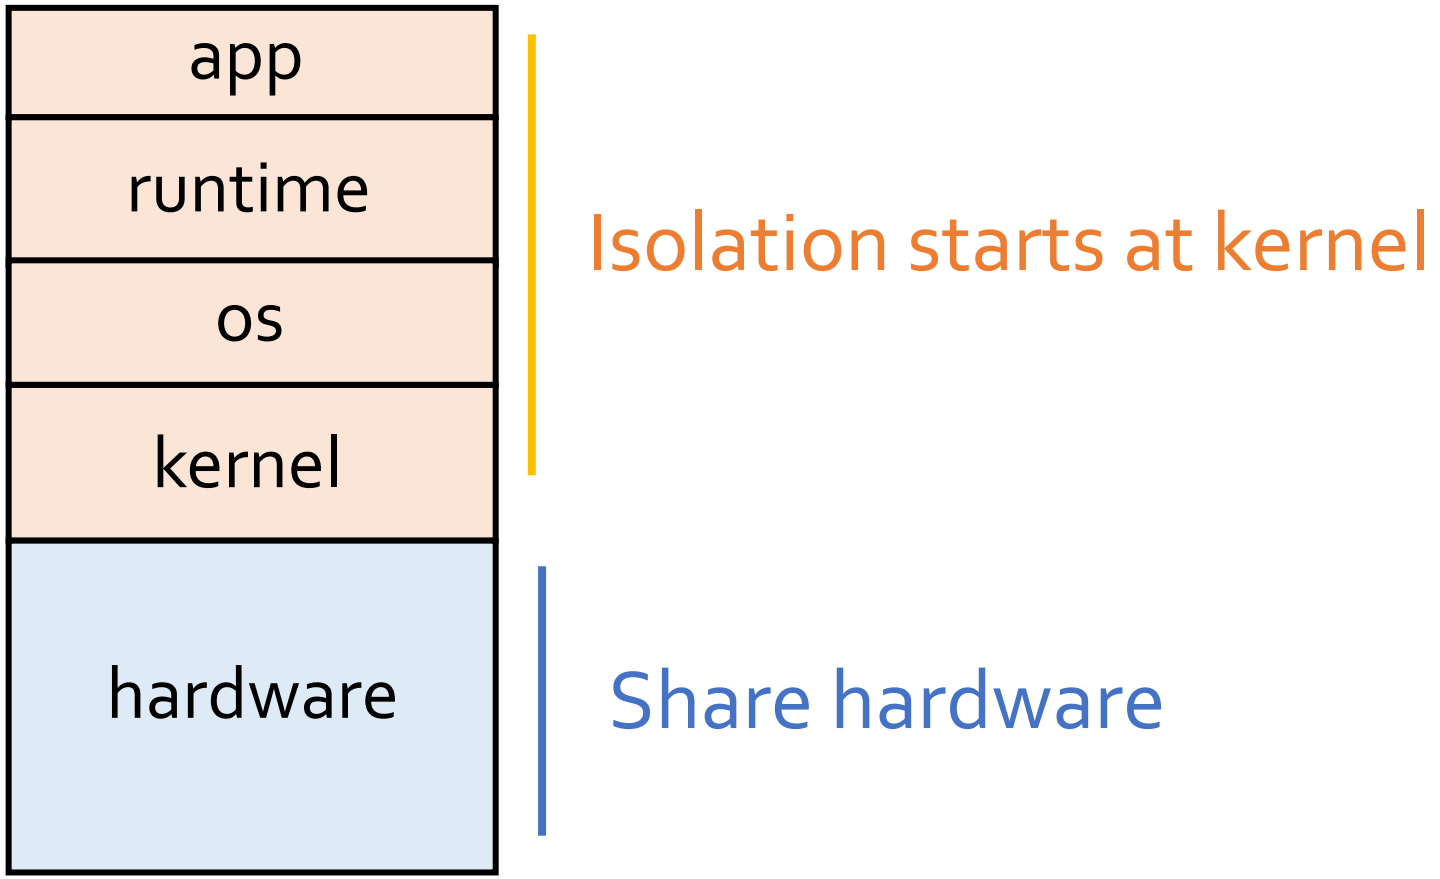

Different Isolation Points

This is a VM Image

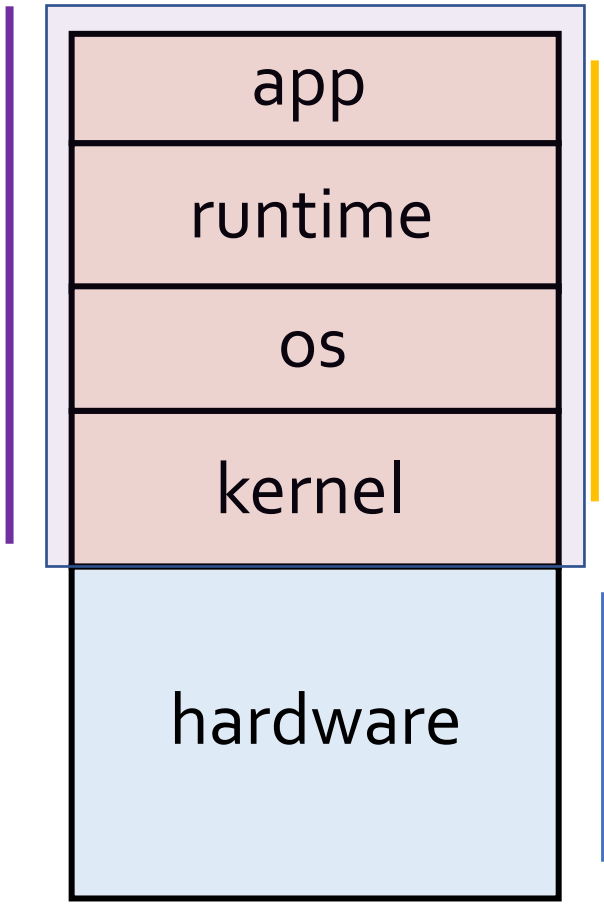

Isolation starts at kernel

Share hardware

Different Isolation Points

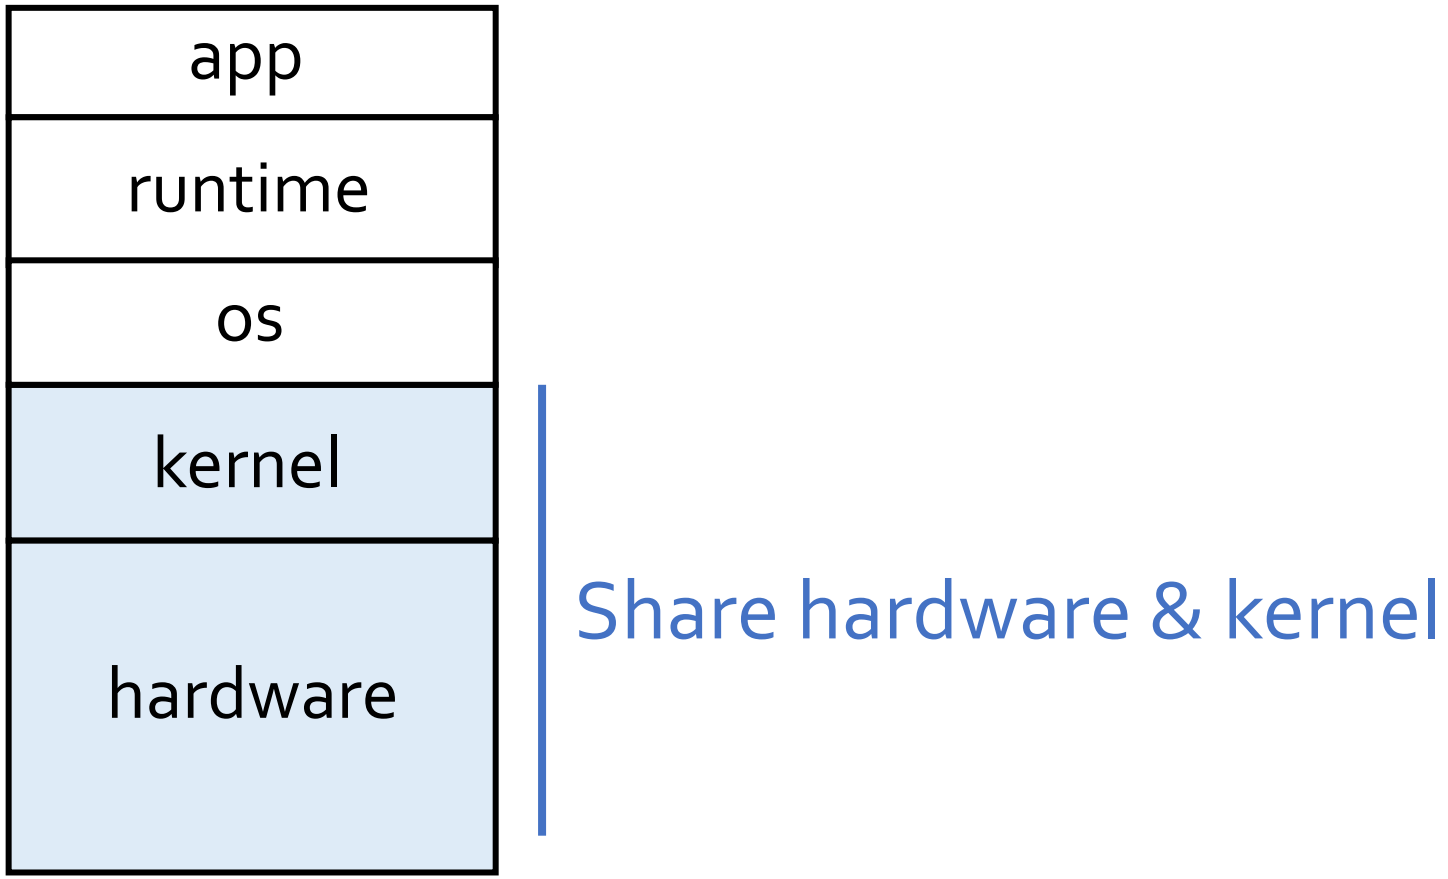

Different Isolation Points

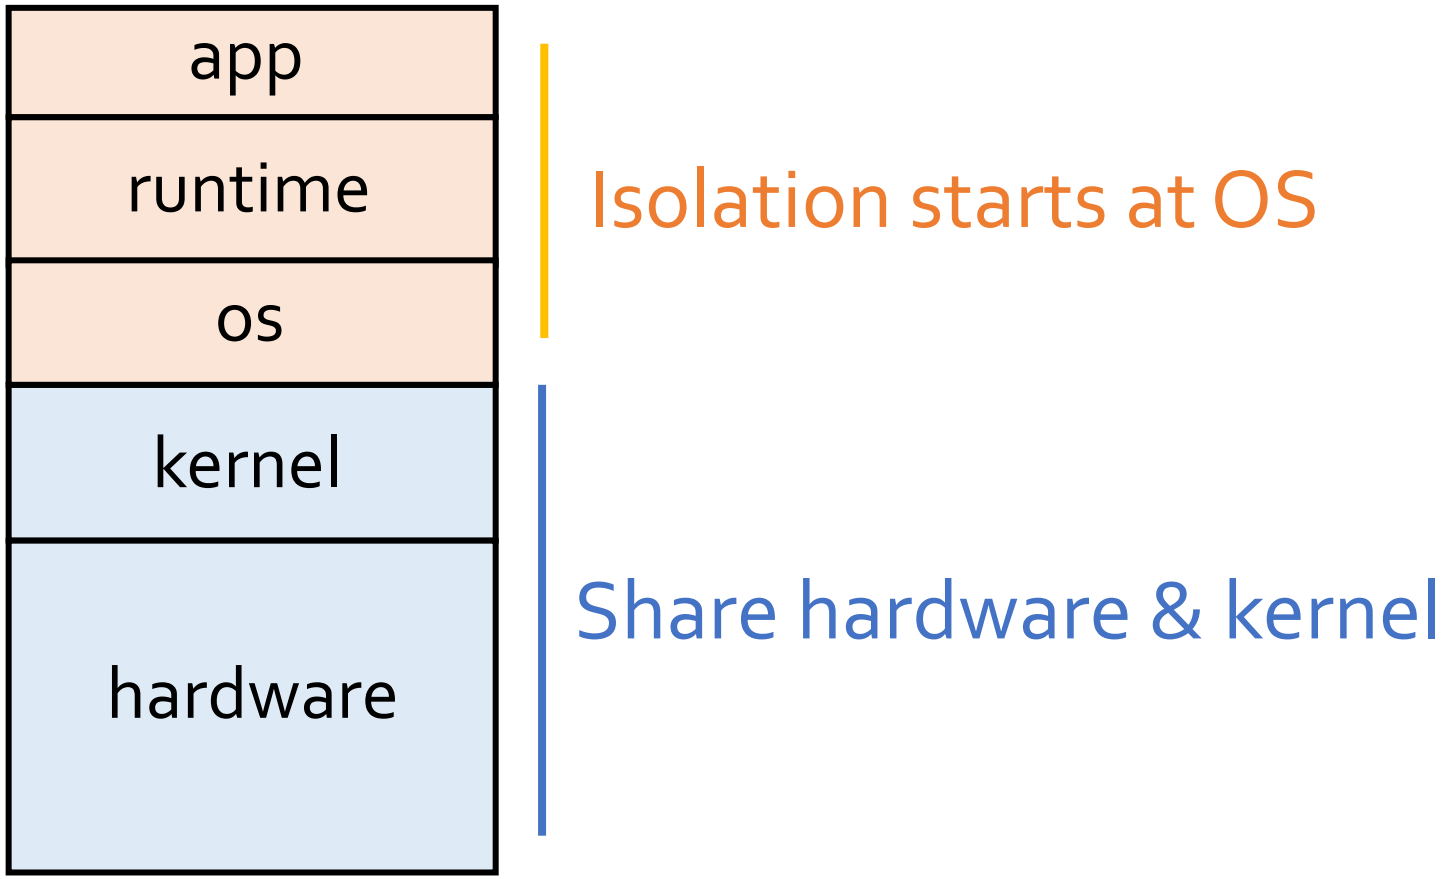

Different Isolation Points

This is a container image

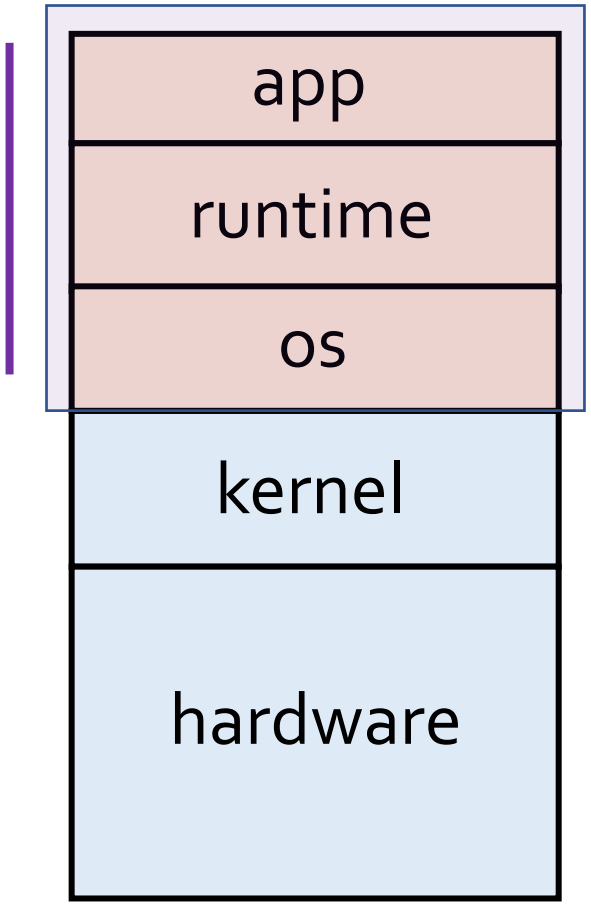

Isolation starts at OS

Share hardware & kernel

Container Benefits

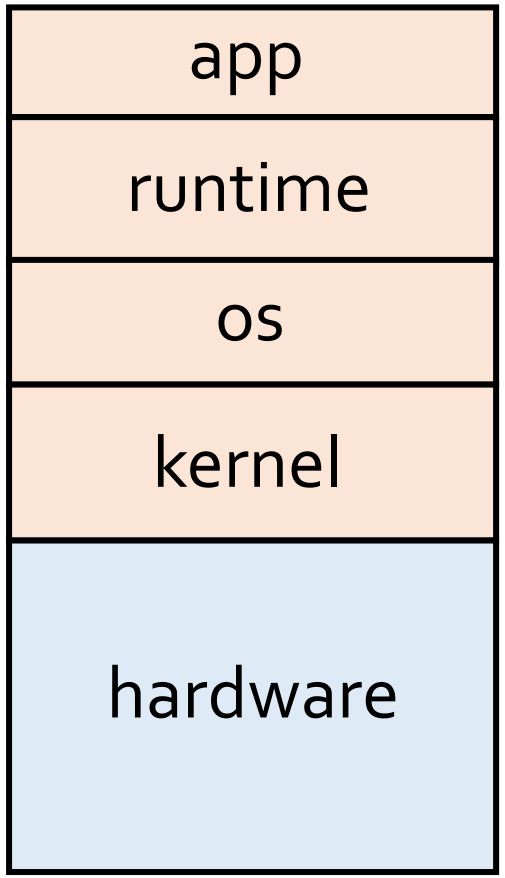

virtual machine

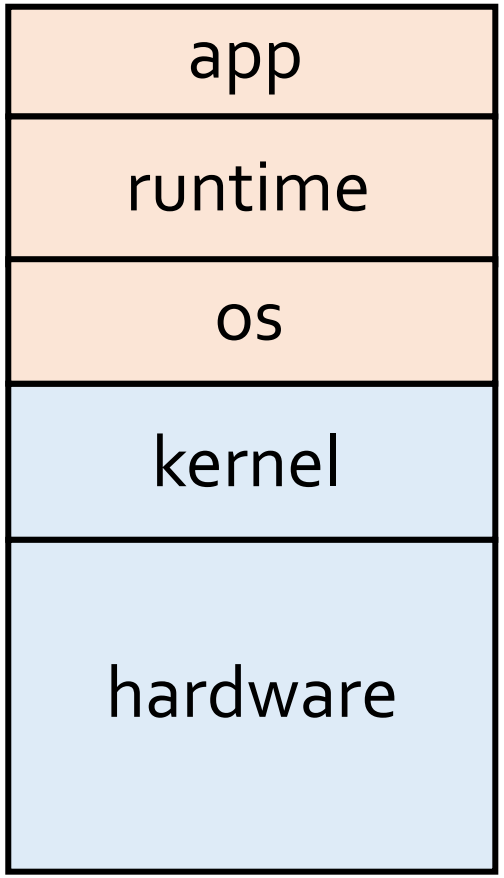

container

- Lightweight

  (megabytes vs gigabytes [10x])
- Fast startup

  (milliseconds vs. minutes [100x])
- Simple to build, deploy, send, & maintain

Benefits come from specialization

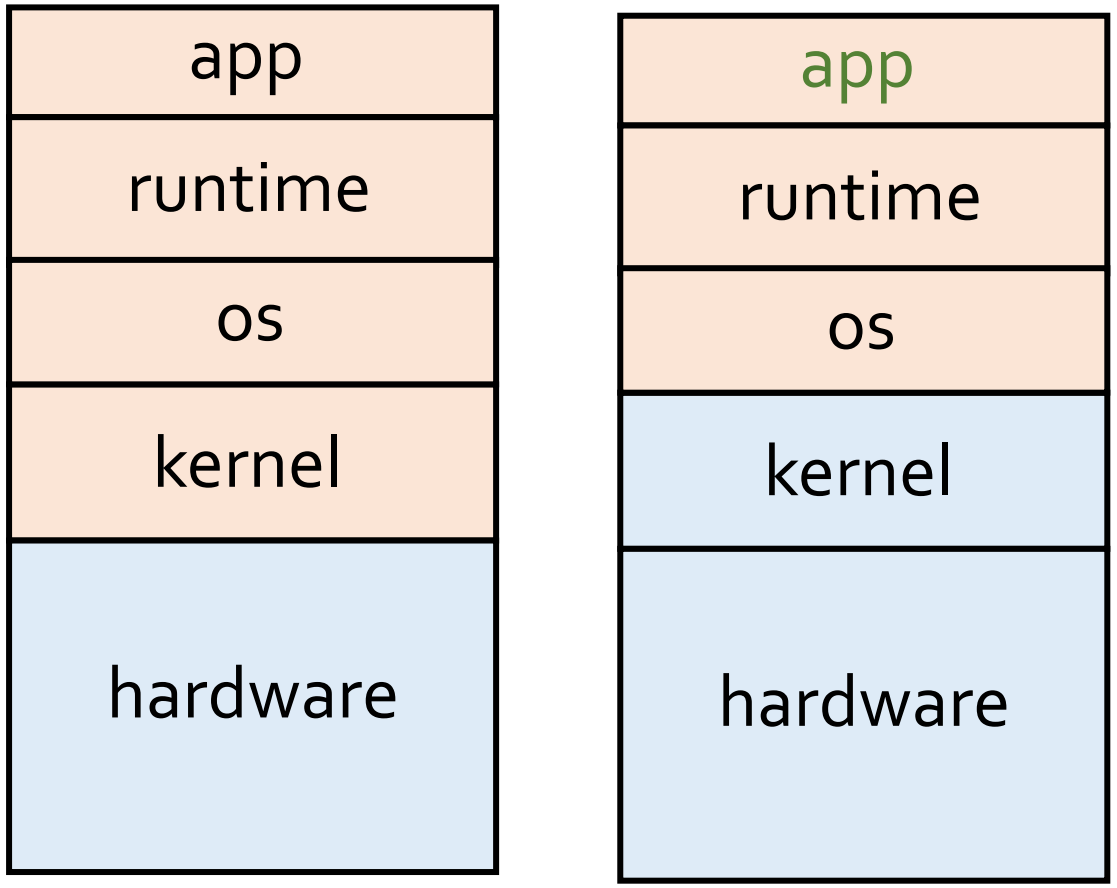

virtual machine

container

- Lightweight

  (megabytes vs gigabytes [10x])

- Fast startup

  (milliseconds vs. minutes [100x])

- Simple to build, deploy, send, & maintain

Benefits come from specialization

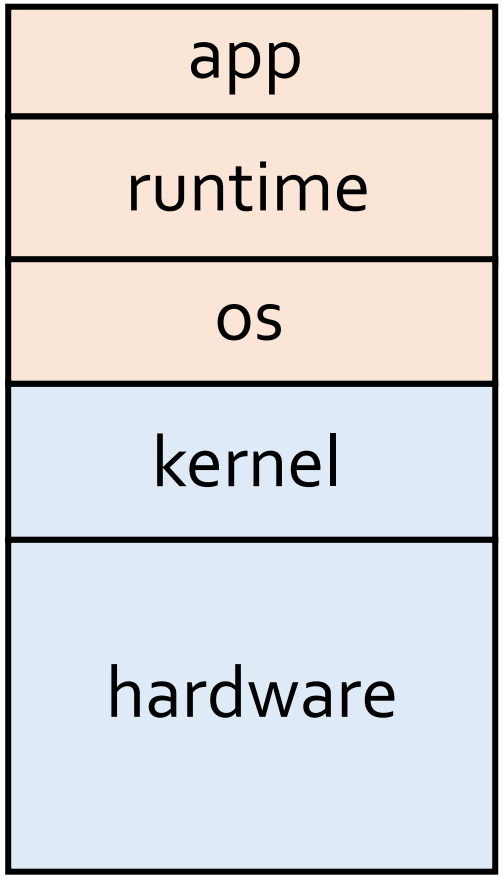

Built to support all applications and users

Benefits come from specialization

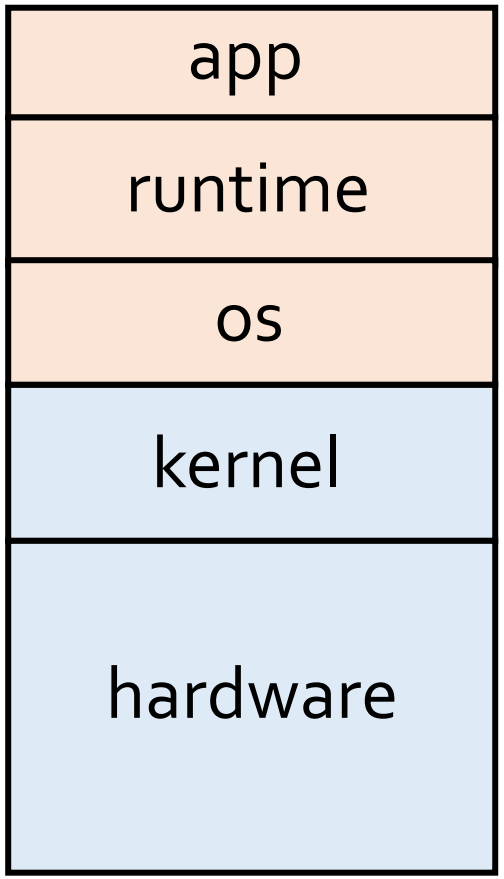

But we only care about a single application

Benefits come from specialization

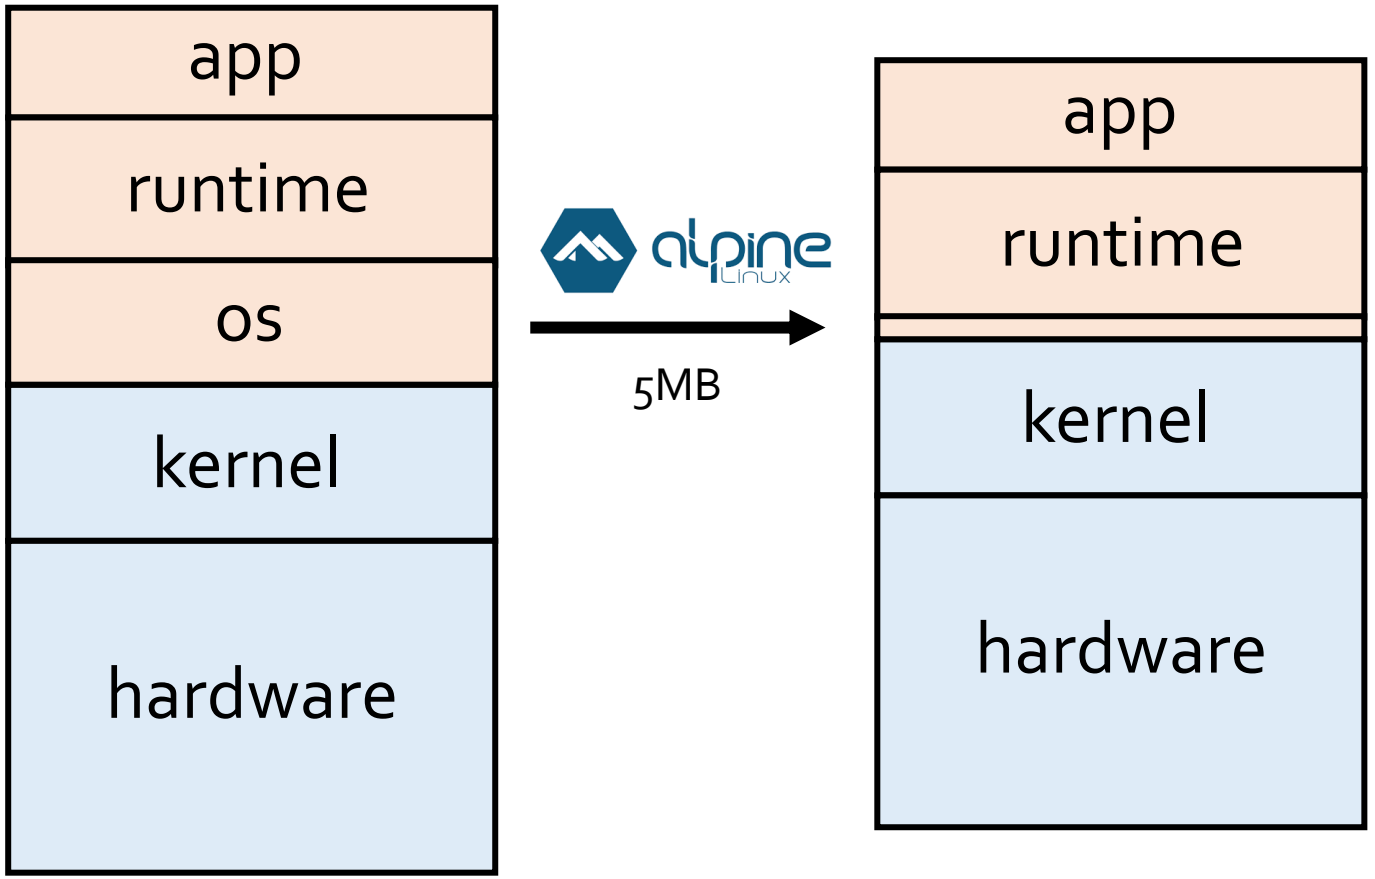

Benefits come from specialization

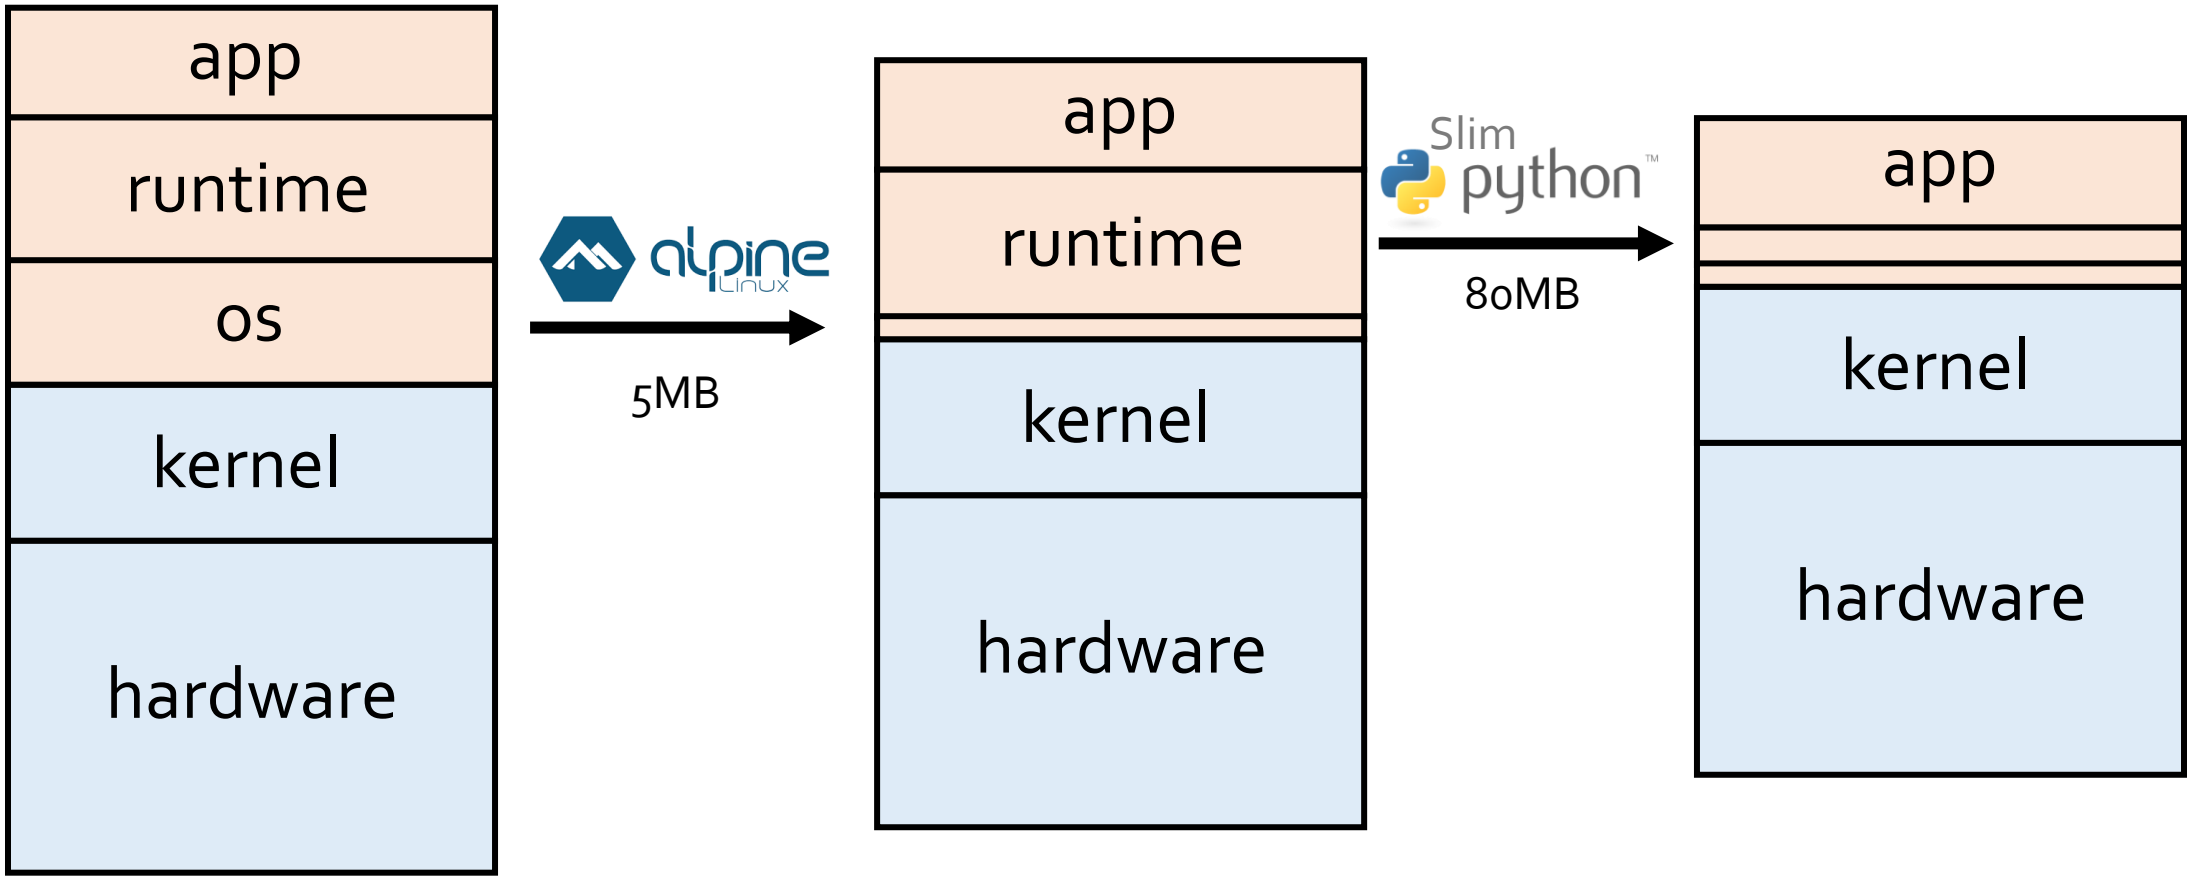

Container image is mostly application

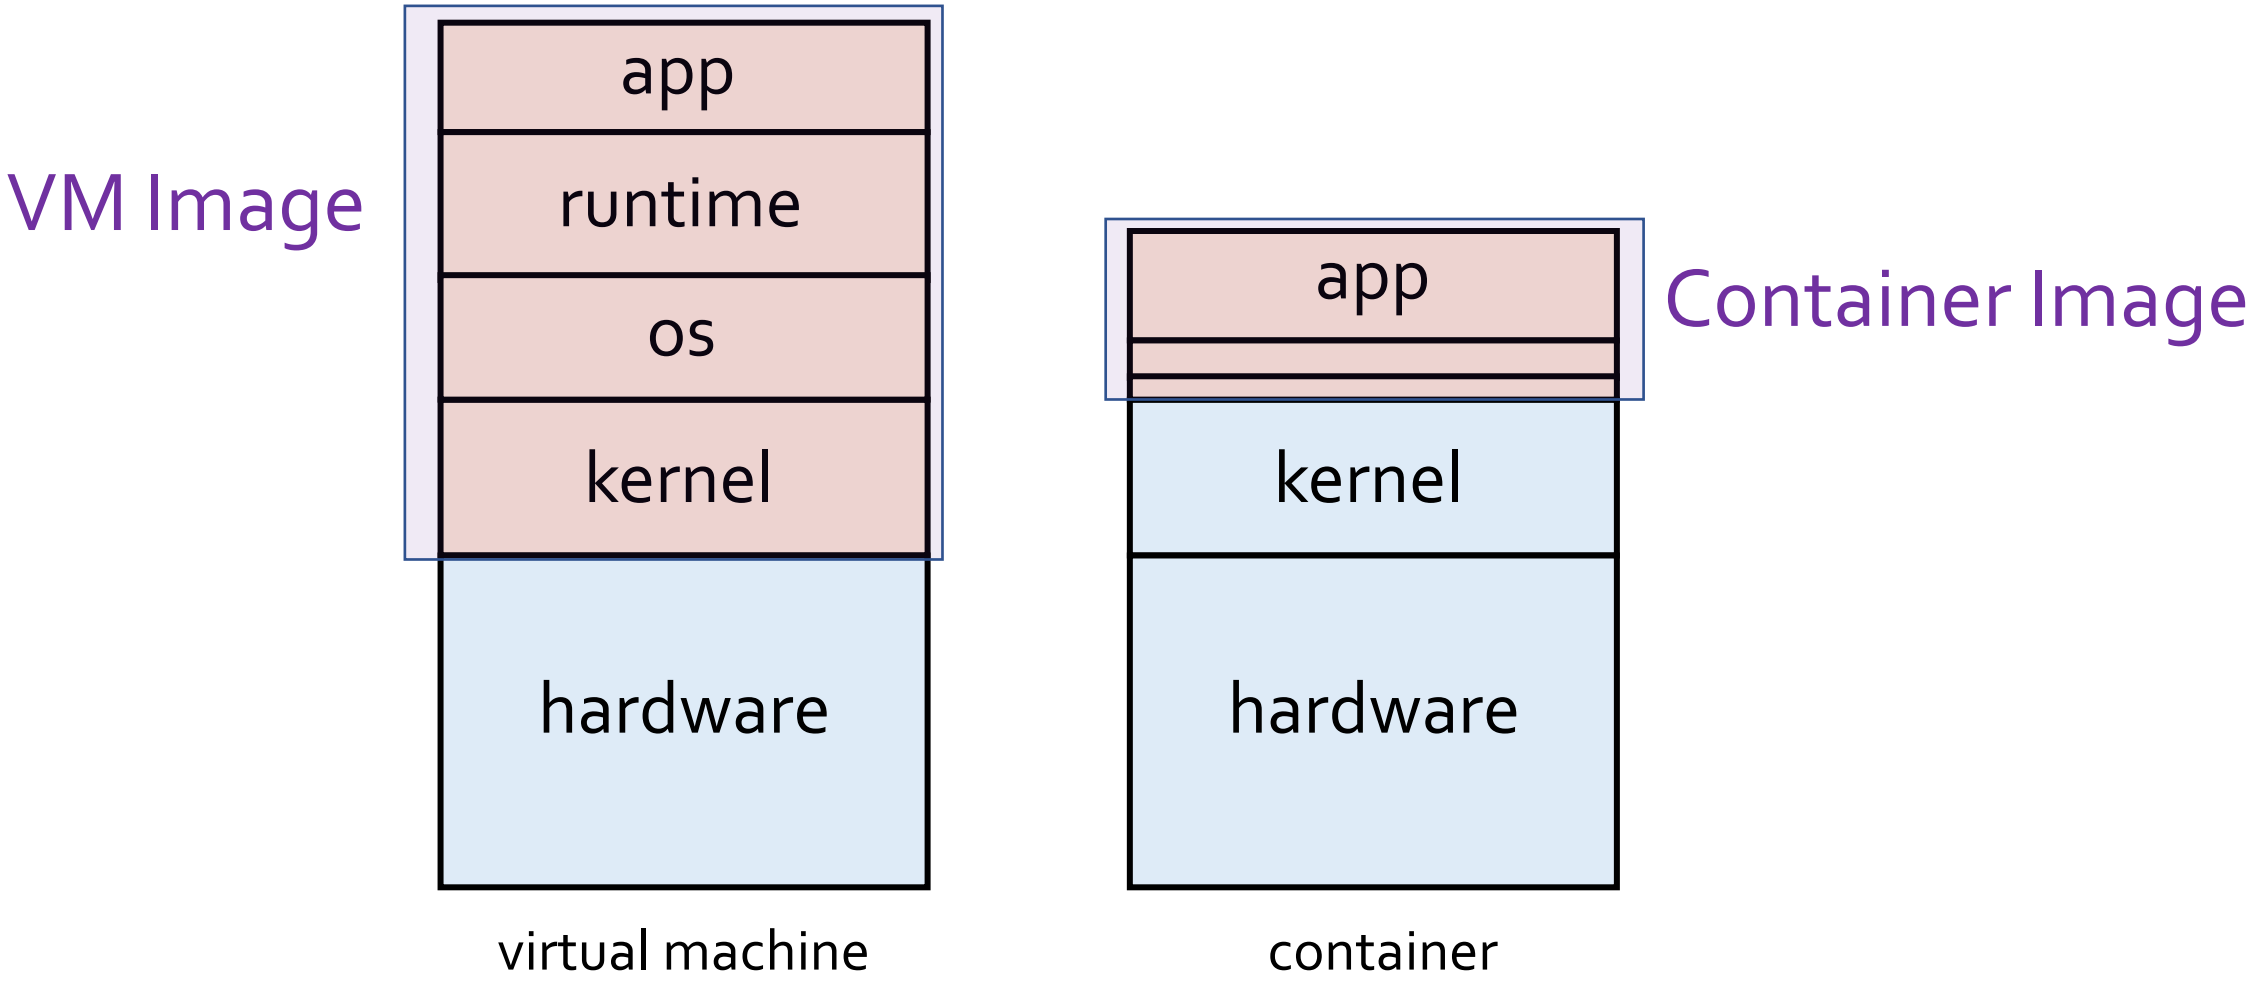

Containers to Serverless Computing

Traditional Cloud Model

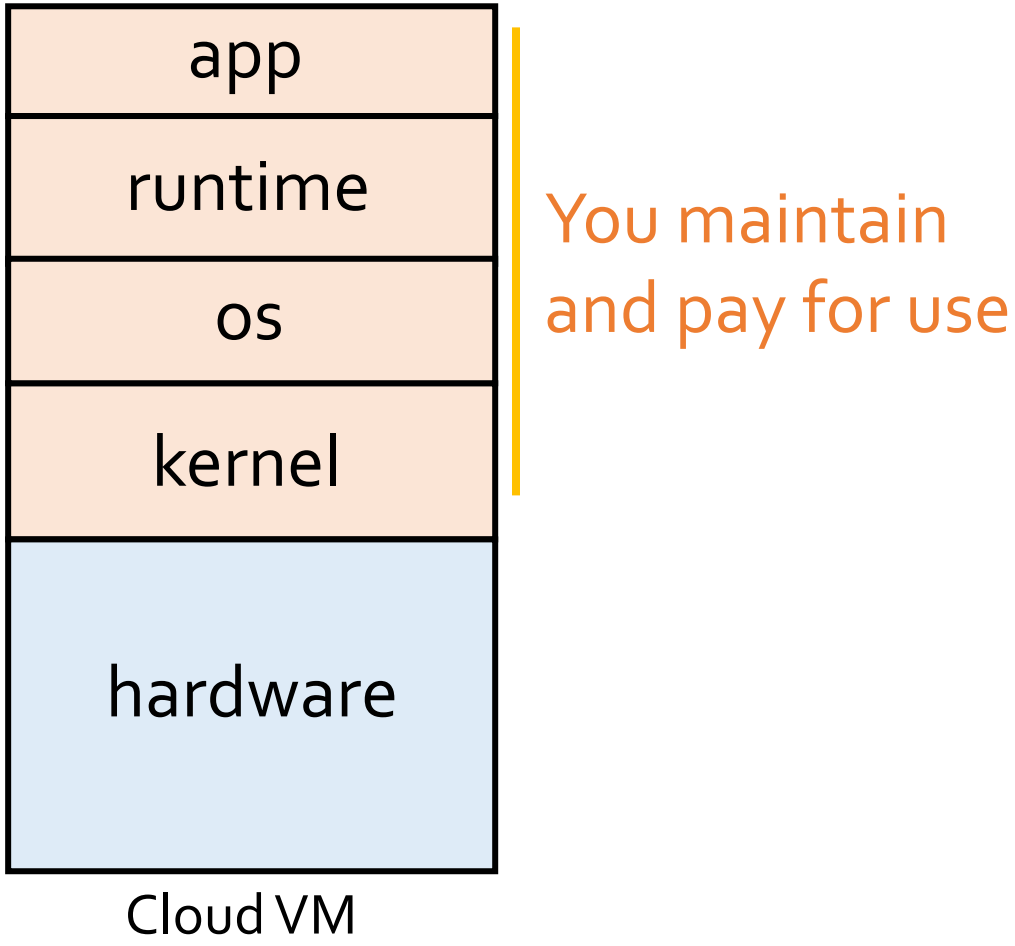

Serverless Cloud Model

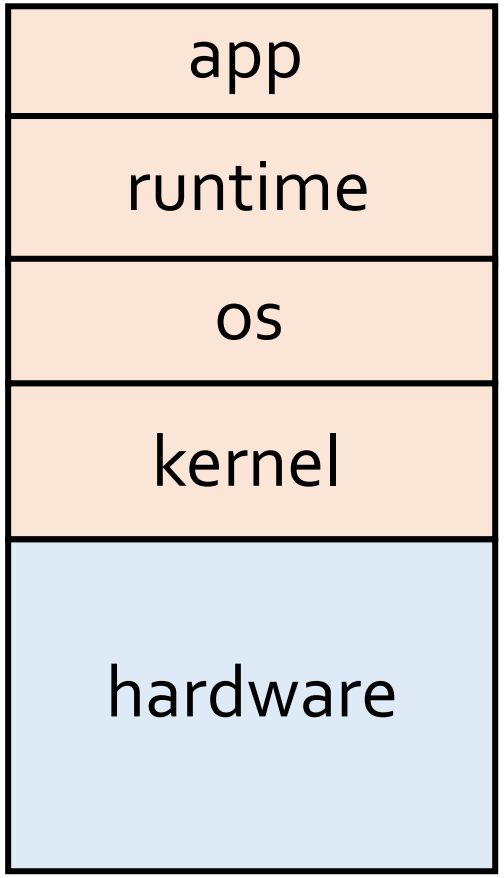

You maintain
and pay for use

Cloud VM

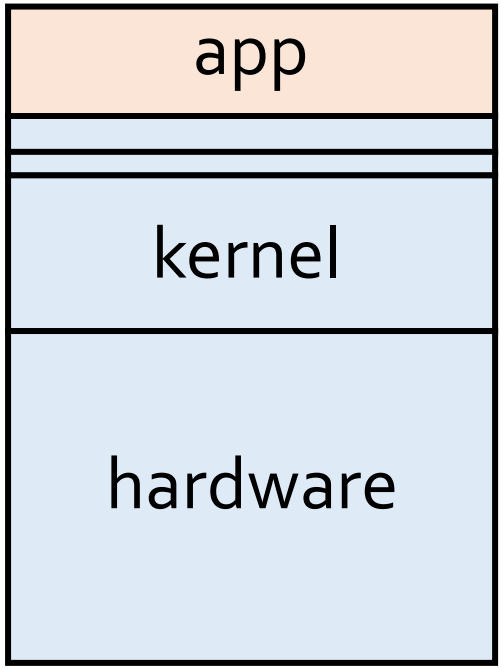

You maintain
and pay for use

Cloud Function

OS and runtime are partially managed by cloud providers

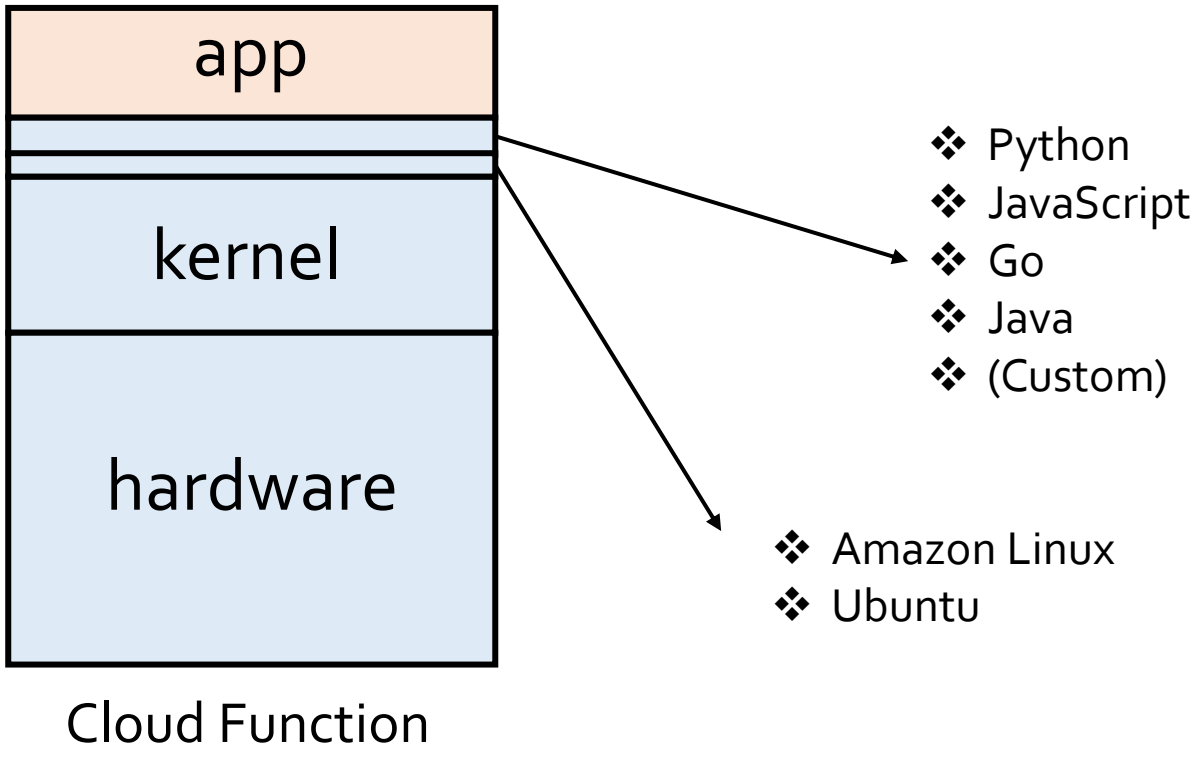

OS and runtime are partially managed by cloud providers

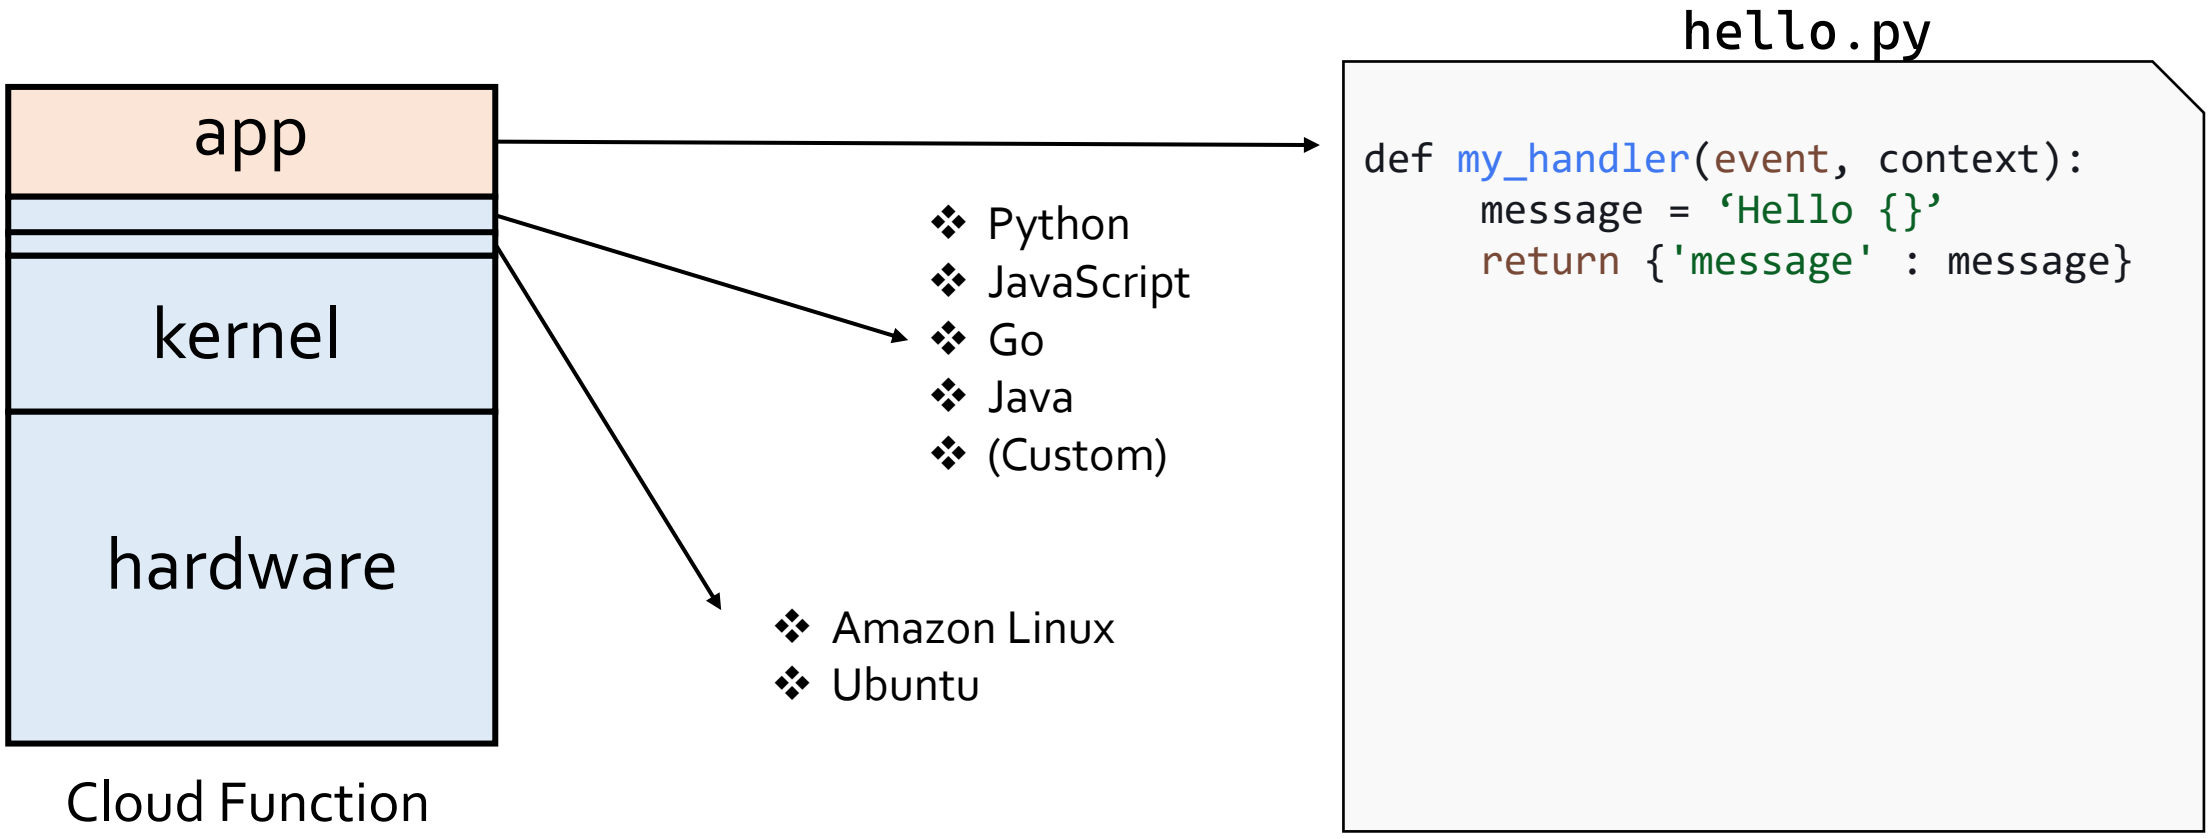

Why would we do this?

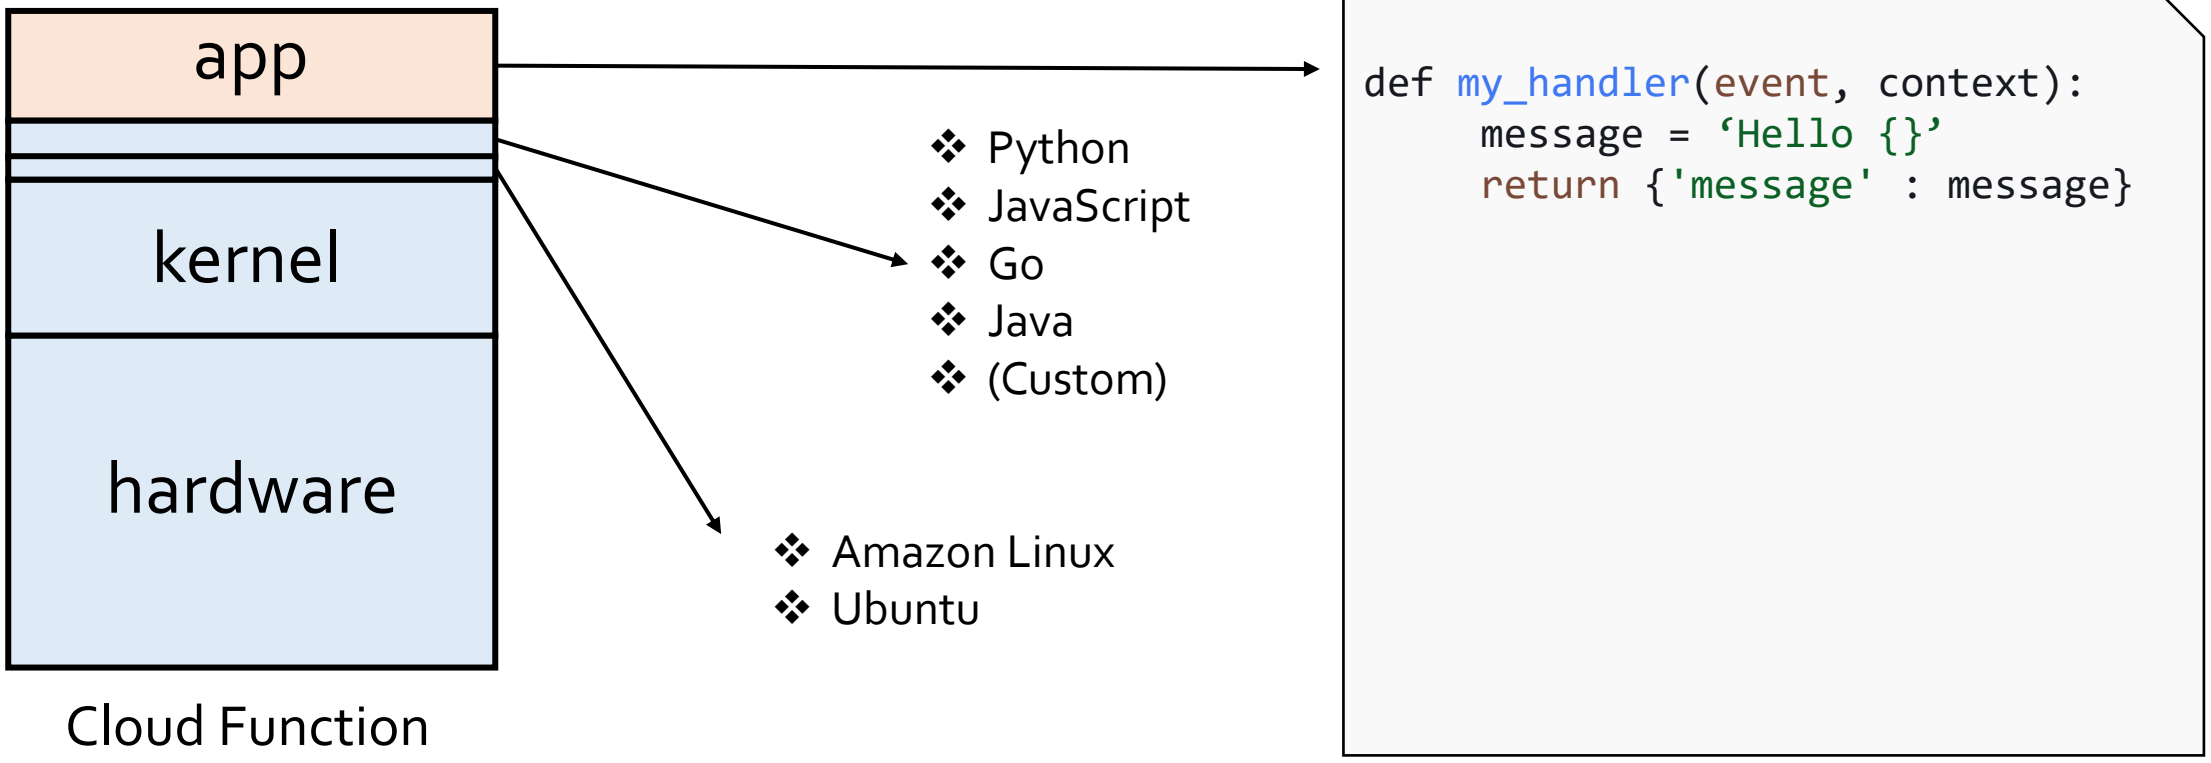

Building a GIF Creator

Building a GIF Creator

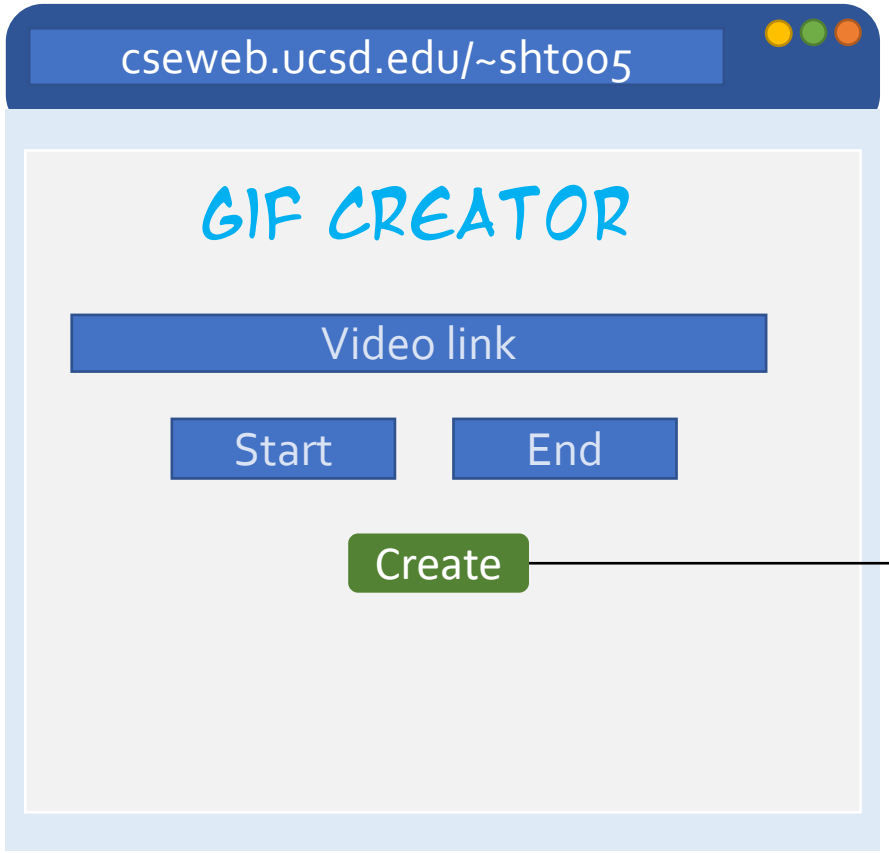

1. Download the video locally
2. Clip the video
3. Compress
4. Convert to gif

Hosting server may be too slow to do this or cannot do this

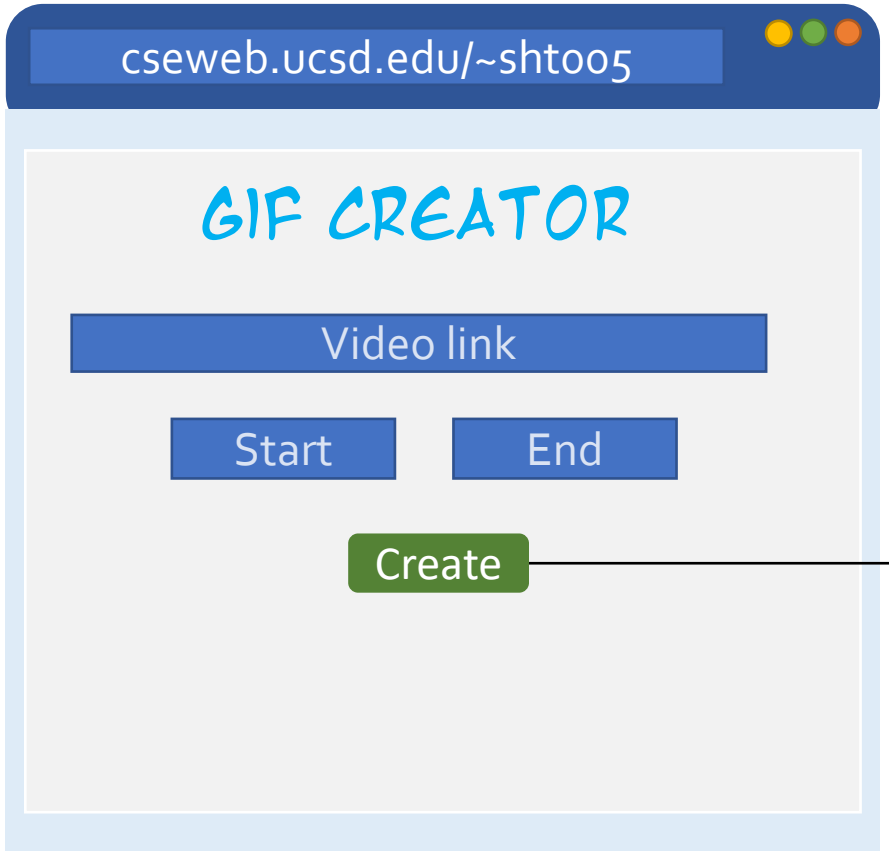

1. Download the video locally
2. Clip the video
3. Compress
4. Convert to gif

With VMs we need to provision a machine

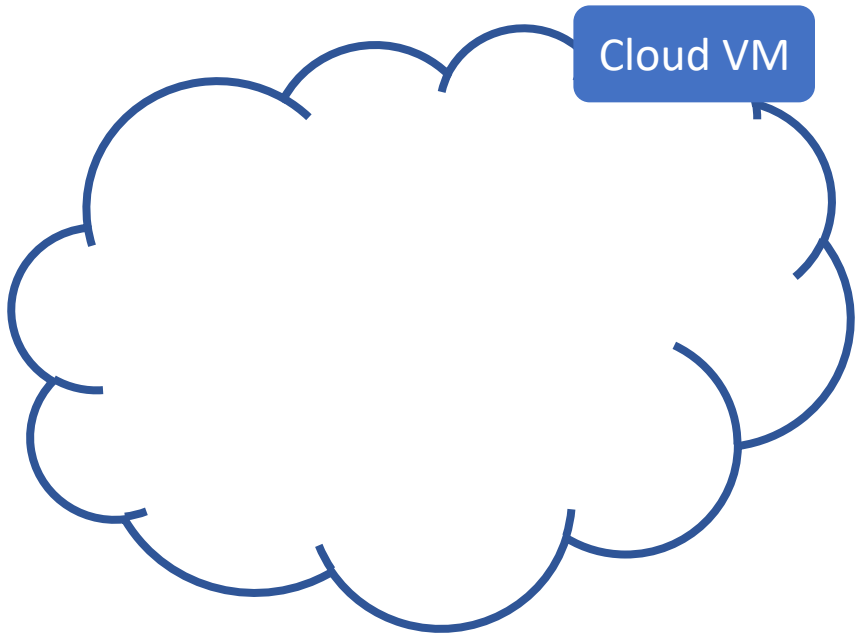

Setup the OS and kernel

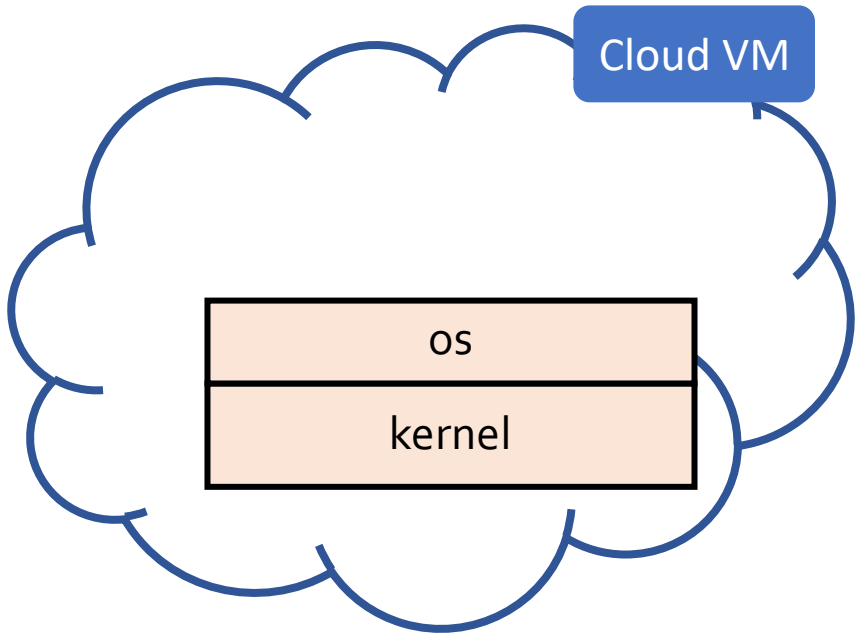

Setup a custom runtime

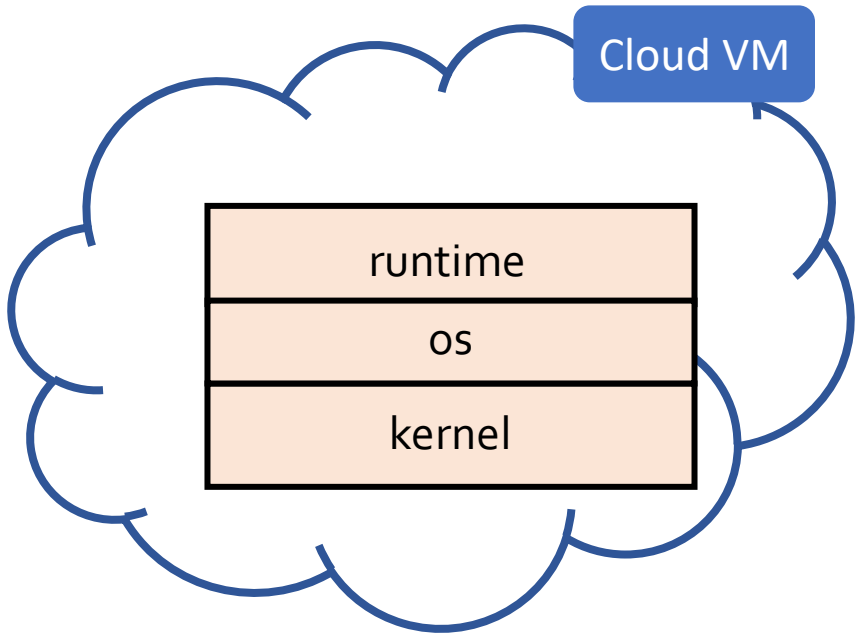

Finally, we write our application

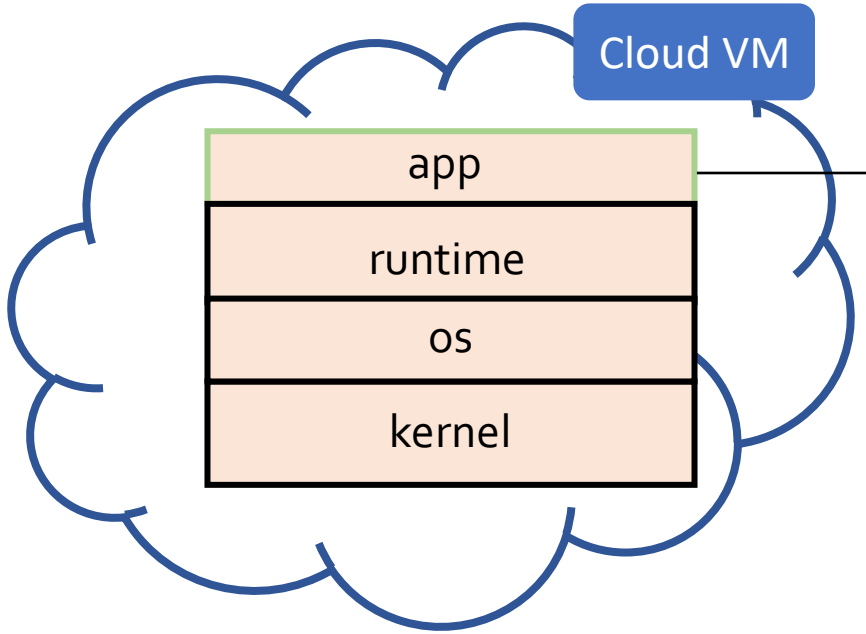

1. Download the video locally
2. Clip the video
3. Compress
4. Convert to gif

But if anything breaks

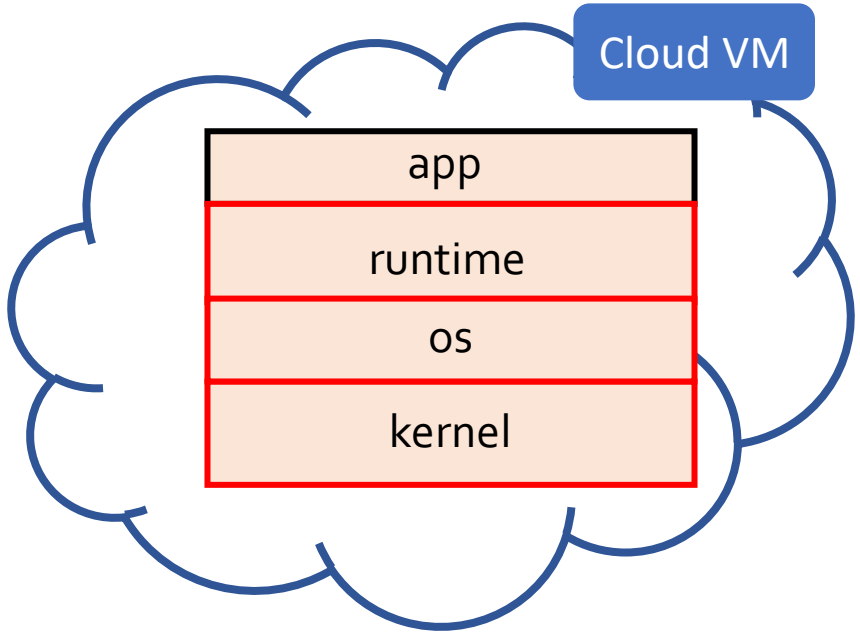

The application stops working

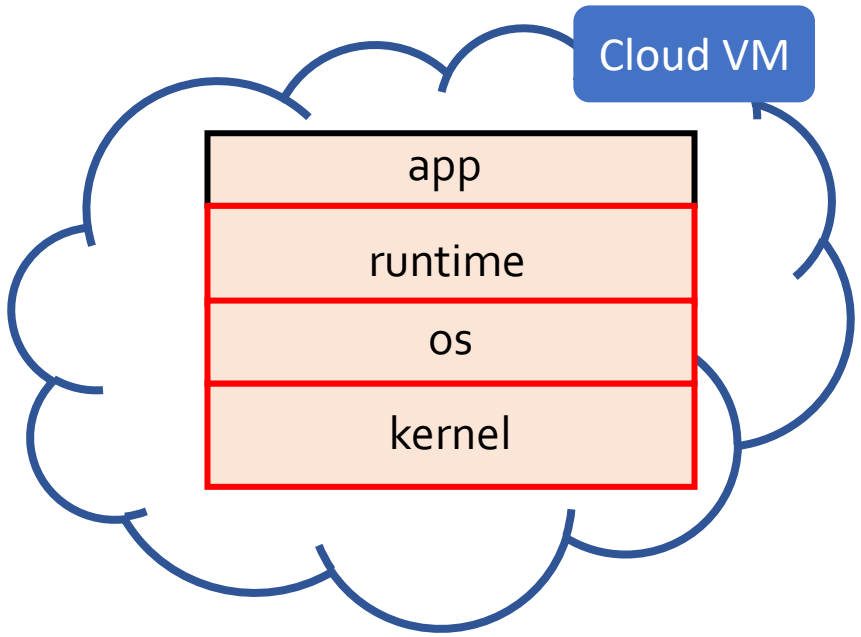

Managing this has high operational overhead

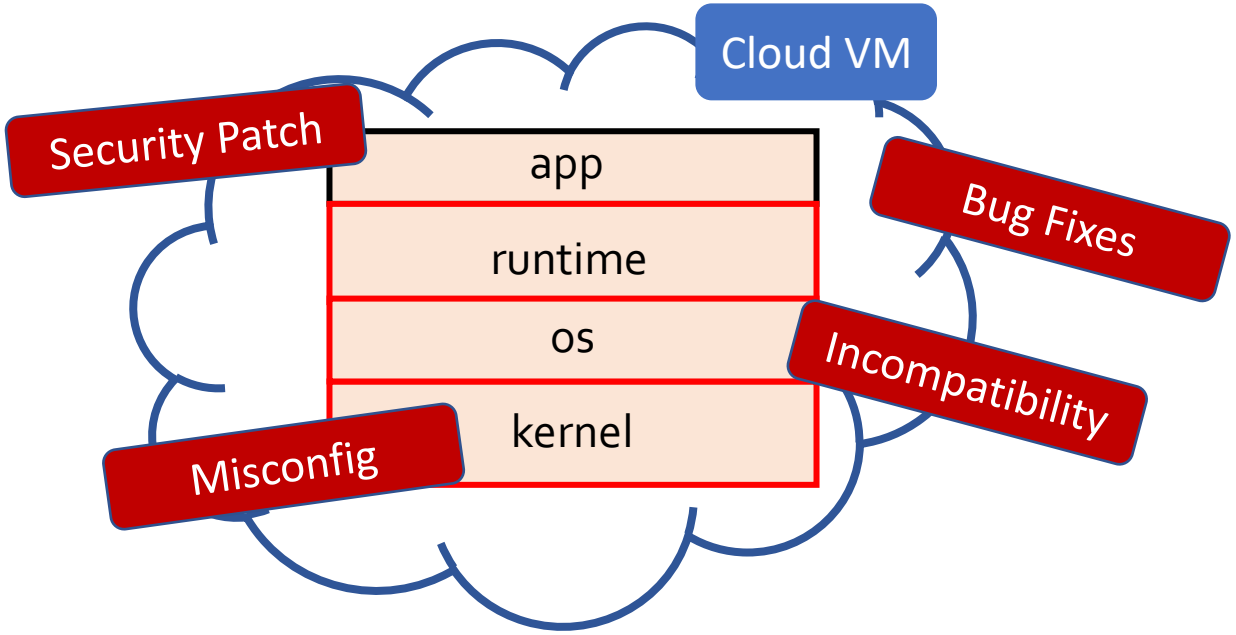

Managing this has high operational overhead

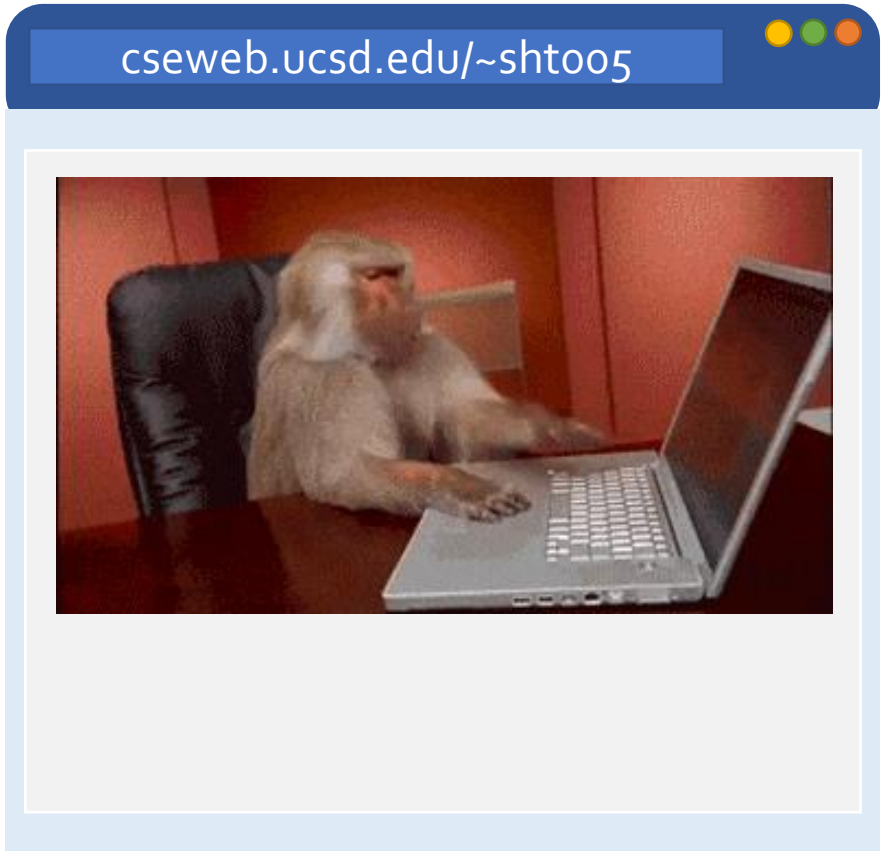

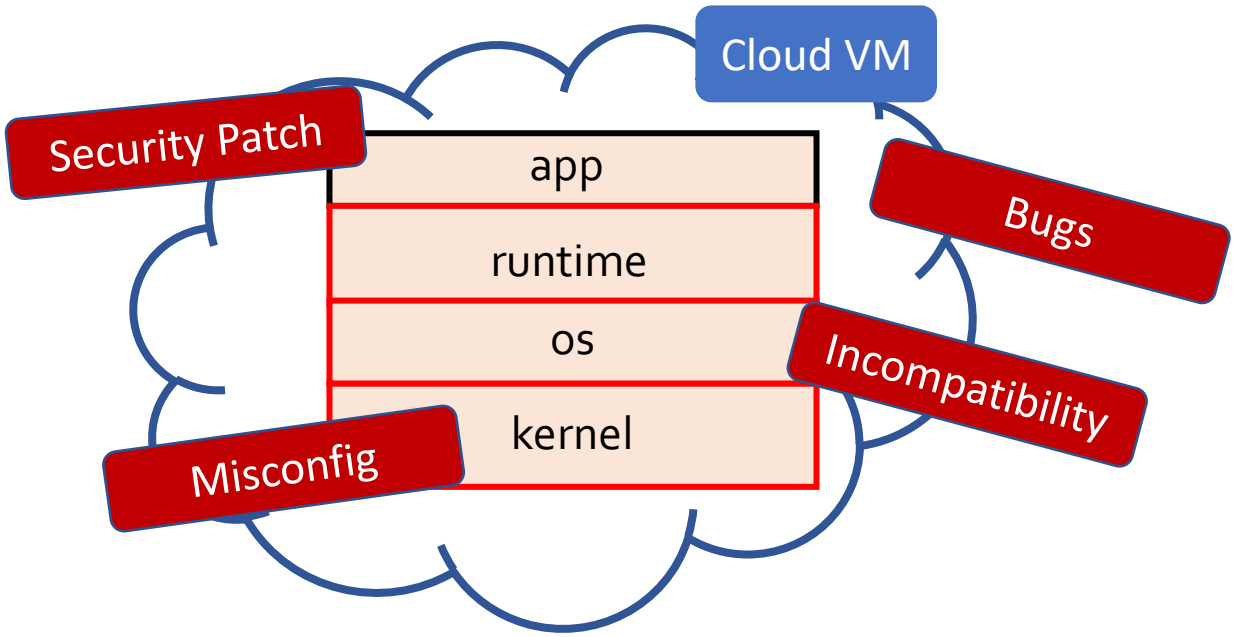

Is there a better way?

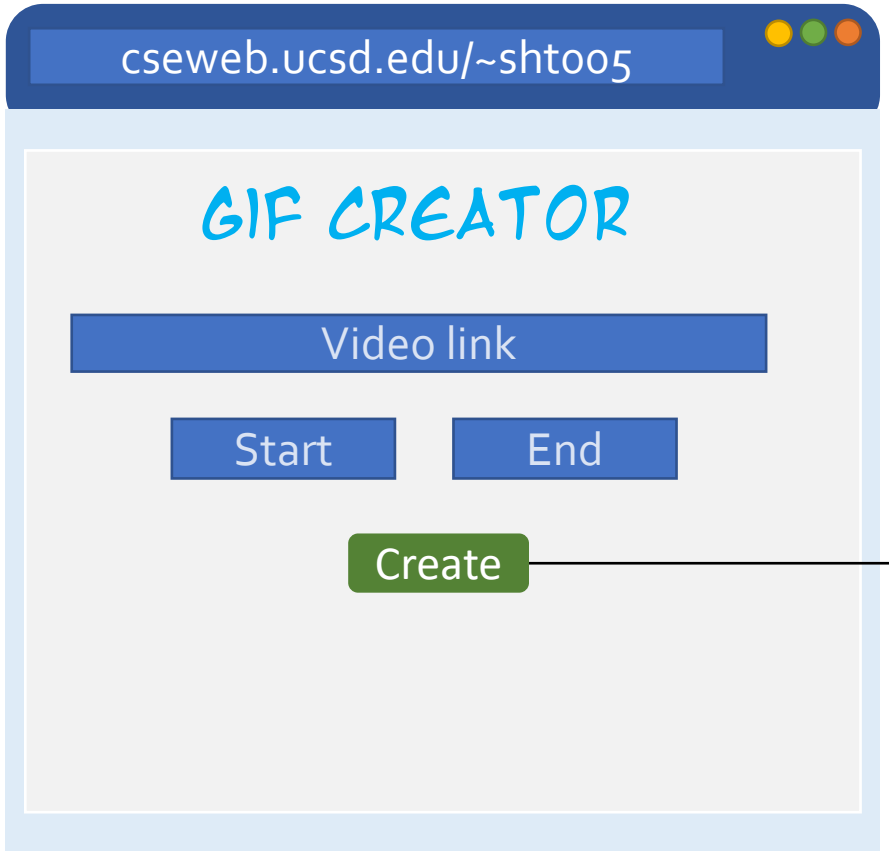

1. Download the video locally
2. Clip the video
3. Compress
4. Convert to gif

Deploy a serverless function

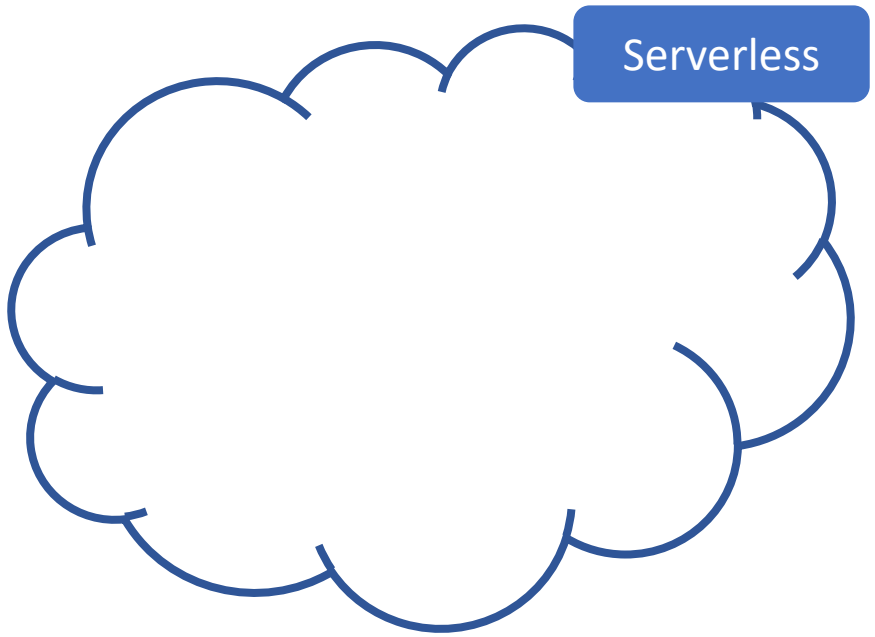

Write the application, specify dependencies

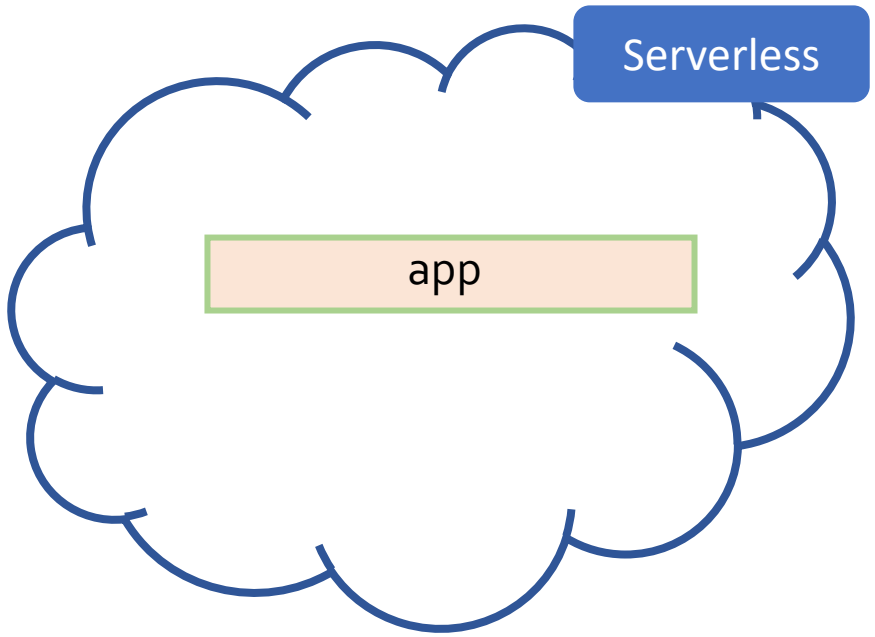

Only need to worry about application bugs

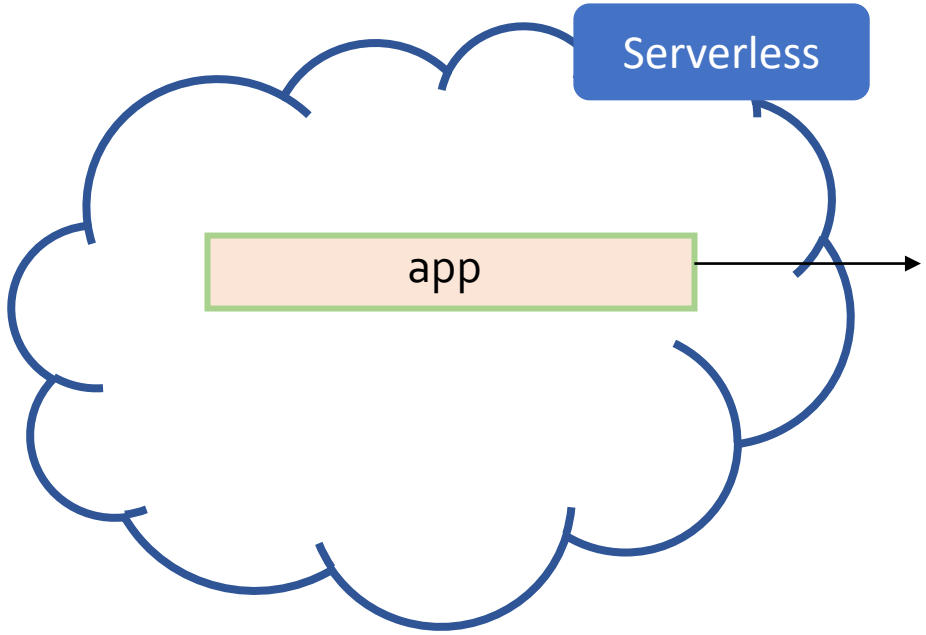

1. Download the video locally
2. Clip the video
3. Compress
4. Convert to gif

Serverless functions run on-demand

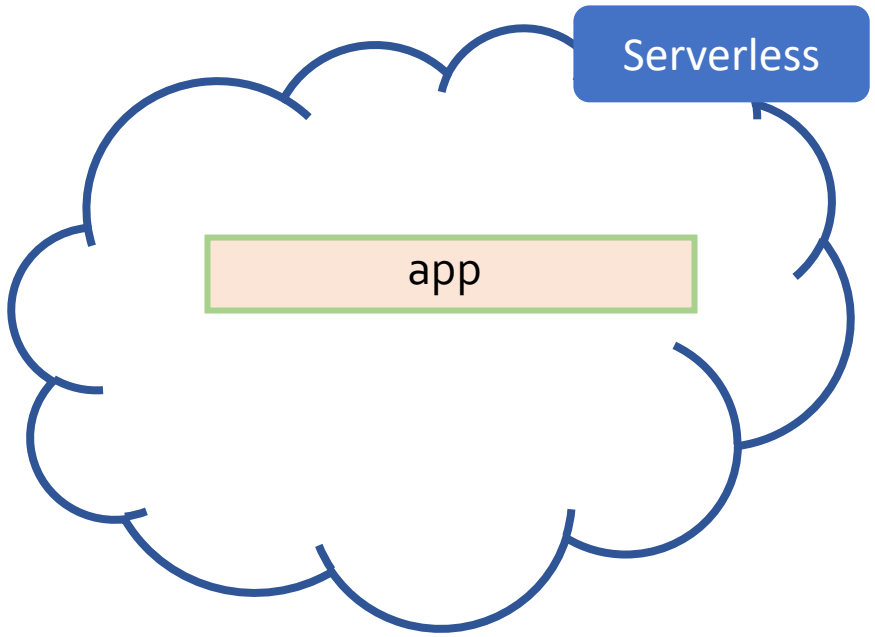

In the Cloud VM the instance is always on

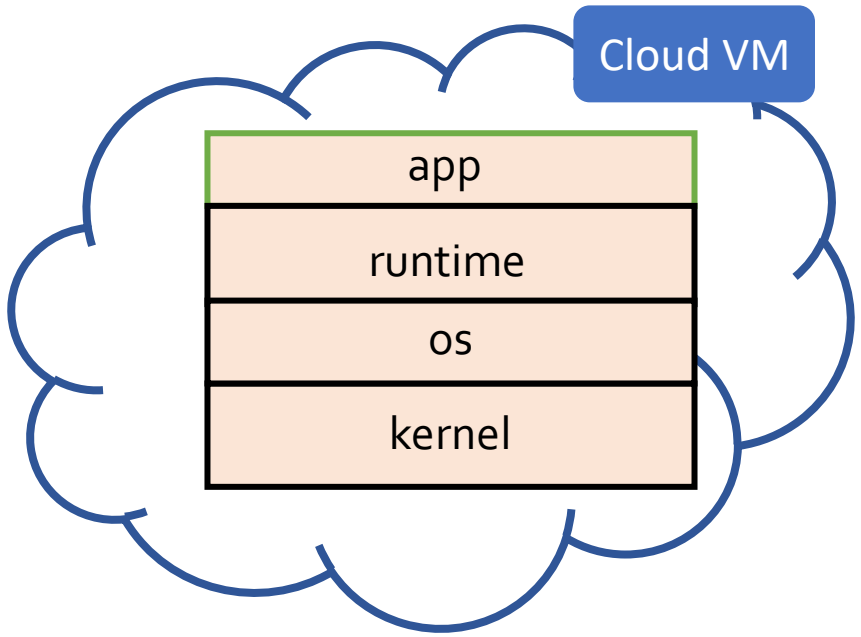

Serverless: New containers are instantiated from scratch every time you click Create

Serverless: New containers are instantiated from scratch every time you click Create

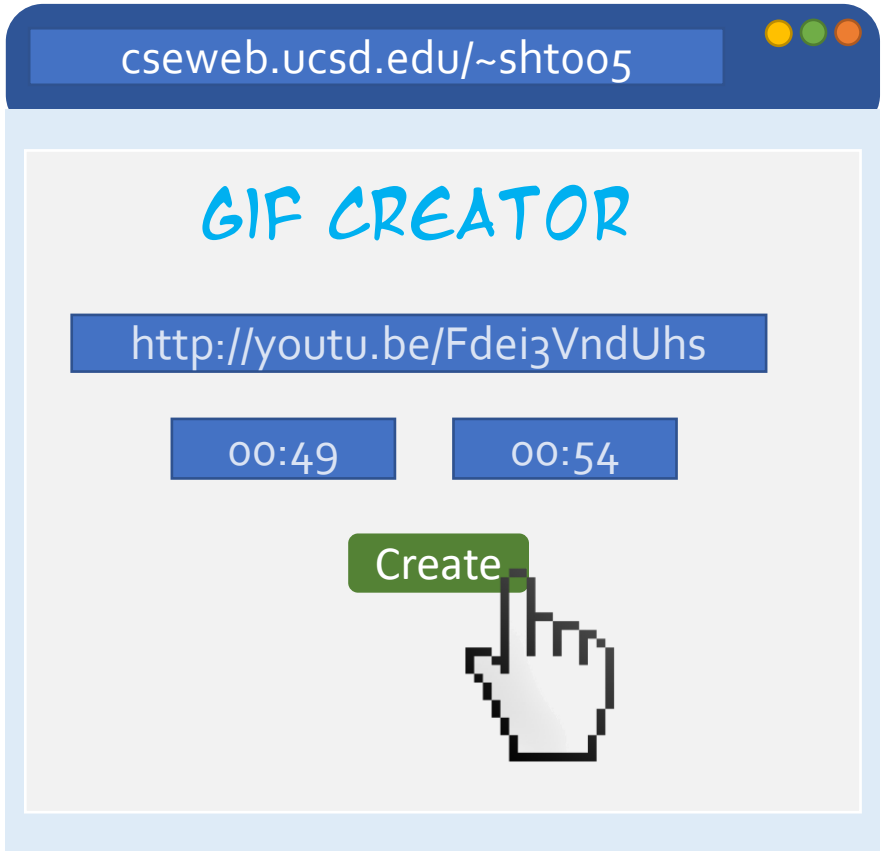

Serverless: New containers are instantiated from scratch every time you click Create

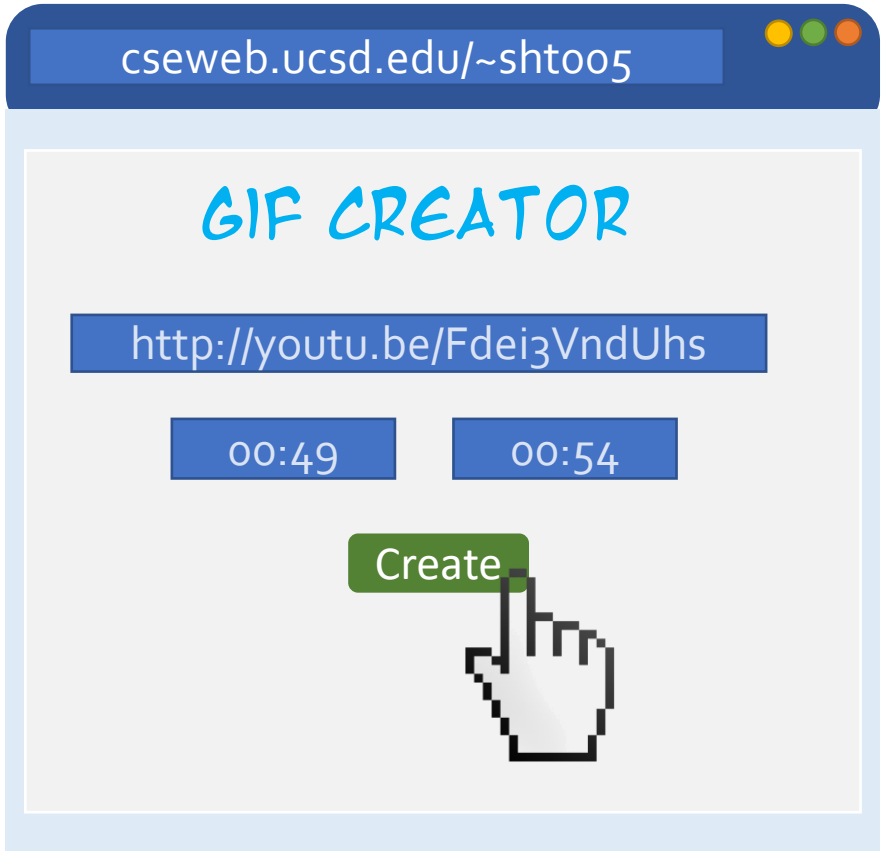

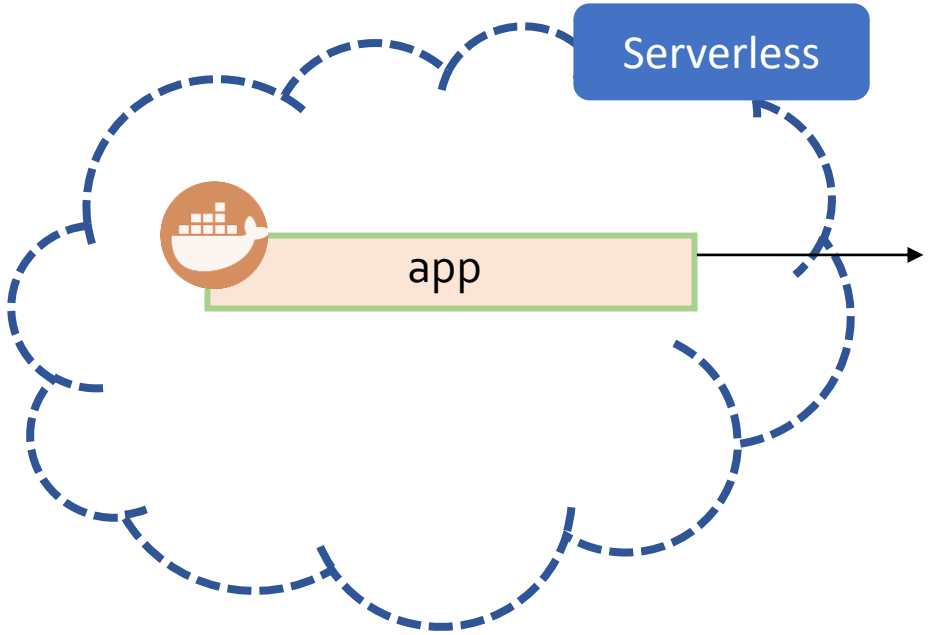

1. **Start container**
2. Download the video locally
3. Clip the video
4. Compress
5. Convert to gif
6. Send result
7. **Stop container**

Serverless: New containers are instantiated from scratch every time you click Create

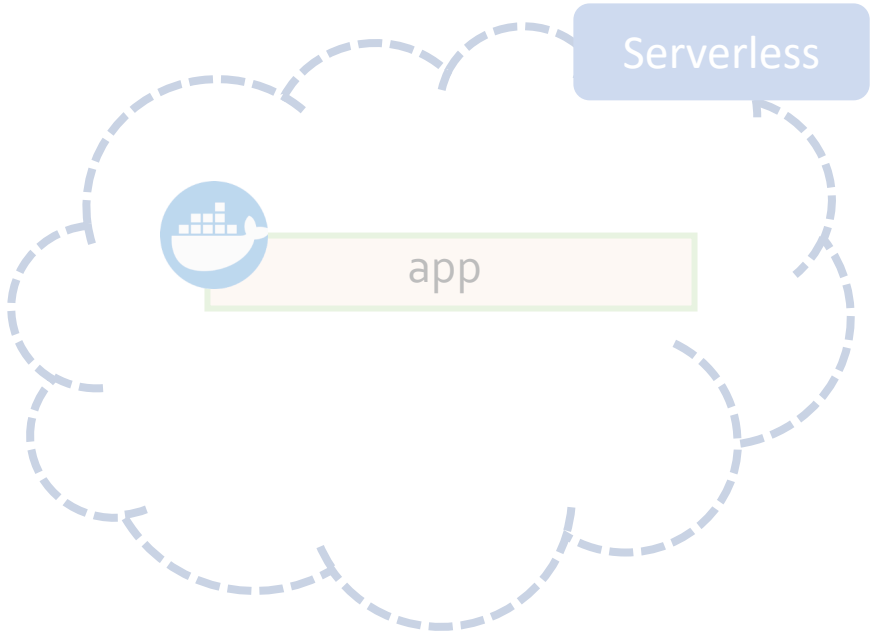

Which one is better?

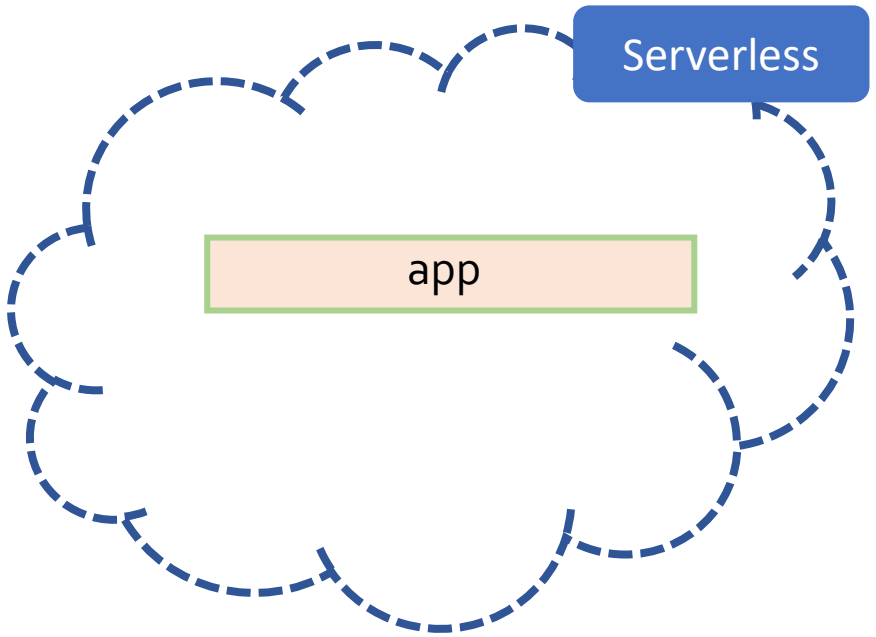

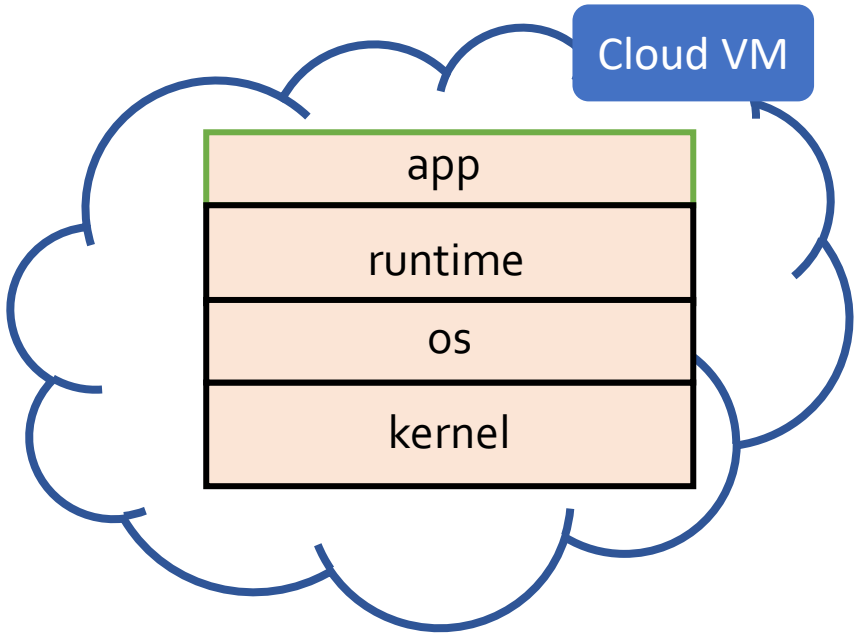

Depends on the use case

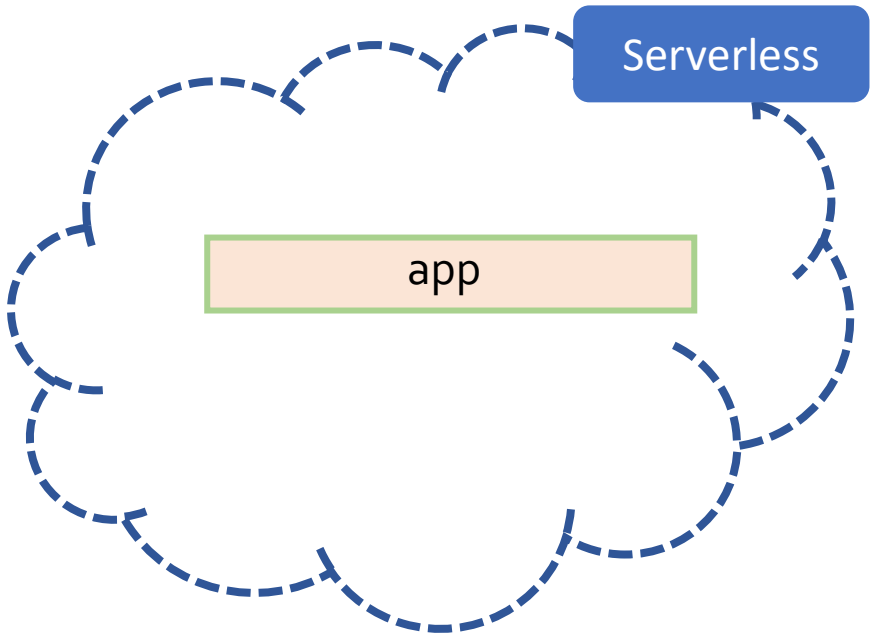

On-demand
Cheaper in many cases
Focus on application
Subject to fluctuations
Questionable security

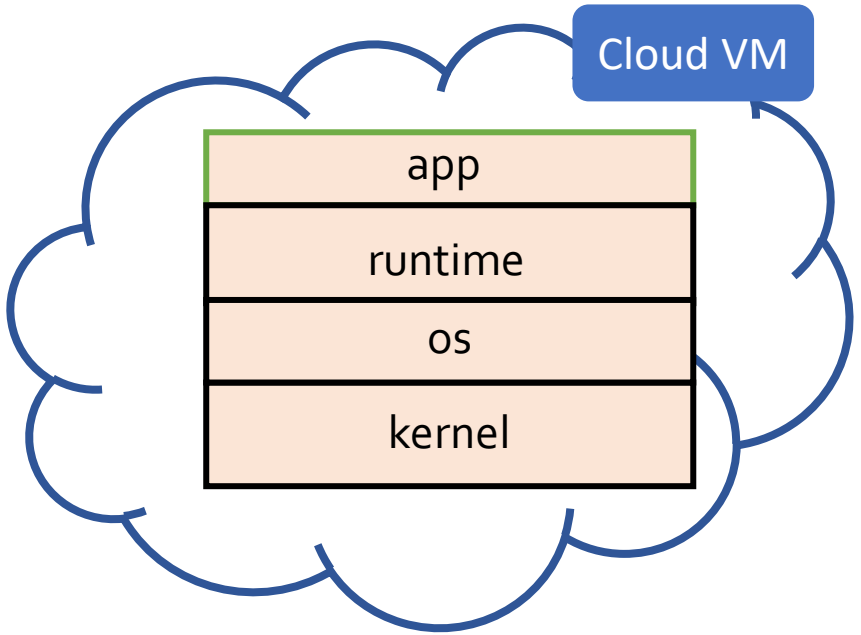

Depends on the use case

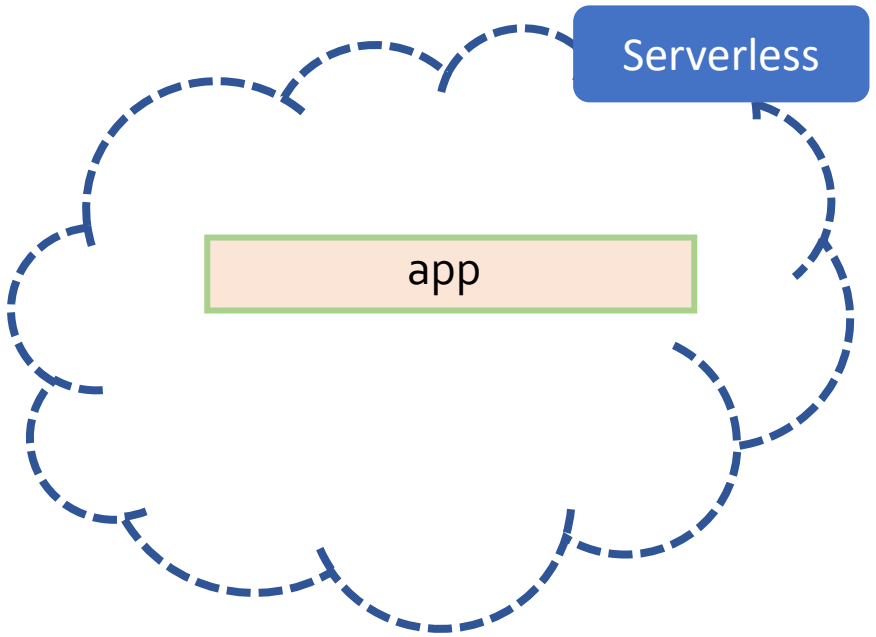

On-demand
Cheaper in many cases
Focus on application
Subject to fluctuations
Questionable security

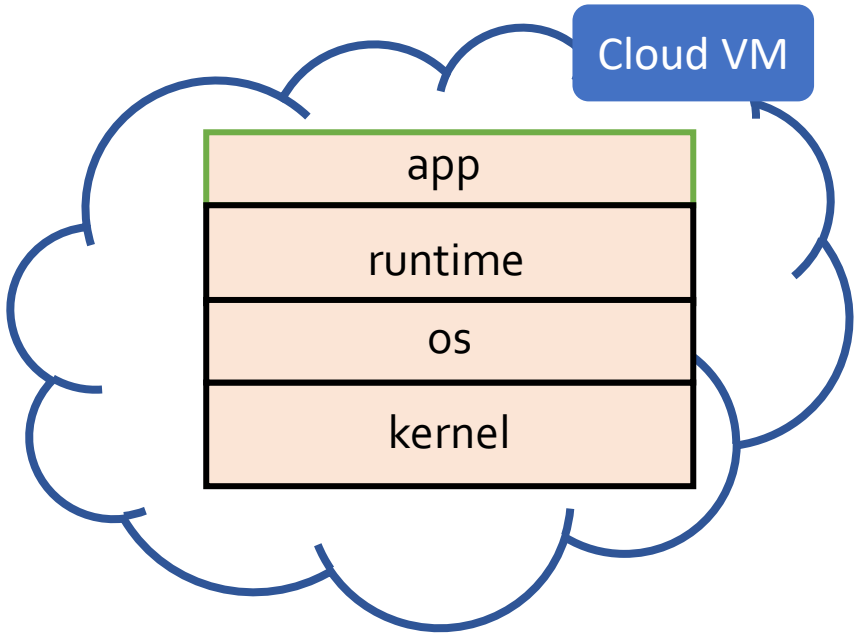

Persistent
More Control
Fast and predictable
Difficult to manage

Depends on the use case

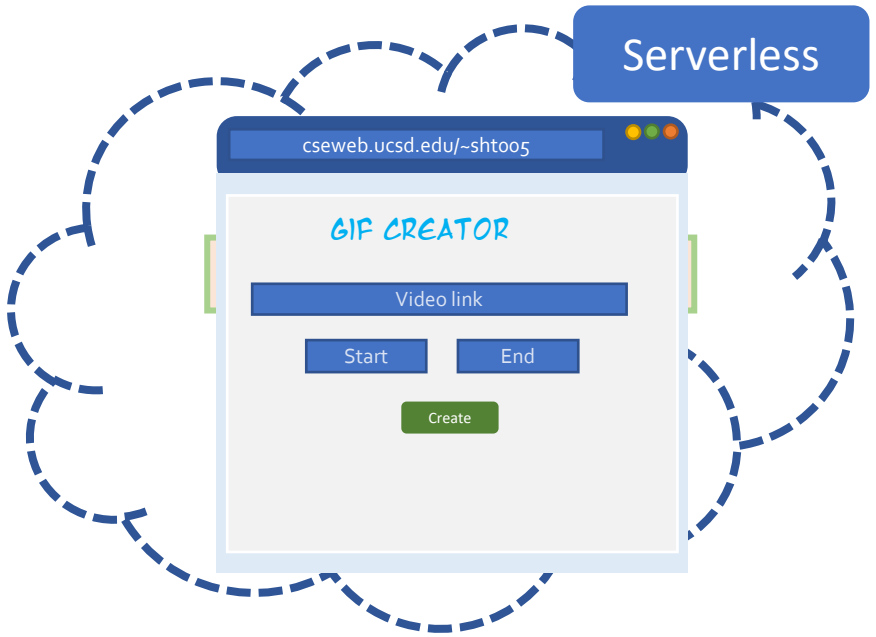

On-demand
Focus on application
Subject to fluctuations
Questionable security

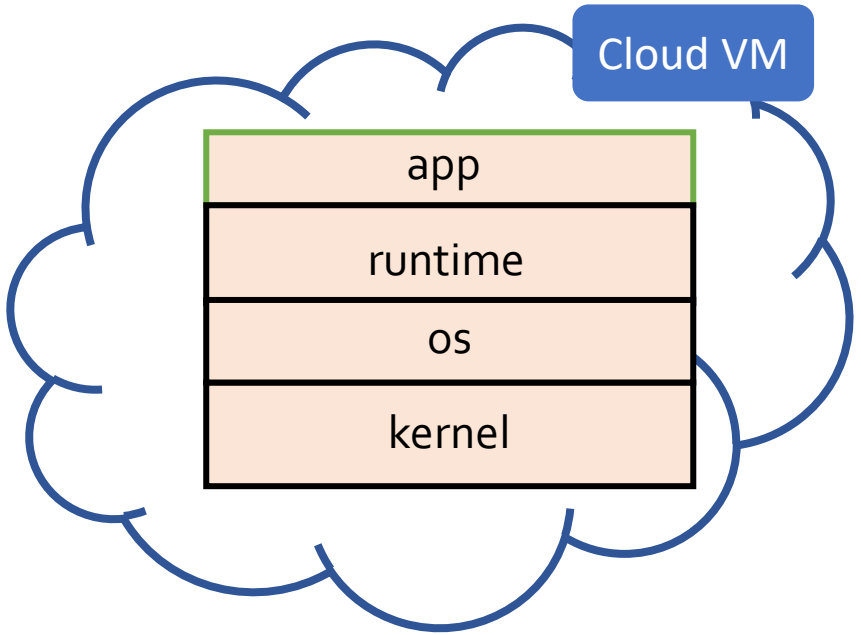

Persistent
More Control
Fast and predictable
Difficult to manage

Review

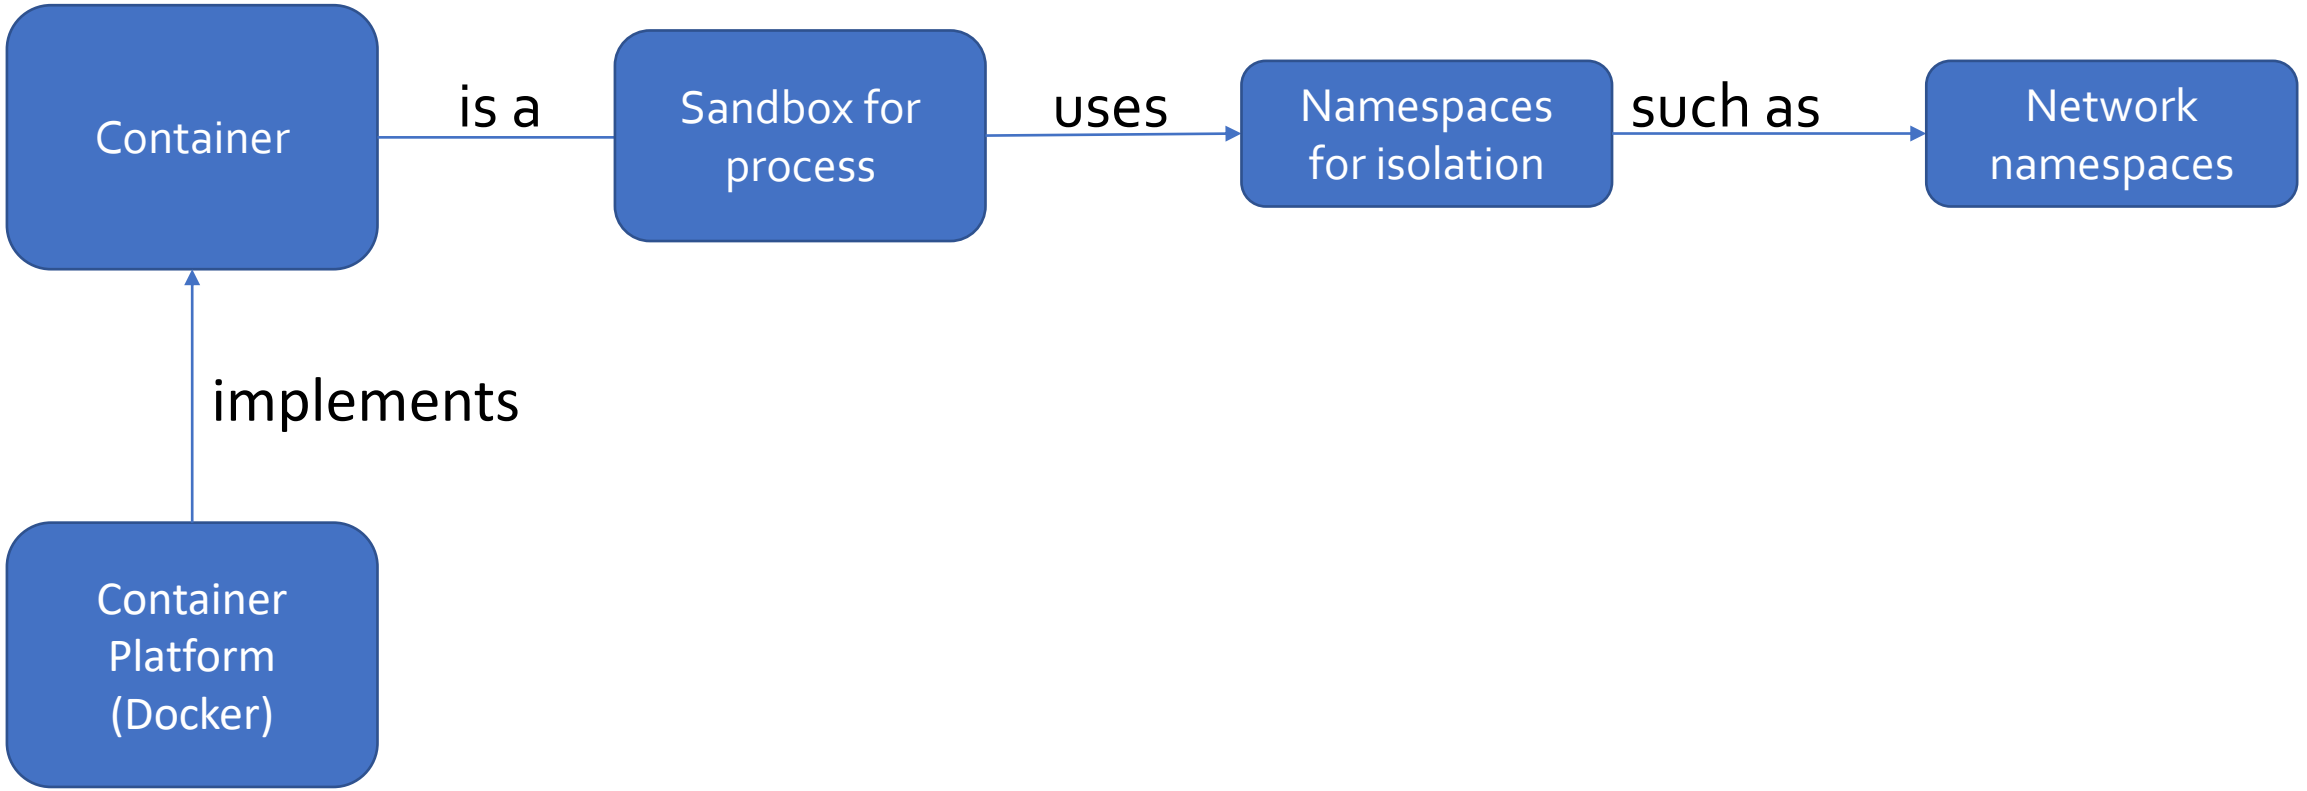
Container
is a
Sandbox for process
uses
Namespaces for isolation
such as
Network namespaces
implements
Container Platform (Docker)

Review

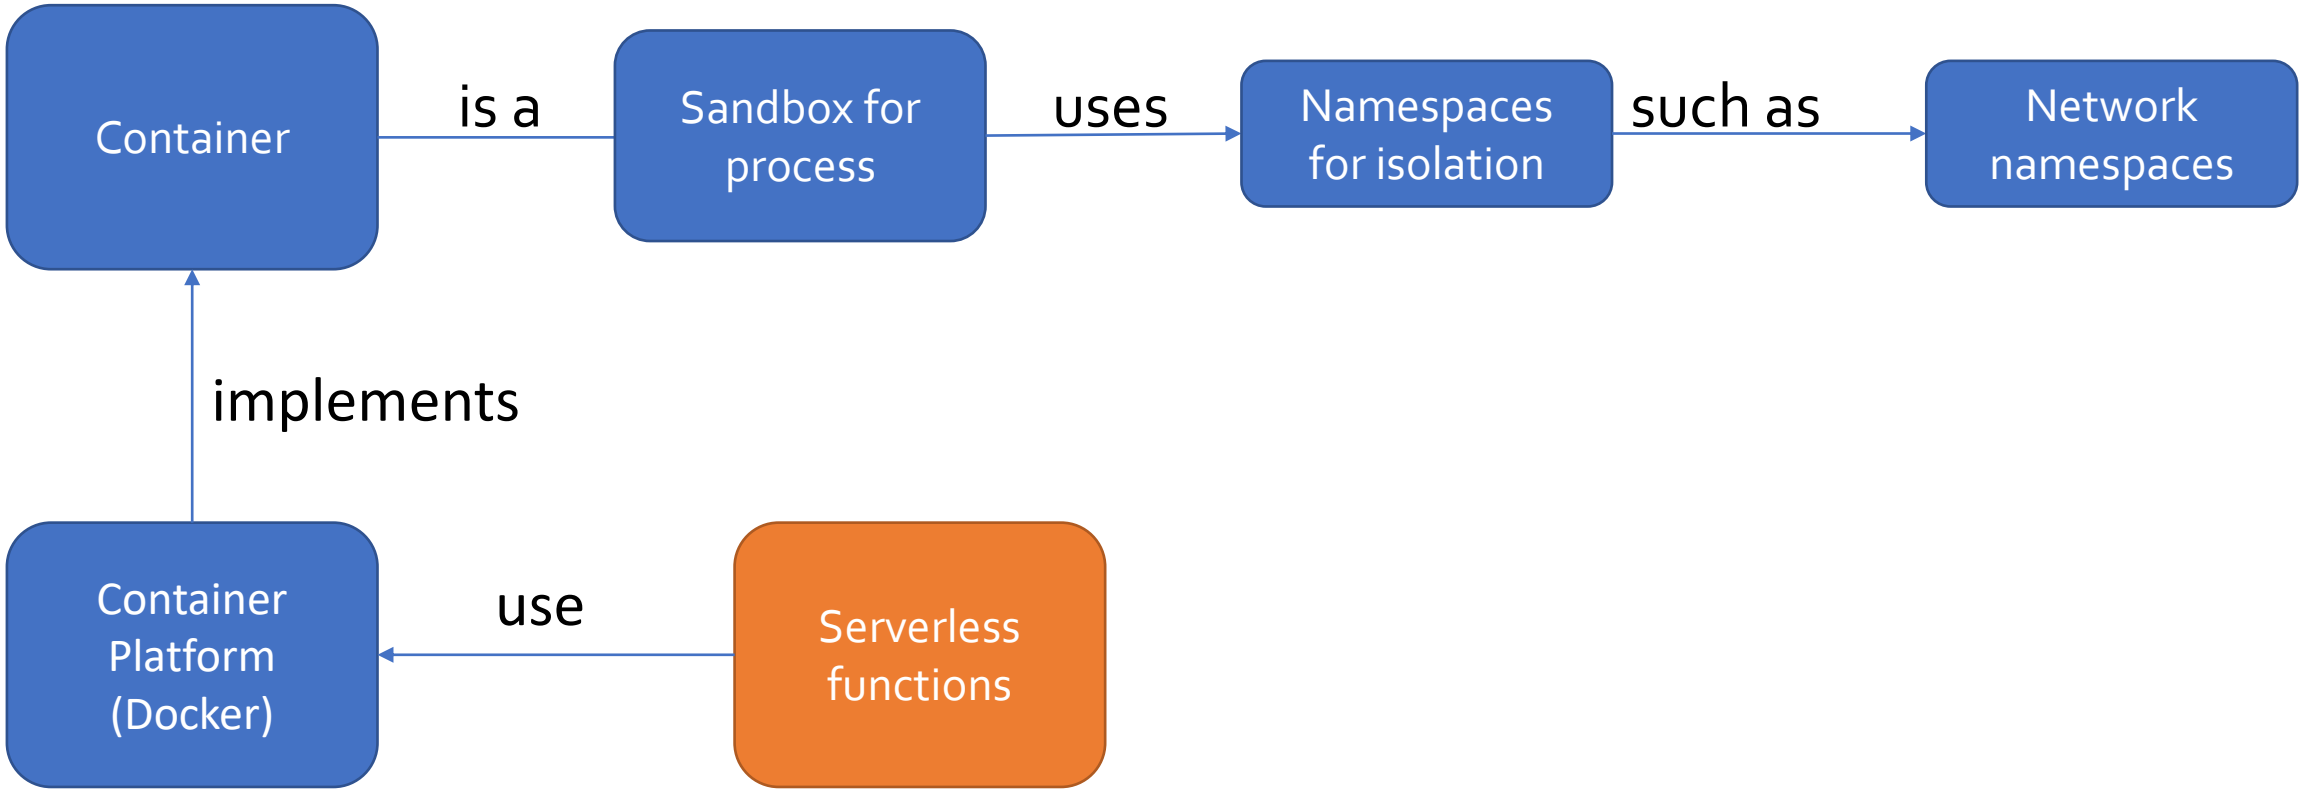
Container
is a
Sandbox for process
uses
Namespaces for isolation
such as
Network namespaces
implements
Container Platform (Docker)
use
Serverless functions

Virtualization Review

Layers of virtualization

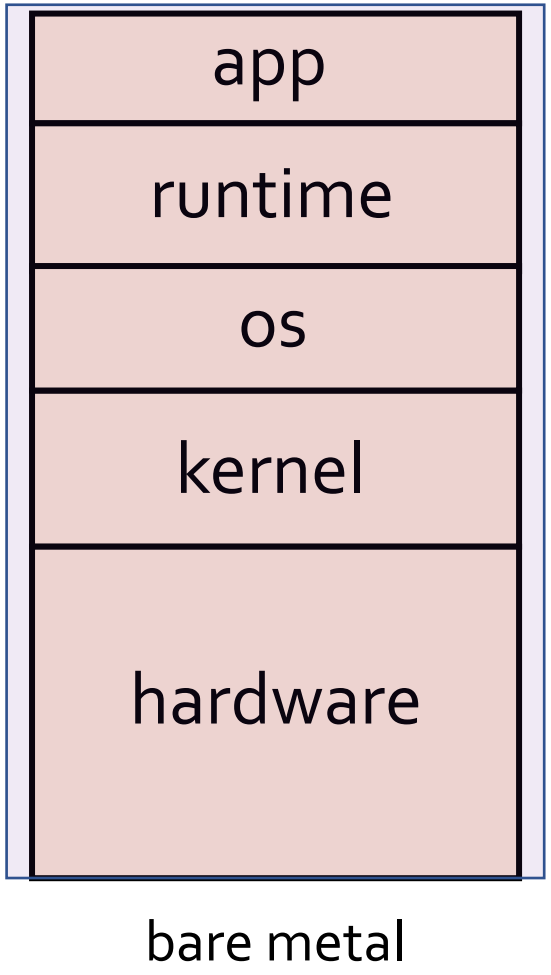

bare metal

Layers of virtualization

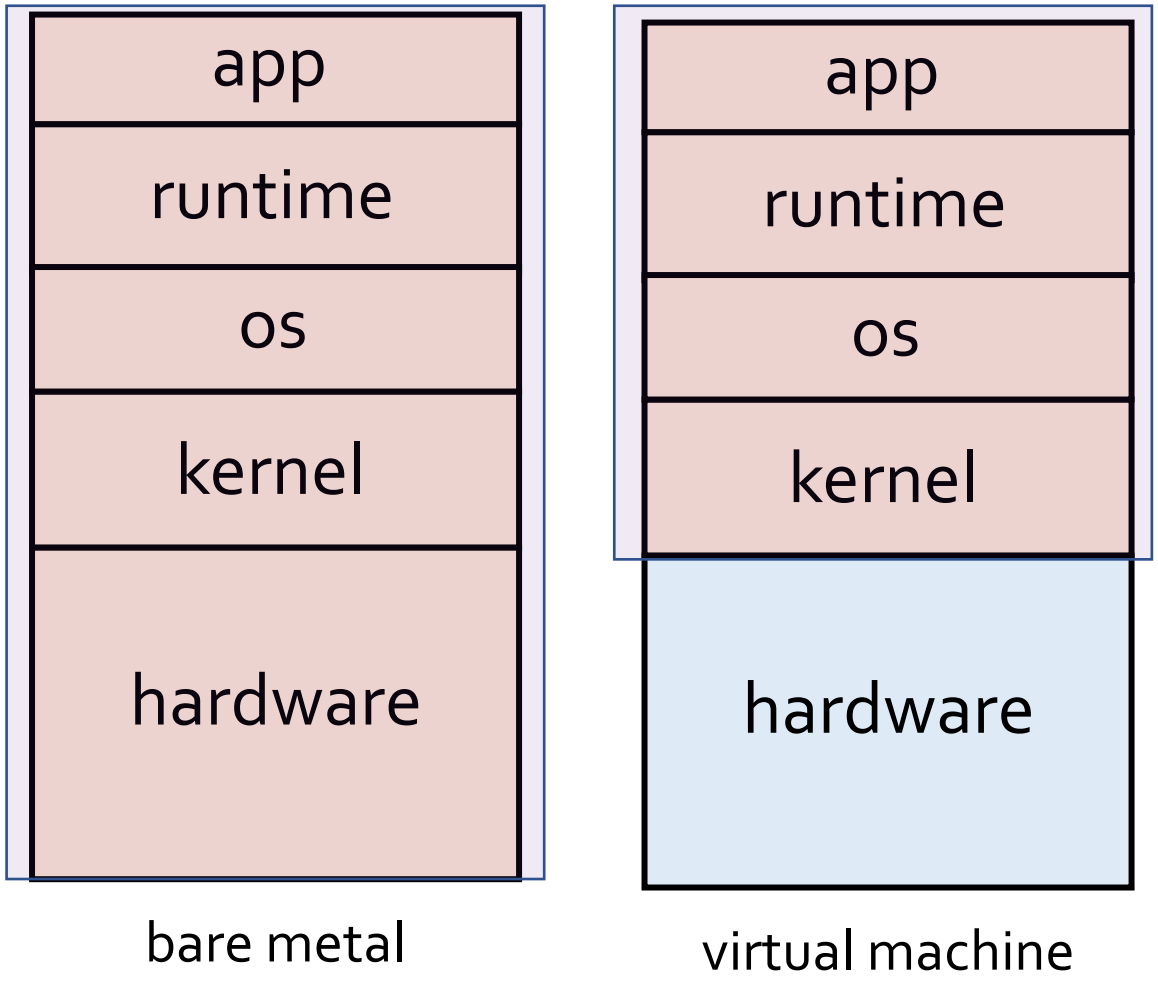

Layers of virtualization

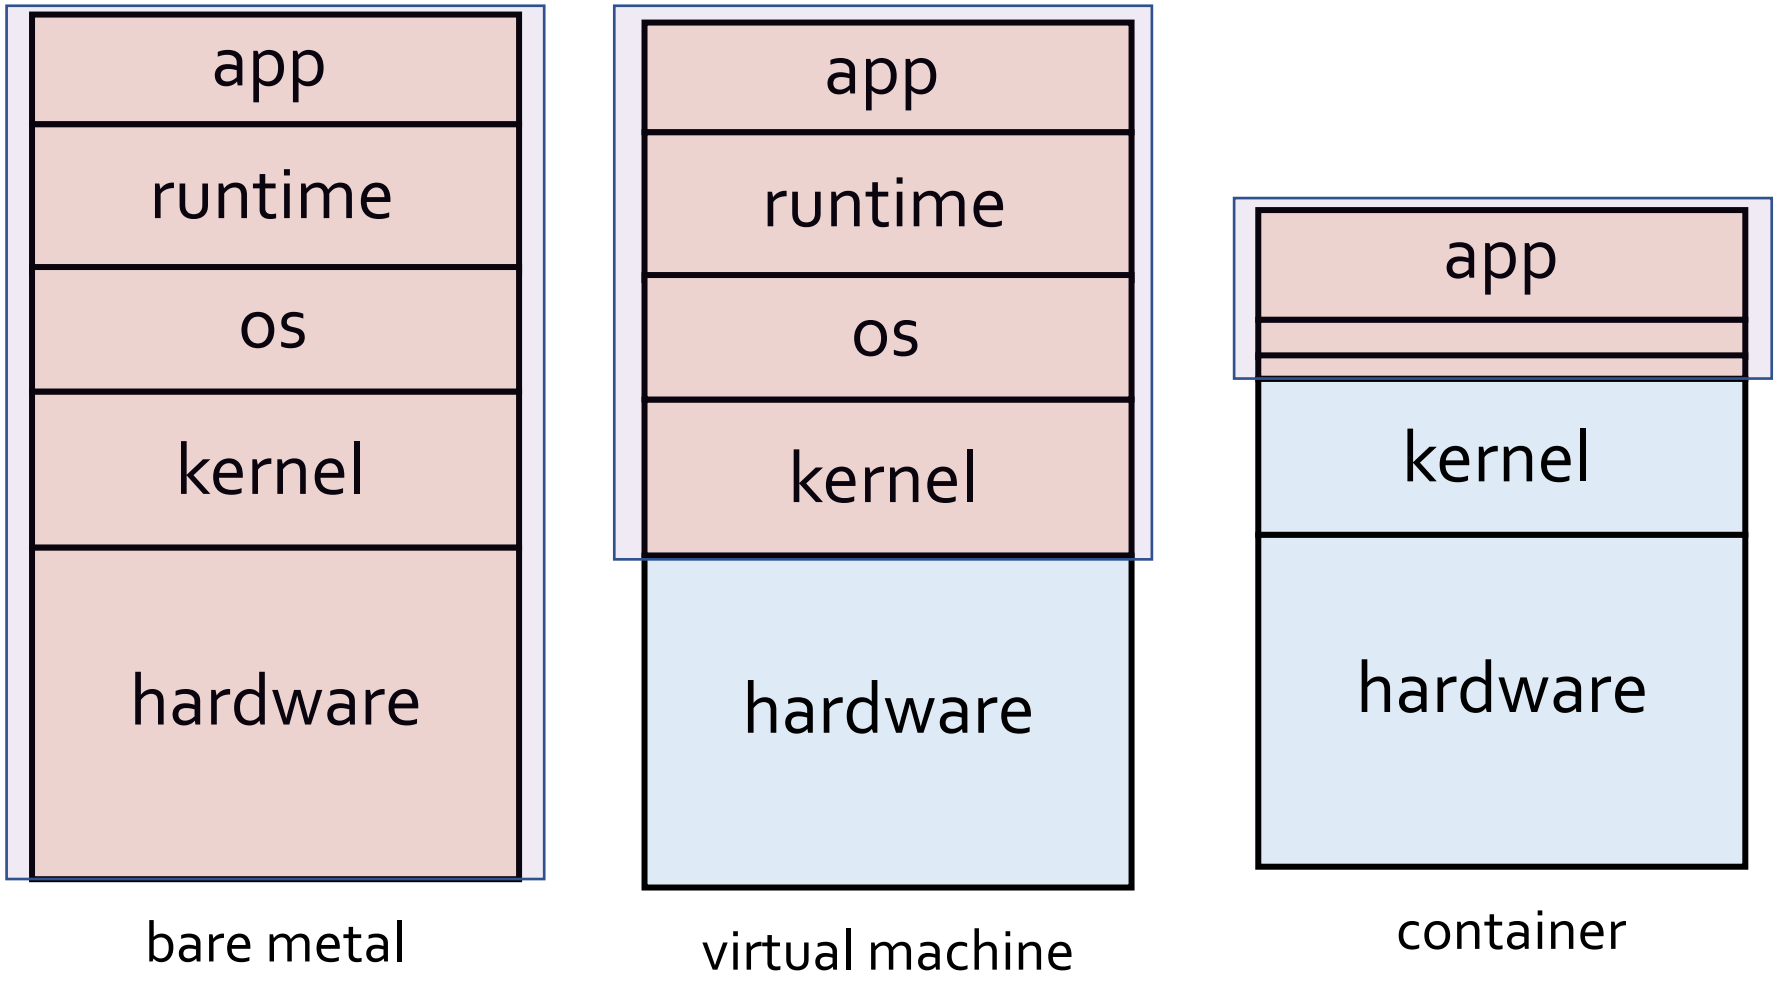

Layers of virtualization

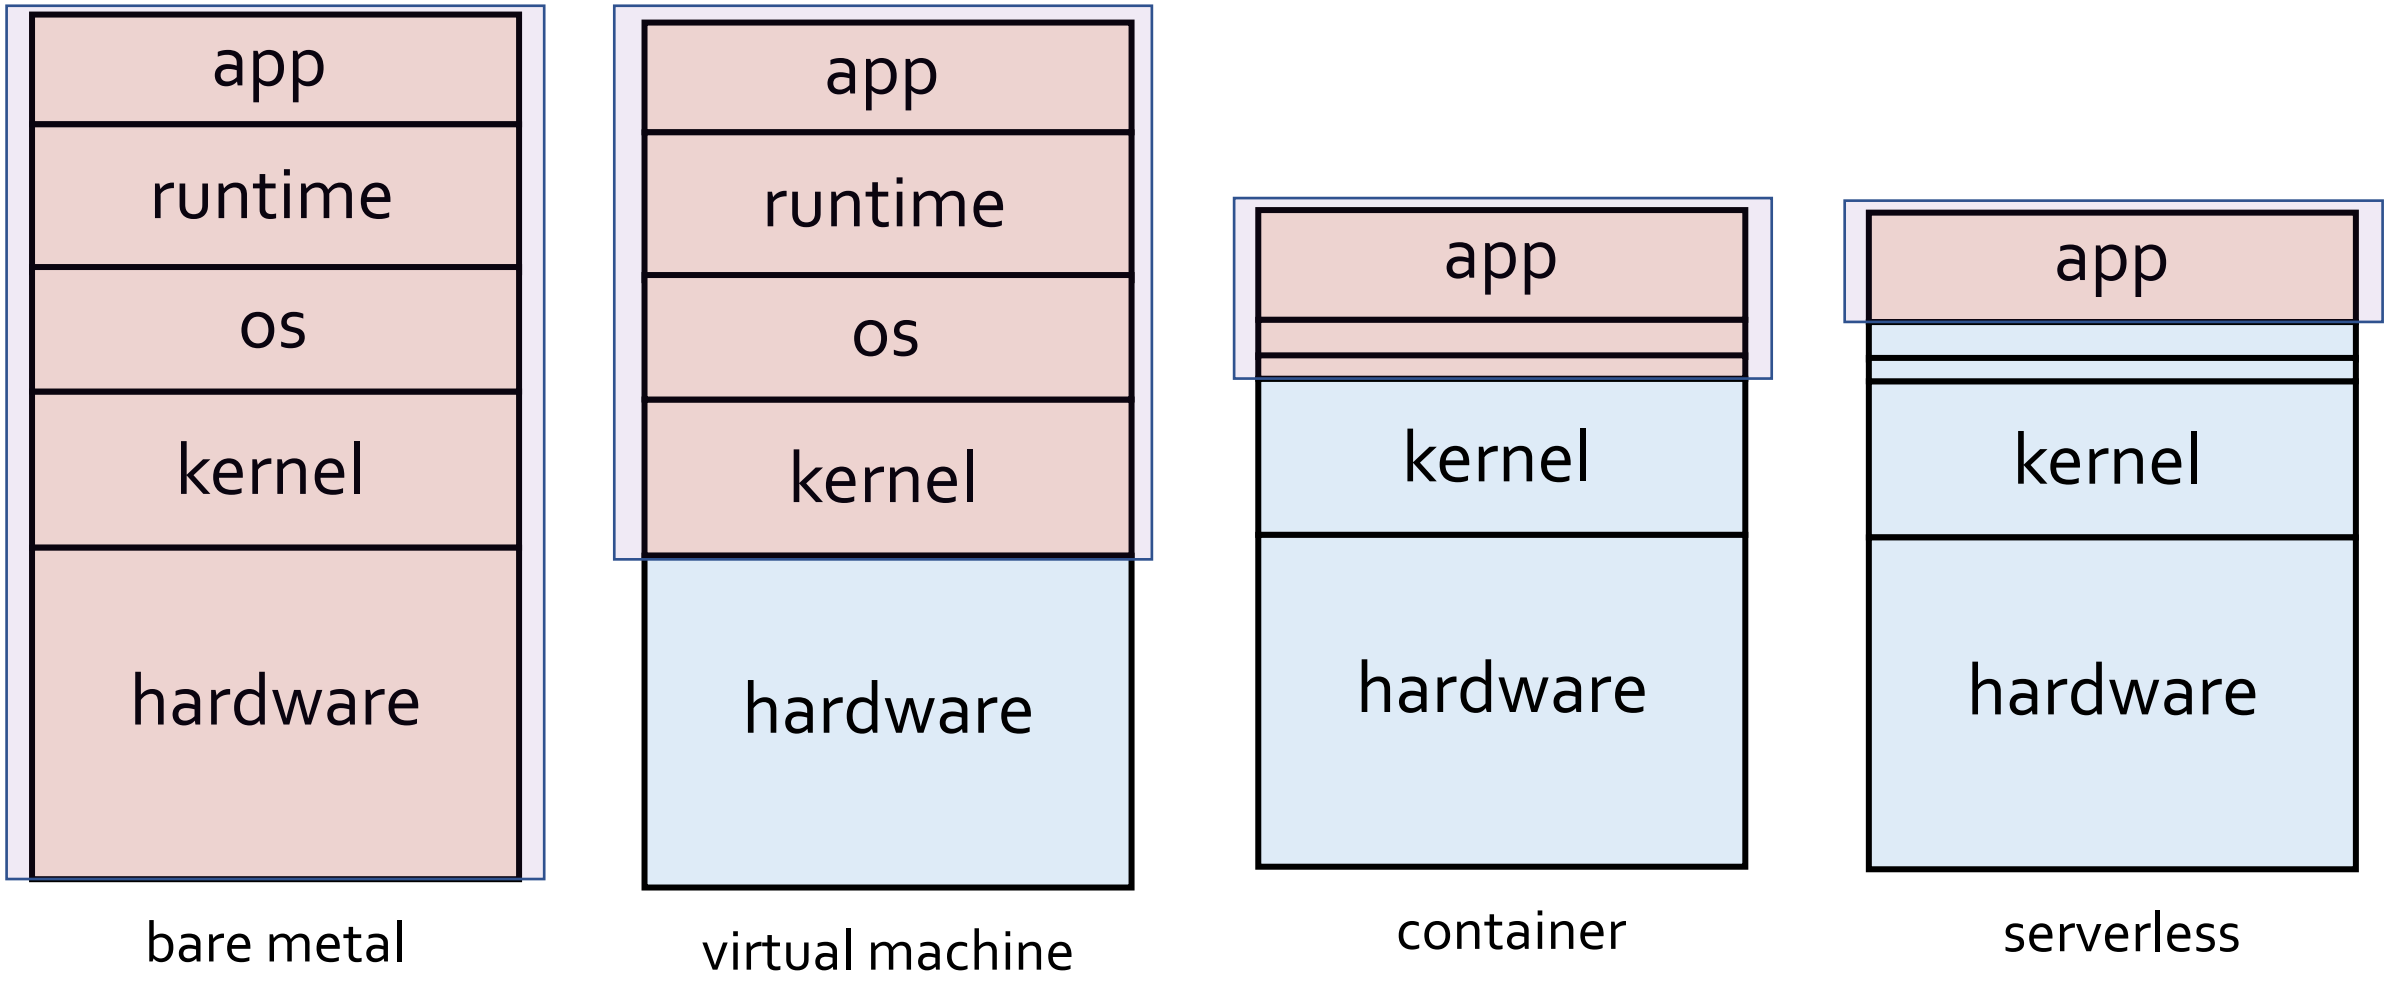

Container and Serverless Research

Topic 1: Reduce startup variance

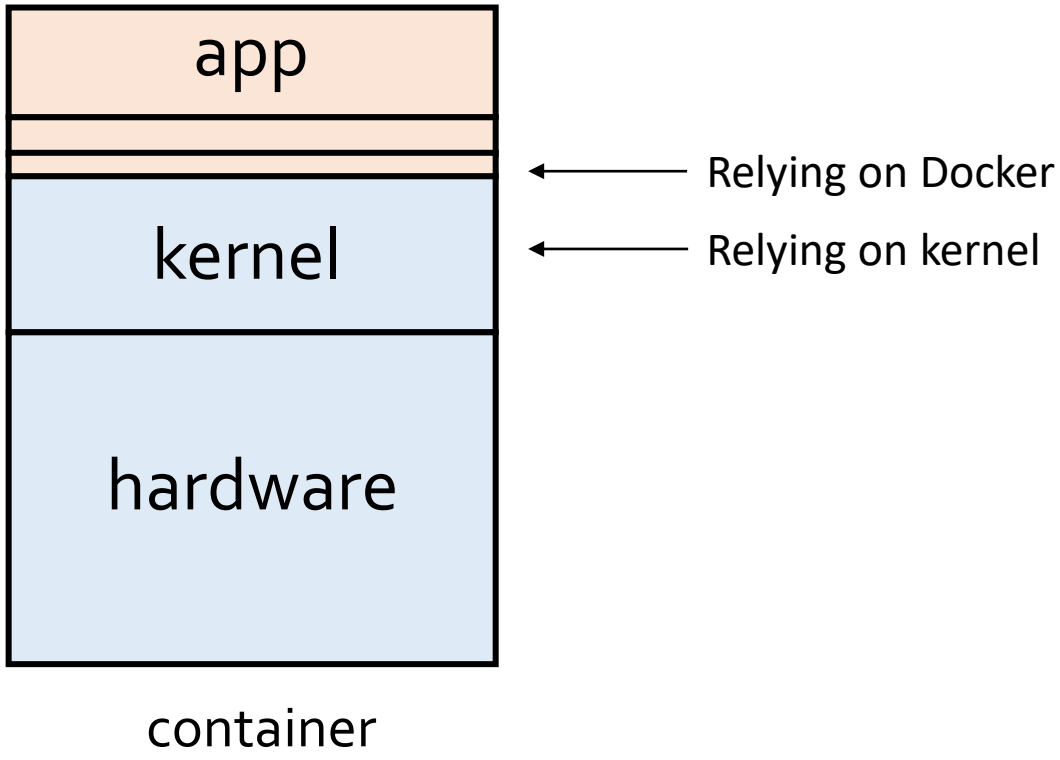

Topic 1: Reduce startup variance

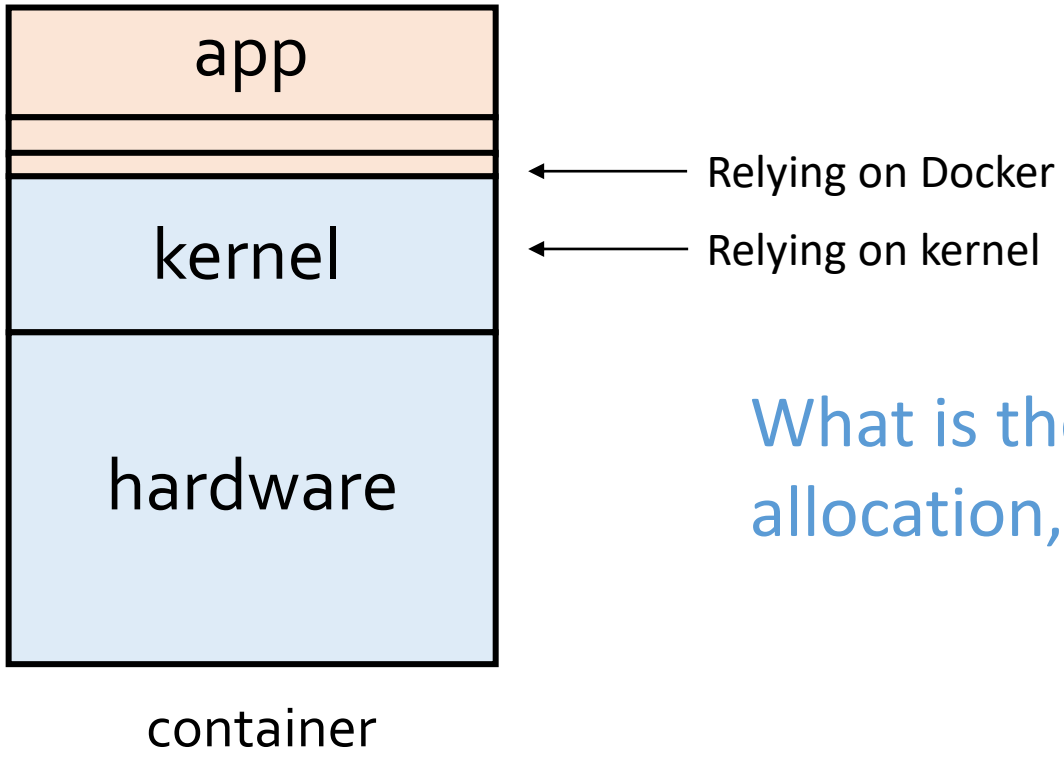

What is the optimal caching, locking, allocation, sharing mechanism?

Topic 1: Reduce serverless startup variance

COLD STARTS

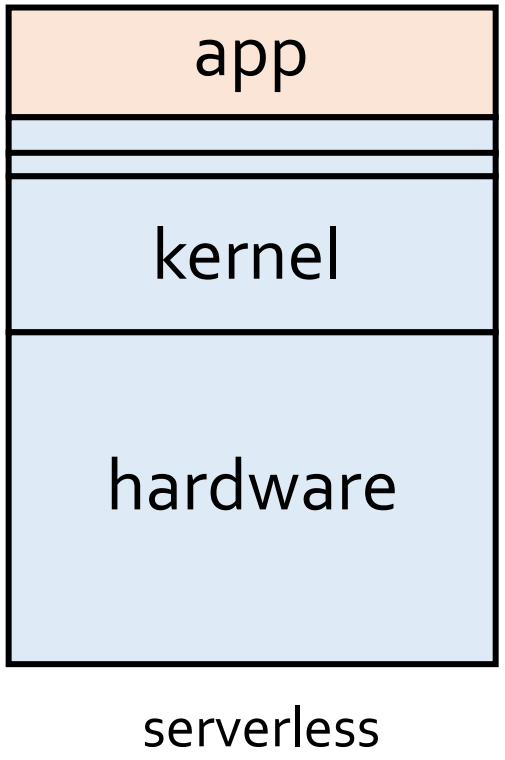

serverless

Topic 1: Reduce serverless startup variance

COLD STARTS

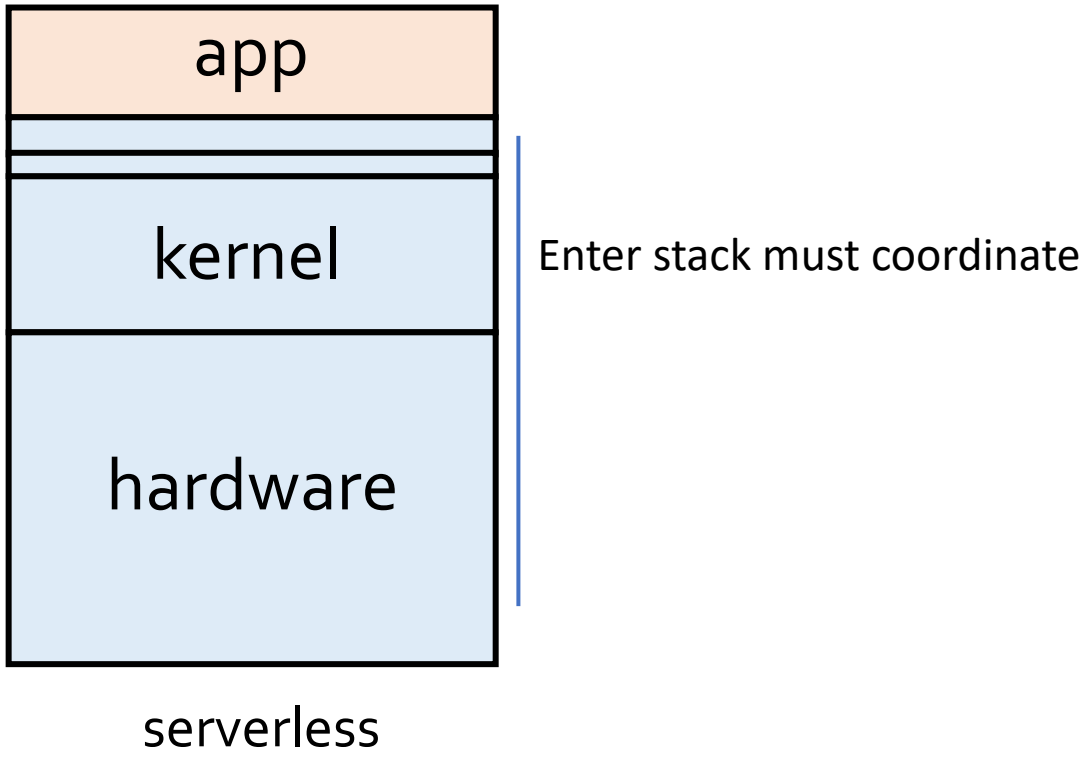

Topic 1: Reduce serverless startup variance

COLD STARTS

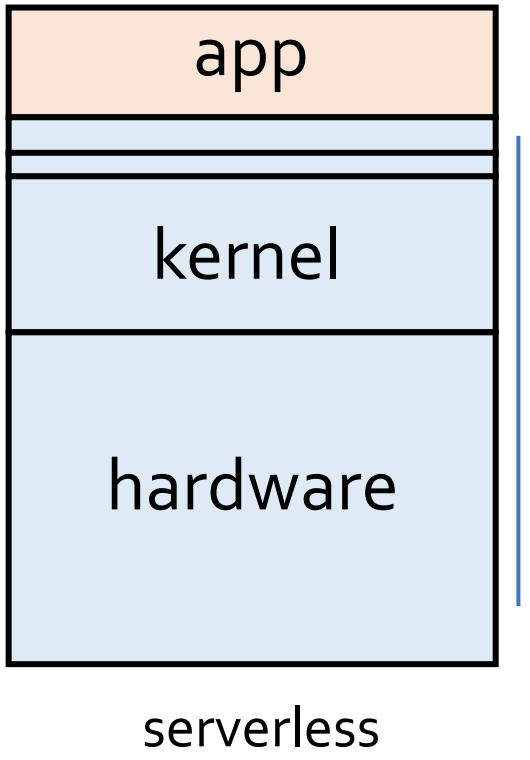

Enter stack must coordinate

Extremely important in serverless context for better response times

Topic 1: Reduce serverless startup variance

COLD STARTS

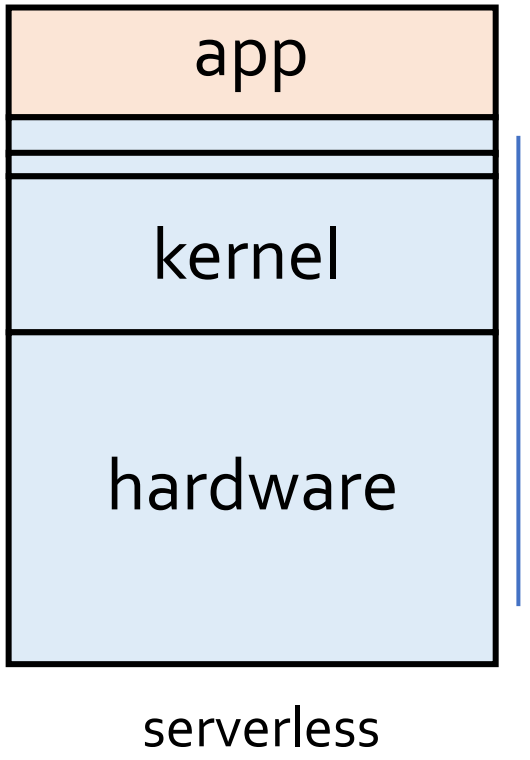

Enter stack must coordinate

Extremely important in serverless context for better response times

A 100-millisecond delay in website load time can hurt conversion rates by 7 percent - Akamai, 2017

Topic 1: Reduce serverless startup variance

COLD STARTS

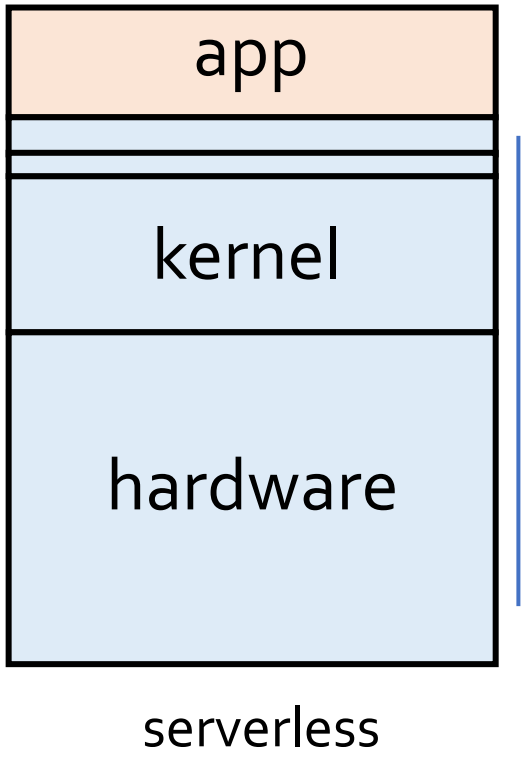

serverless

Particle: Ephemeral Endpoints for Serverless Networking

Shelby Thomas
UC San Diego
shelbyt@ucsd.edu

Lixiang Ao
UC San Diego
liao@eng.ucsd.edu

Geoffrey M. Voelker
UC San Diego
voelker@cs.ucsd.edu

George Porter
UC San Diego
gmporter@cs.ucsd.edu

Extremely important in serverless context for better response times

A 100-millisecond delay in website load time can hurt conversion rates by 7 percent - Akamai, 2017

Topic 2: Improve security and control

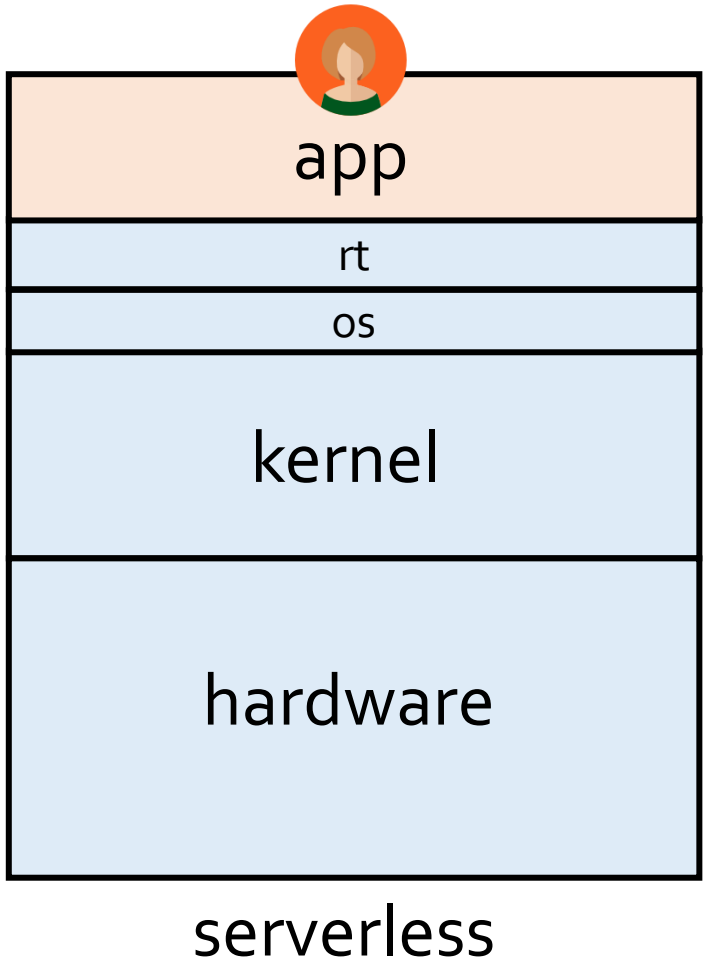

Topic 2: Improve security and control

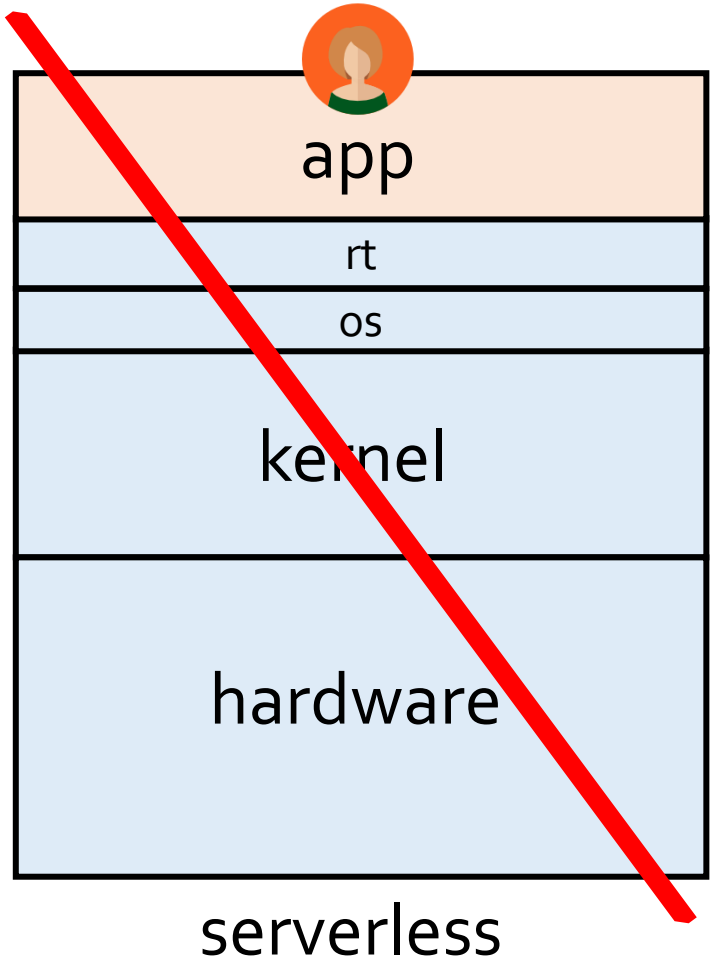

Topic 2: Improve security and control

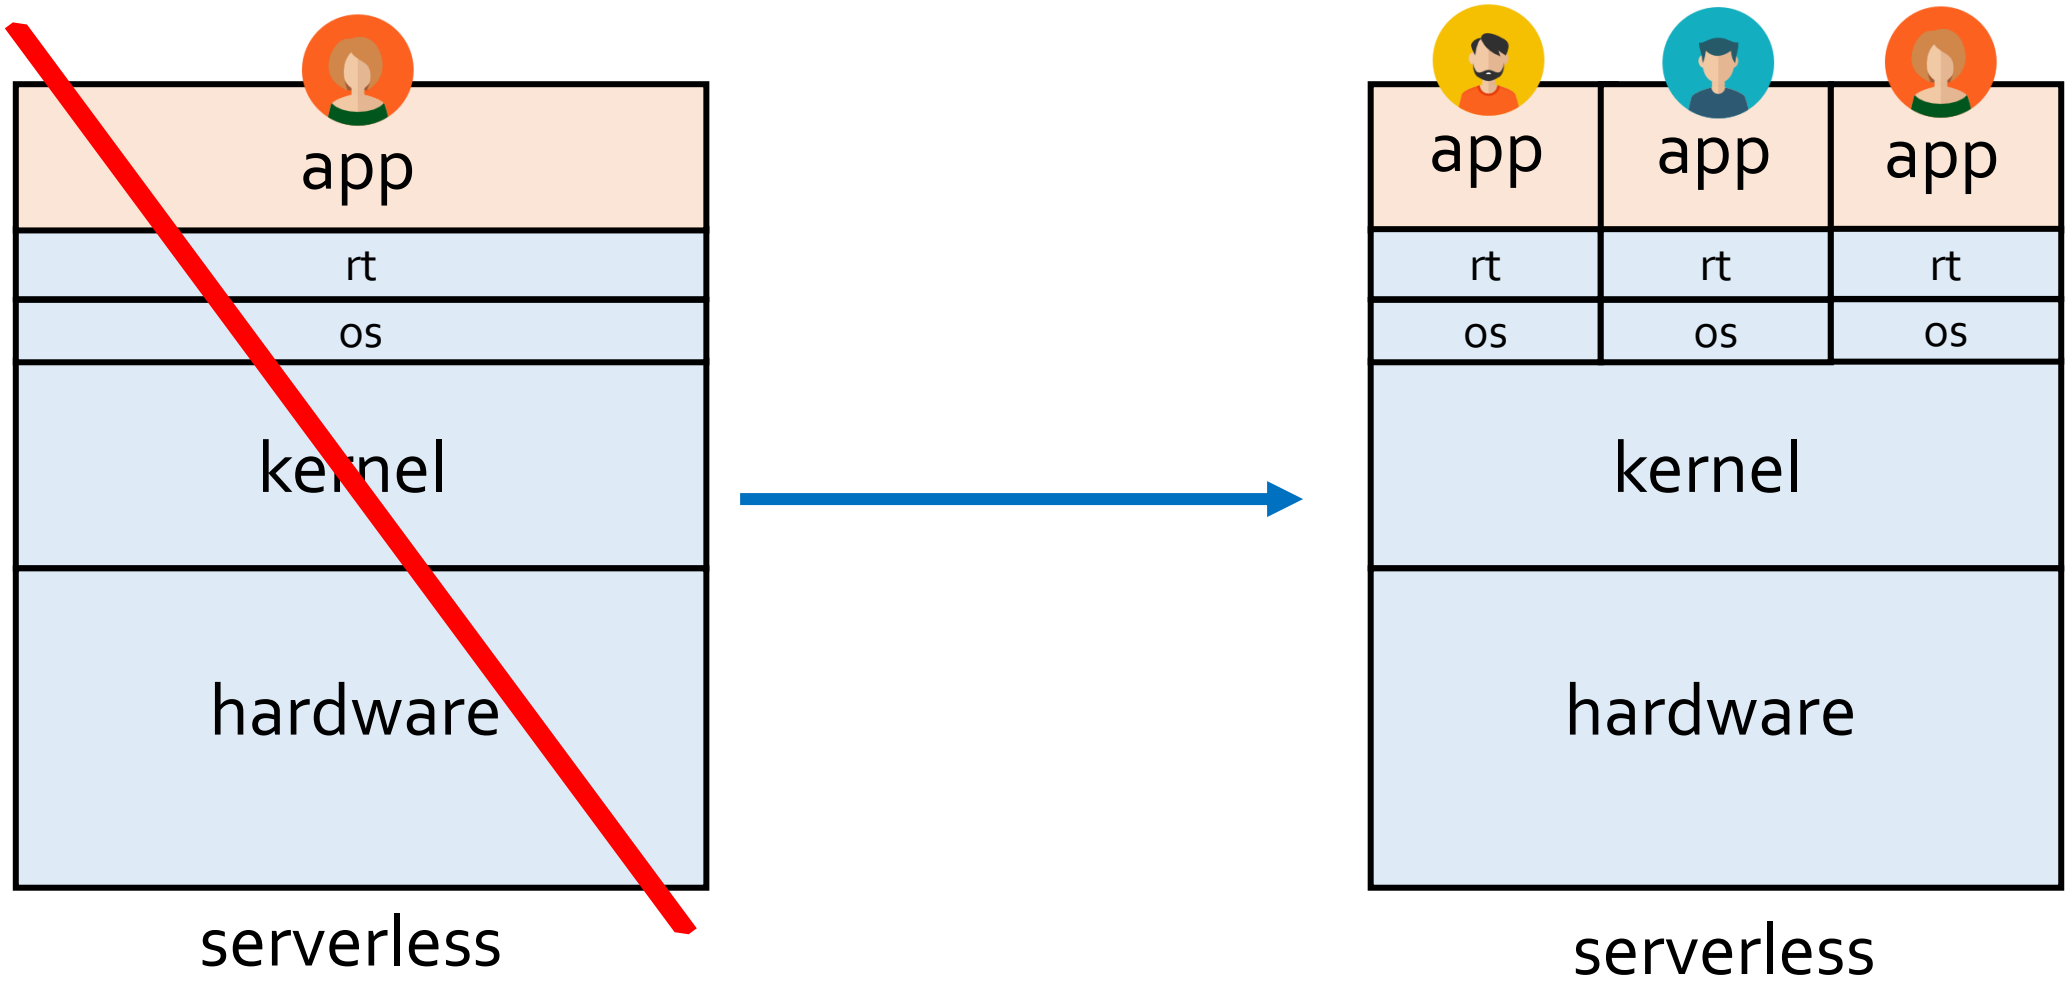

Topic 2: Improve security and control

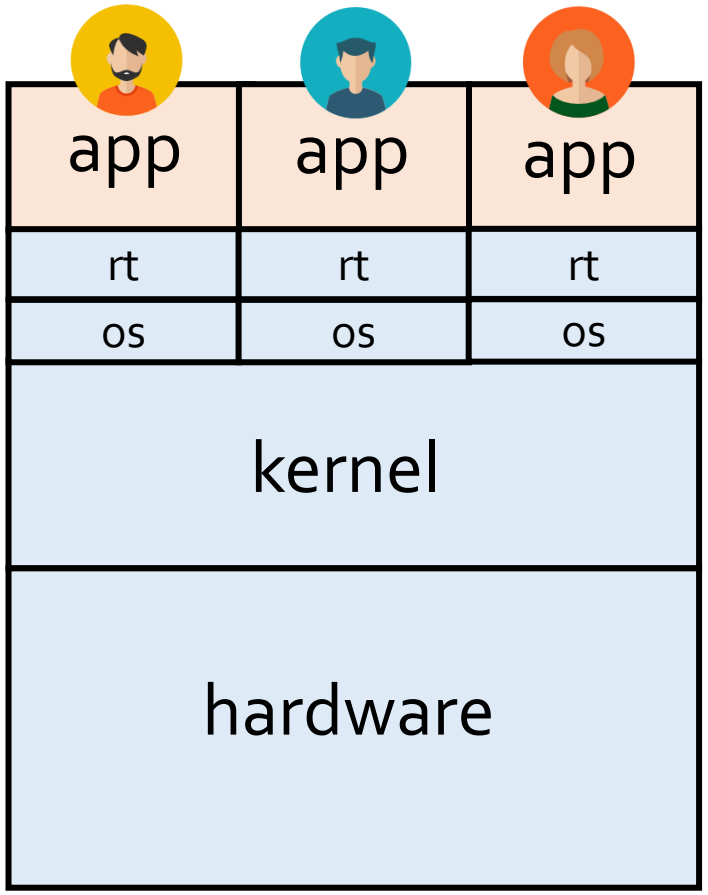

serverless

Topic 2: Improve security and control

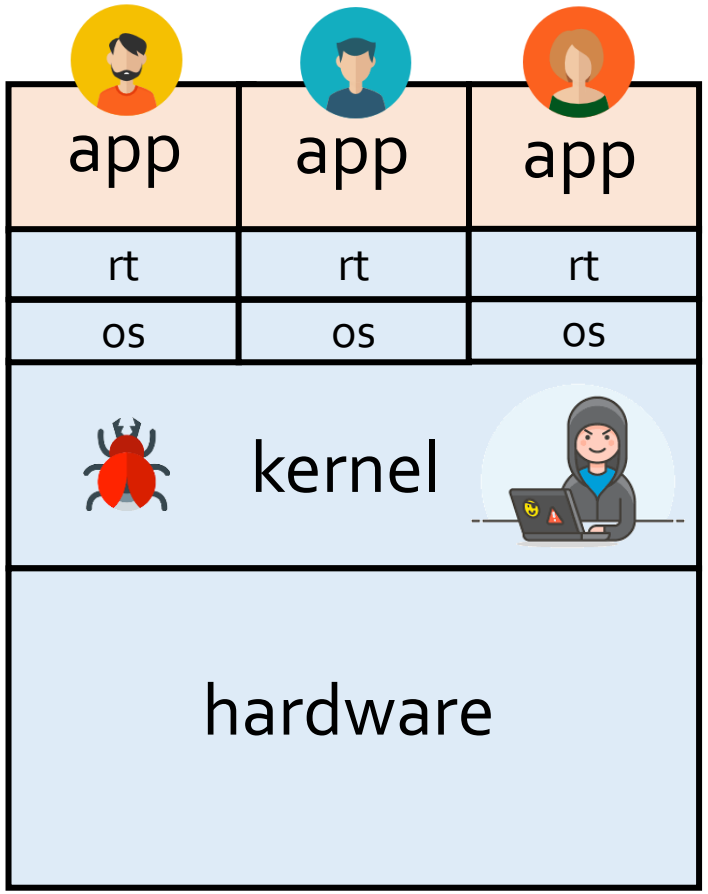

serverless

Topic 2: Improve security and control

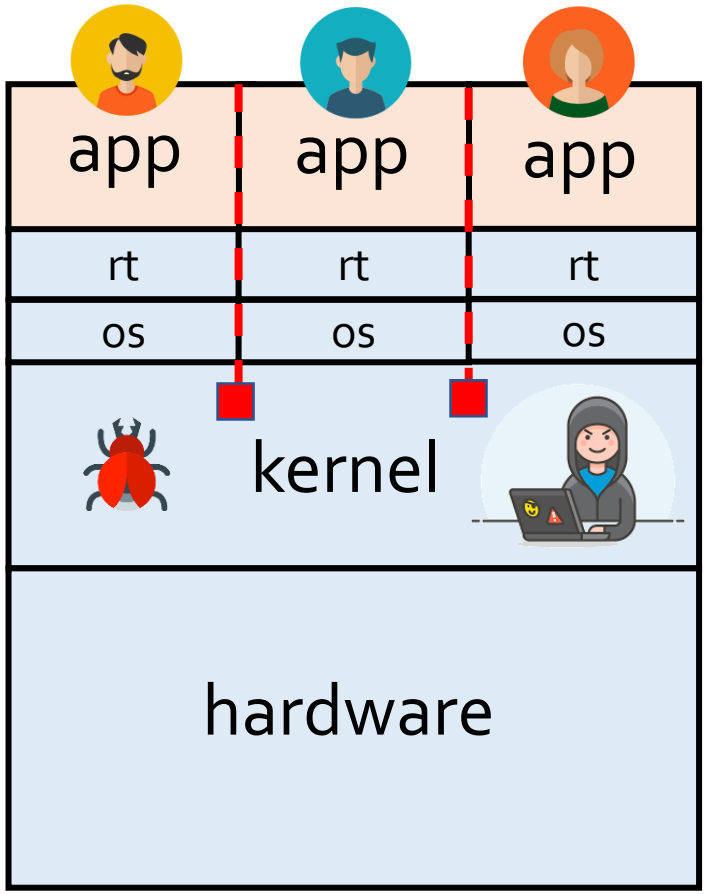

serverless

Topic 2: Improve security and control

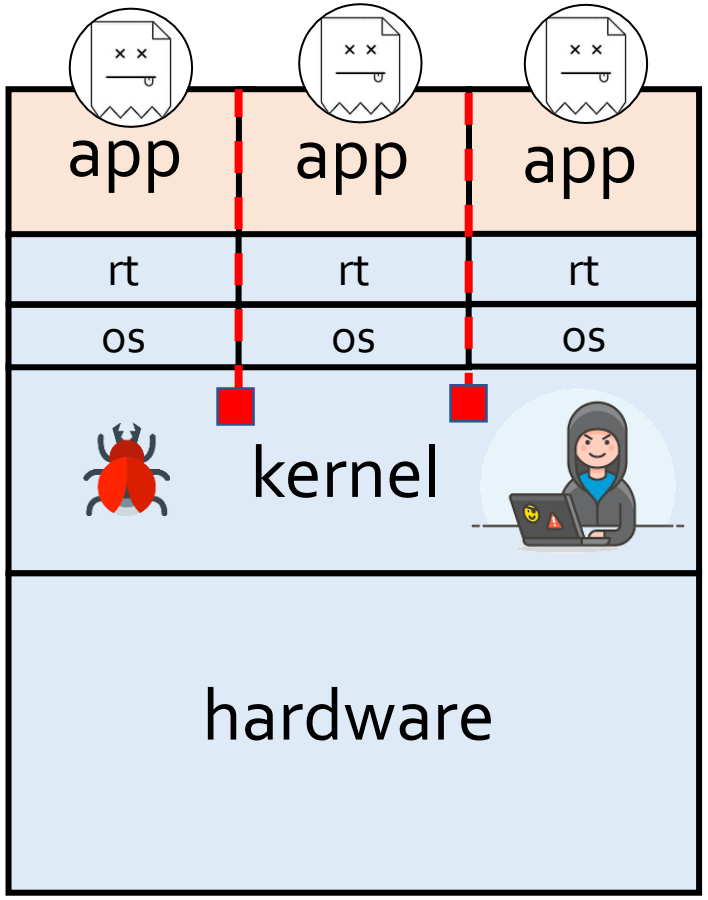

serverless

Topic 2: Improve security and control

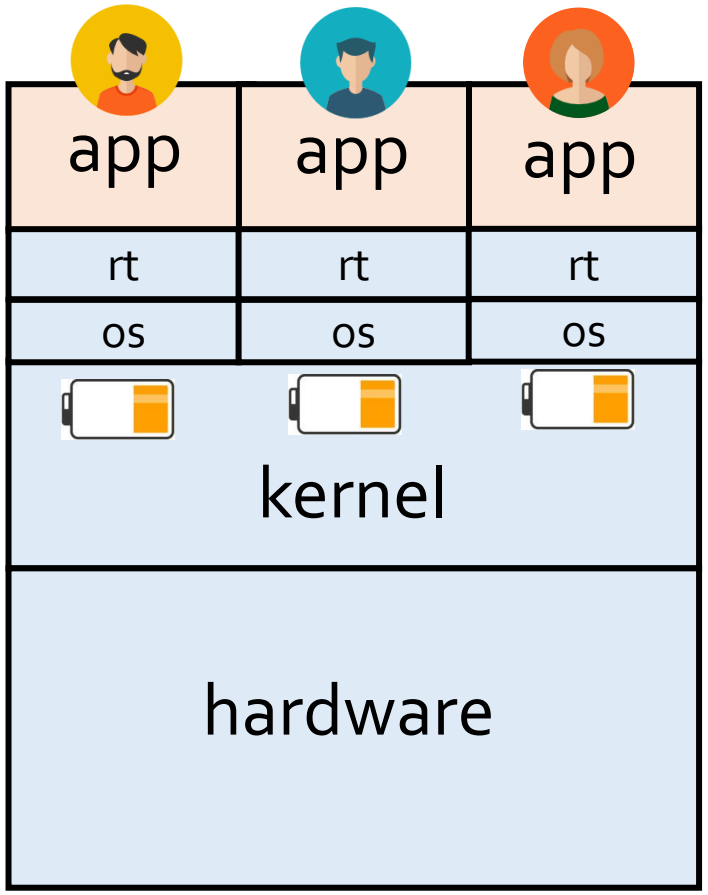

serverless

Topic 2: Improve security and control

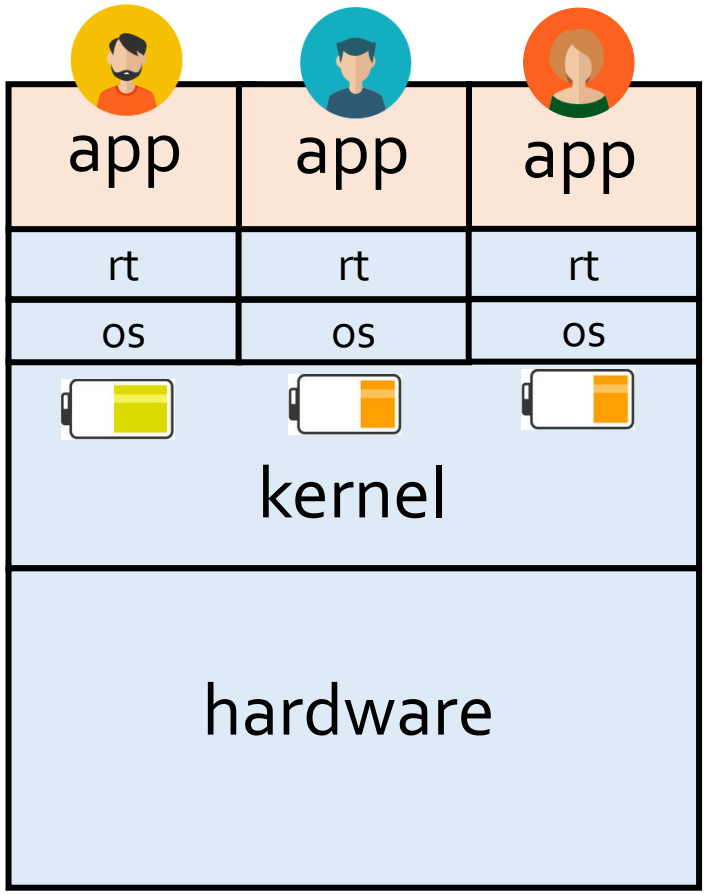

serverless

Topic 2: Improve security and control

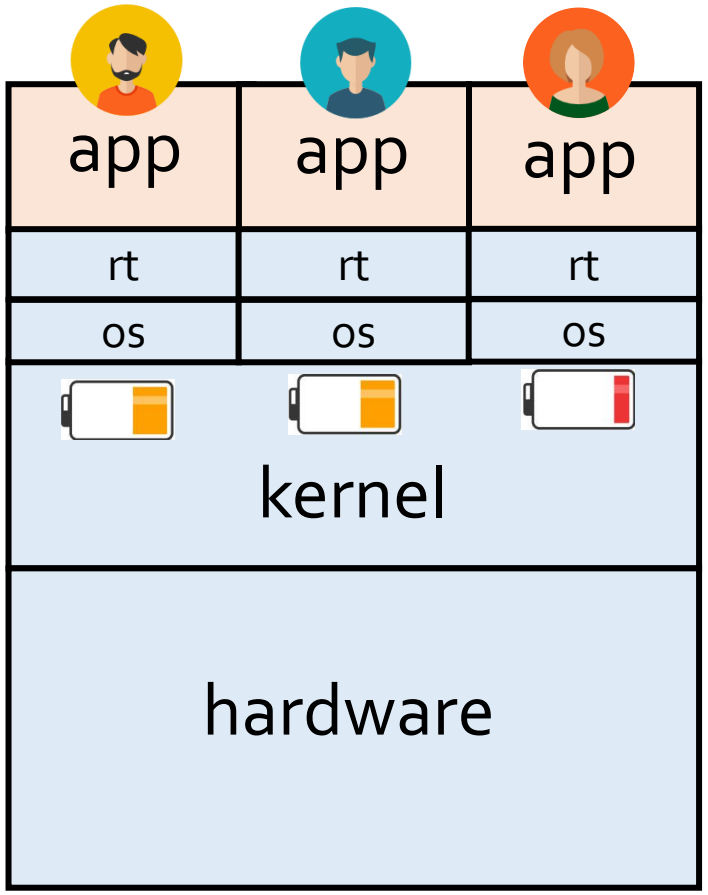

serverless

Topic 2: Improve security and control

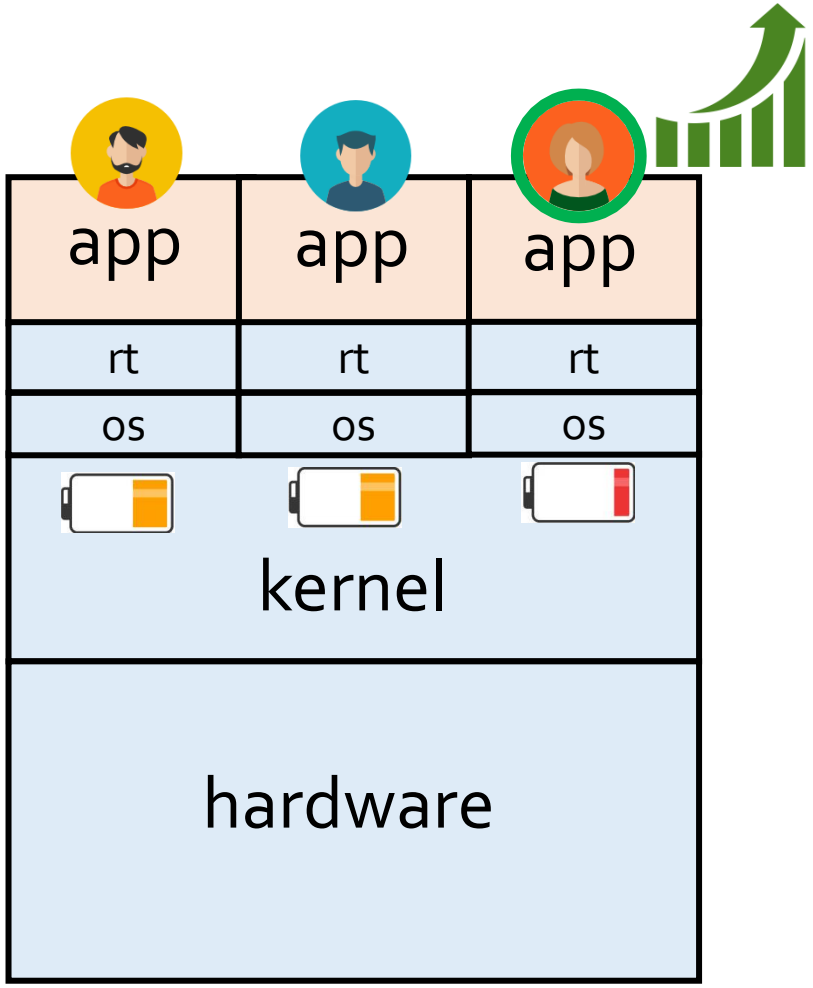

serverless

Topic 2: Improve security and control

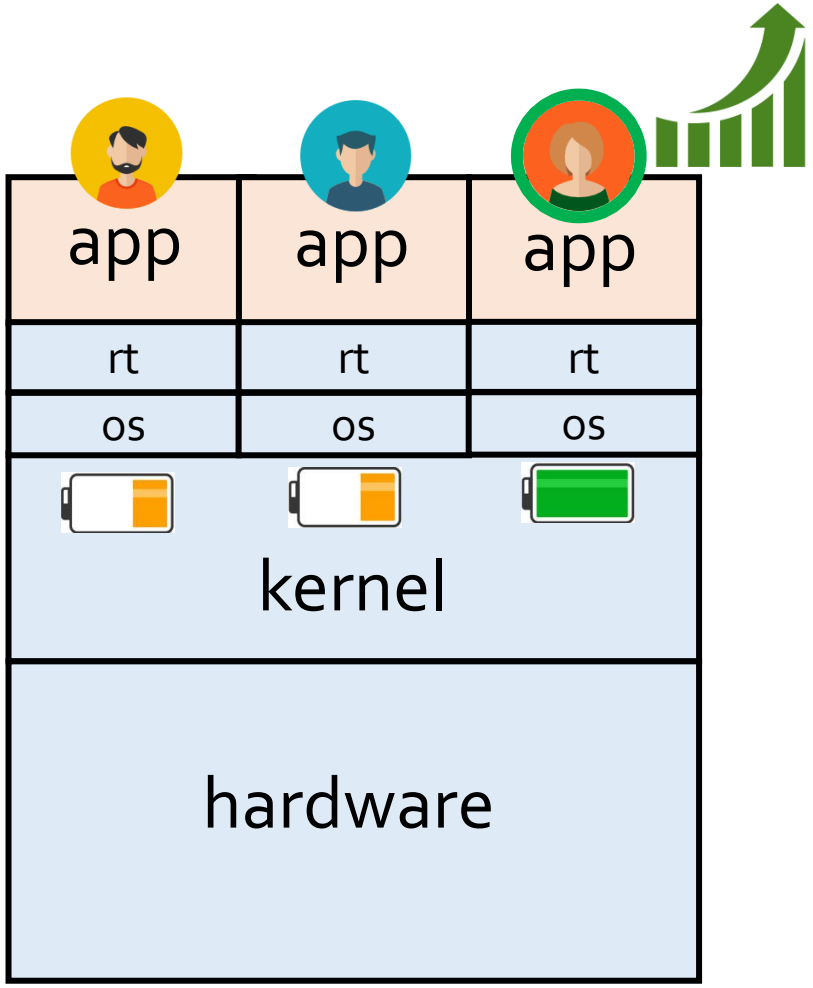

serverless

Topic 2: Improve security and control

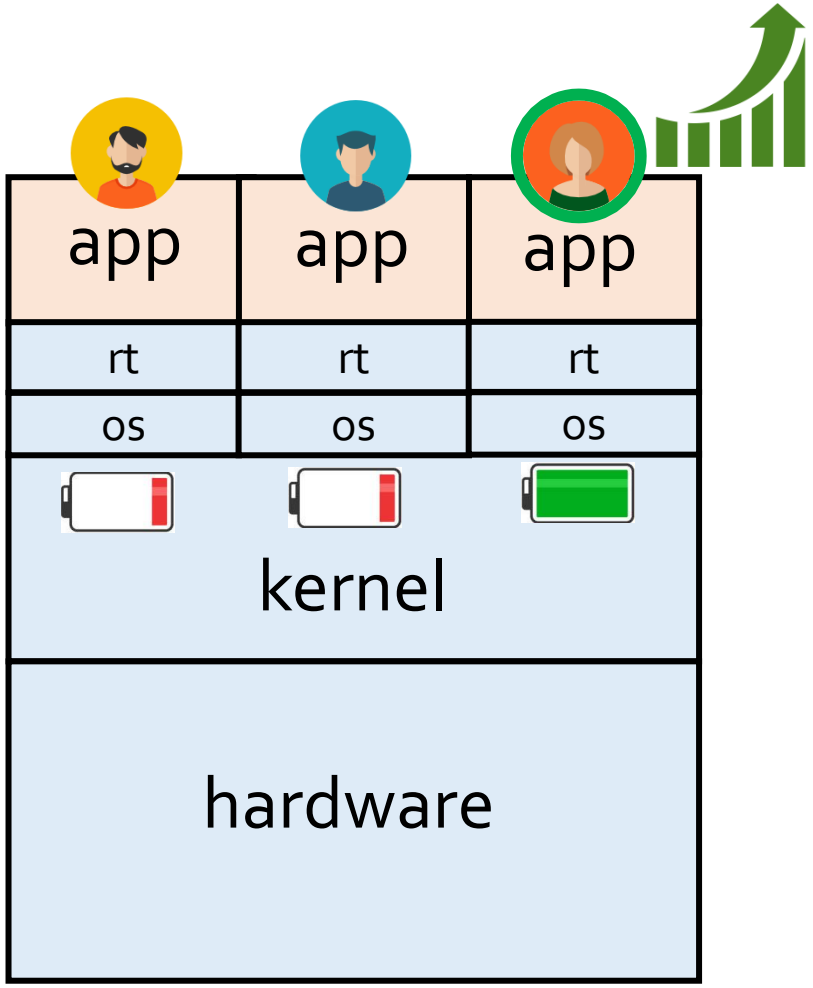

serverless

Topic 2: Improve security and control

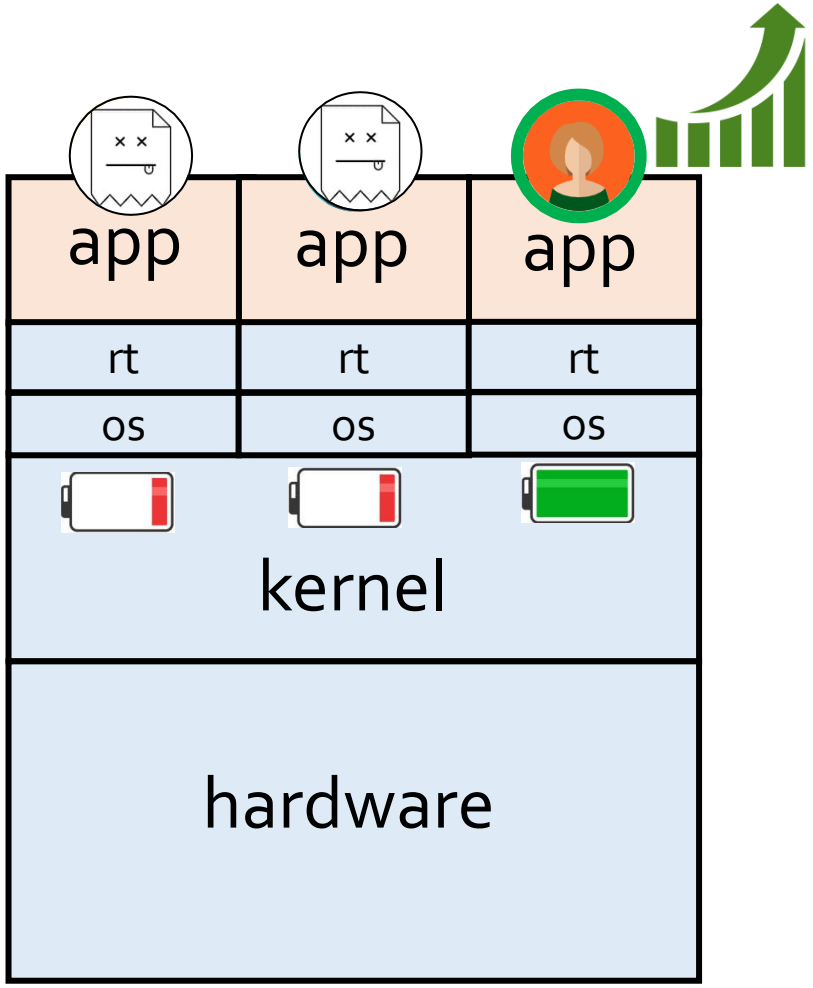

serverless

Topic 2: Improve security and control

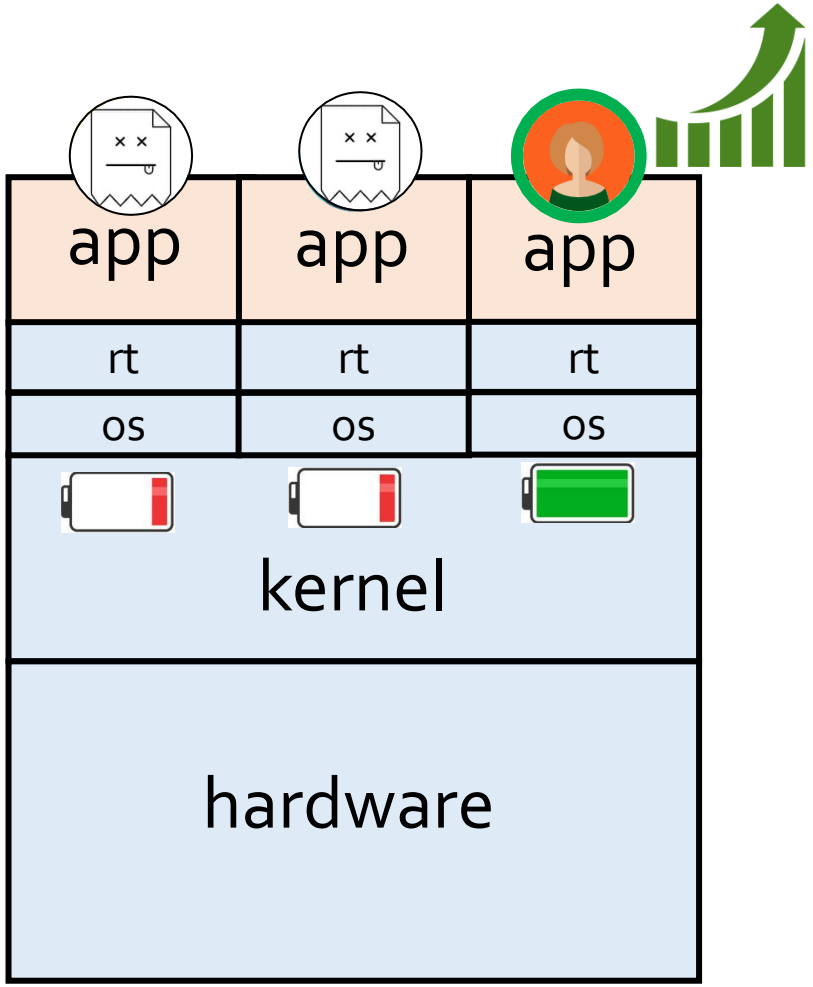

serverless

CPU Problem

Network Problem

Memory Works OK

Topic 2: Improve security and control

Partially solved by running containers inside VMs

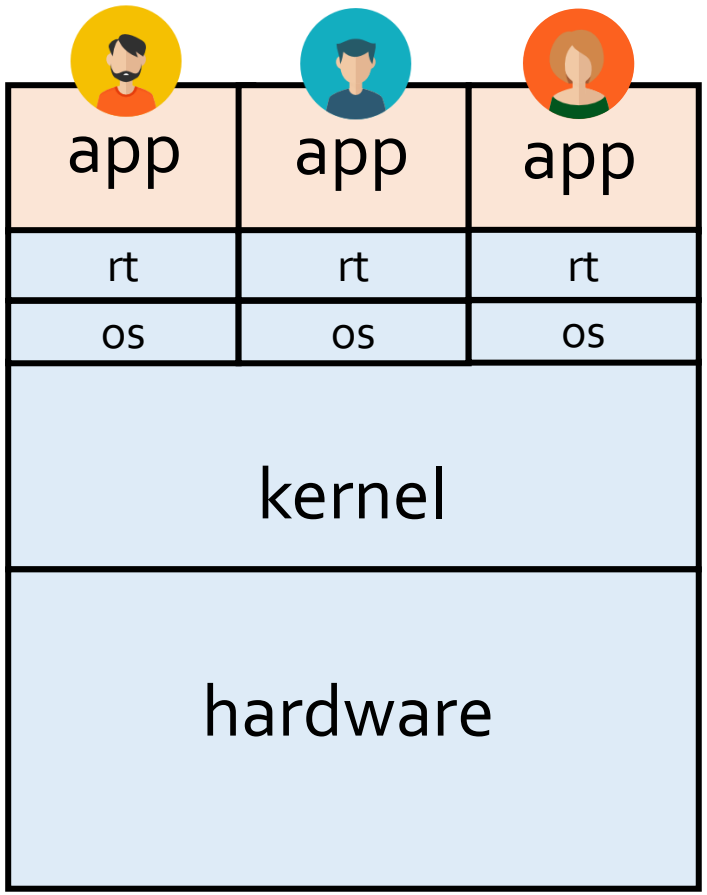

Topic 2: Improve security and control

Partially solved by running containers inside VMs

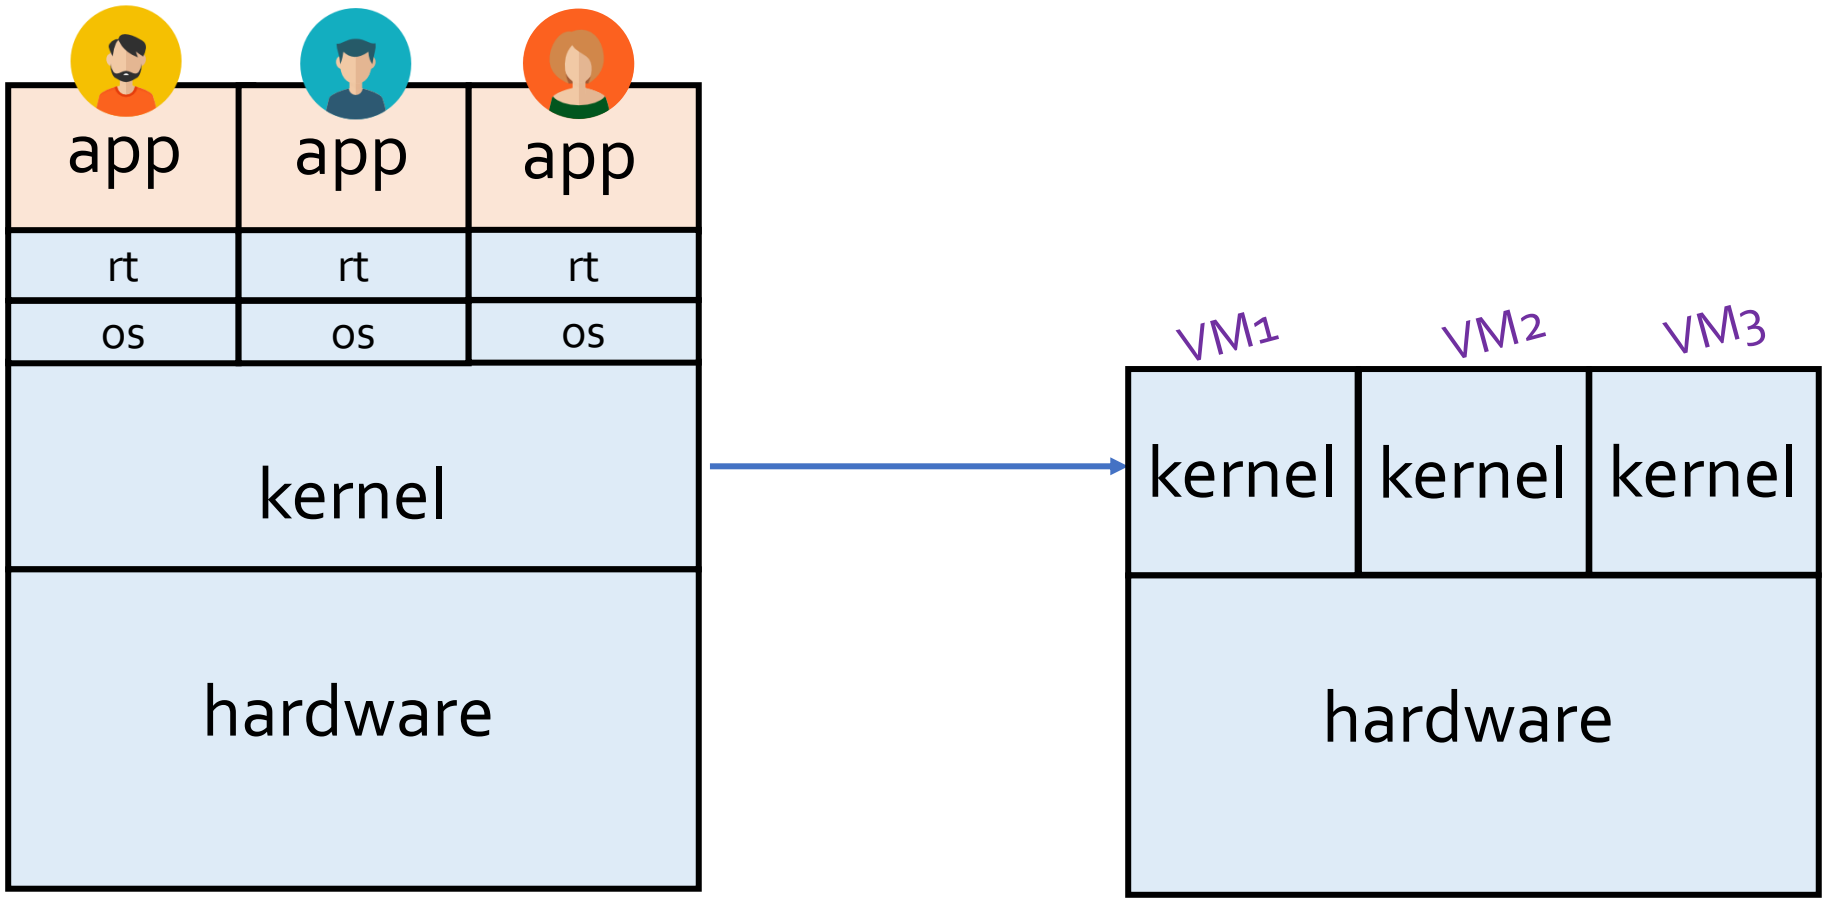

Topic 2: Improve security and control

Partially solved by running containers inside VMs

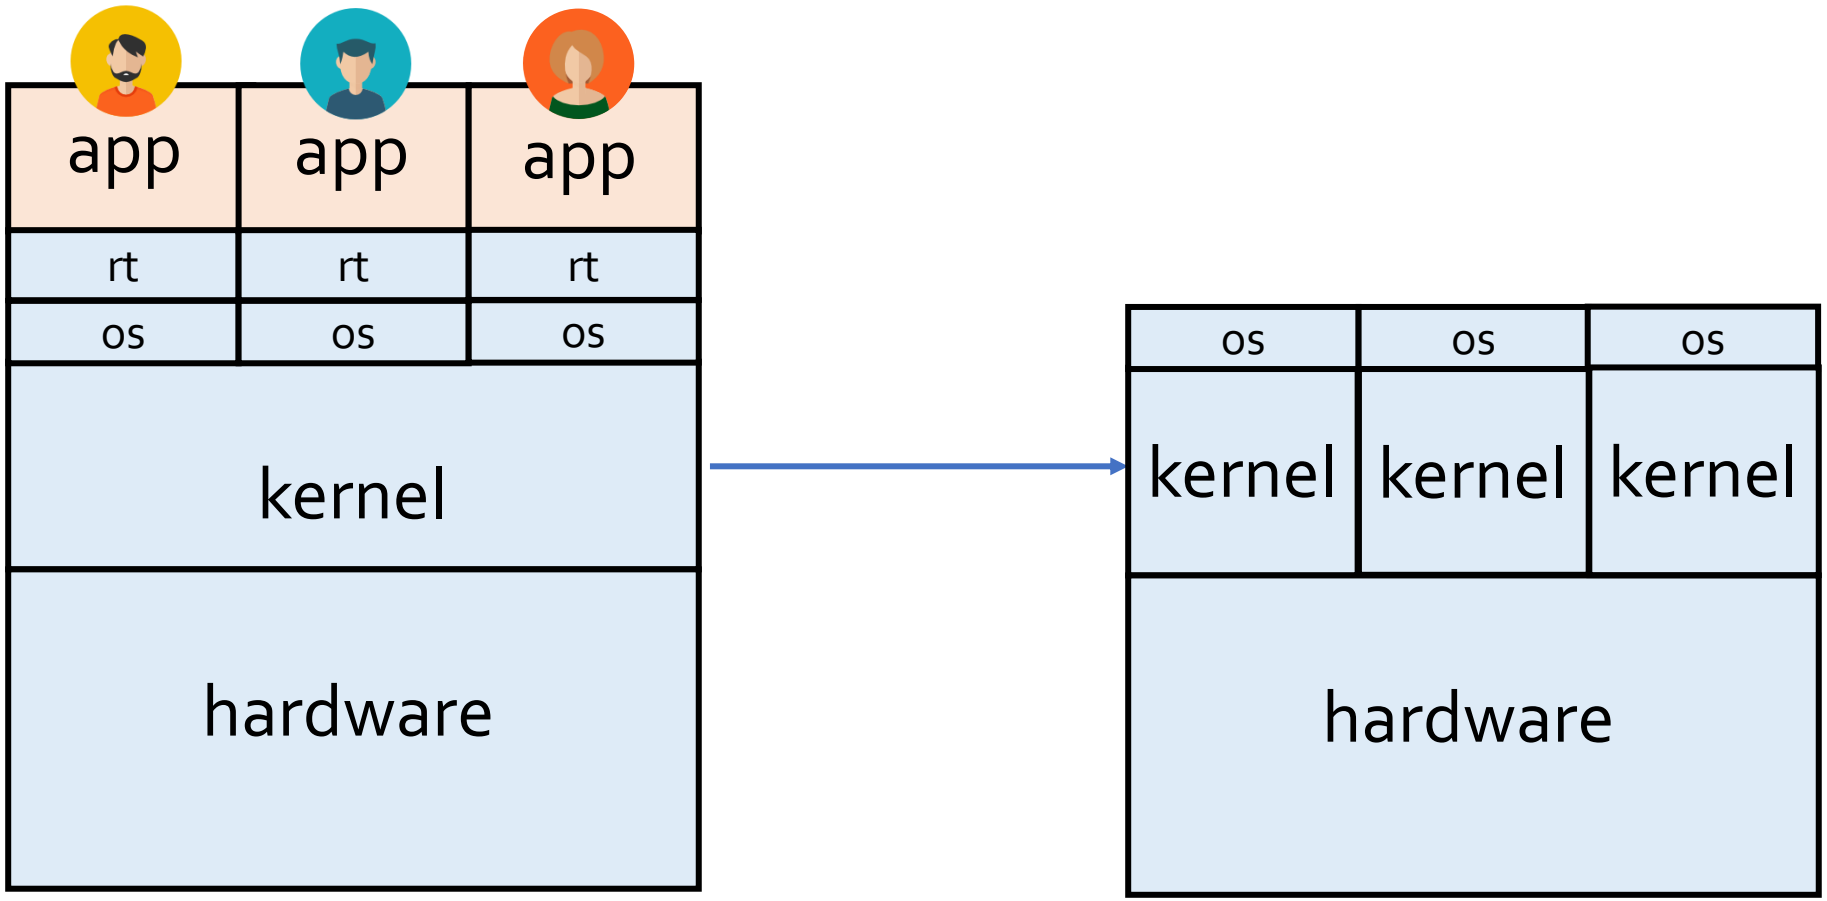

Topic 2: Improve security and control

Partially solved by running containers inside VMs

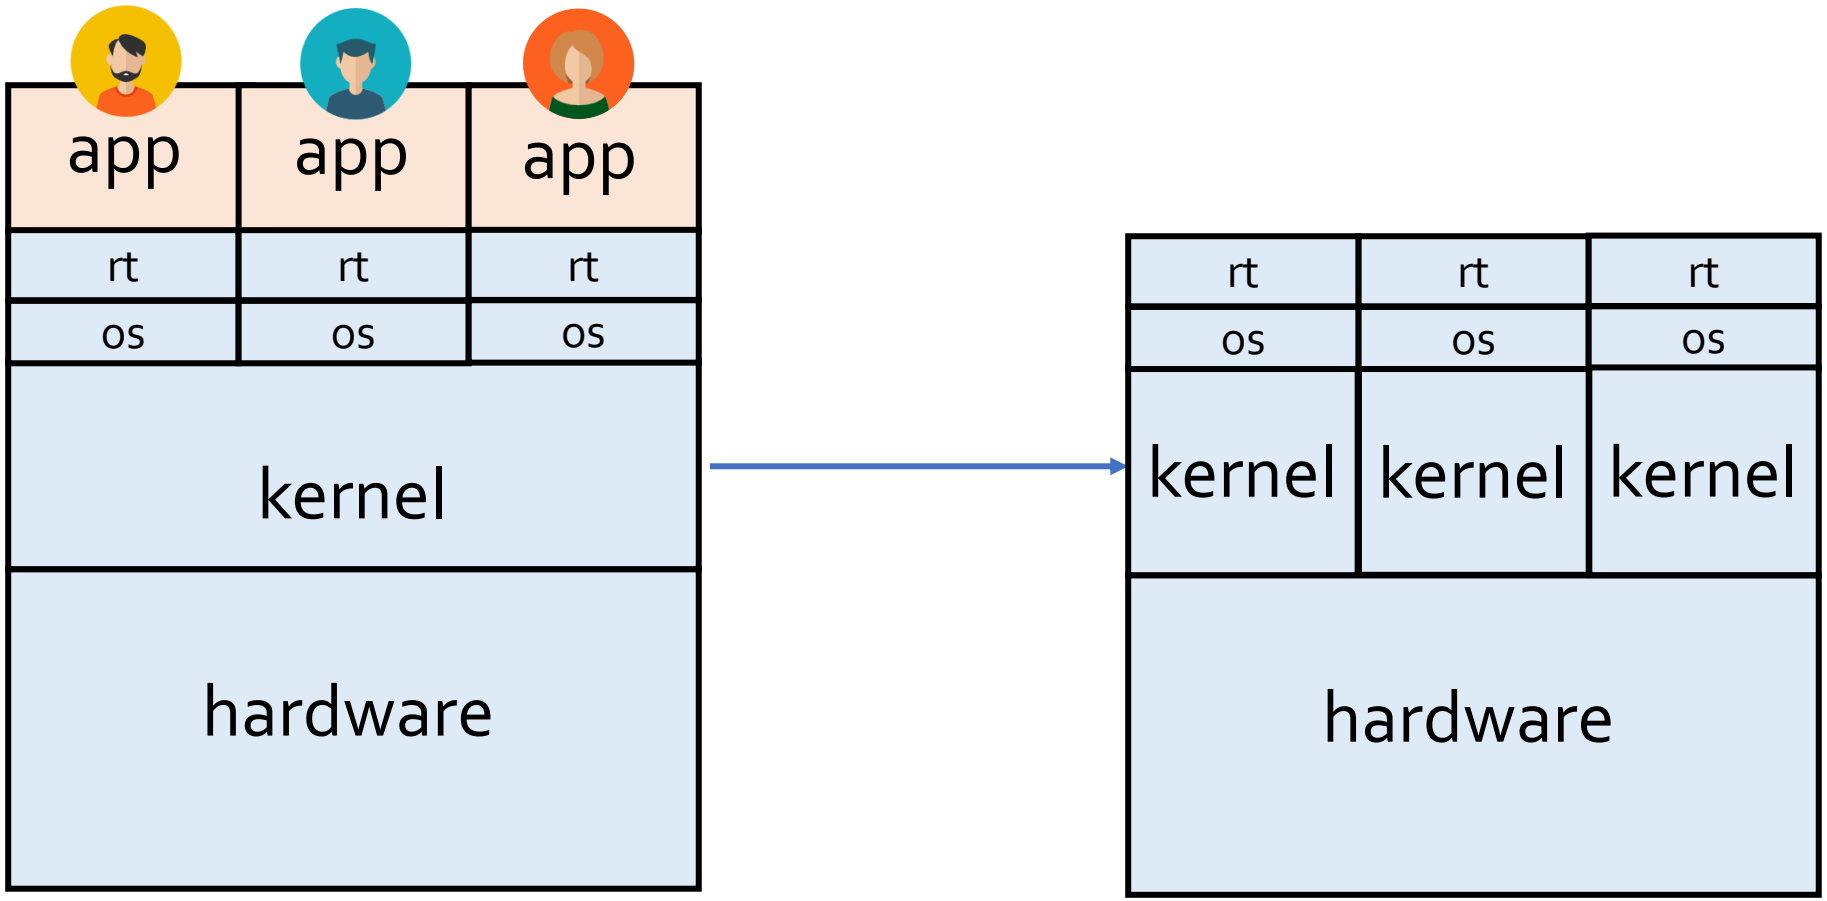

Topic 2: Improve security and control

Partially solved by running containers inside VMs

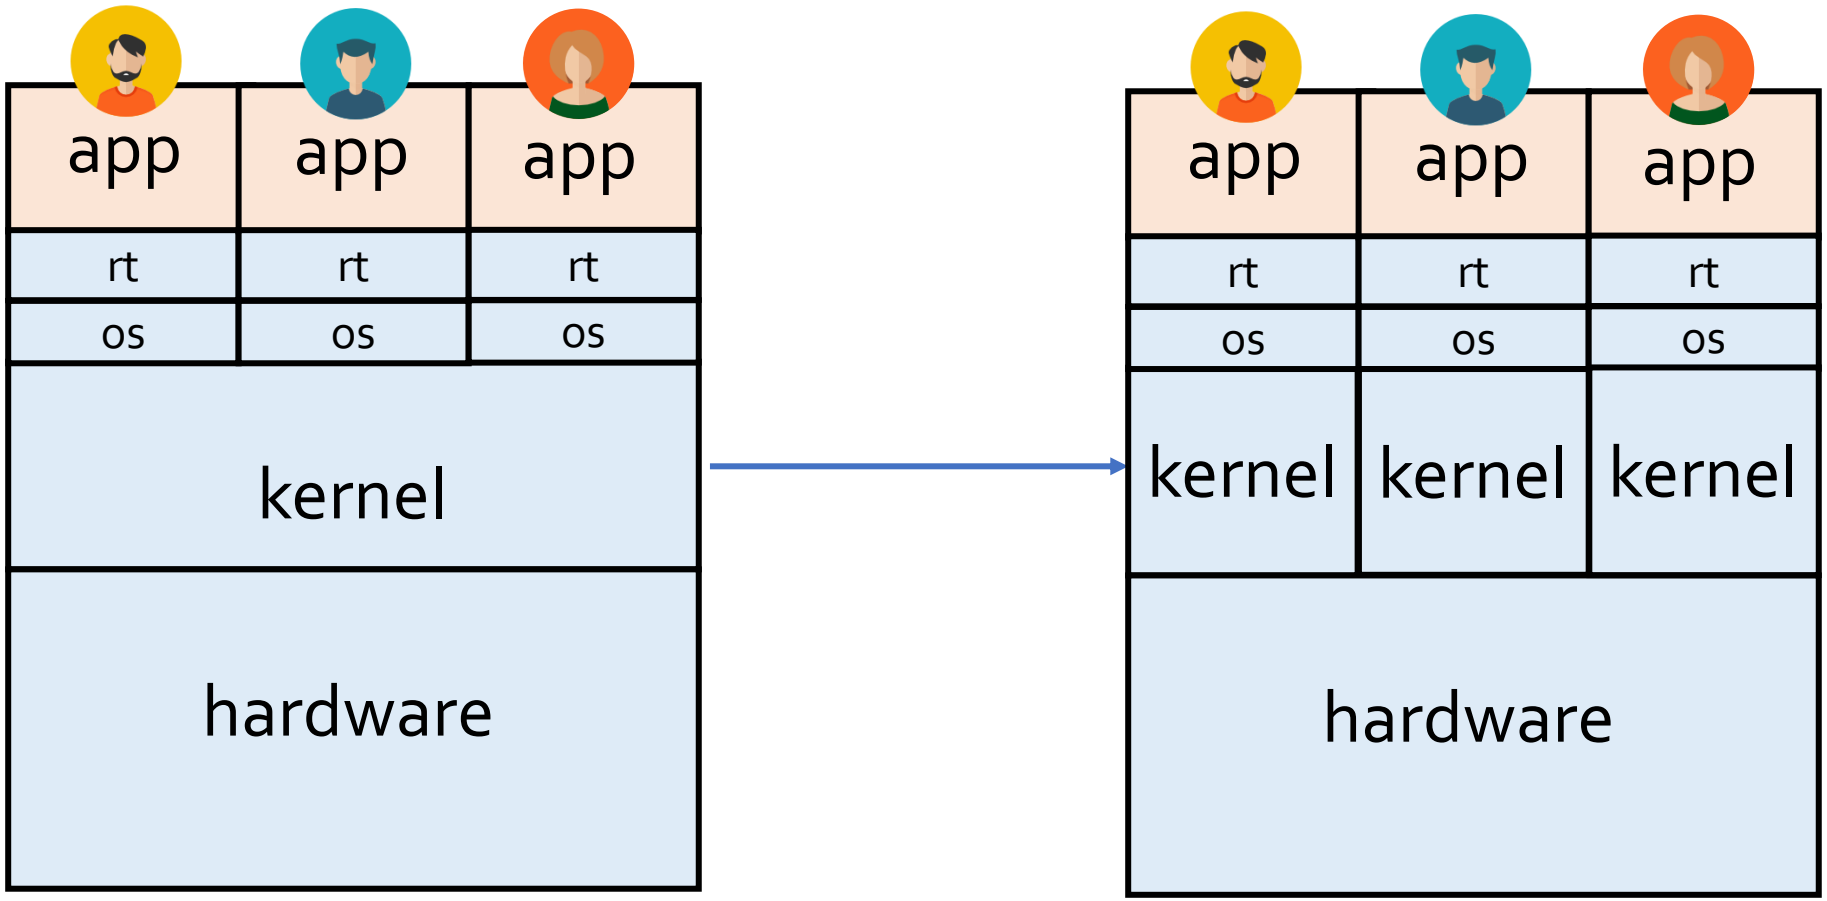

Topic 2: Improve security and control

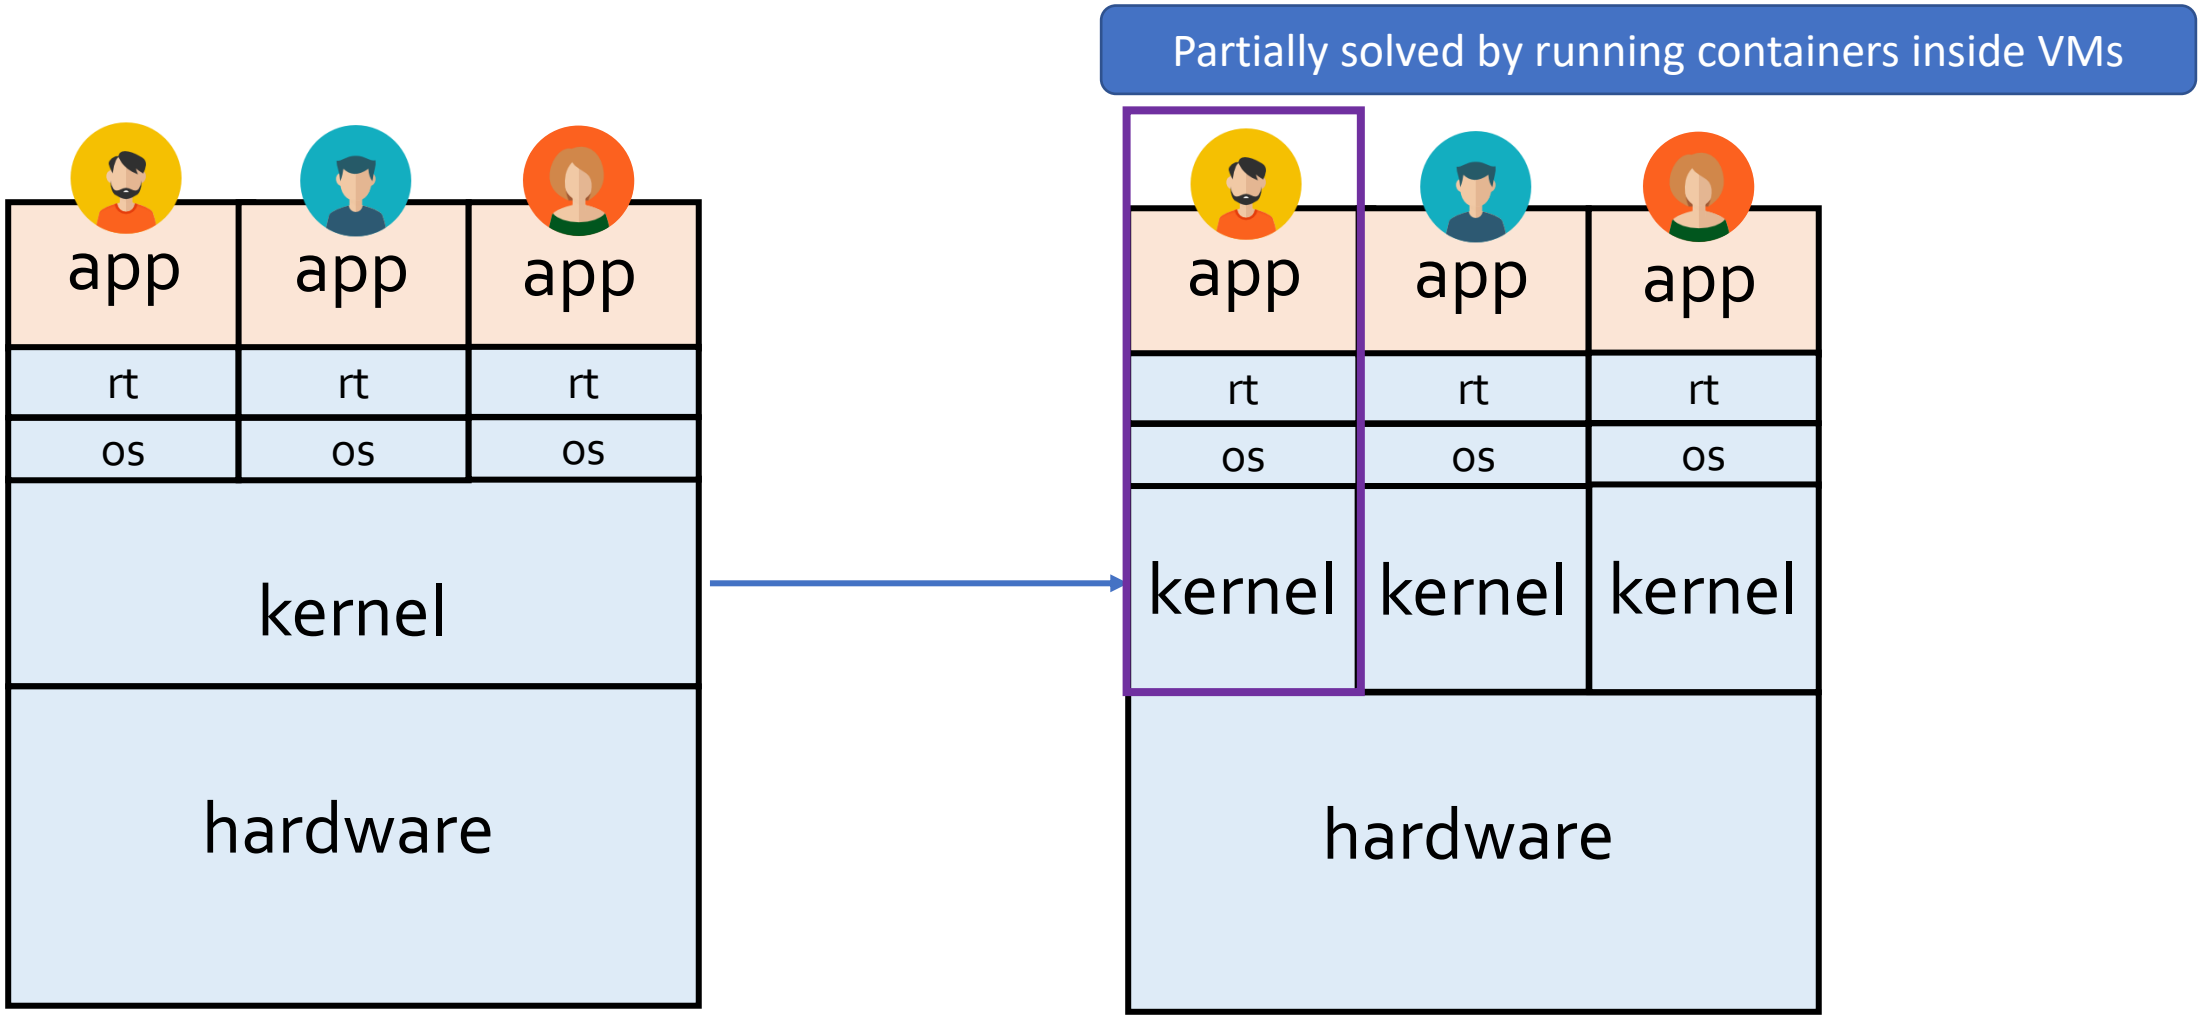

Topic 2: Improve security and control

Partially solved by running containers inside VMs

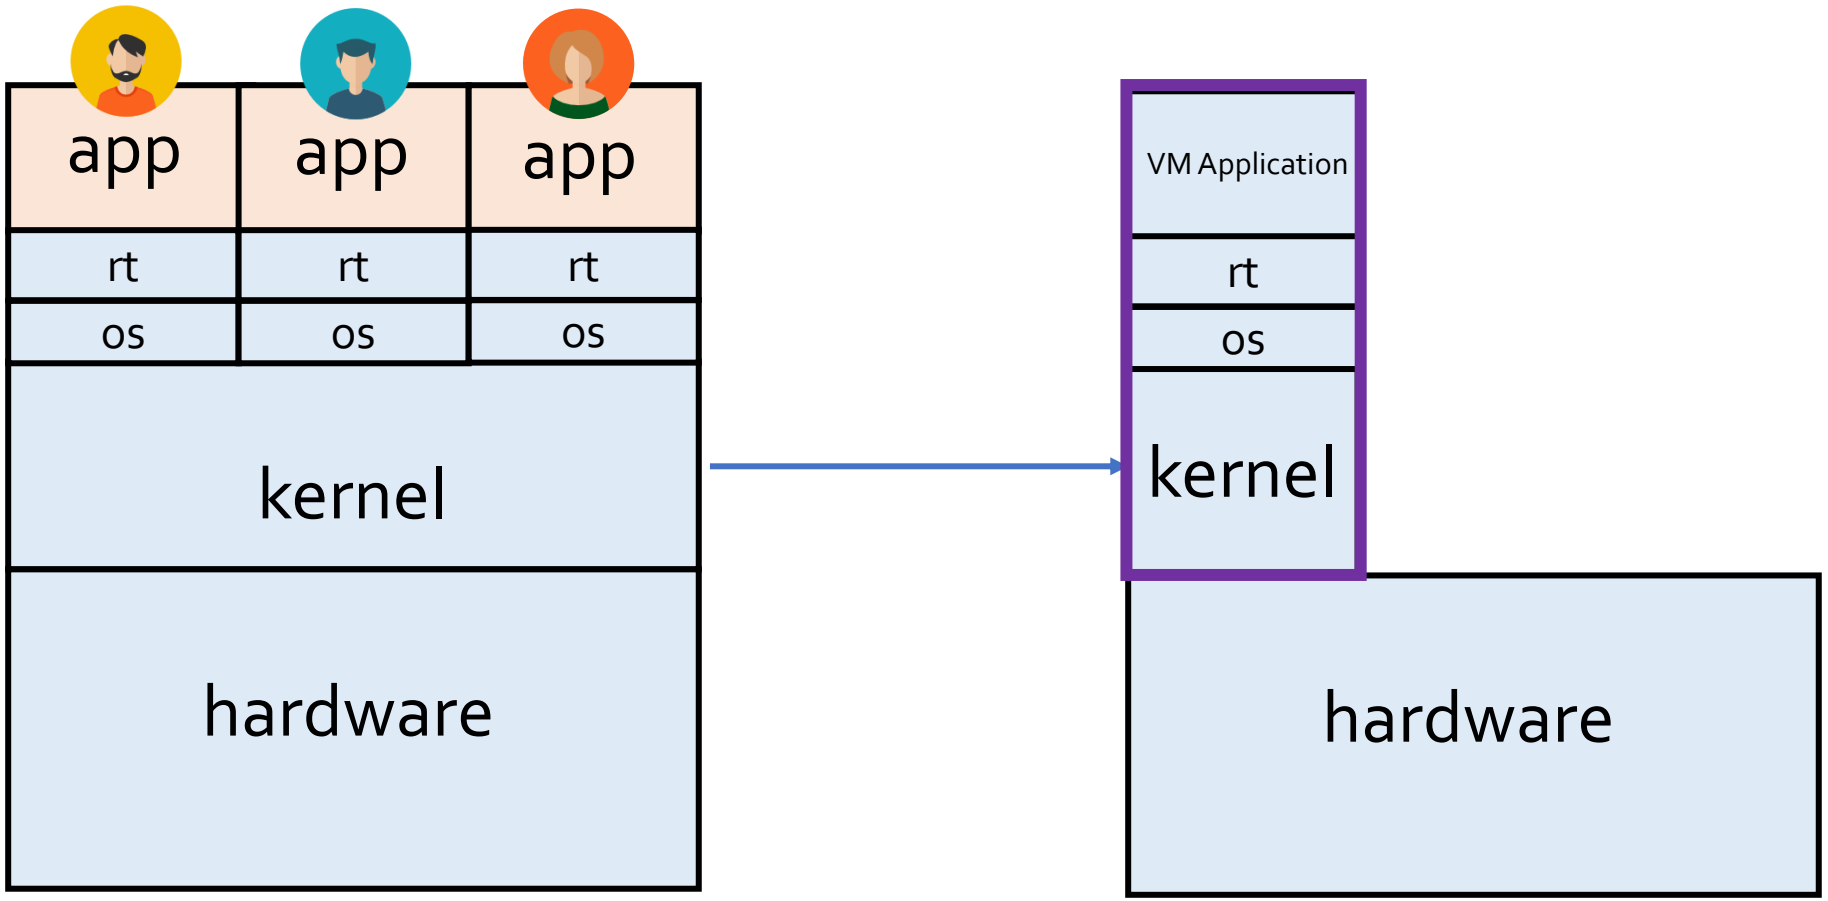

Topic 2: Improve security and control

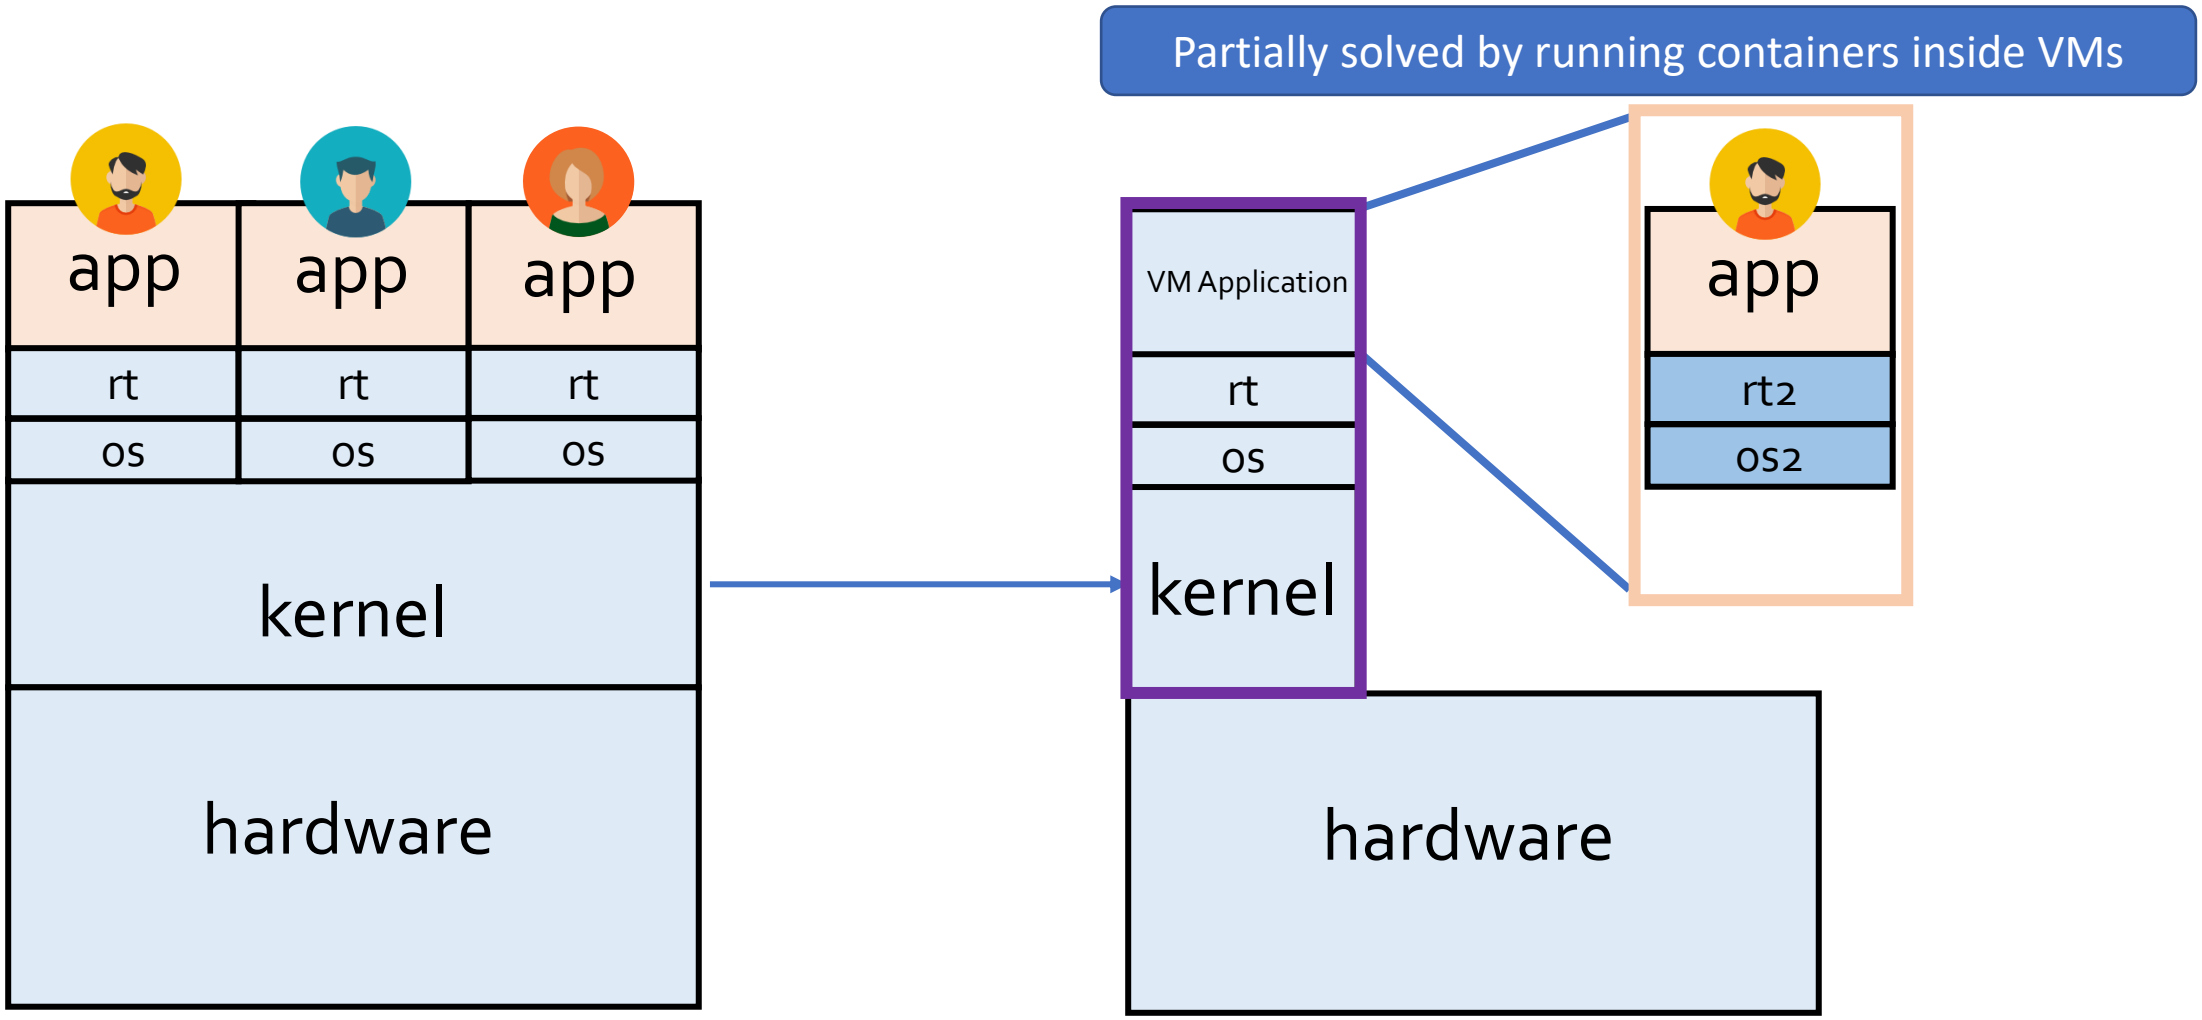

Topic 2: Improve security and control

MicroVMs

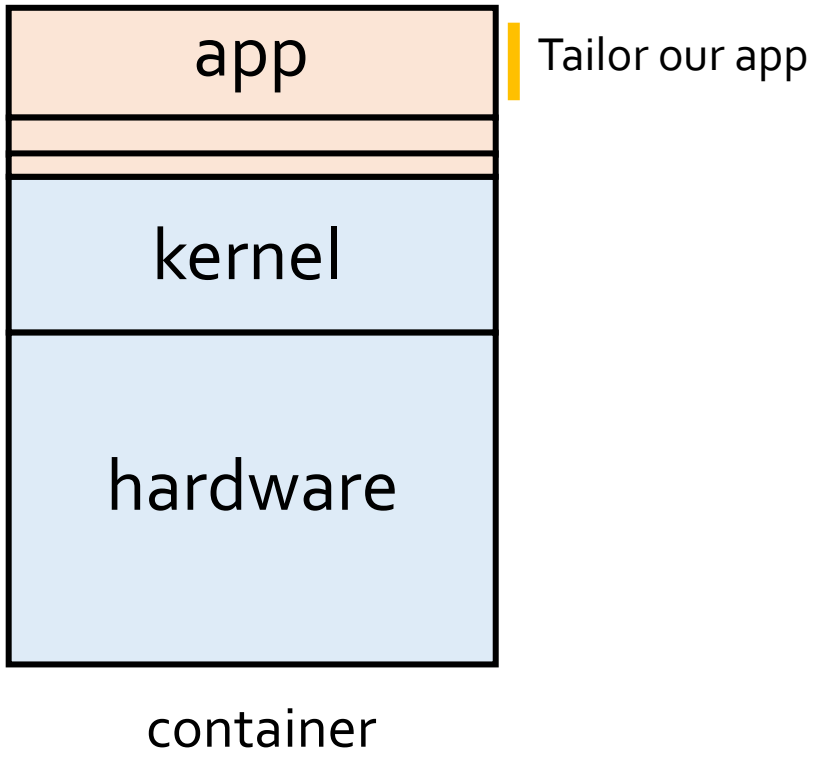

Topic 2: Improve security and control

MicroVMs

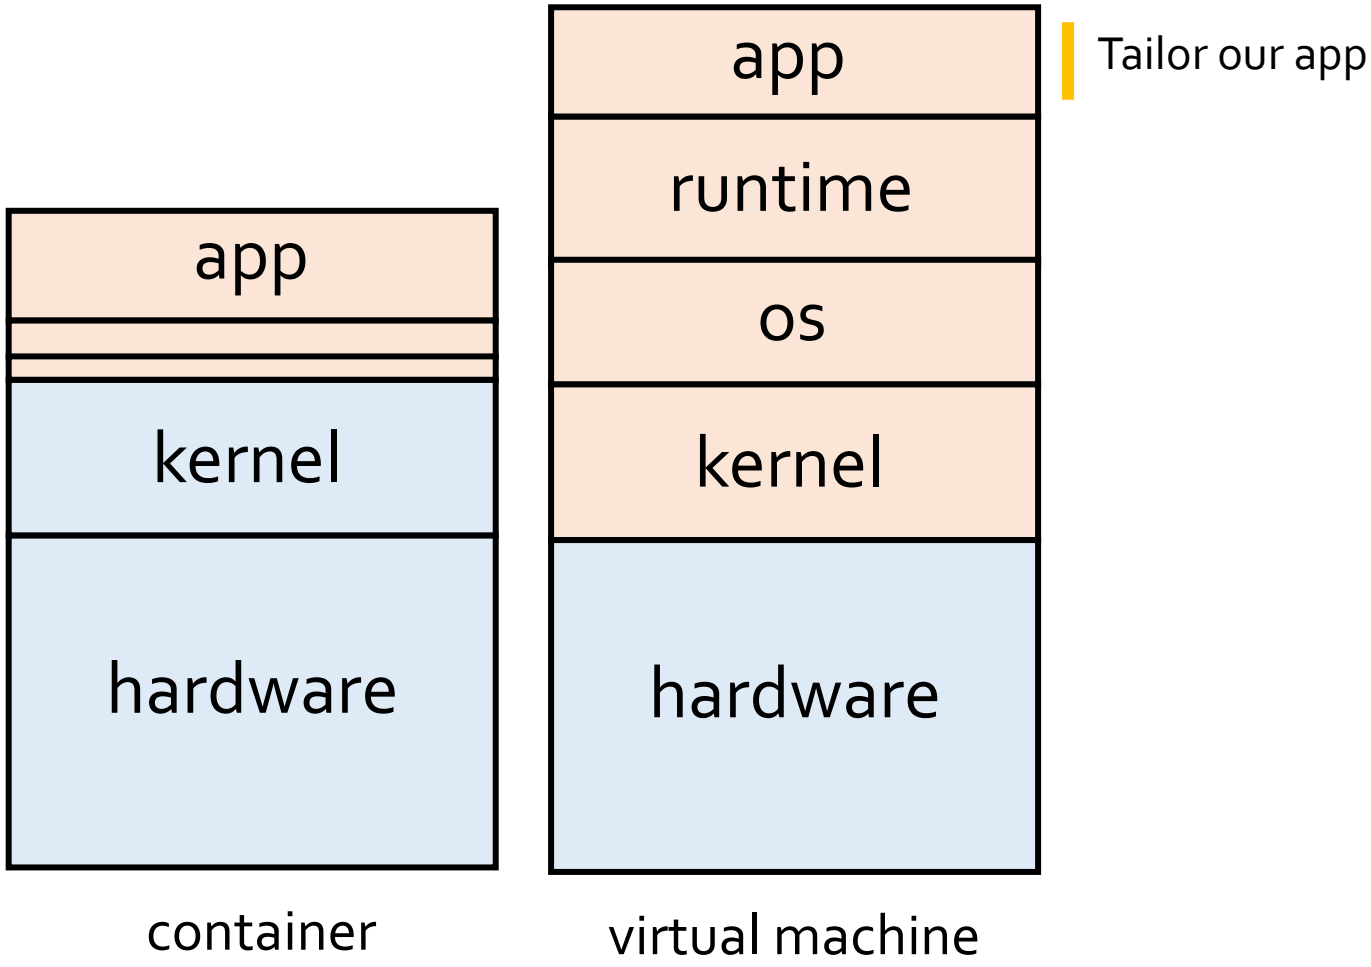

Topic 2: Improve security and control

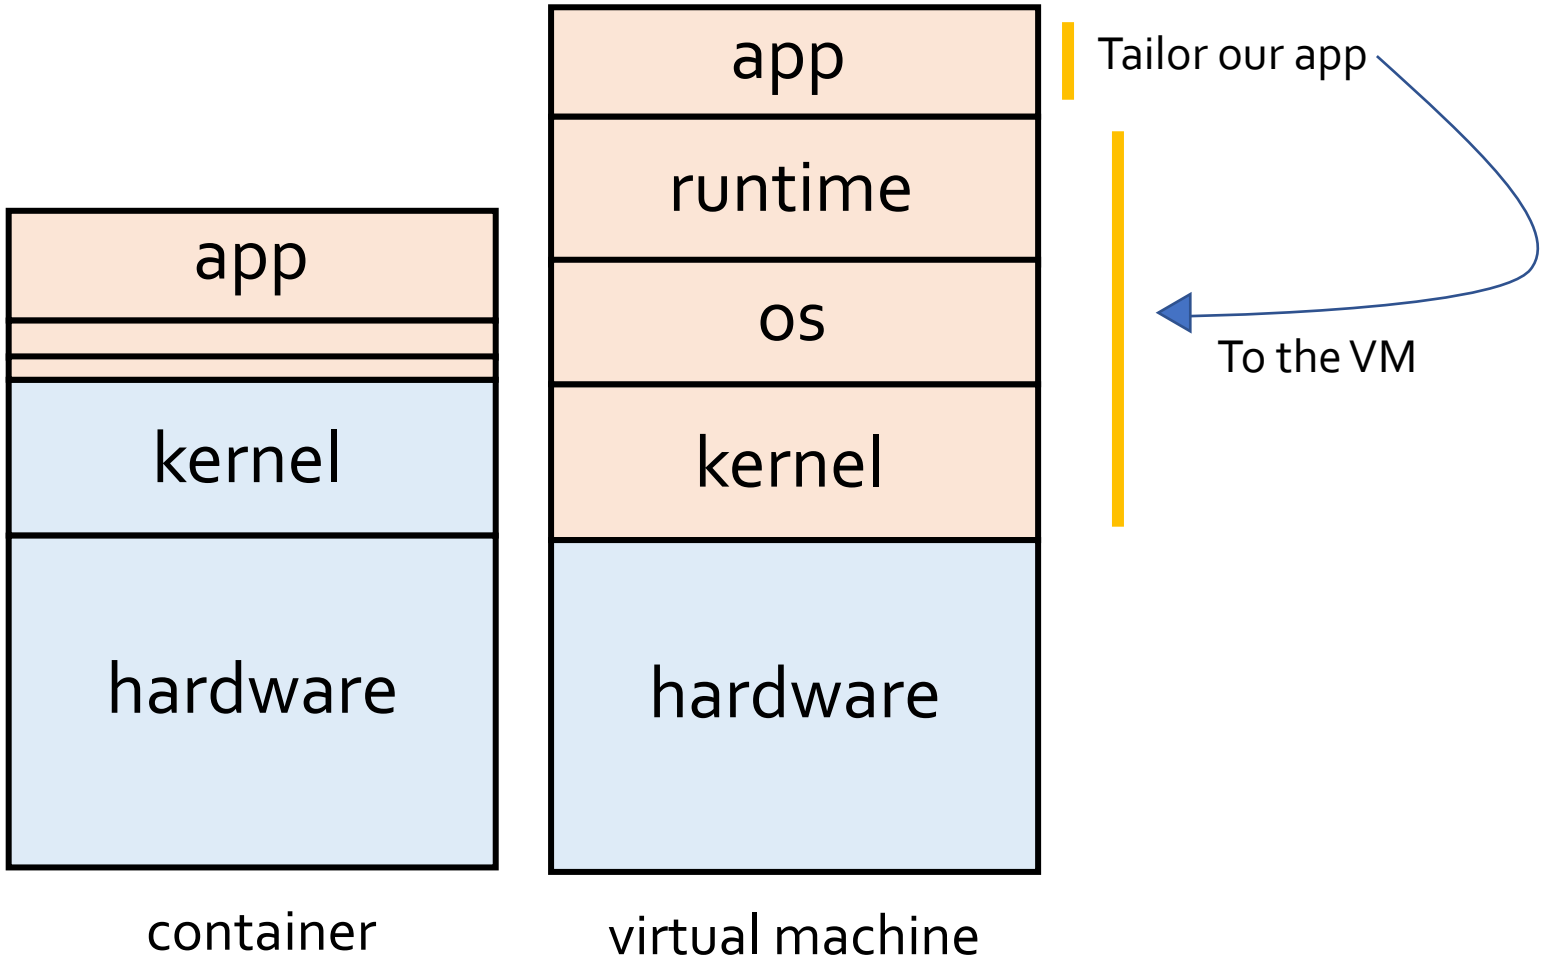

Topic 2: Improve security and control

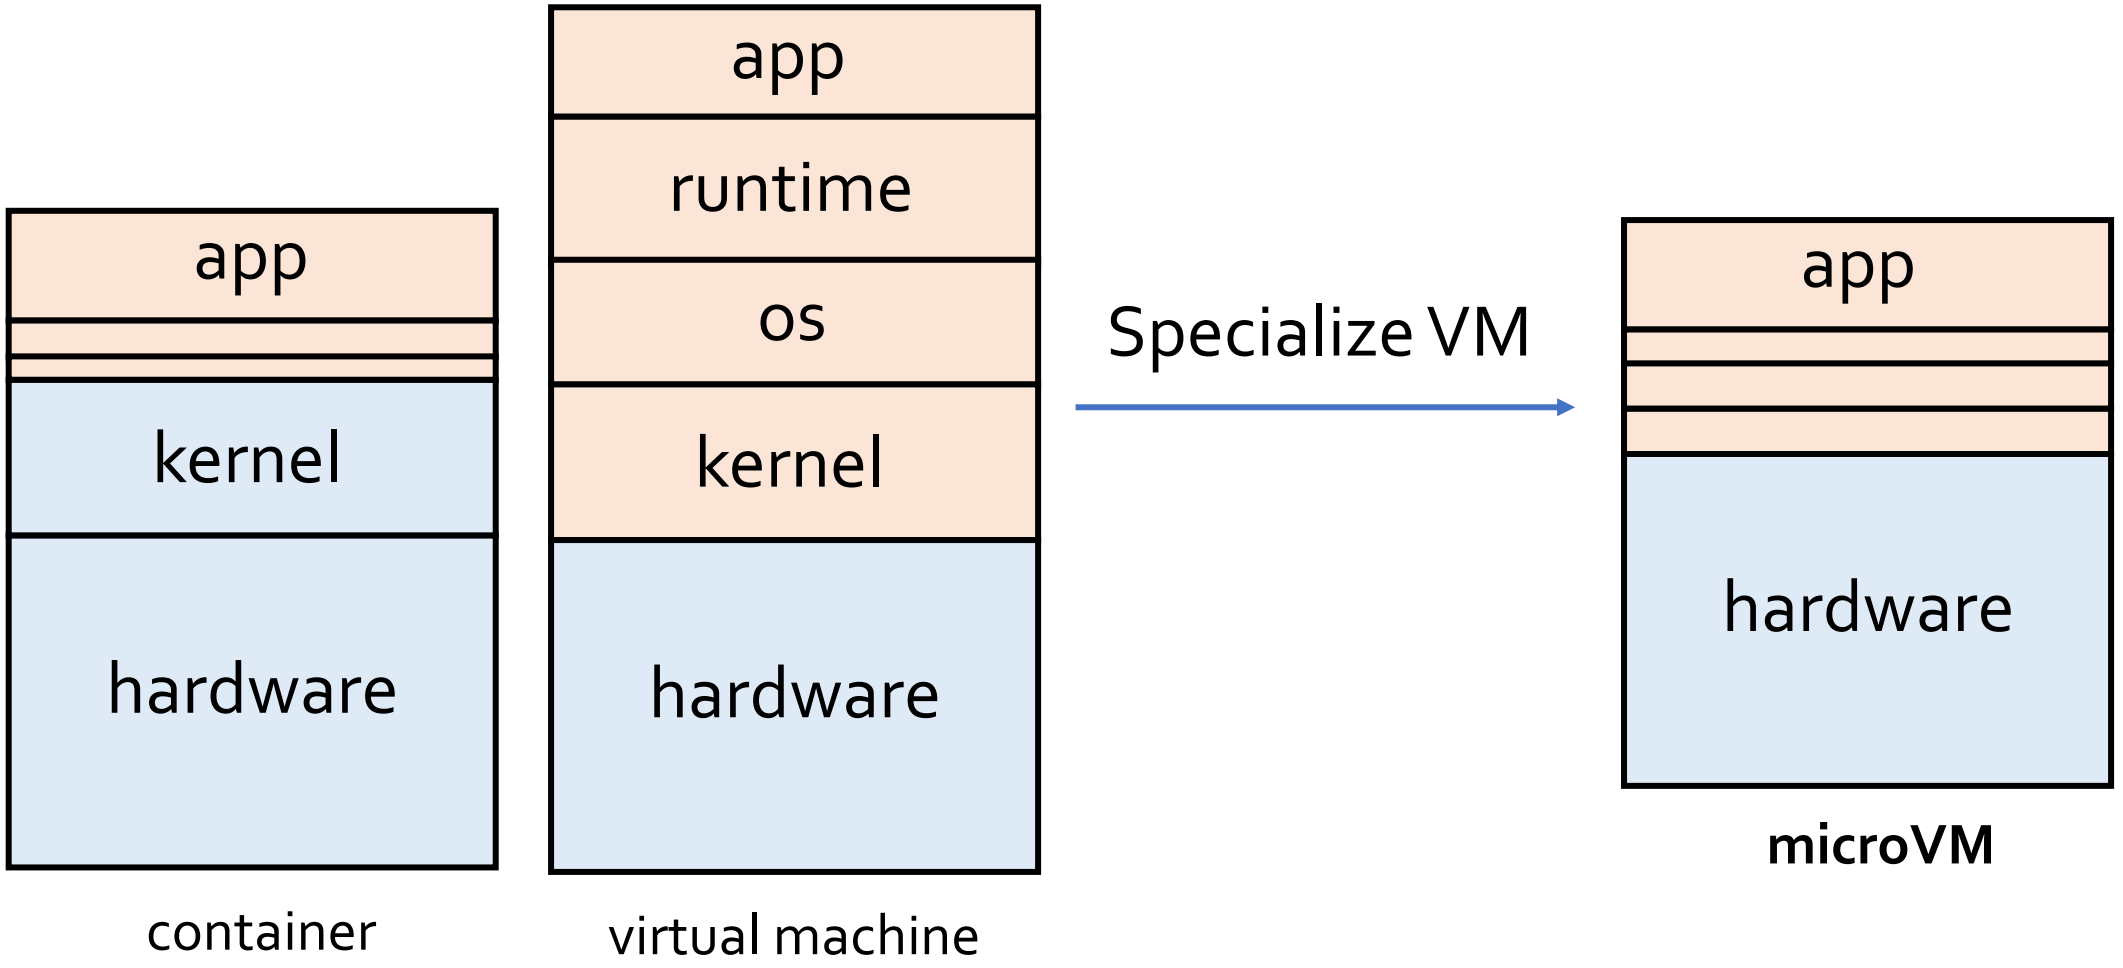

Topic 2: Improve security and control

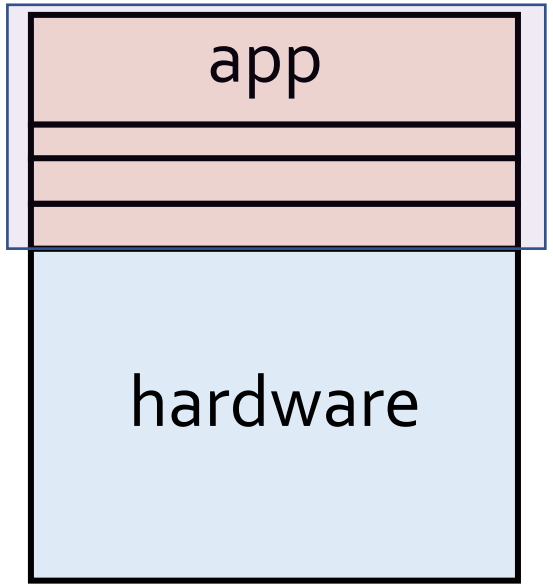

microVM

- Small image size
- Fast startup
- Better security
- Better resource control
- Harder to use/maintain
- Less portable

Conclusion

Review

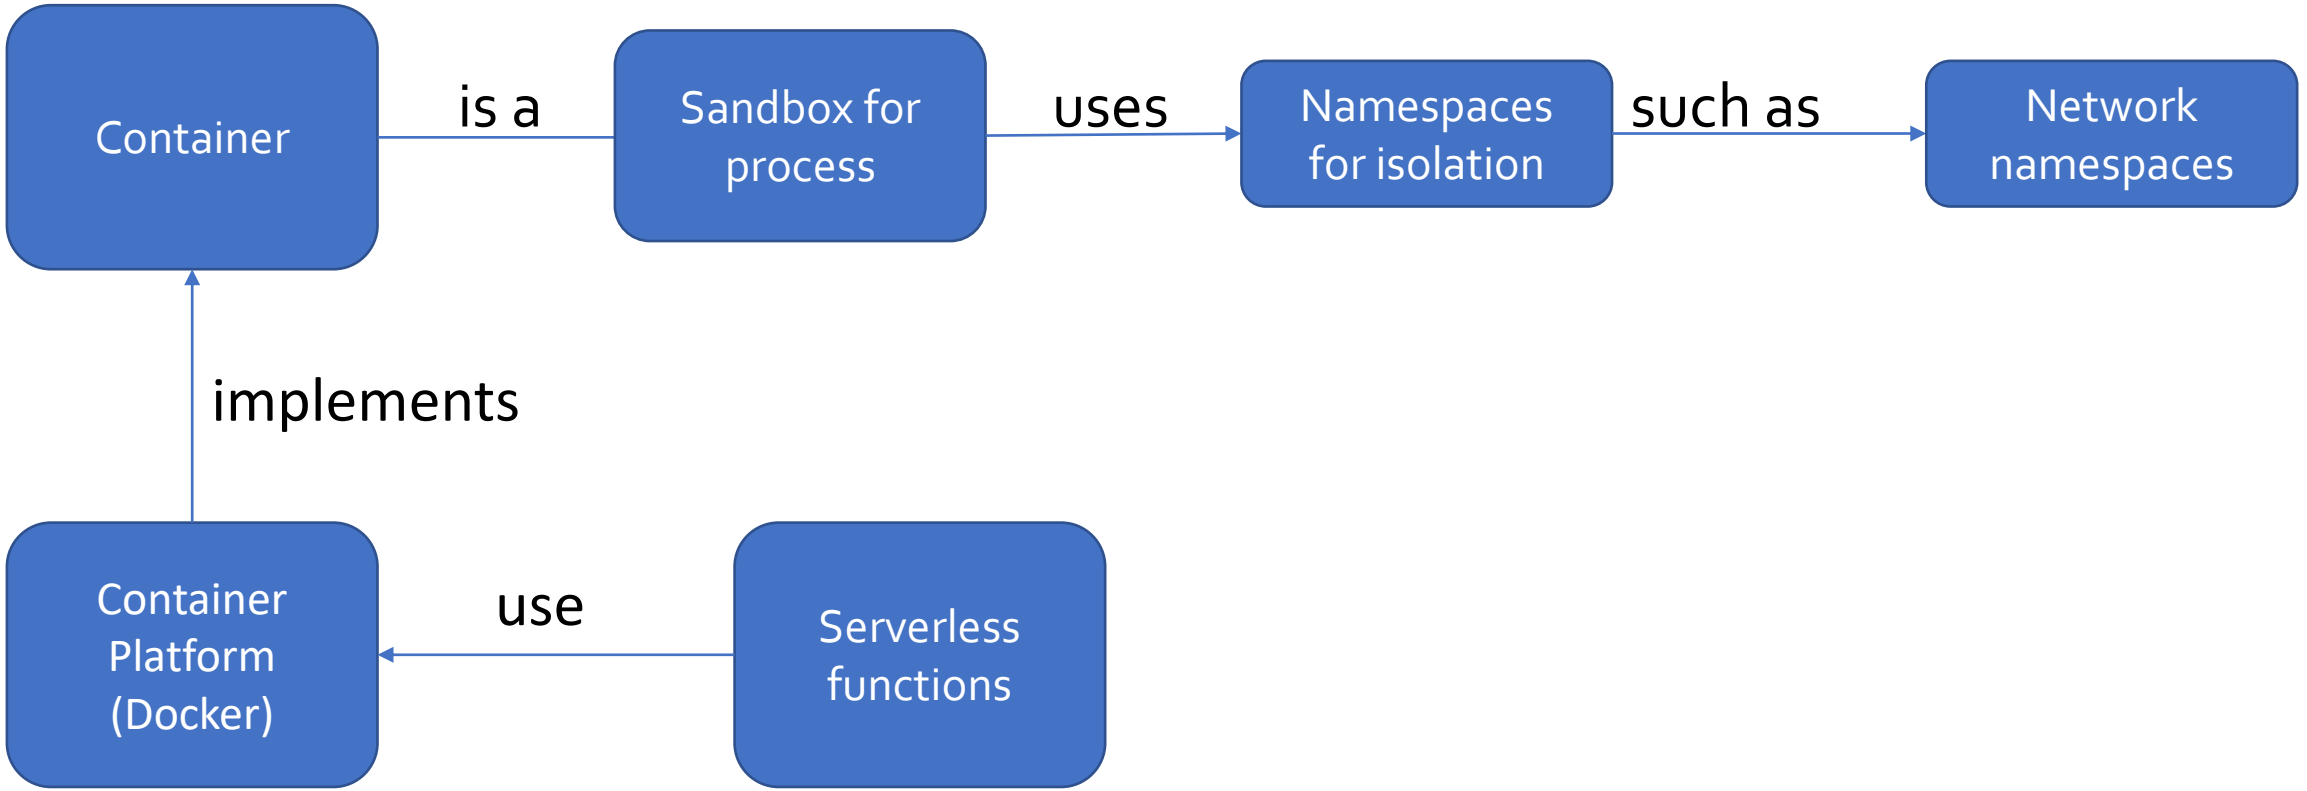
Container
is a
Sandbox for process
uses
Namespaces for isolation
such as
Network namespaces
implements
Container Platform (Docker)
use
Serverless functions

Get exposure to software using containers