

# NETWORKING FUNDAMENTALS

George Porter  
Module 1  
Fall 2020



# ATTRIBUTION

- These slides are released under an Attribution-NonCommercial-ShareAlike 3.0 Unported (CC BY-NC-SA 3.0) Creative Commons license
- These slides incorporate material from:
  - Alex C. Snoeren, UC San Diego
  - Michael Freedman and Kyle Jamieson, Princeton University
  - Internet Society
  - Computer Networking: A Top Down Approach

# Outline

1. Packet switching
2. Addressing



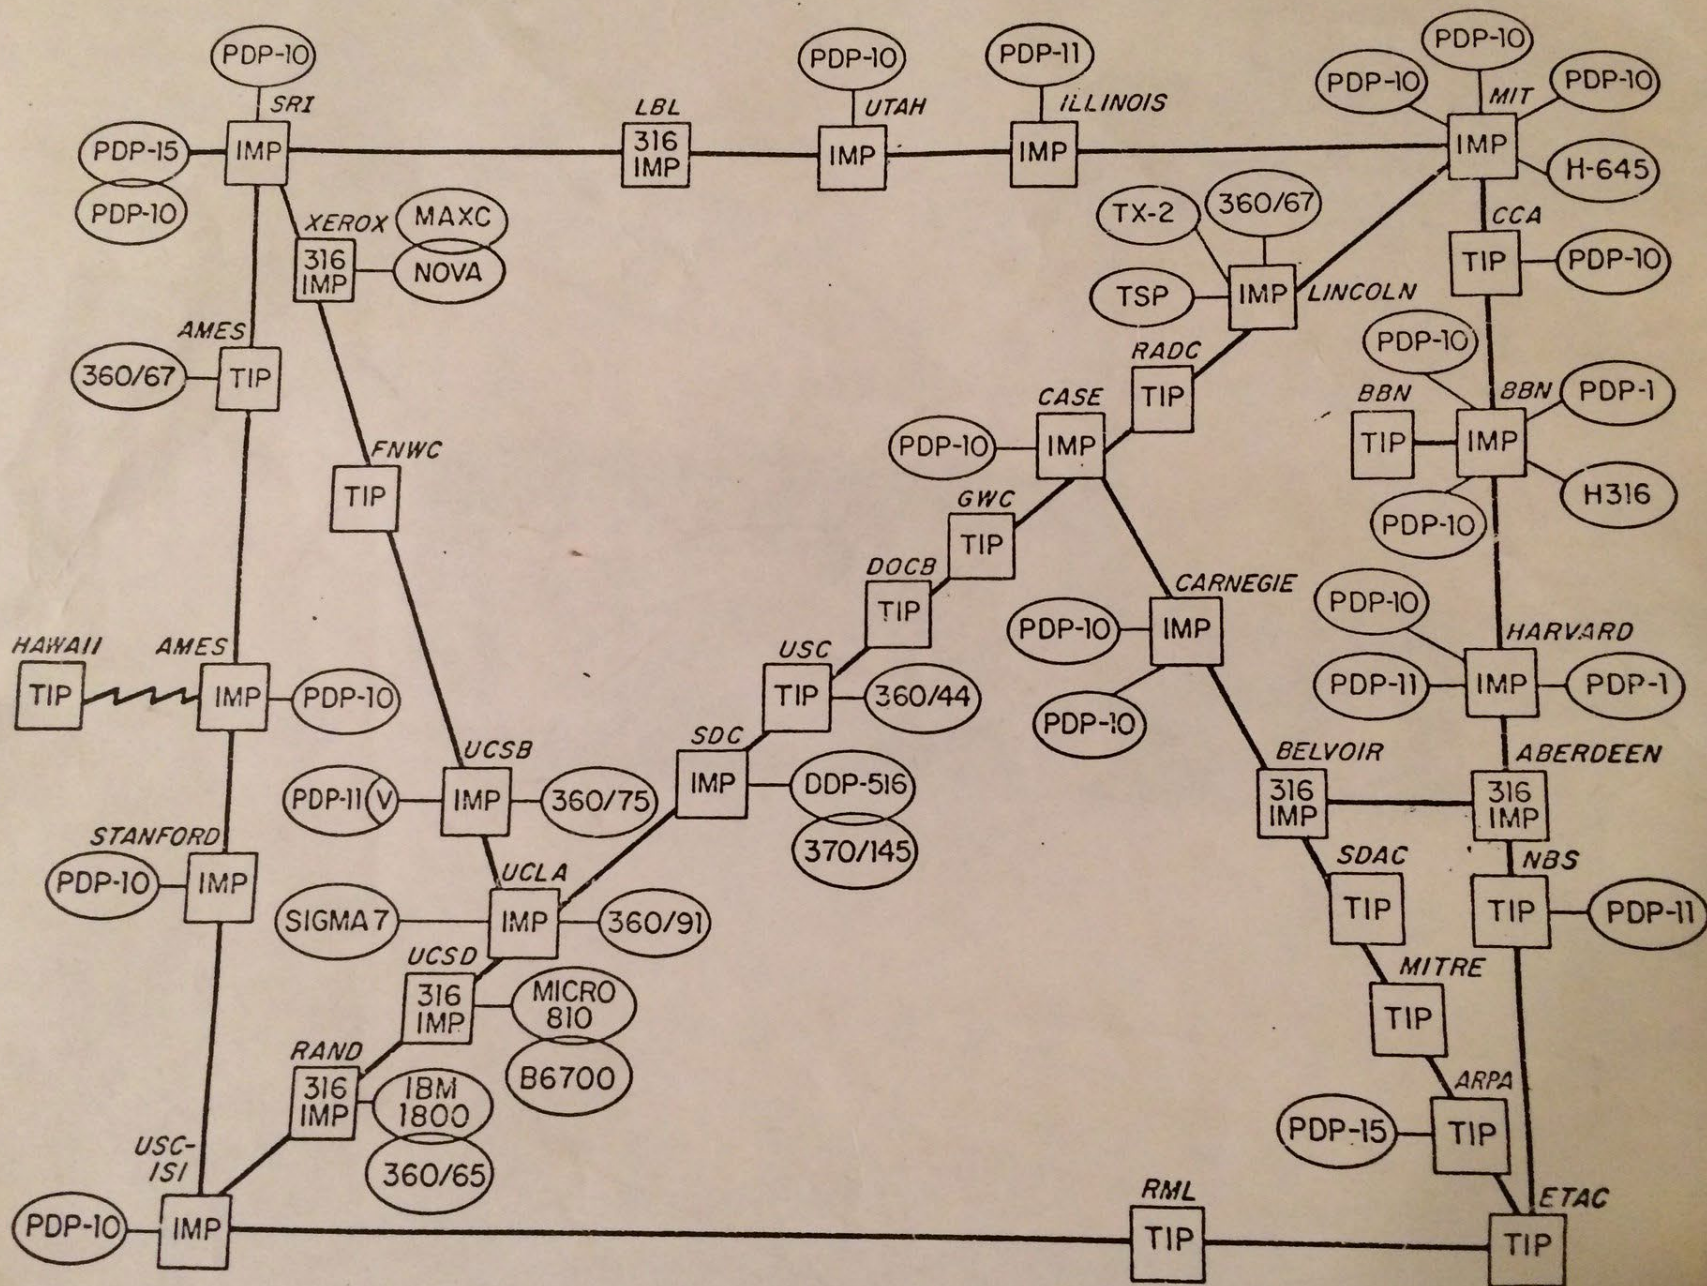
# BRIEF HISTORY OF THE INTERNET

- 1968 - DARPA (Defense Advanced Research Projects Agency) contracts with BBN (Bolt, Beranek & Newman) to create ARPAnet
- 1970 - First five nodes:
  - UCLA
  - Stanford
  - UC Santa Barbara
  - U of Utah, and
  - BBN
- 1974 - TCP specification by Vint Cerf
- 1984 – On January 1, the Internet with its 1000 hosts converts en masse to using TCP/IP for its messaging


*Data from the Internet Society*



## ARPA NETWORK, LOGICAL MAP, MAY 1973



# AN INTER-NETWORK



- billions of connected computing devices:


- *hosts = end systems*
- *running network apps*

❖ *communication links*

- fiber, copper, radio, satellite
- transmission rate: *bandwidth*




- ❖ *Packet switches:* forward packets (chunks of data)
- *routers and switches*




# PACKET SWITCHING

- Data transmitted from source to destination in discrete *packets*
- Variable size, usually a maximum transmission unit (MTU) of ~1500 bytes
- 1 GB file is approx 715,000 packets
- Each packet routed independently
- Has its own source and destination *address* used to find the route to the ultimate destination

# NETWORK TERMINOLOGY



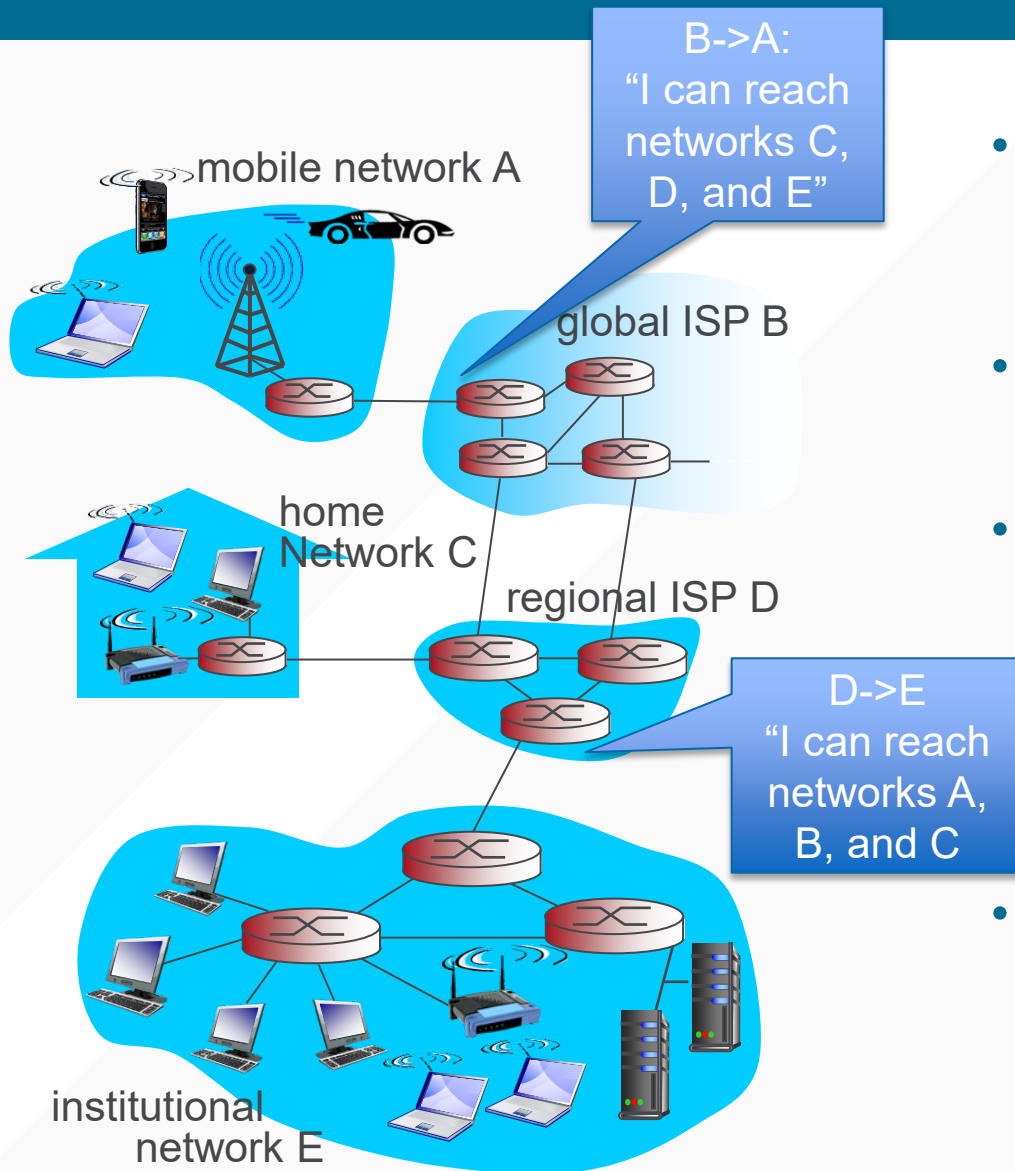
A single packet-switched network



Interconnection of networks



# ROUTING PACKETS BETWEEN NETWORKS




- Networks use *Border Gateway Protocol (BGP)* to announce reachability
- Each network talks just with its neighbors
- Goal is to get a packet to the destination network
- It is up to that destination network to get individual packets to their ultimate destination
- Back-to-back packets from the same "connection" might take different paths!
- Might arrive out of order too





# ROUTING PACKETS WITHIN A SINGLE NETWORK

- E.g. UCSD, San Diego Comcast, a Google datacenter
- Packet switches communicate with each other using an *intra-domain* protocol called a Link-state protocol
- Each packet switch sends a list of its neighbors to every other switch via a *broadcast*
- As a result, each switch has the same “map” of what the network’s *topology* looks like



# ROUTING PACKETS WITHIN A SINGLE NETWORK


- Each packet switch uses its local topology map to choose paths to each destination
- Typically using Dijkstra's shortest path algorithm
  - Take CSE 123/222a to learn more!
- Packets are forwarded along this shortest path *tree*






# RESOURCE SHARING

Multiplexing multiple logical flows over a single physical link



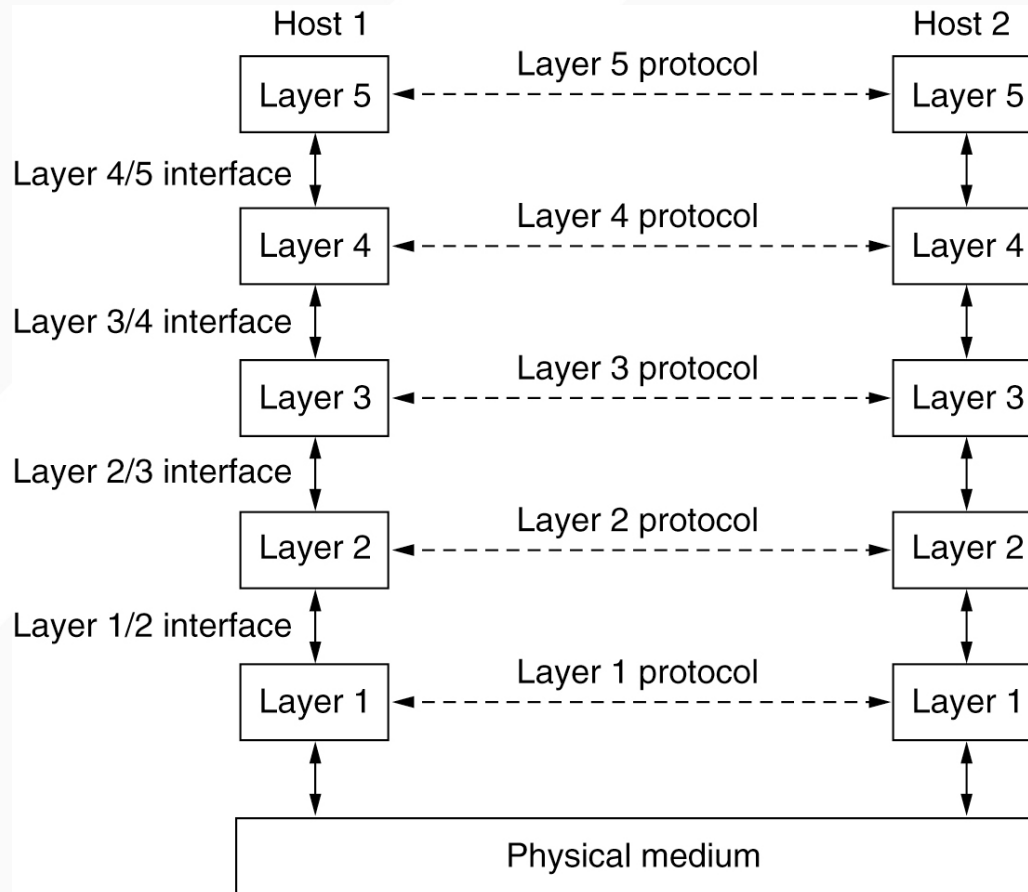
- Resource: links and nodes
- How to share a link?
  - Multiplexing
  - De-multiplexing
  - Queueing



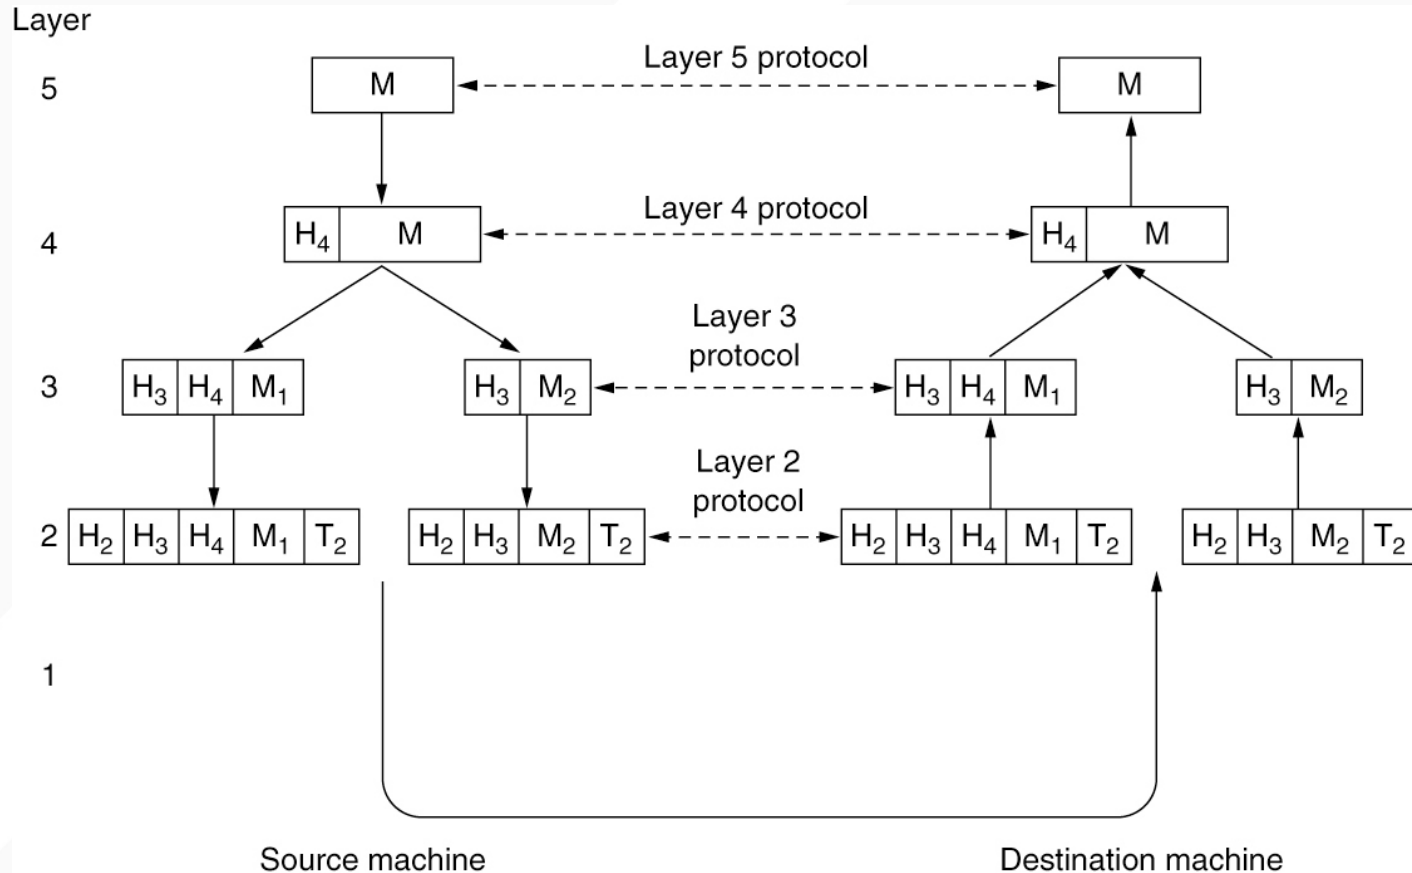
# PROTOCOL LAYERING

- Networks organized as a stack of layers or levels
  - Each layer built upon the one below it
- Communication between corresponding layers
  - Use a common protocol referred to as a “layer n protocol”
  - Below layer 1 is the physical medium through which actual communication occurs
  - Interface lies between each pair of adjacent layers
- Network architecture: a set of layers and protocols
- Protocol stack: a list of the protocols used by a certain system, one protocol per layer

# LAYERED PROTOCOL “STACK”



# INFORMATION FLOW THROUGH LAYERS





# Outline

1. Packet switching
2. Addressing



# ADDRESSING CONSIDERATIONS

- Fixed length or variable length addresses?
- Issues:
  - Flexibility
  - Processing costs
  - Header size
- Engineering choice: IP uses fixed length addresses
- 32-bits in an IPv4 address
  - Dotted decimal format a.b.c.d
  - Each represent 8 bits of address
- Hierarchical: Network part and host part
  - E.g. IP address 128.54.70.238
  - 128.54 refers to the UCSD campus network
  - 70.238 refers to the host ieng6.ucsd.edu
- Which part is network vs. host?

# CLASS-BASED ADDRESSING (NOT REALLY USED ANYMORE)

- Most significant bits determines “class” of address

Class A    

0	Network	Host
---	---------	------

**127 nets, 16M hosts**

Class B    

14		16	
1	0	Network	Host

**16K nets, 64K hosts**

Class C    

21			8	
1	1	0	Network	Host

**2M nets, 254 hosts**

- Special addresses
  - Class D (1110) for multicast, Class E (1111) experimental
  - 127.0.0.1: local host (a.k.a. the loopback address)
  - Host bits all set to 0: network address
  - Host bits all set to 1: broadcast address

# IP ADDRESS PROBLEM (1991)

- Address space depletion
  - In danger of running out of classes A and B
- Why?
  - Class C too small for most organizations (only ~250 addresses)
  - Very few class A – very careful about giving them out (who has 16M hosts anyway?)
  - Class B – greatest problem






# CIDR

- Classless Inter-Domain Routing (1993)
  - Networks described by variable-length prefix and length
  - Allows arbitrary allocation between network and host address

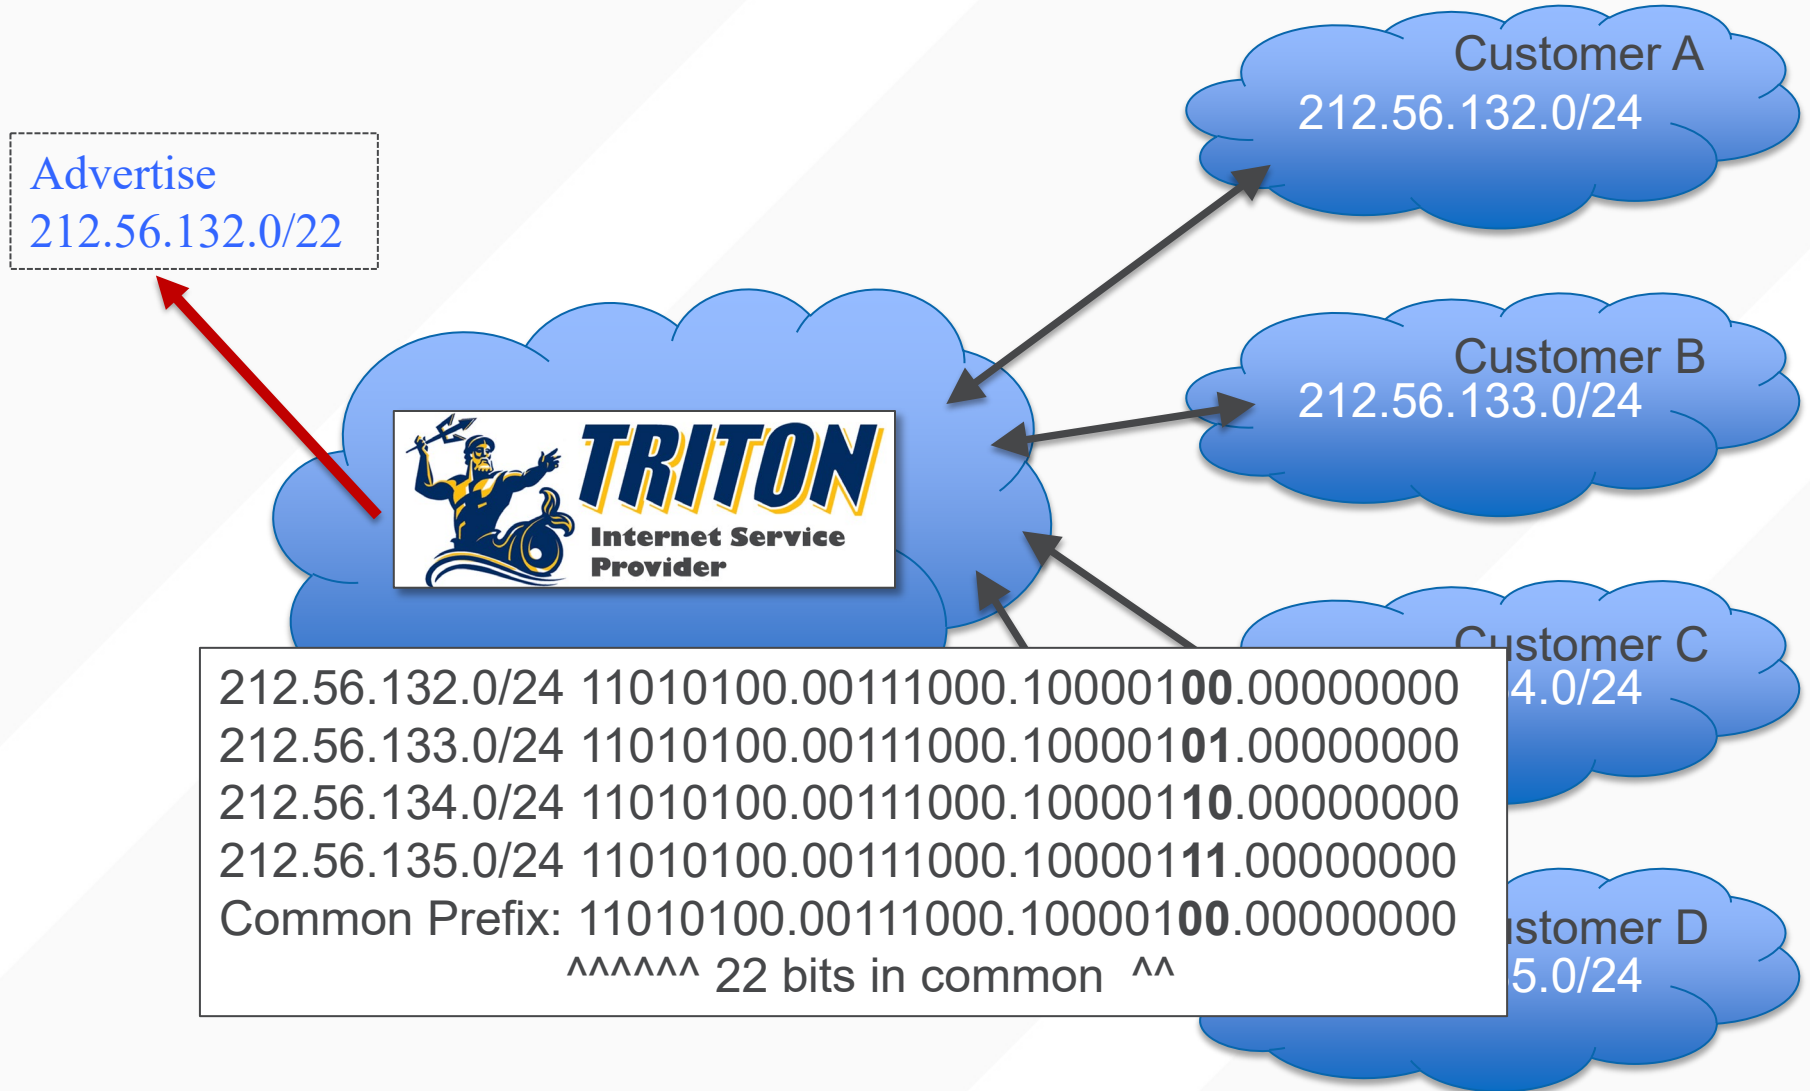


- e.g. 10.95.1.2 contained within 10.0.0.0/8:
  - 10.0.0.0 is network and remainder (95.1.2) is host
- Pro: Finer grained allocation; aggregation
- Con: More expensive lookup: **longest prefix match**

# SUBNETS AND NETMASKS



# ADDRESS AGGREGATION EXAMPLE





# DIVIDING ADDRESSES WITHIN YOUR NETWORK


University	First address	Last address	How many	Prefix
Cambridge	194.24.0.0	194.24.7.255	2048	194.24.0.0/21
Edinburgh	194.24.8.0	194.24.11.255	1024	194.24.8.0/22
(Available)	194.24.12.0	194.24.15.255	1024	194.24.12.0/22
Oxford	194.24.16.0	194.24.31.255	4096	194.24.16.0/20

- Rule: **longest prefix matching**
  - Forward packets to the *sub-network with the longest prefix*

# FORWARDING TABLE EXAMPLE (R2)

- Packet to 10.1.1.6
- Matches 10.1.0.0/23

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.16.1
10.1.8.0/24	interface1
10.1.2.0/23	interface2
<b>10.1.0.0/23</b>	<b>10.1.2.2</b>
10.1.16.0/24	interface3




# FORWARDING TABLE EXAMPLE 2 (R1)

- Packet to 10.1.1.6
- Matches 10.1.1.4/30
  - Longest prefix match

## Routing table at R1

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.2.1
10.1.0.0/24	interface1
<b>10.1.1.0/24</b>	interface2
10.1.2.0/23	interface3
<b>10.1.1.4/30</b>	10.1.1.101




# FORWARDING TABLE EXAMPLE 3 (H1)

- Packet to 10.1.1.6
- Direct route
  - Longest prefix match

## Routing table at H1

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.1.1
<b>10.1.1.0/24</b>	interface1
<b>10.1.1.4/30</b>	interface2



# KEY TAKE-AWAY

```

  _ | _ | _ )
  _ | ( _ | /
  _ | \ _ | _ |
Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
8 package(s) needed for security, out of 16 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-30-72 ~]$ ifconfig -a
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 9001
    inet 172.31.30.72 netmask 255.255.240.0 broadcast 172.31.31.255
    inet6 fe80::65:43ff:fe81:f0 prefixlen 64 scopeid 0x20<link>
    ether 02:65:43:81:00:f0 txqueuelen 1000 (Ethernet)
    RX packets 67181 bytes 5406474 (5.1 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 72005 bytes 5648089 (5.3 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[ec2-user@ip-172-31-30-72 ~]$
```

- To reach a computer, you need:
  - Its IP address
- But to configure that server, you need:
  - Its IP address
  - It's netmask
  - The “default” gateway address

# AMAZON EC2 VIRTUAL MACHINE VIEW

EC2 > Instances > i-0cb7e5abcbec1e200

## Instance summary for i-0cb7e5abcbec1e200 [Info](#)

Updated less than a minute ago



Connect

Actions ▼

### Instance ID

i-0cb7e5abcbec1e200

### Instance state

**Running**

### Instance type

m5a.xlarge

### IAM Role

—

### Public IPv4 address

52.25.99.234 | [open address](#)

### Public IPv4 DNS

ec2-52-25-99-234.us-west-2.compute.amazonaws.com | [open address](#)

### Elastic IP addresses

52.25.99.234 (Public IP)

### Subnet ID

subnet-7baa4c1f

### Private IPv4 addresses

172.31.30.72

### Private IPv4 DNS

ip-172-31-30-72.us-west-2.compute.internal

### VPC ID

vpc-7bfae91e

Details

Security

**Networking**

Storage

Monitoring

Tags

### ▼ Networking details [Info](#)

#### Public IPv4 address

52.25.99.234 | [open address](#)

#### Public IPv4 DNS

ec2-52-25-99-234.us-west-2.compute.amazonaws.com | [open address](#)

#### IPv6 addresses

—

#### Outpost ID

—

#### Private IPv4 addresses

172.31.30.72

#### Private IPv4 DNS

ip-172-31-30-72.us-west-2.compute.internal

#### Secondary private IPv4 addresses

—

#### VPC ID

vpc-7bfae91e

#### Subnet ID

subnet-7baa4c1f

#### Availability zone

us-west-2b

### ▼ Network Interfaces [Info](#)




# ASSIGNING ADDRESSES VIA DHCP

- DHCP server is responsible for providing configuration information to hosts
- There is at least one DHCP server for an administrative domain
- DHCP server maintains a pool of available addresses

# DHCP IN ACTION

- Newly booted or attached host sends DHCPDISCOVER message to a special IP address (255.255.255.255)
- DHCP relay agent unicasts the message to DHCP server and waits for the response



# DNS HOSTNAME VERSUS IP ADDRESS

- **DNS host name** (e.g. www.cs.ucsd.edu)
  - **Mnemonic** name appreciated by humans
  - **Variable length**, full alphabet of characters
  - Provides **little** (if any) information about **location**
- **IP address** (e.g. 128.112.136.35)
  - Numerical address appreciated by **routers**
  - **Fixed length**, decimal number
  - **Hierarchical** address space, related to host **location**

# MAPPING NAMES TO ADDRESSES

GETADDRINFO(3)

Linux Programmer's Manual

GETADDRINFO(3)

## NAME

getaddrinfo, freeaddrinfo, gai\_strerror – network address and service translation

## SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
```

```
int getaddrinfo(const char *node, const char *service,
                const struct addrinfo *hints,
                struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

```
const char *gai_strerror(int errcode);
```

# A FACT ABOUT NAME→ADDRESS MAPPINGS

```
struct addrinfo {  
    int             ai_flags;  
    int             ai_family;  
    int             ai_socktype;  
    int             ai_protocol;  
    socklen_t       ai_addrlen;  
    struct sockaddr *ai_addr;  
    char            *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

- Looking up a name results in a *linked list* of results. Why?
  - Support for newer “IPv6 (128-bit)” and classic “IPv4 (32-bit)”
  - A name can map to many IP addresses. E.g. google.com

