

REMOTE PROCEDURE CALLS

George Porter
Module 3
Fall 2020



Outline

1. RPC fundamentals
2. RPC Implementations
3. Handling failures in RPCs

WHY RPC?

- The typical programmer is trained to write single-threaded code that runs in **one place**
- **Goal:** Easy-to-program network communication that makes client-server communication **transparent**
 - Retains the “feel” of writing centralized code
 - Programmer needn’t think about the network

WHAT'S THE GOAL OF RPC?

- Within a single program, running in a single process, recall the well-known notion of a procedure call:
 - Caller pushes arguments onto stack,
 - jumps to address of callee function
 - Callee reads arguments from stack,
 - executes, puts return value in register,
 - returns to next instruction in caller

RPC's Goal: To make communication appear like a local procedure call: transparency for procedure calls

RPC EXAMPLE (PSEUDOCODE)

Local computing

`X = 3 * 10;`

`print(X)`

`> 30`

Remote computing

`server = connectToServer(S);`

Try:

`X = server.mult(3,10);`

`print(X)`

Except e:

`print "Error!"`

`> 30`

or

`> Error`

RPC ISSUES


- Heterogeneity
 - Client needs to **rendezvous** with the server
 - Server must **dispatch** to the required function
 - What if server is **different** type of machine?
- Failure
 - What if messages get **dropped**?
 - What if client, server, or network **fails**?
- Performance
 - Procedure call takes ≈ 10 cycles ≈ 3 ns
 - RPC in a data center takes ≈ 10 μ s ($10^3\times$ slower)
 - In the wide area, typically $10^6\times$ slower

CLOUD-HOSTED RPC



- Functions that just can't run locally
 - `tags = Facebook.MatchFacesInPhoto(photo)`
 - Print tags: ["Chuck Thacker", "Leslie Lamport"]

MICROSOFT AZURE VISION API

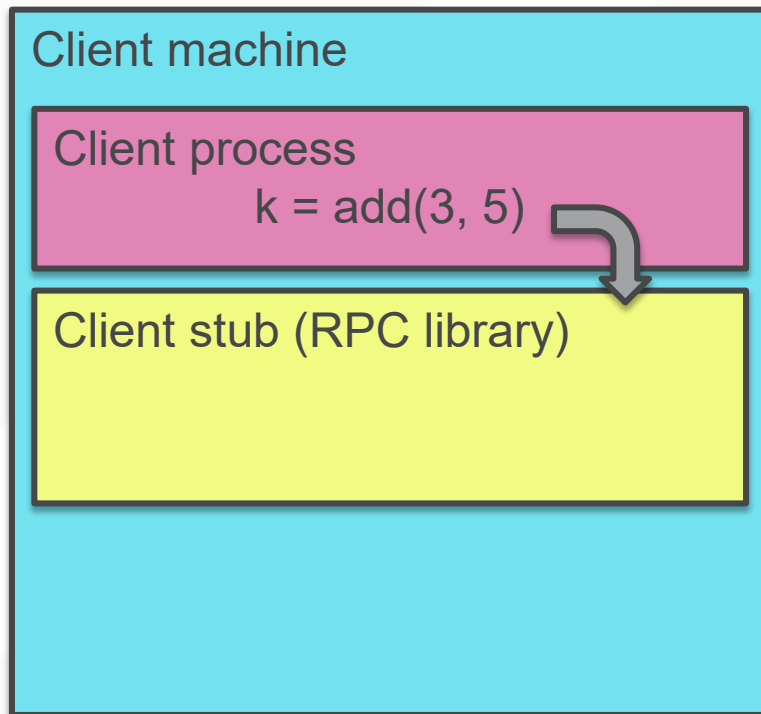


Objects [{ "rectangle": { "x": 238, "y": 299, "w": 177, "h": 117 }, "object": "Skateboard", "confidence": 0.903 }, { "rectangle": { "x": 118, "y": 63, "w": 305, "h": 321 }, "object": "person", "confidence": 0.955 }]

Tags [{ "name": "skating", "confidence": 0.999951541 }, { "name": "snowboarding", "confidence": 0.990067363 }, { "name": "sports equipment", "confidence": 0.9774853 }, { "name": "person", "confidence": 0.9605776 }, { "name": "roller skating", "confidence": 0.945730746 }, { "name": "boardsport", "confidence": 0.9242261 }, { "name": "man", "confidence": 0.9188208 }, { "name": "outdoor", "confidence": 0.9107821 }, { "name": "riding", "confidence": 0.900007248 }, { "name": "skiing", "confidence": 0.894337356 }, { "name": "footwear", "confidence": 0.8788208 }, { "name": "sport", "confidence": 0.86974 }, { "name": "skateboarder", "confidence": 0.840728462 }, { "name": "snowboard", "confidence": 0.834259868 }, { "name": "skateboarding equipment", "confidence": 0.831454 }, { "name": "individual sports", "confidence": 0.824958563 }, { "name": "skateboard", "confidence": 0.7788888 }, { "name": "skateboarder", "confidence": 0.7788888 }]

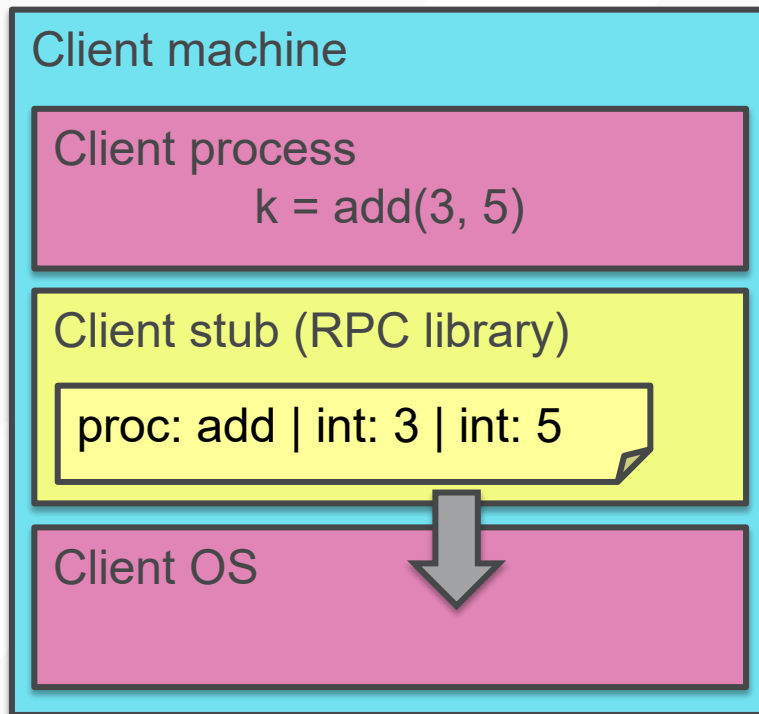
A DAY IN THE LIFE OF AN RPC

1. Client calls stub function (pushes params onto stack)




A DAY IN THE LIFE OF AN RPC

1. Client calls stub function (pushes params onto stack)
- 2. Stub marshals parameters to a network message**



A DAY IN THE LIFE OF AN RPC

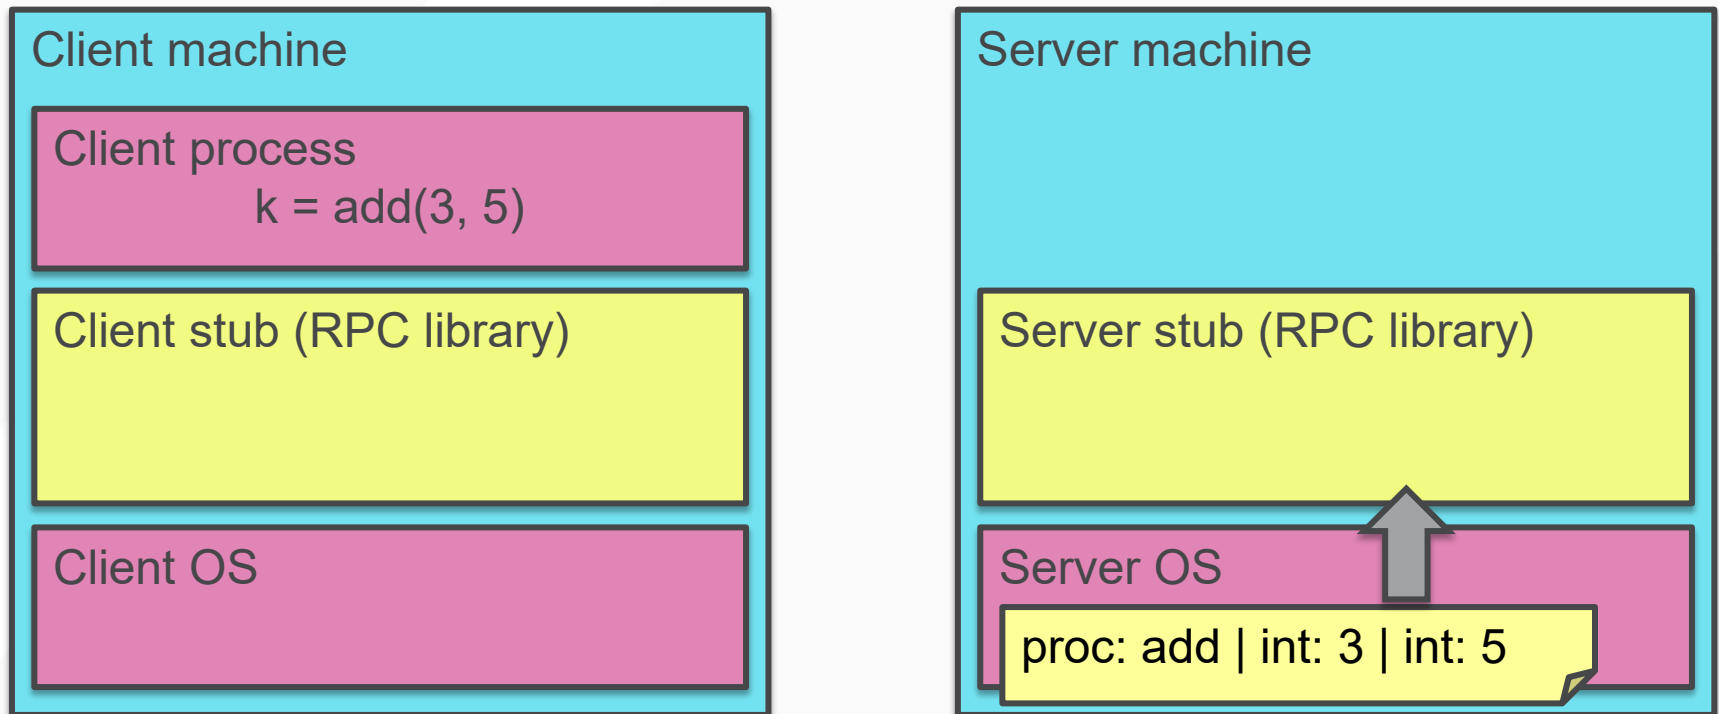
2. Stub marshals parameters to a network message
- 3. OS sends a network message to the server**



A DAY IN THE LIFE OF AN RPC

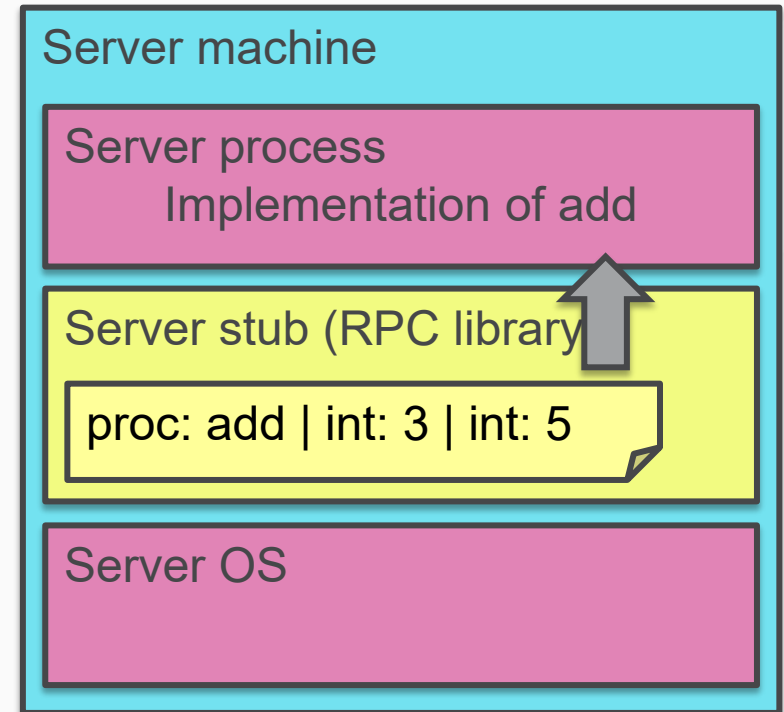
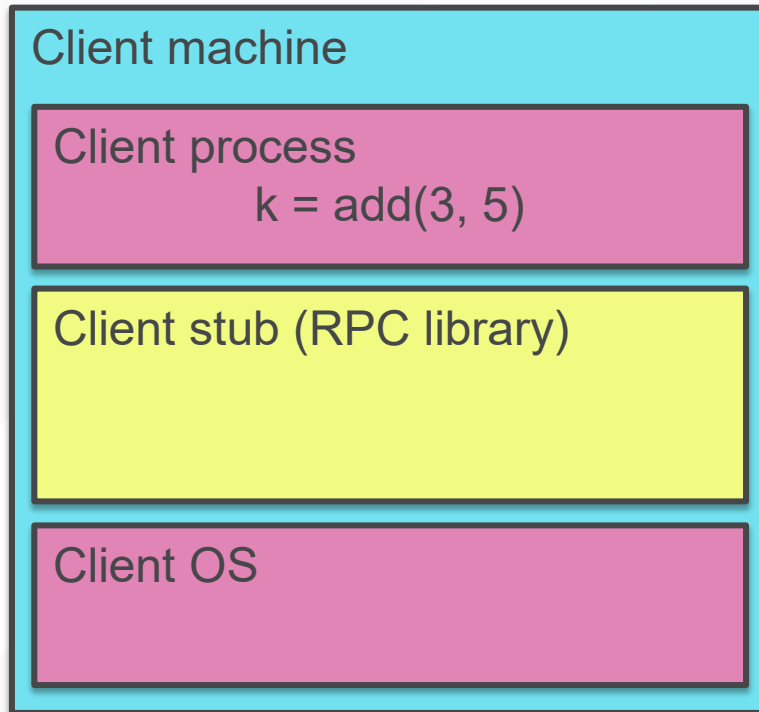
3. OS sends a network message to the server

4. **Server OS receives message, sends it up to stub**



A DAY IN THE LIFE OF AN RPC

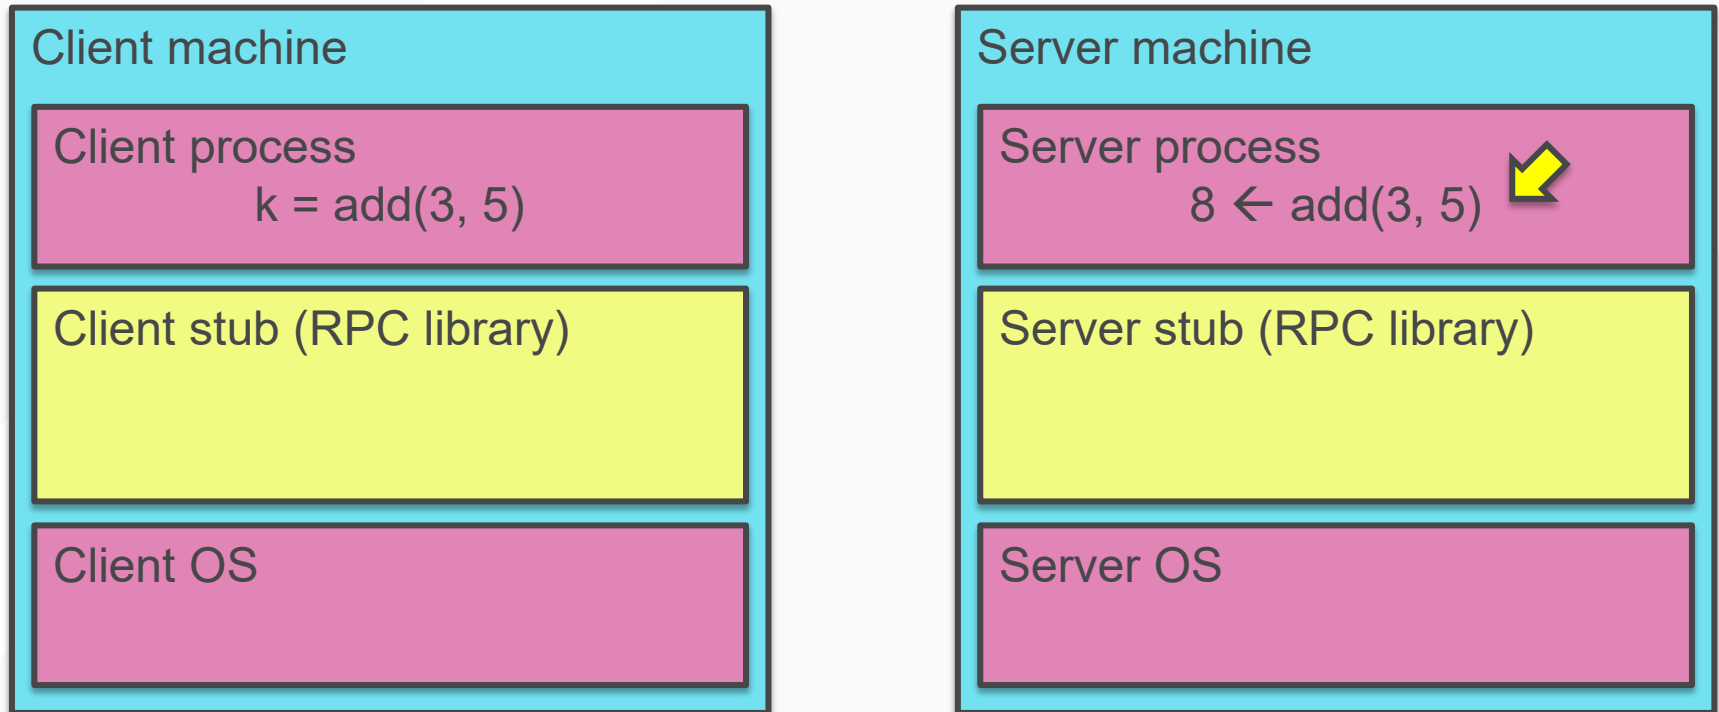
4. Server OS receives message, sends it up to stub
5. **Server stub unmarshals params, calls server function**



A DAY IN THE LIFE OF AN RPC

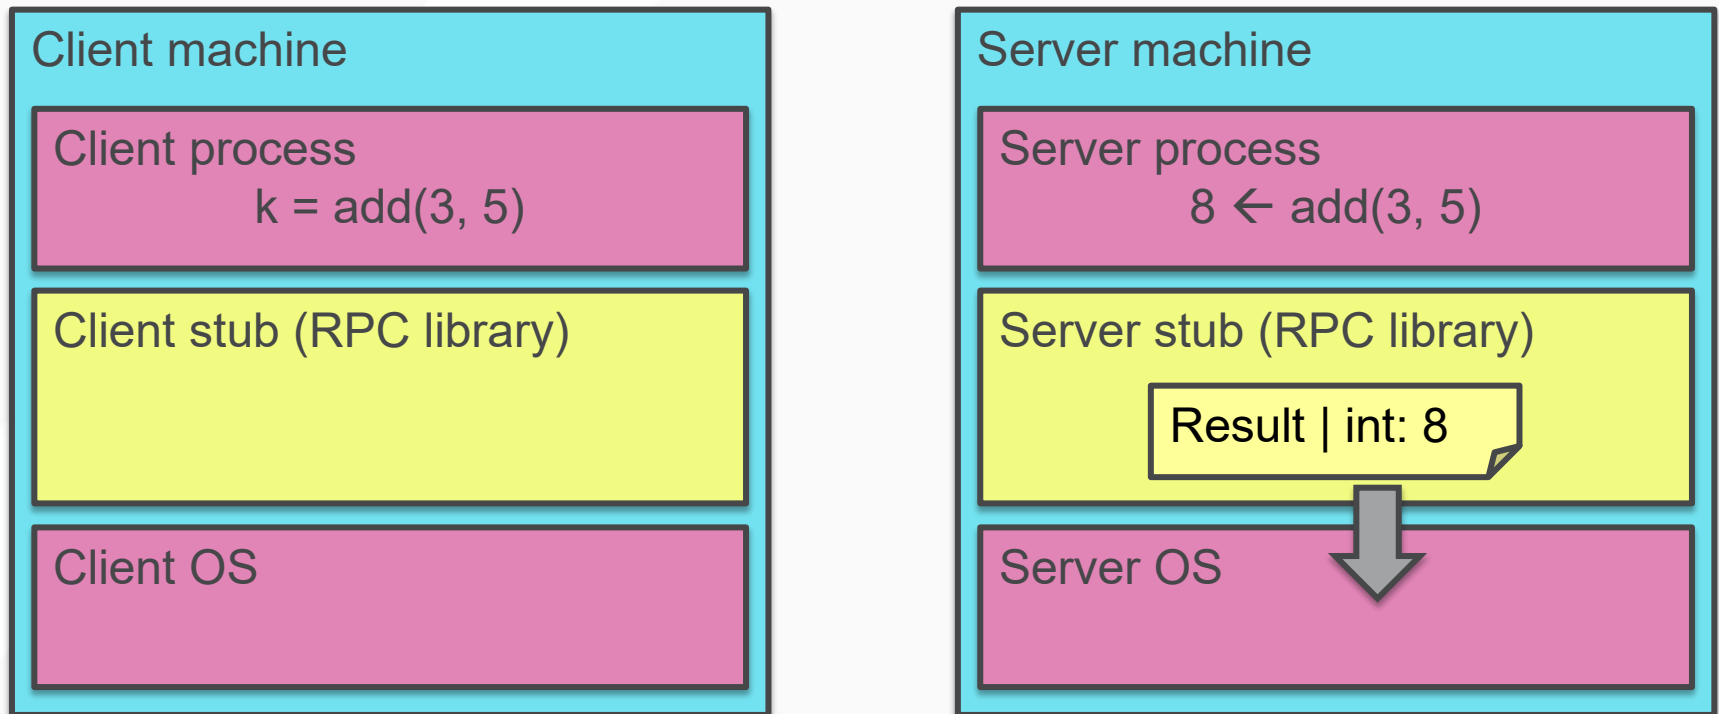
5. Server stub unmarshals params, calls server function

6. **Server function runs, returns a value**



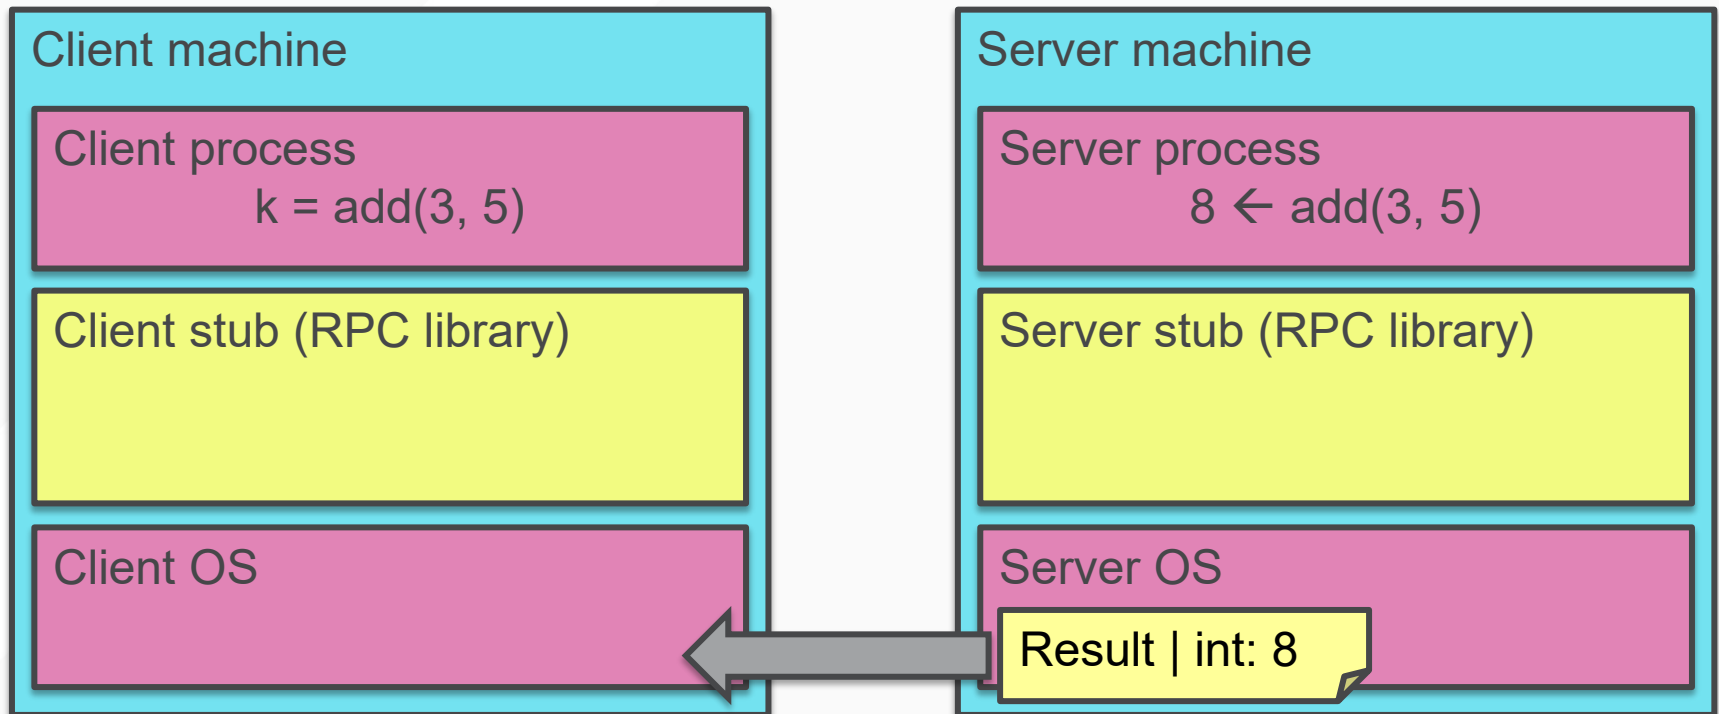
A DAY IN THE LIFE OF AN RPC

6. Server function runs, returns a value
- 7. Server stub marshals the return value, sends msg**



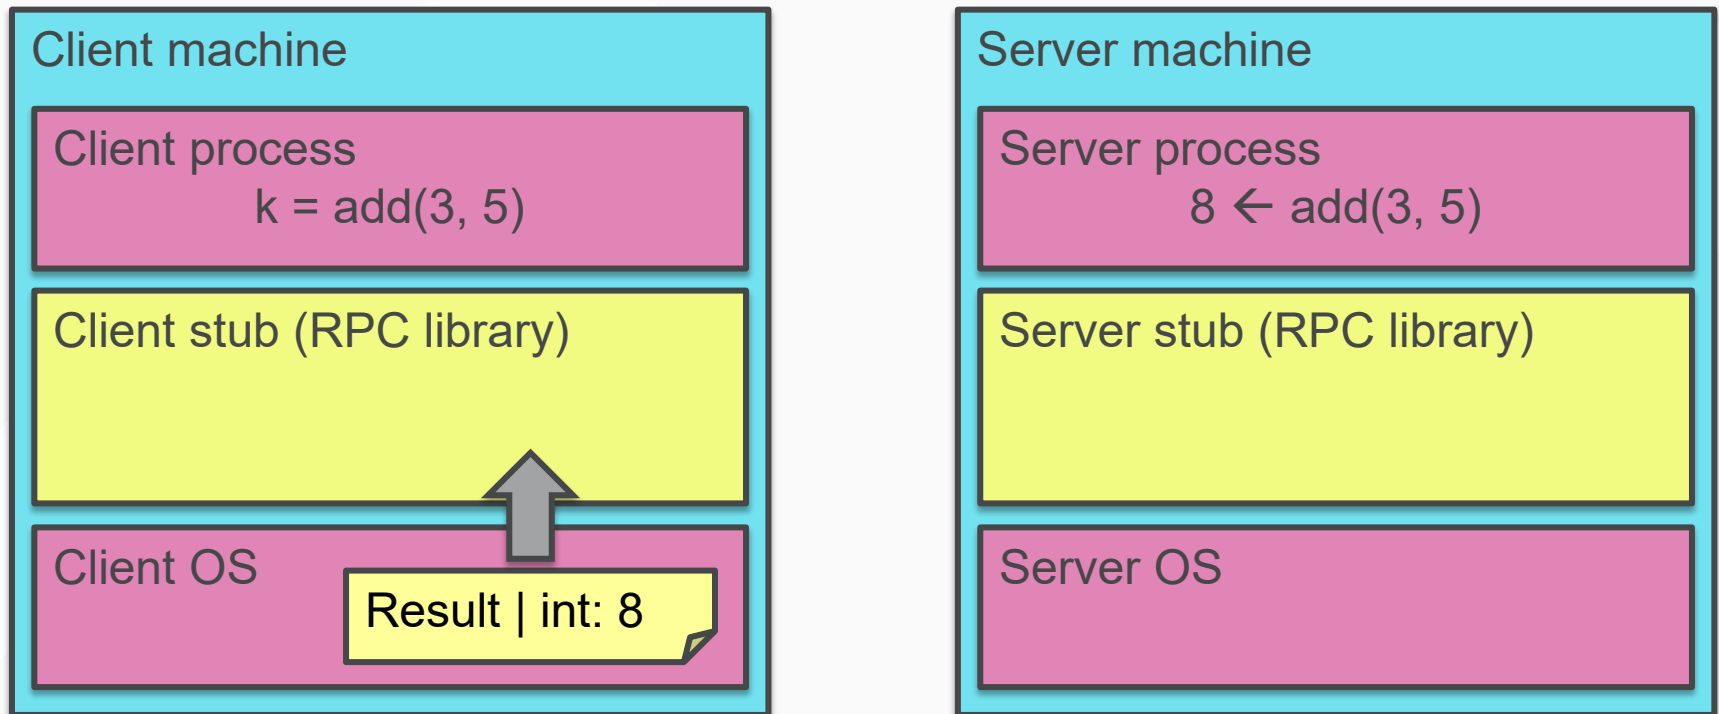
A DAY IN THE LIFE OF AN RPC

7. Server stub marshals the return value, sends msg
8. **Server OS sends the reply back across the network**



A DAY IN THE LIFE OF AN RPC

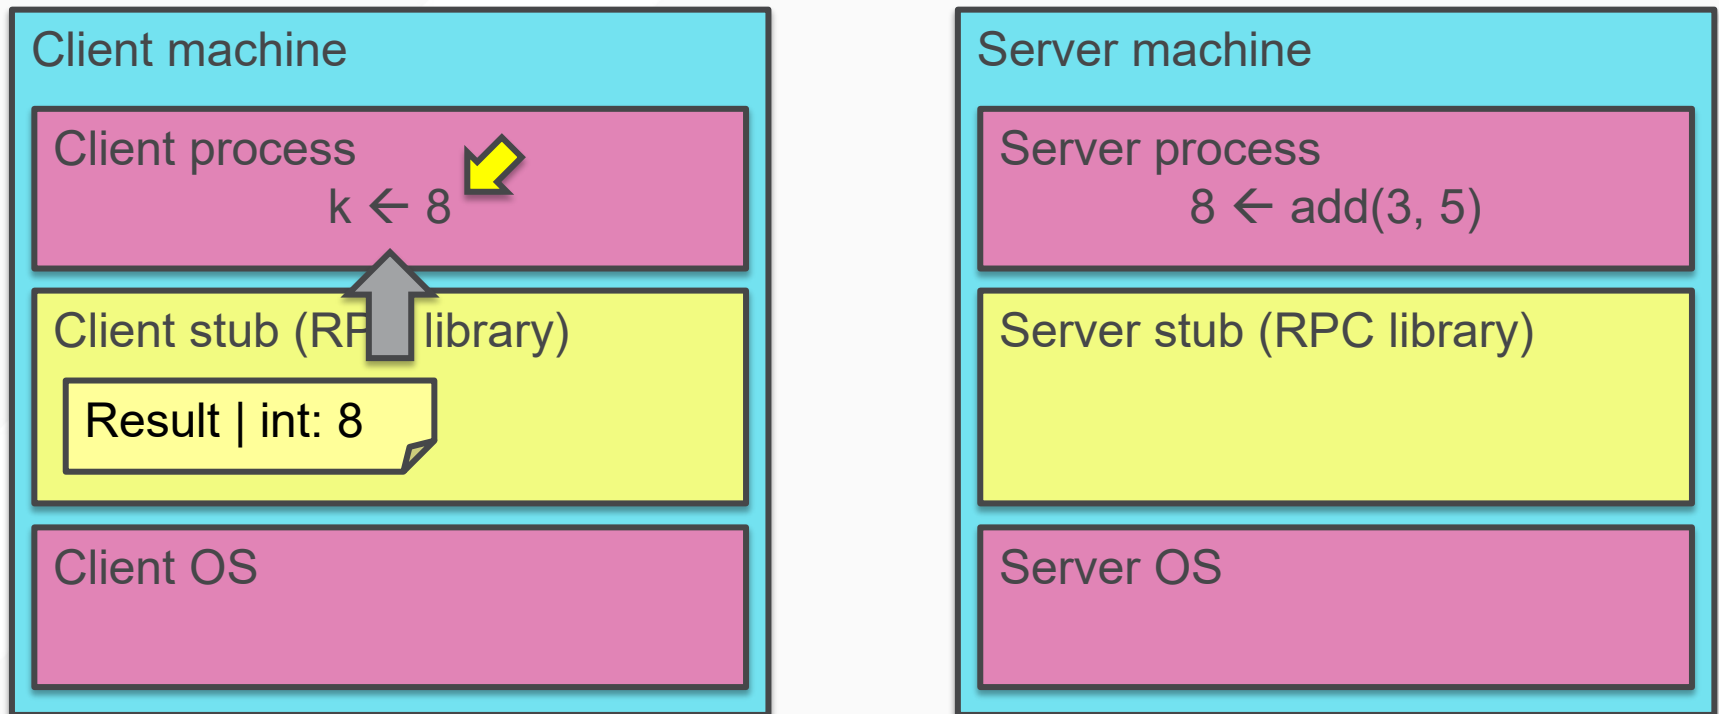
8. Server OS sends the reply back across the network
9. **Client OS receives the reply and passes up to stub**



A DAY IN THE LIFE OF AN RPC

9. Client OS receives the reply and passes up to stub

10. Client stub unmarshals return value, returns to client

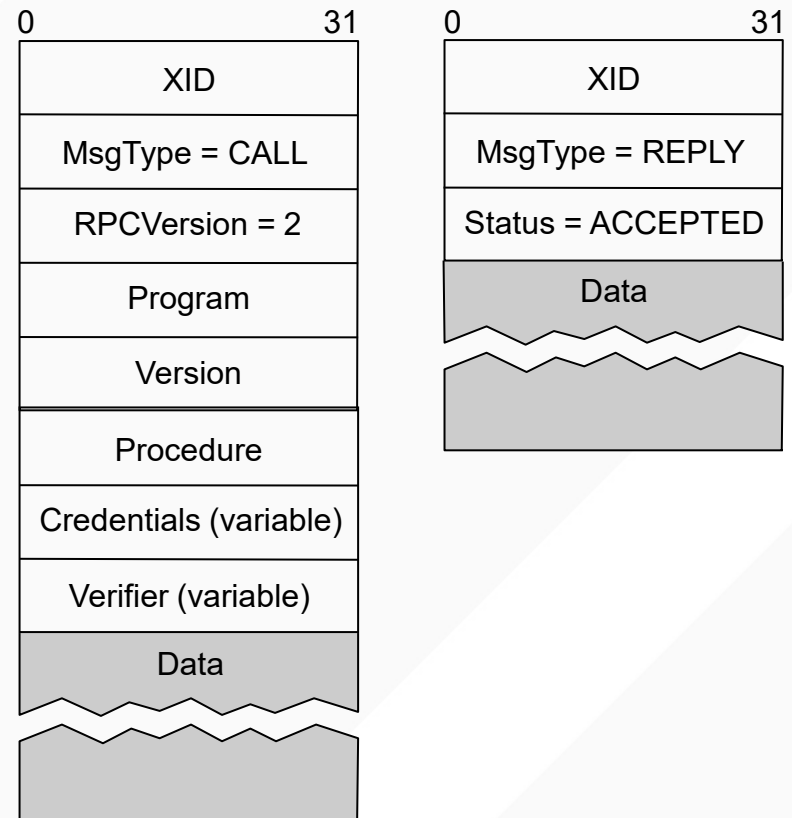




Outline

1. RPC fundamentals
2. RPC Implementations
3. Handling failures in RPCs

SUN RPC (RUNS OVER UDP)



XML-RPC

- XML is a standard for describing structured documents
- Uses tags to define structure: **<tag> ... </tag>** demarcates an element
 - Tags have no predefined semantics ...
 - ... except when document refers to a specific namespace
- Elements can have attributes, which are encoded as name-value pairs
 - A well-formed XML document corresponds to an element tree
- **<?xml version="1.0"?>**
<methodCall>
 <methodName>SumAndDifference</methodName>
 <params>
 <param><value><i4>40</i4></value></param>
 <param><value><i4>10</i4></value></param>
 </params>
</methodCall>

XML-RPC WIRE FORMAT

- Scalar values
 - Represented by a `<value><type> ... </type></value>` block
- Integer
 - `<i4>12</i4>`
- Boolean
 - `<boolean>0</boolean>`
- String
 - `<string>Hello world</string>`
- Double
 - `<double>11.4368</double>`
- Also Base64 (binary), DateTime, etc.

WIRE FORMAT (STRUCT)

- Structures
 - Represented as a set of **<member>**s
 - Each member contains a **<name>** and a **<value>**
- **<struct>**
 - <member>**
 - <name>lowerBound</name>**
 - <value><i4>18</i4></value>**
 - </member>**
 - <member>**
 - <name>upperBound</name>**
 - <value><i4>139</i4></value>**
 - </member>**
- </struct>**

WIRE FORMAT (ARRAYS)

- Arrays
 - A single **<data>** element, which
 - contains any number of **<value>** elements
- **<array>**
 - <data>**
 - <value><i4>12</i4></value>**
 - <value><string>Egypt</string></value>**
 - <value><boolean>0</boolean></value>**
 - <value><i4>-31</i4></value>**
 - </data>**
 - </array>**

XML-RPC REQUEST

- HTTP *POST* message
 - URI interpreted in an implementation-specific fashion
 - Method name passed to the server program

```
POST /RPC2 HTTP/1.1
Content-Type: text/xml
User-Agent: XML-RPC.NET
Content-Length: 278
Expect: 100-continue
Connection: Keep-Alive
Host: localhost:8080
```

```
<?xml version="1.0"?>
<methodCall>
<methodName>SumAndDifference</methodName>
<params>
<param><value><i4>40</i4></value></param>
<param><value><i4>10</i4></value></param>
</params>
</methodCall>
```

XML-RPC RESPONSE

- HTTP Response
 - Lower-level error returned as an HTTP error code
 - Application-level errors returned as a **<fault>** element (next slide)

HTTP/1.1 200 OK
Date: Mon, 22 Sep 2003 21:52:34 GMT
Server: Microsoft-IIS/6.0
Content-Type: text/xml
Content-Length: 467

```
<?xml version="1.0"?>
<methodResponse>
<params><param>
<value><struct>
<member><name>sum</name><value><i4>50</i4></value></member>
<member><name>diff</name><value><i4>30</i4></value></member>
</struct></value>
</param></params>
</methodResponse>
```

XML-RPC FAULT HANDLING

- Another kind of a MethodResponse

```
<?xml version="1.0"?>
<methodResponse>
<fault>
<value><struct>
<member>
<name>faultCode</name>
<value><i4>500</i4></value>
</member>
<member>
<name>faultString</name>
<value><string>Arg `a' out of
range</string></value>
</member>
</struct></value>
</fault>
</methodResponse>
```



Outline

1. RPC fundamentals
2. RPC Implementations
3. Handling failures in RPCs

HANDLING FAILURES IN REMOTE PROCEDURE CALLS

when ur trying your best but nothing is going right

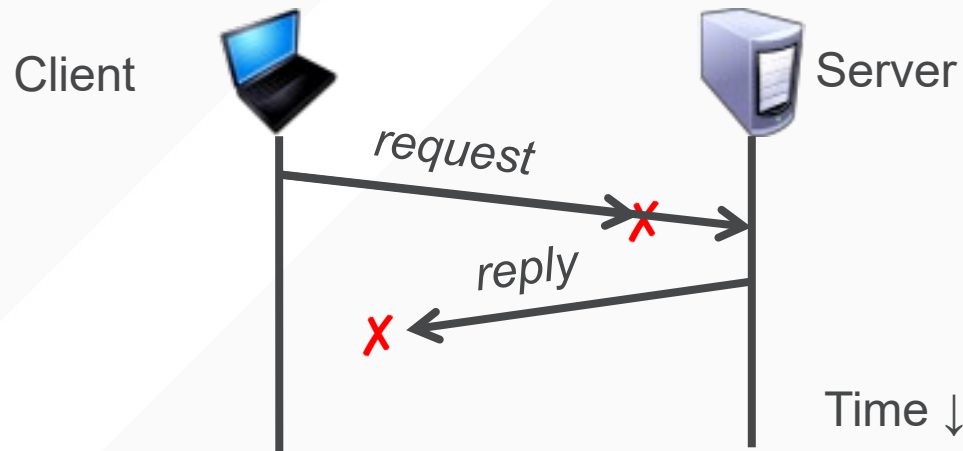


WHAT COULD POSSIBLY GO WRONG?

1. Client may **crash and reboot**
2. Packets may be **dropped**
 - Some individual **packet loss** in the Internet
 - **Broken routing** results in many lost packets
3. Server may **crash and reboot**
4. Network or server might just be **very slow**

All these may **look the same** to the client...

FAILURES, FROM CLIENT'S PERSPECTIVE



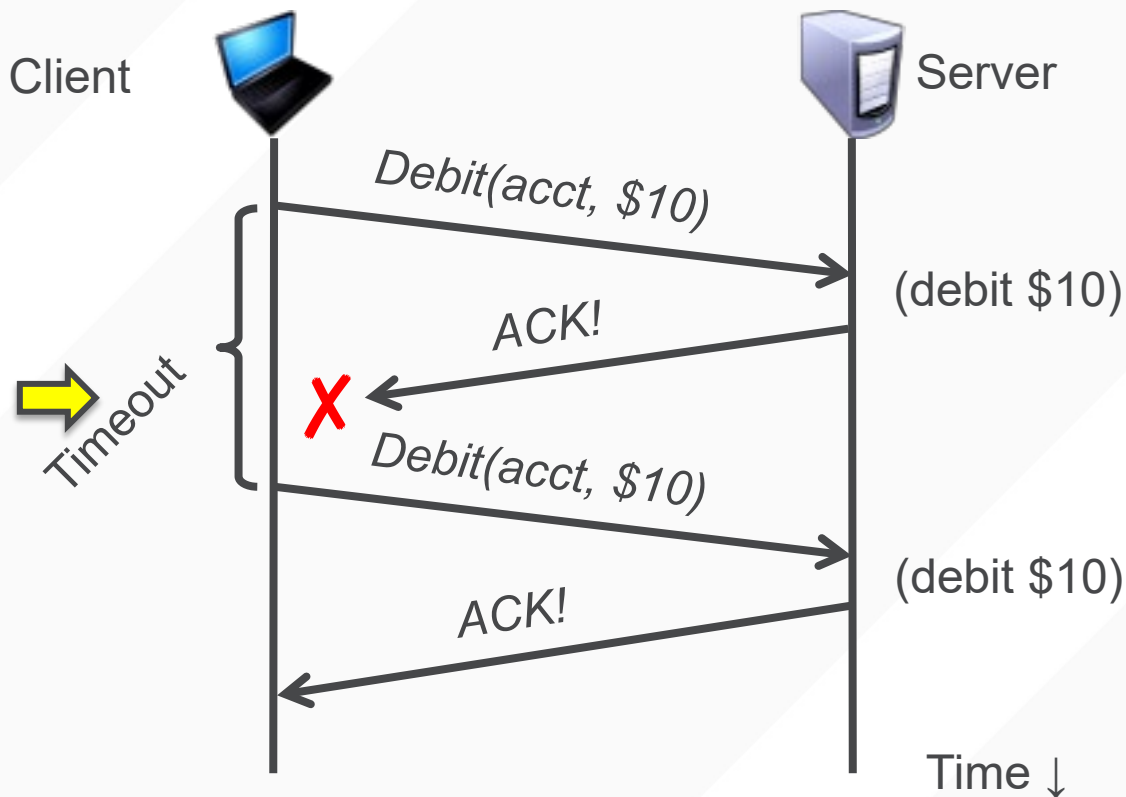
The cause of the failure is **hidden** from the **client**!

AT-LEAST-ONCE SCHEME

- **Simplest** scheme for handling failures
 1. Client stub **waits for a response**, for a while
 - Response takes the form of an **acknowledgement** message from the server stub
 2. If no response arrives after a fixed **timeout** time period, then client stub **re-sends the request**
- Repeat the above a few times
 - *Still no response?* Return an error to the application

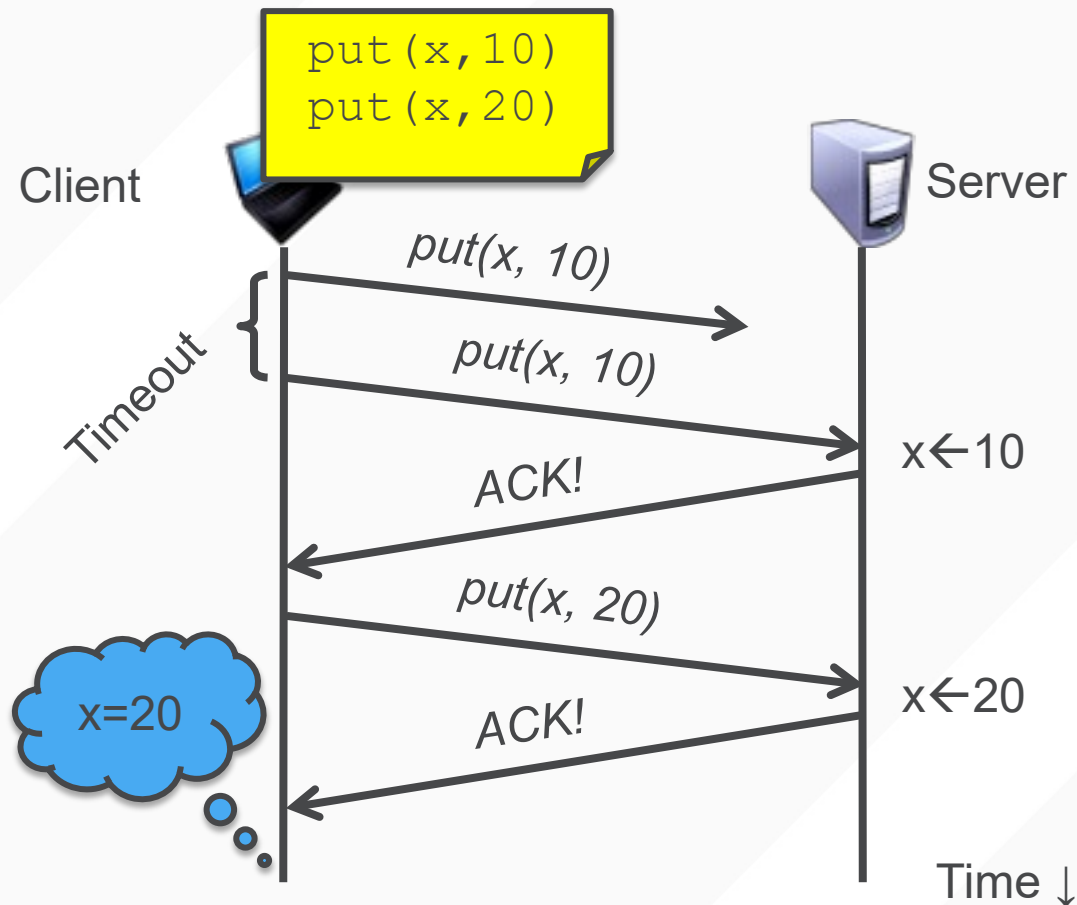
AT-LEAST-ONCE AND SIDE EFFECTS

- Client sends a “debit \$10 from bank account” RPC



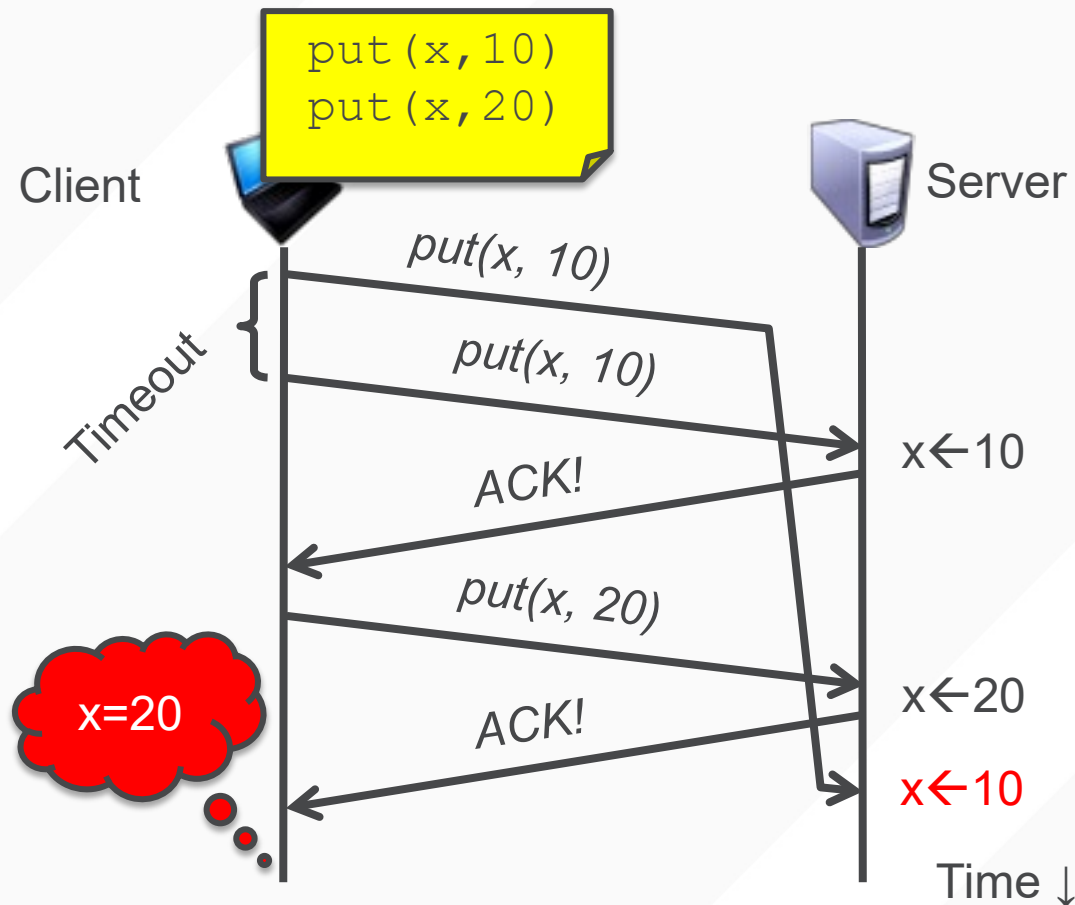
AT-LEAST-ONCE AND WRITES

- `put(x, value)`, then `get(x)`: expect answer to be *value*



AT-LEAST-ONCE AND WRITES

- Consider a client storing **key-value pairs** in a **database**
 - $\text{put}(x, \text{value})$, then $\text{get}(x)$: expect answer to be *value*



SO IS AT-LEAST-ONCE *EVER* OKAY?

- **Yes:** If they are read-only operations with no side effects
 - *e.g.*, read a key's value in a database
- **Yes:** If the application has its own functionality to cope with duplication and reordering
 - You will implement this in a later project

AT-MOST-ONCE SCHEME

- **Idea:** server RPC code detects duplicate requests
 - Returns previous reply **instead of re-running handler**
- *How to detect a duplicate request?*
 - **Test:** Server sees same function, same arguments twice
 - **No!** Sometimes applications **legitimately** submit the same function with same arguments, twice in a row

AT-MOST-ONCE SCHEME

- *How to detect a duplicate request?*
- Client includes unique **transaction ID (xid)** with each one of its RPC requests
- Client uses **same xid** for retransmitted requests

```
At-Most-Once Server  
if seen[xid]:  
    retval = old[xid]  
else:  
    retval = handler()  
    old[xid] = retval  
    seen[xid] = true  
return retval
```

AT MOST ONCE: ENSURING UNIQUE XIDS

- *How to ensure that the xid is unique?*
 1. Combine a unique client ID (e.g., IP address) with the current time of day
 2. Combine unique client ID with a sequence number
 - Suppose the client crashes and restarts. *Can it reuse the same client ID?*
 3. Big random number

AT-MOST-ONCE: DISCARDING SERVER STATE

- **Problem:** `seen` and `old` arrays will **grow without bound**
- **Observation:** By construction, when the client gets a response to a particular `xid`, it will **never re-send it**
- Client could **tell** server “I’m done with `xid x` – delete it”
 - Have to tell the server about **each and every** retired `xid`
 - Could **piggyback** on subsequent requests

Significant overhead if many RPCs are in flight, in parallel

AT-MOST-ONCE: DISCARDING SERVER STATE

- **Problem:** `seen` and `old` arrays will **grow without bound**
- Suppose `xid` = $\langle \text{unique client id, sequence no.} \rangle$
 - *e.g.* $\langle 42, 1000 \rangle, \langle 42, 1001 \rangle, \langle 42, 1002 \rangle$
- Client includes “seen all replies $\leq X$ ” with every RPC
 - Much like TCP sequence numbers, acks
- *How does the client **know** that the server received the information about retired RPCs?*
 - Each one of these is cumulative: later seen messages subsume earlier ones

AT-MOST-ONCE: CONCURRENT REQUESTS

- **Problem:** How to handle a duplicate request while the original is still executing?
 - Server doesn't know reply yet. Also, we don't want to run the procedure twice
- **Idea:** Add a **pending** flag per executing RPC
 - Server waits for the procedure to finish, or ignores

AT MOST ONCE: SERVER CRASH AND RESTART

- **Problem:** Server may crash and restart
- *Does server need to write its tables to disk?*
- Yes! On **server crash and restart**:
 - If `old[]`, `seen[]` tables are only in memory:
 - Server will forget, **accept duplicate requests**

RPC SEMANTICS

Delivery Guarantees			RPC Call Semantics
Retry Request	Duplicate Filtering	Retransmit Response	
No	NA	NA	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least once</i>
Yes	Yes	Retransmit reply	<i>At-most once</i>

UC San Diego