

# QUORUMS AND FAULT TOLERANCE

Module 4  
Fall 2020

George Porter





# ATTRIBUTION

- These slides are released under an Attribution-NonCommercial-ShareAlike 3.0 Unported (CC BY-NC-SA 3.0) Creative Commons license
- These slides incorporate material from:
  - Tanenbaum and Van Steen, Dist. Systems: Principles and Paradigms
  - Kyle Jamieson, Princeton University (also under a CC BY-NC-SA 3.0 Creative Commons license)

# SURFSTORE METADATA SERVER PROBLEM

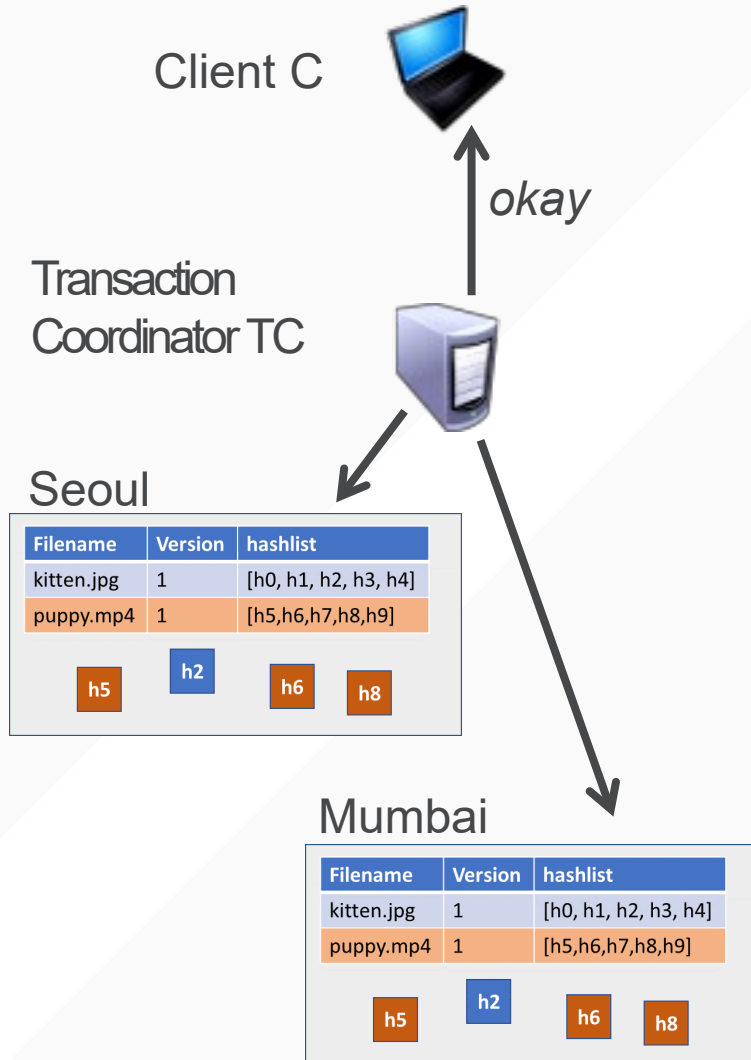
```
UpdateFile(  
  file="kitten.jpg",  
  ver=2,  
  hashlist = {h3,h4}  
);
```



*All data is lost!*

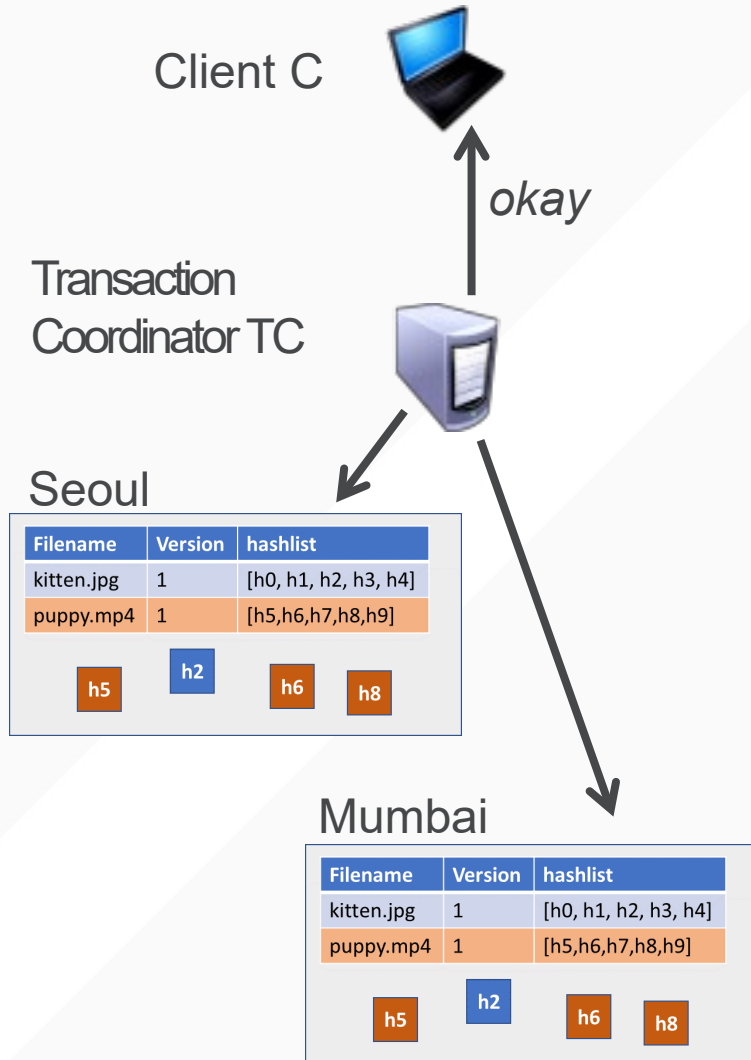
Surfstore  
Client

# IDEA 1: ADAPT TWO-PHASE COMMIT TO SAVE DATA



1.  $C \rightarrow TC$ : “go!”
2.  $TC \rightarrow \text{Seoul (S), Mumbai (M)}$ : “prepare!”
3.  $S, M \rightarrow P$ : “yes” or “wrong\_version”
4.  $TC \rightarrow S, M$ : “commit!” or “abort!”
  - TC sends **commit** if **both** say yes
  - TC sends **abort** if **either** say no
5.  $TC \rightarrow C$ : “okay” or “failed”
  - **S, M** commit on receipt of commit message

## IDEA 2: ASSUME TC DOESN'T FAIL (FOR NOW)



1.  $C \rightarrow TC$ : "go!"
2.  $TC \rightarrow \text{Seoul (S), Mumbai (M)}$ : "prepare!"
3.  $S, M \rightarrow P$ : "yes" [*why always yes?*]
4.  $TC \rightarrow S, M$ : "commit!"
  - TC sends **commit**
5.  $TC \rightarrow C$ : "okay"
  - S, M commit on receipt of commit message
  - *Why do we still need the commit?*

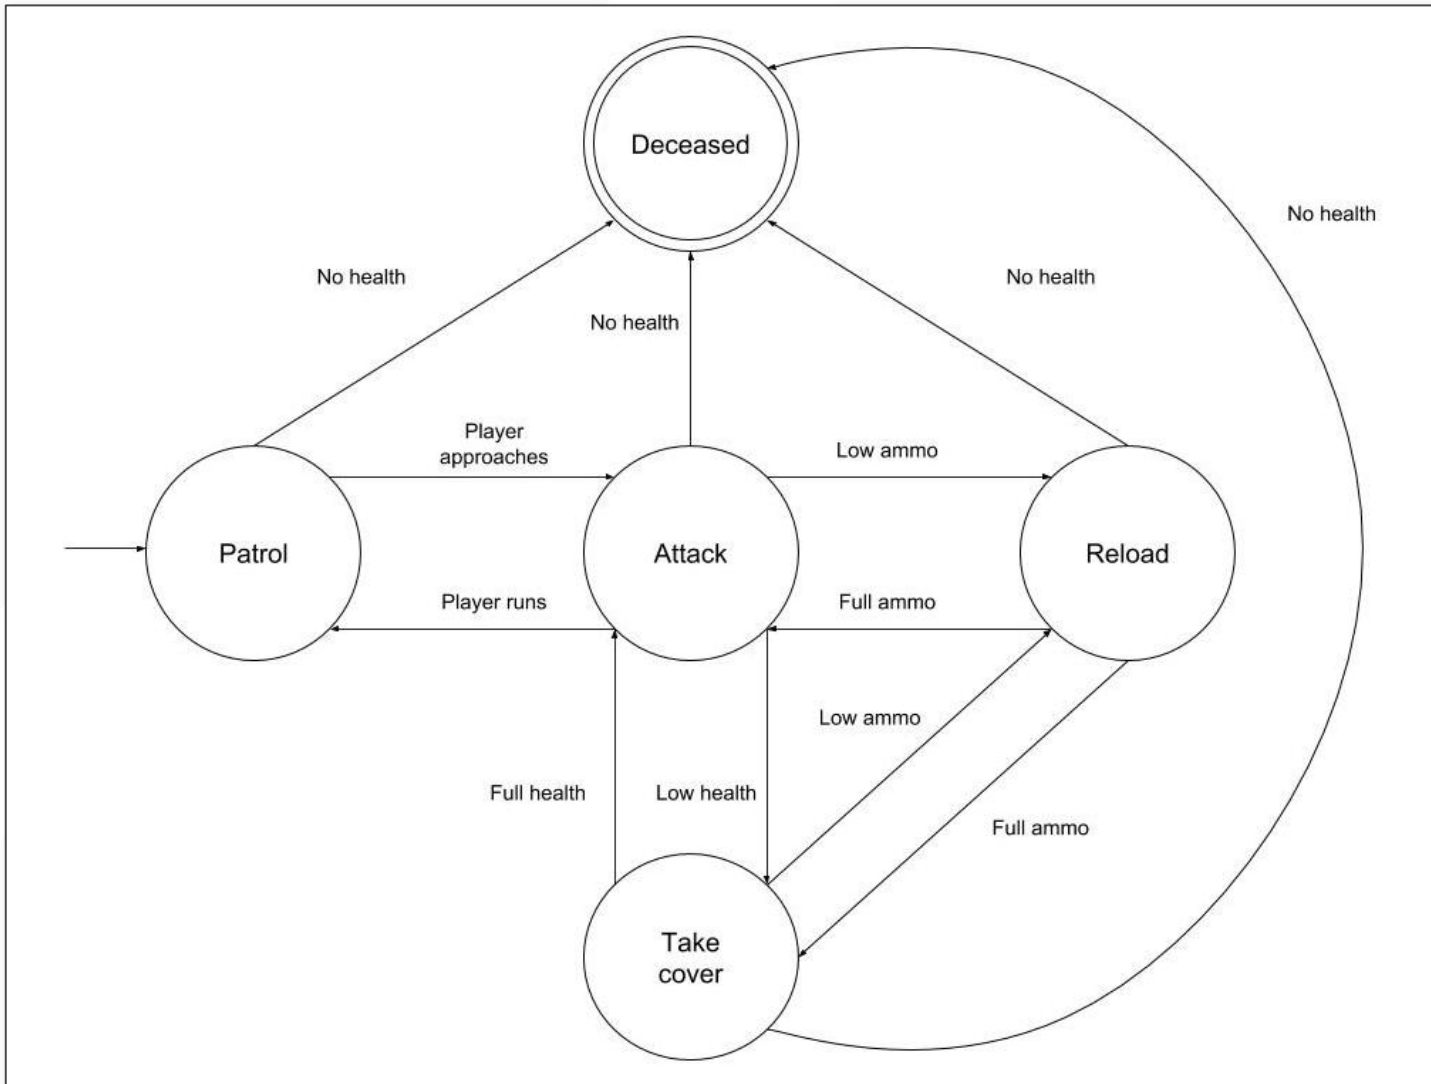
# NETWORK PARTITIONS

- Some failure (either network or host) keeps replicas from communicating with one another
- Two-phase commit (even if we assume all replicas agree) only works if all nodes can be contacted
- How to proceed with read/write transactions in case where not all replicas can be contacted?

# QUORUM-BASED PROTOCOLS

- Idea: Tell client that a file's version is updated after a majority of SurfStoreServers get the update
- Form a “read quorum” of size  $N_R$ 
  - Contact  $N_R$  servers and read all their versions
  - Select highest version as the “correct” version
- Form a “write quorum” of size  $N_W$ 
  - Contact  $N_W$  servers
  - Increment the highest version from that set
  - Write out that new version to the servers in the write quorum

# REPRESENTING A SERVICE AS A FINITE STATE MACHINE






# STATE MACHINE REPLICATION

- Any server is essentially a *state machine*
  - Operations **transition** between states
- Need an op to be executed on all replicas, or none at all
  - *i.e.*, we need **distributed all-or-nothing atomicity**
  - If op is deterministic, replicas will end in same state


# Two phase commit protocol



1.  $C \rightarrow P$ : *"request <op>"*
2.  $P \rightarrow A, B$ : *"prepare <op>"*
3.  $A, B \rightarrow P$ : *"prepared" or "error"*
4.  $P \rightarrow C$ : *"result exec<op>" or "failed"*
5.  $P \rightarrow A, B$ : *"commit <op>"*

What if primary fails?  
Backup fails?



# Two phase commit protocol



1.  $C \rightarrow P$ : *"request <op>"*
2.  $P \rightarrow A, B$ : *"prepare <op>"*
3.  $A, B \rightarrow P$ : *"prepared" or "error"*
4.  $P \rightarrow C$ : *"result exec<op>" or "failed"*
5.  $P \rightarrow A, B$ : *"commit <op>"*

"Okay" (i.e., op is stable) if written to  $> \frac{1}{2}$  backups

# Two phase commit protocol



- Commit sets always overlap  $\geq 1$  node
- Any  $> 1/2$  nodes guaranteed to see committed op


# CONSTANTS AND CONSTRAINTS

- $N$ : Total #Replicas
- $N_R$ : #Replicas in Read Quorum
- $N_W$ : #Replicas in Write Quorum
- Constraints for *strong* consistency:
  1.  $N_R + N_W > N$
  2.  $N_W > N/2$

# QUORUM CONSENSUS

- Write operations can be propagated in background to replicas not in quorum
  - Assumes eventual repair of any network partition
- Operations are slowed by the necessity of first gathering a quorum
  - Though previously, all writes had to go to all replicas
    - With quorum system, must only contact subset of replicas

# QUORUMS IN MICROSOFT ACTIVE DIRECTORY

 **Microsoft**

Windows IT Pro Center

Explore ▾ Docs ▾ Downloads ▾ Scripts Support

Docs / Windows Server / Failover Clustering / Deploy / Manage quorum and witnesses

Feedback Edit Share Da

Filter by title

Failover Clustering

What's New in Failover Clustering

> Understand

> Plan

> Deploy

Create a failover cluster

Deploy a two-node file server

> Prestage a cluster in AD DS

**Manage quorum and witnesses**





Deploy a Cloud Witness

Deploy a file share witness

Cluster operating system rolling upgrades

> Manage

## Configure and manage quorum

01/17/2019 • 20 minutes to read • Contributors    

Applies to: Windows Server 2019, Windows Server 2016, Windows Server 2012 R2, Windows Server 2012

This topic provides background and steps to configure and manage the quorum in a Windows Server failover cluster.

### Understanding quorum

The quorum for a cluster is determined by the number of voting elements that must be part of active cluster membership for that cluster to start properly or continue running. For a more detailed explanation, see the [understanding cluster and pool quorum doc](#).

### Quorum configuration options

The quorum model in Windows Server is flexible. If you need to modify the quorum

In this

Under quoru

Quoru config option

Genera recom for qu config

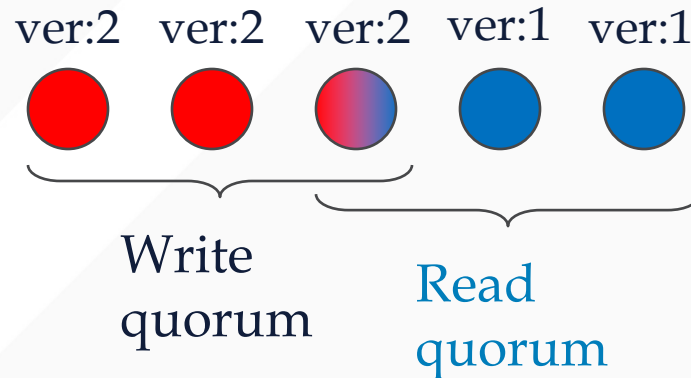
Config cluster

Recover startin quoru

Quoru consid disaste config

More i


# QUORUM EXAMPLE



- 5 replicas, read quorum: 3, write quorum: 3
  - $R+W > 5$  votes ensures overlap between any read/write quorum
- How does this perform for reads?
- How does this perform for writes?




# QUORUM EXAMPLE



- 5 replicas, read quorum: 5, write quorum: 1
  - $R+W > 5$  votes ensures overlap between any read/write quorum
- How does this perform for reads?
- How does this perform for writes?

# QUORUM EXAMPLE



- 5 replicas, read quorum: 1, write quorum: 5
  - $R+W > 5$  votes ensures overlap between any read/write quorum
  - Also called ROWA (read one, write all)
- How does this perform for reads?
- How does this perform for writes?

# STRONGLY CONSISTENT AND EVENTUALLY CONSISTENT EXAMPLES

UC San Diego