Module 2 project milestone

In this project, you are going to build a simple web server that implements a subset of the HTTP/1.1 protocol specification called TritonHTTP (located in Canvas).

TritonHTTP Specification

The TritonHTTP spec is located in Canvas

Project overview

Basic web server functionality

At a high level, a web server listens for connections on a socket (bound to a specific port on a host machine). Clients connect to this socket and use the TritonHTTP protocol to retrieve files from the server. Your server will read data from the client, using the framing and parsing techniques discussed in class to interpret one or more requests (if the client is using pipelined requests). Every time your server reads in a full request, you will service that request and send back a response back to the client. After sending back one (or more) responses, your server will either close the connection (if instructed to do so by the client via the "Connection: close" header, described below, or after an appropriate timeout occurs (also described below). Your web server will then continue waiting for future client connections. Your client should be implemented in a concurrent manner, so that it can process multiple client requests overlapping in time.

Project details

Below are some of the details on implementing your project.

Mapping relative URLs to absolute file paths

Clients make requests to files using a Uniform Resource Locator, such as /images/cyrpto/enigma.jpg. One of the key things to keep in mind in building your web server is that the server must translate that relative URL into an absolute filename on the local filesystem. For example, you might decide to keep all the files for your server in ~aturing/cse101/server/www-files/, which we call the document root. When your server gets a request for the above-mentioned enigma.jpg file, it will prepend the document root to the specified file to get an absolute file name of ~aturing/cse101/server/www-files/images/crypto/enigma.jpg. You need to ensure that malformed or malicious URLs cannot "escape" your document root to access other files. For

example, if a client submits the URL /images/../../../.ssh/id_dsa, they should not be able to download the ~aturing/.ssh/id_dsa file. If a client uses one or more .. directories in such a way that the server would "escape" the document root, you should return a 404 Not Found error back to the client. However, it is valid if the client have URL like ~aturing/cse101/server/../../cse101/server/www-files/images/crypto/enigma.jpg, where the URL does not escape from the document root.

Supporting MIME types

Most applications interpret the contents of files based on their extensions (e.g. homework.txt represents an ASCII text file, whereas icon.jpg represents a JPEG image). Historically, file extensions have not played much of a role with the HTTP protocol, in part because sometimes web servers will dynamically generate content on demand. So, for example, a request line of the form

```
GET /cgi-bin/genhaiku.pl HTTP/1.1
....
```

Might invoke an external helper program (in this case, written in Perl), to generate a haiku poem on demand. But how to interpret the data that is returned? Is it ASCII text? A Word file? An image with the poem written in the center?

Web servers explicitly indicate the type of file via the "Content-Type" header. Examples include:

```
Content-Type: image/jpeg
Content-Type: application/msword
Content-Type: text/plain
```

When serving files out of the document root, your web server will need to convert the file's extension to one of these "MIME types". There is a mapping file called mime.types included in your starter code that you should use to determine the mime type for local files so you can properly set the Content-Type response header. If a file is requested with a file extension not included in the mime.types file, you should return "application/octet-stream".

Executable

Your server binary should be called with ./run-server.sh script and should take one argument which references the configuration file (which may be stored anywhere on the filesystem, so do not assume it is in the current directory). For example:

```
$ ./run-server.sh /home/aturing/myconfig.ini
```

```
$ ./run-server.sh ~aturing/myconfig.ini

$ ./run-server.sh ../configs/myconfig.ini
```

Experimentation

We have provided you with starter code that relies on Go's own internal http web server. You can use this to experiment and explore with what a "real" webserver does under a variety of conditions. Note that the full HTTP specification is several *thousand* pages long, and we certainly do not expect you to implement all of that! We're only focusing on a relatively small subset of the overall protocol. As a result, you may see Go's implementation of the web server doing things we didn't ask you to (for example, Go's web server appends a character set specifier to the Content-Type header and returns a header called "X-Content-Type-Options" which we don't cover in our class. You can just ignore those differences and focus on the subset of HTTP described in the TritonHTTP specification, linked above.

If an aspect of what we're expecting from you is unclear, please ask. We do encourage you to experiment with Go's in-built webserver first before asking about whether your web server "should" act in a certain way or not.

Testing strategies

We have written up a separate document to guide you in developing a testing strategy for your project. This document is located in Canvas.

Grading

Correctness/functionality: 75%

Basic functionality for 200 error code responses (30%)

- This category represents error-free, valid requests that result in a 200 status code.
  - The response headers should be set correctly
  - The response body should match the content
- You should support directories and subdirectories
- "http://server:port/" should be mapped to "http://server:port/index.html"
- "http://server:port/subdirectory/" should be mapped to "http://server:port/subdirectory/index.html" (for any level of nesting)
- You should correctly support the MIME types specified in mime.types

Basic functionality for non-200 error code responses (22.5%):

- Handles 404 for files that aren't found
- Handles 404 for URLs that escape the doc root
- Correctly handles malformed HTTP requests by issuing a 400 error

Concurrency (7.5%):

- Your server should be able to handle concurrent clients using goroutines

Pipelining (15%):

- Your server should be able to handle two or more requests that are pipelined together in the same TCP connection
- Your server should handle the Connection: close header correctly
- Your server should implement a 5-second timeout (refreshed after a successful read) correctly for requests that are not explicitly closed via the Connection: close header.

Testing strategy and completeness: 25%

One quarter of your grade will be determined by the completeness of your testing strategy. For each of the rubric items above, you are to create one or more test cases that you use to determine whether your web server is working correctly for that particular rubric item. The specific form your tests take can vary. On one extreme is developing a complete unit test framework that is fully automated. While that would be the most complete, it is certainly not necessary for this class project. Other examples of testing strategies include shell scripts or other scripts (written in Go, Python, or other languages), or even a text file write up of the command-line tests that you ran along with the output that you saw so that you knew it was correct. Note that it is possible to get full points on the testing strategy even if your actual code doesn't work, as long as it is complete and fully tests the above rubric items. Likewise, it is possible to submit a fully working project and yet not receive full points on the testing strategy if the testing strategy is incomplete.

You should document your testing strategy using a file or files in the *testing* subdirectory. If you used command line tools (e.g. printf and nc), list each of your tests and the output you found and expected in a clearly marked file (e.g. TESTING.md or TESTING.txt, etc). If you have scripts or other code to test your server, put it there and make sure you clearly mark which rubric items are associated with the different tests.

Submitting your work

Log into gradescope.com and upload your code. This assignment is to be done in groups of 1 (solo) or 2. Make sure to log into gradescope by clicking the link in Canvas! This will ensure that your grade "syncs" correctly.

Submission guidelines:

You'll be turning in your project on Gradescope. If you have worked in groups, please add your group members by clicking on "Group Members" option. There are two options to submit your code on Gradescope:

- **GitHub**:
  Choose your GitHub repository and upload the right branch.

- **Upload**:
  Your submission should be a *single ZIP file* named **module2.zip** which should have files with the following structure:
  |-- module2.zip
         |-- README.md
         |-- src
         |-- testing
         |-- sample_htdocs

Make sure that the directory structure remains the same as provided in the starter code.

Do not include large files(video, audio files) in your submission.

Due date/time

Listed on the course calendar/schedule on the Canvas site

Starter code

- [Link to the starter code github invitation](#)